

Open Watcom FORTRAN 77

Programmer's Guide



Version 2.0

Open Watcom

Notice of Copyright

Copyright © 2002-2026 the Open Watcom Contributors. Portions Copyright © 1984-2002 Sybase, Inc. and its subsidiaries. All rights reserved.

Any part of this publication may be reproduced, transmitted, or translated in any form or by any means, electronic, mechanical, manual, optical, or otherwise, without the prior written permission of anyone.

For more information please visit <https://github.com/open-watcom/open-watcom-v2> .

Portions of this manual are reprinted with permission from Tenberry Software, Inc.

Preface

The *Open Watcom FORTRAN 77 Programmer's Guide* includes the following major components:

- DOS Programming Guide
- The DOS/4GW DOS Extender
- Windows 3.x Programming Guide
- Windows NT Programming Guide
- OS/2 Programming Guide
- Novell NLM Programming Guide
- Mixed Language Programming
- Common Problems

Acknowledgements

This book was produced with the Open Watcom GML electronic publishing system, a software tool developed by WATCOM. In this system, writers use an ASCII text editor to create source files containing text annotated with tags. These tags label the structural elements of the document, such as chapters, sections, paragraphs, and lists. The Open Watcom GML software, which runs on a variety of operating systems, interprets the tags to format the text into a form such as you see here. Writers can produce output for a variety of printers, including laser printers, using separately specified layout directives for such things as font selection, column width and height, number of columns, etc. The result is type-set quality copy containing integrated text and graphics.

Many users have provided valuable feedback on earlier versions of the Open Watcom FORTRAN 77 compilers and related tools. Their comments were greatly appreciated. If you find problems in the documentation or have some good suggestions, we would like to hear from you.

July, 1997.

Trademarks Used in this Manual

DOS/4G and DOS/16M are trademarks of Tenberry Software, Inc.

OS/2 is a trademark of International Business Machines Corp. IBM Developer's Toolkit, Presentation Manager, and OS/2 are trademarks of International Business Machines Corp. IBM is a registered trademark of International Business Machines Corp.

Intel and Pentium are registered trademarks of Intel Corp.

Microsoft, Windows and Windows 95 are registered trademarks of Microsoft Corp. Windows NT is a trademark of Microsoft Corp.

NetWare, NetWare 386, and Novell are registered trademarks of Novell, Inc.

Phar Lap, 286|DOS-Extender and 386|DOS-Extender are trademarks of Phar Lap Software, Inc.

UNIX is a registered trademark of The Open Group.

WATCOM is a trademark of Sybase, Inc. and its subsidiaries.

Table of Contents

1 Open Watcom FORTRAN 77 Application Development	1
DOS Programming Guide	3
2 Creating 16-bit DOS Applications	5
2.1 The Sample Application	5
2.2 Building and Running the Sample DOS Application	5
2.3 Debugging the Sample DOS Application	6
3 Creating 32-bit Phar Lap 386 DOS-Extender Applications	9
3.1 The Sample Application	9
3.2 Building and Running the Sample 386 DOS-Extender Application	10
3.3 Debugging the Sample 386 DOS-Extender Application	10
4 Creating 32-bit DOS/4GW Applications	13
4.1 The Sample Application	13
4.2 Building and Running the Sample DOS/4GW Application	13
4.3 Debugging the Sample DOS/4GW Application	14
5 32-bit Extended DOS Application Development	17
5.1 Introduction	17
5.2 How can I write directly to video memory using a DOS extender?	17
5.3 How do I issue interrupts in a DOS/4GW application?	18
5.4 How do I get information about free memory in the 32-bit environment?	19
The DOS/4GW DOS Extender	21
6 The Tenberry Software DOS/4GW DOS Extender	23
7 Linear Executables	25
7.1 The Linear Executable Format	25
7.1.1 The Stub Program	25
7.2 Memory Use	26
8 Configuring DOS/4GW	29
8.1 The DOS4G Environment Variable	29
8.2 Changing the Switch Mode Setting	30
8.3 Fine Control of Memory Usage	31
8.3.1 Specifying a Range of Extended Memory	31
8.3.2 Using Extra Memory	32
8.4 Setting Runtime Options	33
8.5 Controlling Address Line 20	34
9 VMM	35
9.1 VMM Default Parameters	35
9.2 Changing the Defaults	36
9.2.1 The .VMC File	36
10 Interrupt 21H Functions	37
10.1 Functions 25H and 35H: Interrupt Handling in Protected Mode	40

Table of Contents

10.1.1 32-Bit Gates	40
10.1.2 Chaining 16-bit and 32-bit Handlers	41
10.1.3 Getting the Address of the Interrupt Handler	41
11 Interrupt 31H DPMI Functions	43
11.1 Using Interrupt 31H Function Calls	43
11.2 Int31H Function Calls	44
11.2.1 Local Descriptor Table (LDT) Management Services	44
11.2.2 DOS Memory Management Services	49
11.2.3 Interrupt Services	51
11.2.4 Translation Services	53
11.2.5 DPMI Version	60
11.2.6 Memory Management Services	61
11.2.7 Page Locking Services	62
11.2.8 Demand Paging Performance Tuning Services	63
11.2.9 Physical Address Mapping	64
11.2.10 Virtual Interrupt State Functions	65
11.2.11 Vendor Specific Extensions	67
11.2.12 Coprocessor Status	67
12 Utilities	69
12.1 DOS4GW	70
12.2 PMINFO	71
12.3 PRIVATXM	73
12.4 RMINFO	74
13 Error Messages	77
13.1 Kernel Error Messages	77
13.2 DOS/4G Errors	80
14 DOS/4GW Commonly Asked Questions	85
14.1 Differences Within the DOS/4G Product Line	85
14.2 Addressing	88
14.3 Interrupt and Exception Handling	89
14.4 Memory Management	92
14.5 DOS, BIOS, and Mouse Services	92
14.6 Virtual Memory	93
14.7 Debugging	95
14.8 Compatibility	99
Windows 3.x Programming Guide	101
15 Creating 16-bit Windows 3.x Applications	103
15.1 The Sample GUI Application	103
15.2 Building and Running the GUI Application	104
15.3 Debugging the GUI Application	104
16 Porting Non-GUI Applications to 16-bit Windows 3.x	107
16.1 Console Device in a Windowed Environment	107
16.2 The Sample Non-GUI Application	108
16.3 Building and Running the Non-GUI Application	108

Table of Contents

16.4 Debugging the Non-GUI Application	109
16.5 Default Windowing Library Functions	110
17 Creating 32-bit Windows 3.x Applications	113
17.1 The Sample GUI Application	113
17.2 Building and Running the GUI Application	114
17.3 Debugging the GUI Application	115
18 Porting Non-GUI Applications to 32-bit Windows 3.x	117
18.1 Console Device in a Windowed Environment	117
18.2 The Sample Non-GUI Application	118
18.3 Building and Running the Non-GUI Application	118
18.4 Debugging the Non-GUI Application	120
18.5 Default Windowing Library Functions	121
19 The Open Watcom 32-bit Windows 3.x Extender	123
19.1 Pointers	123
19.2 Implementation Overview	124
19.3 System Structure	125
19.4 System Overview	126
19.5 Steps to Obtaining a 32-bit Application	127
20 Windows 3.x 32-bit Programming Overview	129
20.1 WINAPI.FI	129
20.2 Environment Notes	130
20.3 Floating-point Emulation	130
20.4 Multiple Instances	130
20.5 Pointer Handling	131
20.5.1 When To Convert Incoming Pointers	132
20.5.2 When To Convert Outgoing Pointers	132
20.5.2.1 SendMessage and SendDlgItemMessage	133
20.5.3 GlobalAlloc and LocalAlloc	134
20.5.4 Callback Function Pointers	134
20.5.4.1 Window Sub-classing	136
20.6 Calling 16-bit DLLs	137
20.7 _16 Functions	138
21 Windows 32-Bit Dynamic Link Libraries	139
21.1 Introduction to 32-Bit DLLs	139
21.2 A Sample 32-bit DLL	140
21.3 Calling Functions in a 32-bit DLL from a 16-bit Application	141
21.4 Calling Functions in a 32-bit DLL from a 32-bit Application	143
21.5 A Sample 32-bit DLL Using a Structure	144
21.6 Creating and Debugging Dynamic Link Libraries	147
21.6.1 Building the Applications	148
21.6.2 Installing the Examples under Windows	148
21.6.3 Running the Examples	149
21.6.4 Debugging a 32-bit DLL	149
21.6.5 Summary	150
22 Interfacing Visual Basic and Open Watcom FORTRAN 77 DLLs	151
22.1 Introduction to Visual Basic and DLLs	151

Table of Contents

22.2 A Working Example	152
22.3 Sample Visual Basic DLL Programs	154
22.3.1 Source Code for VBDLL32.DLL	154
22.3.2 Source code for COVER16.DLL	155
22.4 Compiling and Linking the Examples	156
23 WIN386 Library Subprograms	159
AllocAlias16	160
AllocHugeAlias16	161
_Call16	162
DefineDLLEntry	164
DefineUserProc16	166
FreeAlias16	168
FreeHugeAlias16	169
FreeIndirectFunctionHandle	170
GetIndirectFunctionHandle	172
GetProcAddress	174
InvokeIndirectFunction	176
MapAliasToFlat	178
PASS_WORD_AS_POINTER	179
ReleaseProc16	180
24 32-bit Extended Windows Application Development	181
24.1 Can you call 16-bit code from a 32-bit code?	181
24.2 How do I add my Windows resources?	181
24.3 What size of function pointers passed to Windows?	182
24.4 Why are 32-bit callback routines FAR?	182
24.5 Why use the _16 API functions?	182
25 Special Windows API Functions	183
Windows NT Programming Guide	187
26 Windows NT Programming Overview	189
26.1 Windows NT Character-mode Versus GUI	189
27 Creating Windows NT Character-mode Applications	191
27.1 The Sample Character-mode Application	191
27.2 Building and Running the Character-mode Application	192
27.3 Debugging the Character-mode Application	192
28 Windows NT Multi-threaded Applications	195
28.1 Programming Considerations	195
28.2 Creating Threads	195
28.2.1 Creating a New Thread	195
28.2.2 Terminating the Current Thread	196
28.2.3 Getting the Current Thread Identifier	196
28.3 A Multi-threaded Example	196
29 Windows NT Dynamic Link Libraries	199
29.1 Creating Dynamic Link Libraries	199

Table of Contents

29.2 Creating a Sample Dynamic Link Library	200
29.3 Using Dynamic Link Libraries	202
29.4 The Dynamic Link Library Data Area	203
29.5 Dynamic Link Library Initialization/Termination	204
OS/2 Programming Guide	207
30 Creating 16-bit OS/2 1.x Applications	209
30.1 The Sample Application	209
30.2 Building and Running the Sample OS/2 1.x Application	210
30.3 Debugging the Sample OS/2 1.x Application	210
31 Creating 32-bit OS/2 Applications	213
31.1 The Sample Application	213
31.2 Building and Running the Sample OS/2 Application	214
31.3 Debugging the Sample OS/2 Application	214
32 OS/2 2.x Multi-threaded Applications	217
32.1 Programming Considerations	217
32.2 Creating Threads	217
32.2.1 Creating a New Thread	217
32.2.2 Terminating the Current Thread	218
32.2.3 Getting the Current Thread Identifier	218
32.3 A Multi-threaded Example	218
32.4 Thread Limits	220
33 OS/2 2.x Dynamic Link Libraries	221
33.1 Creating Dynamic Link Libraries	221
33.2 Creating a Sample Dynamic Link Library	222
33.3 Using Dynamic Link Libraries	223
33.4 The Dynamic Link Library Data Area	224
33.5 Dynamic Link Library Initialization/Termination	225
34 Programming for OS/2 Presentation Manager	227
34.1 Porting Existing FORTRAN 77 Applications	227
34.1.1 An Example	227
34.2 Calling Presentation Manager API Functions	228
Novell NLM Programming Guide	233
35 Creating NetWare 386 NLM Applications	235
Mixed Language Programming	237
36 Inter-Language calls: C and FORTRAN	239
36.1 Symbol Naming Convention	239
36.2 Argument Passing Convention	240
36.3 Memory Model Compatibility	240
36.4 Linking Considerations	241

Table of Contents

36.5 Integer Type Compatibility	241
36.6 How do I pass integers from C to a FORTRAN function?	241
36.7 How do I pass integers from FORTRAN to a C function?	242
36.8 How do I pass a string from a C function to FORTRAN?	243
36.9 How do I pass a string from FORTRAN to a C function?	244
36.10 How do I access a FORTRAN common block from within C?	245
36.11 How do I call a C function that accepts a variable number of arguments?	246
 Common Problems	 247
 37 Commonly Asked Questions and Answers	 249
37.1 Converting to Open Watcom F77	249
37.2 What you should know about optimization	250
37.3 Reading a stream of binary data from a file	251
37.4 Redefining math error handling with Open Watcom F77	251
37.5 The compiler cannot find my include files	257
37.6 The linker reports a "stack segment not found" error	257
37.7 Resolving an "Undefined Reference" linker error	257
37.8 Why local variable values are not maintained between subprogram calls	258
37.9 What "Stack Overflow!" means	258
37.10 What are the probable causes of a General Protection Fault in 32-bit applications?	259
37.11 Which floating-point compiler option should I use for my application?	260
37.12 How more than 20 files at a time can be opened	261
37.13 How source files can be seen in the debugger	262
37.14 The difference between the "d1" and "d2" compiler options	264
37.15 The difference between the "debug" and "d2" compiler options	264

List of Figures

Figure 1. Basic Memory Layout	27
Figure 2. Physical Memory/Linear Address Space	28
Figure 3. Access Rights/Type	47
Figure 4. Extended Access Rights/Type	48
Figure 5. WIN386 Structure	125
Figure 6. 32-bit Application Structure	126

1 Open Watcom FORTRAN 77 Application Development

This document contains guides to application development for several environments including 16-bit DOS, 32-bit extended DOS, Windows 3.x, 32-bit extended Windows 3.x, Windows NT/2000/XP, Win9x, OS/2, and Novell NLMS. It also describes mixed language (C, FORTRAN) application development. It concludes with a chapter on some general questions and the answers to them.

This document covers the following topics:

- DOS Programming Guide

Creating 16-bit DOS Applications
Creating 32-bit Phar Lap 386 DOS-Extender Applications
Creating 32-bit DOS/4GW Applications
32-bit Extended DOS Application Development

- The DOS/4GW DOS Extender

The Tenberry Software DOS/4GW DOS Extender
Linear Executables
Configuring DOS/4GW
VMM
Interrupt 21H Functions
Interrupt 31H DPMI Functions
Utilities
Error Messages
DOS/4GW Commonly Asked Questions

- Windows 3.x Programming Guide

Creating 16-bit Windows 3.x Applications
Porting Non-GUI Applications to 16-bit Windows 3.x
Creating 32-bit Windows 3.x Applications
Porting Non-GUI Applications to 32-bit Windows 3.x
The Open Watcom 32-bit Windows Extender
Windows 3.x 32-bit Programming Overview
Windows 32-Bit Dynamic Link Libraries
Interfacing Visual Basic and Open Watcom FORTRAN 77 DLLs
WIN386 Library Subprograms
32-bit Extended Windows Application Development
Special Windows API Functions

- Windows NT Programming Guide

Windows NT Programming Overview
Creating Windows NT GUI Applications
Porting Non-GUI Applications to Windows NT GUI
Windows NT Multi-threaded Applications
Windows NT Dynamic Link Libraries

- OS/2 Programming Guide

Creating 16-bit OS/2 1.x Applications
Creating 32-bit OS/2 Applications
OS/2 Multi-threaded Applications
OS/2 Dynamic Link Libraries
Programming for OS/2 Presentation Manager

- Novell NLM Programming Guide

Creating NetWare 386 NLM Applications

- Mixed Language Programming

Inter-Language calls: C and FORTRAN

- Common Problems

Commonly Asked Questions and Answers

DOS Programming Guide

2 Creating 16-bit DOS Applications

This chapter describes how to compile and link 16-bit DOS applications simply and quickly.

We will illustrate the steps to creating 16-bit DOS applications by taking a small sample application and showing you how to compile, link, run and debug it.

2.1 The Sample Application

To demonstrate the creation of 16-bit DOS applications using command-line oriented tools, we introduce a simple sample program. For our example, we are going to use the "sieve" program.

- * This program computes the prime numbers between 1 and 10,000
- * using the Sieve of Eratosthenes algorithm.

```
IMPLICIT NONE
INTEGER UPBOUND
PARAMETER (UPBOUND=10000)
INTEGER I, K, PRIMES
LOGICAL*1 NUMBERS(2:UPBOUND)
CHARACTER*11 FORM
PARAMETER (FORM='(A,I5,A,I5)')
DO I = 2, UPBOUND
    NUMBERS(I) = .TRUE.
ENDDO
PRIMES = 0
DO I = 2, UPBOUND
    IF( NUMBERS(I) ) THEN
        PRIMES = PRIMES + 1
        DO K = I + I, UPBOUND, I
            NUMBERS(K) = .FALSE.
        ENDDO
    ENDIF
ENDDO
PRINT FORM, 'The Number of Primes between 1 and ', UPBOUND,
1      ' are: ', PRIMES
END
```

The goal of this program is to count the prime numbers between 1 and 10,000. It uses the famous *Sieve of Eratosthenes* algorithm to accomplish this task. We will take you through the steps necessary to produce this result.

2.2 Building and Running the Sample DOS Application

To compile and link our example program which is stored in the file `sieve.for`, enter the following command:

```
C>wfl -l=dos sieve.for
```

The typical messages that appear on the screen are shown in the following illustration.

```
C>wfl -l=dos sieve.for
Open Watcom F77/16 Compile and Link Utility
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1990-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
    wfc sieve.for
Open Watcom FORTRAN 77/16 Optimizing Compiler
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1984-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
sieve.for: 21 statements, 311 bytes, 6 extensions, 0 warnings, 0 errors

Open Watcom Linker
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1985-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
loading object files
searching libraries
creating a DOS executable
```

Provided that no errors were encountered during the compile or link phases, the "sieve" program may now be run.

```
C>sieve
The Number of Primes between 1 and 10000 are: 1229
```

If you examine the current directory, you will find that two files have been created. These are `sieve.obj` (the result of compiling `sieve.for`) and `sieve.exe` (the result of linking `sieve.obj` with the appropriate Open Watcom FORTRAN 77 libraries). It is `sieve.exe` that is run by DOS when you enter the "sieve" command.

2.3 Debugging the Sample DOS Application

Let us assume that you wish to debug your application in order to locate an error in programming. In the previous section, the "sieve" program was compiled with default compile and link options. When debugging an application, it is useful to refer to the symbolic names of routines and variables. It is also convenient to debug at the source line level rather than the machine language level. To do this, we must direct both the compiler and linker to include additional debugging information in the object and executable files. Using the **WFL** command, this is fairly straightforward. **WFL** recognizes the Open Watcom F77 compiler "debug" options and will create the appropriate debug directives for the Open Watcom Linker.

For example, to compile and link the "sieve" program with debugging information, the following command may be issued.

```
C>wfl -l=dos -d2 sieve.for
```

The typical messages that appear on the screen are shown in the following illustration.

```
C>wfl -l=dos -d2 sieve.for
Open Watcom F77/16 Compile and Link Utility
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1990-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
    wfc sieve.for -d2
Open Watcom FORTRAN 77/16 Optimizing Compiler
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1984-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
sieve.for: 21 statements, 392 bytes, 6 extensions, 0 warnings, 0 errors

Open Watcom Linker
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1985-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
loading object files
searching libraries
creating a DOS executable
```

The "d2" option requests the maximum amount of debugging information that can be provided by the Open Watcom F77 compiler. **WFL** will make sure that this debugging information is included in the executable file that is produced by the linker.

The "bytes" value is larger than in the previous example since selection of the "d2" option results in fewer code optimizations by default. You can request more optimization by specifying the appropriate options. However, you do so at the risk of making it more difficult for yourself to determine the relationship between the object code and the original source language code.

To request the Open Watcom Debugger to assist in debugging the application, the following command may be issued.

```
C>wd sieve
```

It would be too ambitious to describe the debugger in this introductory chapter so we refer you to the book entitled *Open Watcom Debugger User's Guide*.

3 Creating 32-bit Phar Lap 386/DOS-Extender Applications

This chapter describes how to compile and link 32-bit Phar Lap 386/DOS-Extender applications simply and quickly.

We will illustrate the steps to creating 32-bit Phar Lap 386/DOS-Extender applications by taking a small sample application and showing you how to compile, link, run and debug it.

3.1 The Sample Application

To demonstrate the creation of 32-bit Phar Lap 386/DOS-Extender applications using command-line oriented tools, we introduce a simple sample program. For our example, we are going to use the "sieve" program.

```
* This program computes the prime numbers between 1 and 10,000
* using the Sieve of Eratosthenes algorithm.

      IMPLICIT NONE
      INTEGER UBOUND
      PARAMETER (UBOUND=10000)
      INTEGER I, K, PRIMES
      LOGICAL*1 NUMBERS(2:UBOUND)
      CHARACTER*11 FORM
      PARAMETER (FORM=' (A,I5,A,I5)')
      DO I = 2, UBOUND
         NUMBERS(I) = .TRUE.
      ENDDO
      PRIMES = 0
      DO I = 2, UBOUND
         IF( NUMBERS(I) )THEN
            PRIMES = PRIMES + 1
            DO K = I + I, UBOUND, I
               NUMBERS(K) = .FALSE.
            ENDDO
         ENDIF
      ENDDO
      PRINT FORM, 'The Number of Primes between 1 and ', UBOUND,
1      ' are: ', PRIMES
      END
```

The goal of this program is to count the prime numbers between 1 and 10,000. It uses the famous *Sieve of Eratosthenes* algorithm to accomplish this task. We will take you through the steps necessary to produce this result.

3.2 Building and Running the Sample 386/DOS-Extender Application

To compile and link our example program which is stored in the file `sieve.for`, enter the following command:

```
C>wfl386 -l=pharlap sieve.for
```

The typical messages that appear on the screen are shown in the following illustration.

```
C>wfl386 -l=pharlap sieve.for
Open Watcom F77/32 Compile and Link Utility
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1990-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
    wfc386 sieve.for
Open Watcom FORTRAN 77/32 Optimizing Compiler
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1984-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
sieve.for: 21 statements, 172 bytes, 6 extensions, 0 warnings, 0 errors

Open Watcom Linker
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1985-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
loading object files
searching libraries
creating a Phar Lap simple executable
```

Provided that no errors were encountered during the compile or link phases, the "sieve" program may now be run.

```
C>run386 sieve
The Number of Primes between 1 and 10000 are: 1229
```

If you examine the current directory, you will find that two files have been created. These are `sieve.obj` (the result of compiling `sieve.for`) and `sieve.exp` (the result of linking `sieve.obj` with the appropriate Open Watcom FORTRAN 77 libraries). It is `sieve.exp` that is run by DOS when you enter the "run386 sieve" command.

3.3 Debugging the Sample 386/DOS-Extender Application

Let us assume that you wish to debug your application in order to locate an error in programming. In the previous section, the "sieve" program was compiled with default compile and link options. When debugging an application, it is useful to refer to the symbolic names of routines and variables. It is also convenient to debug at the source line level rather than the machine language level. To do this, we must direct both the compiler and linker to include additional debugging information in the object and executable files. Using the **WFL386** command, this is fairly straightforward. **WFL386** recognizes the Open Watcom F77 compiler "debug" options and will create the appropriate debug directives for the Open Watcom Linker.

For example, to compile and link the "sieve" program with debugging information, the following command may be issued.

```
C>wfl386 -l=pharlap -d2 sieve.for
```

The typical messages that appear on the screen are shown in the following illustration.

```
C>wfl386 -l=pharlap -d2 sieve.for
Open Watcom F77/32 Compile and Link Utility
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1990-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
      wfc386 sieve.for  -d2
Open Watcom FORTRAN 77/32 Optimizing Compiler
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1984-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
sieve.for: 21 statements, 237 bytes, 6 extensions, 0 warnings, 0 errors

Open Watcom Linker
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1985-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
loading object files
searching libraries
creating a Phar Lap simple executable
```

The "d2" option requests the maximum amount of debugging information that can be provided by the Open Watcom F77 compiler. **WFL386** will make sure that this debugging information is included in the executable file that is produced by the linker.

The "bytes" value is larger than in the previous example since selection of the "d2" option results in fewer code optimizations by default. You can request more optimization by specifying the appropriate options. However, you do so at the risk of making it more difficult for yourself to determine the relationship between the object code and the original source language code.

To request the Open Watcom Debugger to assist in debugging the application, the following command may be issued.

```
C>wd /trap=pls sieve
```

It would be too ambitious to describe the debugger in this introductory chapter so we refer you to the book entitled *Open Watcom Debugger User's Guide*.

4 Creating 32-bit DOS/4GW Applications

This chapter describes how to compile and link 32-bit DOS/4GW applications simply and quickly.

We will illustrate the steps to creating 32-bit DOS/4GW applications by taking a small sample application and showing you how to compile, link, run and debug it.

4.1 The Sample Application

To demonstrate the creation of 32-bit DOS/4GW applications using command-line oriented tools, we introduce a simple sample program. For our example, we are going to use the "sieve" program.

- * This program computes the prime numbers between 1 and 10,000
- * using the Sieve of Eratosthenes algorithm.

```
IMPLICIT NONE
INTEGER UPBOUND
PARAMETER (UPBOUND=10000)
INTEGER I, K, PRIMES
LOGICAL*1 NUMBERS(2:UPBOUND)
CHARACTER*11 FORM
PARAMETER (FORM='(A,I5,A,I5)')
DO I = 2, UPBOUND
    NUMBERS(I) = .TRUE.
ENDDO
PRIMES = 0
DO I = 2, UPBOUND
    IF( NUMBERS(I) ) THEN
        PRIMES = PRIMES + 1
        DO K = I + I, UPBOUND, I
            NUMBERS(K) = .FALSE.
        ENDDO
    ENDIF
ENDDO
PRINT FORM, 'The Number of Primes between 1 and ', UPBOUND,
1      ' are: ', PRIMES
END
```

The goal of this program is to count the prime numbers between 1 and 10,000. It uses the famous *Sieve of Eratosthenes* algorithm to accomplish this task. We will take you through the steps necessary to produce this result.

4.2 Building and Running the Sample DOS/4GW Application

To compile and link our example program which is stored in the file `sieve.for`, enter the following command:

```
C>wfl386 -l=dos4g sieve.for
```

The typical messages that appear on the screen are shown in the following illustration.

```
C>wfl386 -l=dos4g sieve.for
Open Watcom F77/32 Compile and Link Utility
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1990-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
    wfc386 sieve.for
Open Watcom FORTRAN 77/32 Optimizing Compiler
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1984-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
sieve.for: 21 statements, 172 bytes, 6 extensions, 0 warnings, 0 errors

Open Watcom Linker
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1985-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
loading object files
searching libraries
creating a DOS/4G executable
```

Provided that no errors were encountered during the compile or link phases, the "sieve" program may now be run.

```
C>sieve
The Number of Primes between 1 and 10000 are: 1229
```

If you examine the current directory, you will find that two files have been created. These are `sieve.obj` (the result of compiling `sieve.for`) and `sieve.exe` (the result of linking `sieve.obj` with the appropriate Open Watcom FORTRAN 77 libraries). It is `sieve.exe` that is run by DOS when you enter the "sieve" command.

4.3 Debugging the Sample DOS/4GW Application

Let us assume that you wish to debug your application in order to locate an error in programming. In the previous section, the "sieve" program was compiled with default compile and link options. When debugging an application, it is useful to refer to the symbolic names of routines and variables. It is also convenient to debug at the source line level rather than the machine language level. To do this, we must direct both the compiler and linker to include additional debugging information in the object and executable files. Using the **WFL386** command, this is fairly straightforward. **WFL386** recognizes the Open Watcom F77 compiler "debug" options and will create the appropriate debug directives for the Open Watcom Linker.

For example, to compile and link the "sieve" program with debugging information, the following command may be issued.

```
C>wfl386 -l=dos4g -d2 sieve.for
```

The typical messages that appear on the screen are shown in the following illustration.

```
C>wfl386 -l=dos4g -d2 sieve.for
Open Watcom F77/32 Compile and Link Utility
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1990-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
    wfc386 sieve.for -d2
Open Watcom FORTRAN 77/32 Optimizing Compiler
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1984-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
sieve.for: 21 statements, 237 bytes, 6 extensions, 0 warnings, 0 errors

Open Watcom Linker
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1985-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
loading object files
searching libraries
creating a DOS/4G executable
```

The "d2" option requests the maximum amount of debugging information that can be provided by the Open Watcom F77 compiler. **WFL386** will make sure that this debugging information is included in the executable file that is produced by the linker.

The "bytes" value is larger than in the previous example since selection of the "d2" option results in fewer code optimizations by default. You can request more optimization by specifying the appropriate options. However, you do so at the risk of making it more difficult for yourself to determine the relationship between the object code and the original source language code.

To request the Open Watcom Debugger to assist in debugging the application, the following command may be issued.

```
C>wd /trap=rsi sieve
```

It would be too ambitious to describe the debugger in this introductory chapter so we refer you to the book entitled *Open Watcom Debugger User's Guide*.

5 32-bit Extended DOS Application Development

5.1 Introduction

The purpose of this chapter is to anticipate common programming questions for 32-bit extended DOS application development. Note that these programming solutions may be DOS-extender specific and therefore may not work for other DOS extenders.

The following topics are discussed in this chapter:

- How can I write directly to video memory using DOS/4GW?
- How do I issue interrupts in a DOS/4GW application?
- How do I get information about free memory with DOS/4GW?

Please refer to the *DOS Protected-Mode Interface (DPMI) Specification* for information on DPMI services.

The DPMI 1.0 specification can be obtained from the Open Watcom website <https://open-watcom.github.io/open-watcom-1.9/devel/dpmi10.pdf>. This file contains a PDF version of the DPMI 1.0 specification.

Alternatively the DPMI 1.0 specification can be obtained at <http://www.delorie.com/djgpp/doc/dpmi/dpmi-specs.tar.gz>. This file contains online HTML version of the DPMI 1.0 specification.

5.2 How can I write directly to video memory using a DOS extender?

Many programmers require access to video RAM in order to directly manipulate data on the screen. Under DOS, it was standard practice to use a far pointer, with the segment part of the far pointer set to the screen segment. Under DOS extenders, this practice is not so standard. Each DOS extender provides its own method for accessing video memory. The following program demonstrates the method used with DOS/4GW.

```
* FSCREEN.FOR
* The following program shows how to access screen memory
* from a FORTRAN program under the DOS/4GW DOS extender.

* Compile & Link: wfl386 -l=dos4g fscreen

      program screen

* Allocatable arrays must be declared by specifying their
* dimensions using colons only (see Open Watcom FORTRAN 77
* Language Reference on the ALLOCATE statement for details).

      character*1 screen(:, :)
      integer SCRSIZE, i

      parameter ( SCRSIZE = 80*25 )
* Under DOS/4GW, the first megabyte of physical memory - the
* real memory - is mapped as a shared linear address space.
* This allows your application to access video RAM using its
* linear address. The DOS segment:offset of B800:0000
* corresponds to a linear address of B8000.

      allocate( screen(0:1,0:SCRSIZE-1), location='B8000'x )

      do i = 0, SCRSIZE - 1
         screen(0,i) = '*'
      enddo

      end
```

5.3 How do I issue interrupts in a DOS/4GW application?

The Open Watcom F77 library files contain the FINTR and the FINTRF subroutines which allows the user to perform interrupt calls within a FORTRAN 77 program. This subroutine is described in the Subprogram Library section of the *Open Watcom FORTRAN 77 User's Guide*.

The following sample program illustrates the use of the FINTR subroutine to set up the register information required for Interrupt 21h. The register information is loaded into the regs structure. This structure is defined in the DOS.FI file located in the \WATCOM\SRC\FORTRAN\DOS directory. Assign values to the register elements according to the interrupt call requirements. For example, Interrupt 21h, function 4Eh needs valid values for the AH, ECX, DS and EDX to set up the registers for the Interrupt 21h call. This procedure can be used to perform any interrupt calls that are supported in protected mode by DOS/4GW.

```
* DTA.FOR
* This program demonstrates the use of the FINTR
* function to list the files of the current directory.
* Interrupt 21 Functions for FIND FIRST, FIND NEXT,
* and GET DTA are used.

* Compile & Link: set finclude=\watcom\src\fortran\dos
*                wfl386 -l=dos4g dta

*$pragma aux GetDS = "mov ax,ds" value [ax]

      program dta
      implicit integer*2 (i-n)
      integer*2 res
      integer*2 GetDS
      integer*4 dir,addr
      integer*1 dta(:)
      character fname*1(12), fname2*12
      equivalence (fname, fname2)
```

```

* DTA is declared as a FAR array.  When referencing an array
* element, the pointer to the array is a FAR pointer.  With a
* character variable, the result is a pointer to a string
* control block (SCB).  The run-time library expects the SCB
* to contain a near pointer.  To get around the problem, we
* define the DTA as a byte array, then use the CHAR function
* to get the character equivalent for printing a filename.

*$pragma array dta far
    include 'dos.fi'
*
* Listing of current directory
*
    call fsystem( 'dir/w *.*'//char(0) )
    dir = loc( '*.*/'//char(0) )

    i = 0
10    i = i + 1
    if( i .eq. 1 )then
*
* Find first file
*
        AH = '4E'x
        ECX = 0
        DS = GetDS()
        EDX = dir
    else
*
* Find next file
*
        AH = '4F'x
    endif
    call fintr( '21'x, regs )
    res = AX
    if( res .eq. 0 )then
*
* Extract filename from DTA
*
        AH = '2F'x
        call fintr( '21'x, regs )

        addr = ISHL( IAND( INT( ES ), '0000FFFF'x ), 16 )
        addr = IOR( addr, IAND( INT( BX ), '0000FFFF'x ) )
        allocate( dta(0:42), location=addr )
        fname2 = ' '
        do j = 30, 41
            if( dta(j) .eq. 0 ) goto 20
            fname(j - 29) = char( dta(j) )
        enddo
20    print *, fname2
        deallocate( dta )
        goto 10
    endif
end
end

```

5.4 How do I get information about free memory in the 32-bit environment?

Under a virtual memory system, programmers are often interested in the amount of physical memory they can allocate. Information about the amount of free memory that is available is always provided under a DPMI host, however, the manner in which this information is provided may differ under various environments. Keep in mind that in a multi-tasking environment, the information returned to your task from the DPMI host can easily become obsolete if other tasks allocate memory independently of your task.

DOS/4GW provides a DPMI interface through interrupt 0x31. This allows you to use DPMI service 0x0500 to get free memory information. The following program illustrates this procedure.

```
* FMEMORY.FOR
* This example shows how to get information about free
* memory using DPMI call 0500h under DOS/4GW using Open Watcom
* FORTRAN 77. Note that only the first field of the
* structure is guaranteed to contain a valid value; any
* field not returned by DOS/4GW is set to -1 (0FFFFFFFh).

* Compile & Link:  set finclude=\watcom\src\fortran\dos
*                  wfl386 -l=dos4g fmemory

* Pragma to get the default data segment

*$pragma aux GetDS = "mov ax,ds" value [ax] modify exact [ax]

      program memory
      implicit none
      include 'dos.fi'

      structure /meminfo/
         integer*4 LargestBlockAvail
         integer*4 MaxUnlockedPage
         integer*4 LargestLockablePage
         integer*4 LinAddrSpace
         integer*4 NumFreePagesAvail
         integer*4 NumPhysicalPagesFree
         integer*4 TotalPhysicalPages
         integer*4 FreeLinAddrSpace
         integer*4 SizeOfPageFile
         integer*4 Reserved1
         integer*4 Reserved2
      end structure

* Set up the register information for the interrupt call

      record /meminfo/ MemInfo
         integer interrupt_no
         integer*2 GetDS

      parameter( interrupt_no='31'x)
      DS = FS = GS = 0
      EAX = '00000500'x
      ES = GetDS()
      EDI = loc(MemInfo)

      call fintr( interrupt_no, regs)
* Report the information returned by the DPMI host

      print *, '-----'
      print *, 'Largest available block (in bytes): ',
&      MemInfo.LargestBlockAvail
      print *, 'Maximum unlocked page allocation: ',
&      MemInfo.MaxUnlockedPage
      print *, 'Pages that can be allocated and locked: ',
&      MemInfo.LargestLockablePage
      print *, 'Total linear address space including' //
&      ' allocated pages:', MemInfo.LinAddrSpace
      print *, 'Number of free pages available: ',
&      MemInfo.NumFreePagesAvail
      print *, 'Number of physical pages not in use: ',
&      MemInfo.NumPhysicalPagesFree
      print *, 'Total physical pages managed by host: ',
&      MemInfo.TotalPhysicalPages
      print *, 'Free linear address space (pages): ',
&      MemInfo.FreeLinAddrSpace
      print *, 'Size of paging/file partition (pages): ',
&      MemInfo.SizeOfPageFile

      end
```


The DOS/4GW DOS Extender

6 The Tenberry Software *DOS/4GW* DOS Extender

The chapters in this section describe the 32-bit Tenberry Software *DOS/4GW* DOS Extender which is provided with the Open Watcom F77 package. *DOS/4GW* is a subset of Tenberry Software's *DOS/4G* product. *DOS/4GW* is customized for use with the Open Watcom F77 package. Key differences are:

- *DOS/4GW* will only execute programs built with a Open Watcom 32-bit compiler such as Open Watcom F77 and linked with its run-time libraries.
- The *DOS/4GW* virtual memory manager (VMM), included in the package, is restricted to 32MB of memory.
- *DOS/4GW* does not provide extra functionality such as TSR capability and VMM performance tuning enhancements.

If your application has requirements beyond those provided by *DOS/4GW*, you may wish to acquire *DOS/4GW* Professional or *DOS/4G* from:

Tenberry Software, Inc.
PO Box 20050
Fountain Hills, Arizona
U.S.A 85269-0050

WWW: <http://www.tenberry.com/dos4g/>
Email: info@tenberry.com
Phone: 1.480.767.8868
Fax: 1.480.767.8709

Programs developed to use the restricted version of *DOS/4GW* which is included in the Open Watcom F77 package can be distributed on a royalty-free basis, subject to the licensing terms of the product.

7 Linear Executables

To build a linear executable, compile and link it as described in the chapter entitled "Creating 32-bit DOS/4GW Executables". The resulting file will not run independently: you can run it under the Open Watcom Debugger, Tenberry Software Instant-D debugger, or with the standalone "DOS4GW.EXE".

7.1 The Linear Executable Format

DOS/4GW works with files that use the Linear Executable (LE) file format. The format represents a protected-mode program in the context of a 32-bit 386 runtime environment with linear to physical address translation hardware enabled. It uses a flat address space.

This file format is similar to the Segmented Executable (NE) format used in OS/2 1.x and MS Windows. Both support Dynamic Linking, Resources, and are geared toward protected-mode programs. Both formats use tables of "counted ASCII" names, and they use similar relocation formats.

Both formats begin with a DOS style stub program that sophisticated loaders skip. This stub program executes when the *DOS/4GW* loader is not present, displaying the message, *This program cannot run in DOS mode*.

When the Open Watcom Linker is used to link a *DOS/4GW* application, it automatically replaces the default stub program with one that calls DOS4GW.

7.1.1 The Stub Program

The stub at the beginning of a linear executable is a real-mode program that you can modify as you like. For example, you can:

- make the stub program do a checksum on the "DOS4GW.EXE" file to make sure it's the correct version.
- copy protect your program.
- specify a search path for the "DOS4GW.EXE" file.
- add command line arguments.

The SRC directory contains source code for a sample stub program. "WSTUB.C" is a simple example, a good base to start from when you construct your own stub. Please note that you will require a 16-bit C compiler to compile a new stub program. Following is the code in "WSTUB.C":

```
#include <stdio.h>
#include <stdlib.h>
#include <process.h>
#include <errno.h>
#include <string.h>

/* Add environment strings to be searched here */
char *paths_to_check[] = {
    "DOS4GPATH",
    "PATH"};

char *dos4g_path()
{
    static char fullpath[80];
    int i;

    for( i = 0;
        i < sizeof( paths_to_check ) / sizeof( paths_to_check[0] );
        i++ ) {
        _searchenv( "dos4gw.exe", paths_to_check[i], fullpath );
        if( fullpath[0] ) return( &fullpath );
    }
    for( i = 0;
        i < sizeof( paths_to_check ) / sizeof( paths_to_check[0] );
        i++ ) {
        _searchenv( "dos4g.exe", paths_to_check[i], fullpath );
        if( fullpath[0] ) return( &fullpath );
    }
    return( "dos4gw.exe" );
}

main( int argc, char *argv[] )
{
    char          *av[4];
    auto char      cmdline[128];

    av[0] = dos4g_path();           /* Locate the DOS/4G loader */
    av[1] = argv[0];                /* name of executable to run */
    av[2] = getcmd( cmdline );      /* command line */
    av[3] = NULL;                  /* end of list */
#ifdef QUIET
    putenv( "DOS4G=QUIET" );       /* disables DOS/4G Copyright banner */
#endif
    execvp( av[0], av );
    puts( "Stub exec failed:" );
    puts( av[0] );
    puts( strerror( errno ) );
    exit( 1 );                     /* indicate error */
}
```

If you do not have a C compiler, you can create an assembly language version of the above sample stub program and use it to create your own version of the stub program.

7.2 Memory Use

This section explains how a *DOS/4GW* application uses the memory on a 386-based PC/AT. The basic memory layout of an AT machine consists of 640KB of DOS memory, 384KB of upper memory, and an undetermined amount of extended memory. DOS memory and upper memory together compose real memory, the memory that can be addressed when the processor is running in real mode.




Figure 1. Basic Memory Layout

Under *DOS/4GW*, the first megabyte of physical memory — the real memory — is mapped as a shared linear address space. This allows your application to use absolute addresses in real memory, to access video RAM or BIOS ROM, for example. Because the real memory is available to all processes, you are not guaranteed to be able to allocate a particular area in real memory: another process may have allocated it already.

Most code and data is placed in a paged linear address space starting at 4MB. The linear address space starts at 4MB, the first address in the second page table, to avoid conflicts with VCPI system software.

This split mapping — an executable that is linked to start at 4MB in the linear address space, with the first MB in the address space mapped to the first MB of physical memory — is called a *split flat model*.

The illustration below shows the layout of physical memory on the left, and the layout of the linear address space on the right.




Figure 2. *Physical Memory/Linear Address Space*

The 1KB label in the diagram indicates the top of the real-mode interrupt vectors. 4KB marks the end of the first page.

8 Configuring *DOS/4GW*

This chapter explains various options that can be specified with the **DOS4G** environment variable including how to suppress the banner that is displayed by *DOS/4GW* at startup. It also explains how to use the **DOS16M** environment variable to select the switch mode setting, if necessary, and to specify the range of extended memory in which *DOS/4GW* will operate. *DOS/4GW* is based on Tenberry Software's DOS/16M 16-bit Protected-Mode support; hence the **DOS16M** environment variable name remains unchanged.

8.1 The *DOS4G* Environment Variable

A number of options can be selected by setting the **DOS4G** environment variable. The syntax for setting options is:

```
set DOS4G=option1,option2,...
```

Do not insert a space between **DOS4G** and the equal sign. A space to the right of the equal sign is optional.

Options:

QUIET Use this option to suppress the *DOS/4GW* banner.

The banner that is displayed by *DOS/4GW* at startup can be suppressed by issuing the following command:

```
set DOS4G=quiet
```

Note: Use of the quiet switch is only permitted pursuant to the terms and conditions of the WATCOM Software License Agreement and the additional redistribution rights described in the *Getting Started* manual. Under these terms, suppression of the copyright by using the quiet switch is not permitted for applications which you distribute to others.

VERBOSE Use this option to maximize the information available for postmortem debugging.

Before running your application, issue the following command:

```
set DOS4G=verbose
```

Reproduce the crash and record the output.

NULLP Use this option to trap references to the first sixteen bytes of physical memory.

Before running your application, issue the following command:

```
set DOS4G=nullp
```

To select a combination of options, list them with commas as separators.

Example:

```
set DOS4G=nullp,verbose
```

8.2 Changing the Switch Mode Setting

In almost all cases, *DOS/4GW* programs can detect the type of machine that is running and automatically choose an appropriate real- to protected-mode switch technique. For the few cases in which this default setting does not work we provide the **DOS16M** DOS environment variable, which overrides the default setting.

Change the switch mode settings by issuing the following command:

```
set DOS16M=value
```

Do not insert a space between **DOS16M** and the equal sign. A space to the right of the equal sign is optional.

The table below lists the machines and the settings you would use with them. Many settings have mnemonics, listed in the column "Alternate Name", that you can use instead of the number. Settings that you must set with the **DOS16M** variable have the notation *req'd* in the first column. Settings you may use are marked *option*, and settings that will automatically be set are marked *auto*.

Status	Machine	Setting	Alternate Name	Comment
auto	386/486 w/ DPMI	0	None	Set automatically if DPMI is active
req'd	NEC 98-series	1	9801	Must be set for NEC 98-series
auto	PS/2	2	None	Set automatically for PS/2
auto	386/486	3	386, 80386	Set automatically for 386 or 486
auto	386	INBOARD	None	386 with Intel Inboard
req'd	Fujitsu FMR-70	5	None	Must be set for Fujitsu FMR-70
auto	386/486 w/ VCPI	11	None	Set automatically if VCPI detected
req'd	Hitachi B32	14	None	Must be set for Hitachi B32
req'd	OKI if800	15	None	Must be set for OKI if800
option	IBM PS/55	16	None	May be needed for some PS/55s

The following procedure shows you how to test the switch mode setting.

1. If you have one of the machines listed below, set the **DOS16M** environment variable to the value shown for that machine and specify a range of extended memory. For example, if your machine is a NEC 98-series, set `DOS16M=1 @2M-4M`. See the section entitled "Fine Control of Memory Usage" on page 31 in this chapter for more information about setting the memory range.

Machine	Setting
NEC 98-series	1
Fujitsu FMR-60,-70	5
Hitachi B32	14
OKI if800	15

Before running *DOS/4GW* applications, check the switch mode setting by following this procedure:

2. Run PMINFO and note the switch setting reported on the last line of the display. (PMINFO, which reports on the protected-mode resources available to your programs, is described in more detail in the chapter entitled "Utilities" on page 69)

If PMINFO runs, the setting is usable on your machine.

3. If you changed the switch setting, add the new setting to your AUTOEXEC.BAT file.

Note: PMINFO will run successfully on 286 machines. If your *DOS/4GW* application does not run, and PMINFO does, check the CPU type reported on the first line of the display.

You are authorized (and encouraged) to distribute PMINFO to your customers. You may also include a copy of this section in your documentation.

8.3 Fine Control of Memory Usage

In addition to setting the switch mode as described above, the **DOS16M** environment variable enables you to specify which portion of extended memory *DOS/4GW* will use. The variable also allows you to instruct *DOS/4GW* to search for extra memory and use it if it is present.

8.3.1 Specifying a Range of Extended Memory

Normally, you don't need to specify a range of memory with the **DOS16M** variable. You must use the variable, however, in the following cases:

- You are running on a Fujitsu FMR-series, NEC 98-series, OKI if800-series or Hitachi B-series machine.
- You have older programs that use extended memory but don't follow one of the standard disciplines.
- You want to shell out of *DOS/4GW* to use another program that requires extended memory.

If none of these conditions applies to you, you can skip this section.

The general syntax is:

```
set DOS16M= [switch_mode] [@start_address [- end_address]] [:size]
```

In the syntax shown above, *start_address*, *end_address* and *size* represent numbers, expressed in decimal or in hexadecimal (hex requires a 0x prefix). The number may end with a K to indicate an

address or size in kilobytes, or an M to indicate megabytes. If no suffix is given, the address or size is assumed to be in kilobytes. If both a size and a range are specified, the more restrictive interpretation is used.

The most flexible strategy is to specify only a size. However, if you are running with other software that does not follow a convention for indicating its use of extended memory, and these other programs start before *DOS/4GW*, you will need to calculate the range of memory used by the other programs and specify a range for *DOS/4GW* programs to use.

DOS/4GW ignores specifications (or parts of specifications) that conflict with other information about extended memory use. Below are some examples of memory usage control:

<i>set DOS16M= 1 @2m-4m</i>	Mode 1, for NEC 98-series machines, and use extended memory between 2.0 and 4.0MB.
<i>set DOS16M= :1M</i>	Use the last full megabyte of extended memory, or as much as available limited to 1MB.
<i>set DOS16M= @2m</i>	Use any extended memory available above 2MB.
<i>set DOS16M= @ 0 - 5m</i>	Use any available extended memory from 0.0 (really 1.0) to 5.0MB.
<i>set DOS16M= :0</i>	Use no extended memory.

As a default condition *DOS/4GW* applications take all extended memory that is not otherwise in use. Multiple *DOS/4GW* programs that execute simultaneously will share the reserved range of extended memory. Any non-*DOS/4GW* programs started while *DOS/4GW* programs are executing will find that extended memory above the start of the *DOS/4GW* range is unavailable, so they may not be able to run. This is very safe. There will be a conflict only if the other program does not check the BIOS configuration call (Interrupt 15H function 88H, get extended memory size).

To create a private pool of extended memory for your *DOS/4GW* application, use the PRIVATXM program, described in the chapter entitled "Utilities" on page 69.

The default memory allocation strategy is to use extended memory if available, and overflow into DOS (low) memory.

In a VCPI or DPMI environment, the `start_address` and `end_address` arguments are not meaningful. *DOS/4GW* memory under these protocols is not allocated according to specific addresses because VCPI and DPMI automatically prevent address conflicts between extended memory programs. You can specify a `size` for memory managed by VCPI or DPMI, but *DOS/4GW* will not necessarily allocate this memory from the highest available extended memory address, as it does for memory managed under other protocols.

8.3.2 Using Extra Memory

Some machines contain extra non-extended, non-conventional memory just below 16MB. When *DOS/4GW* runs on a Compaq 386, it automatically uses this memory because the memory is allocated according to a certain protocol, which *DOS/4GW* follows. Other machines have no protocol for allocating this memory. To use the extra memory that may exist on these machines, set **DOS16M** with the + option.

```
set DOS16M=+
```

Setting the + option causes *DOS/4GW* to search for memory in the range from FA0000 to FFFFFFFF and determine whether the memory is usable. *DOS/4GW* does this by writing into the extra memory and reading what it has written. In some cases, this memory is mapped for DOS or BIOS usage, or for other system uses. If *DOS/4GW* finds extra memory that is mapped this way, and is not marked read-only, it will write into that memory. This will cause a crash, but won't have any other effect on your system.

8.4 Setting Runtime Options

The **DOS16M** environment variable sets certain runtime options for all *DOS/4GW* programs running on the same system.

To set the environment variable, the syntax is:

```
set DOS16M=[switch_mode_setting]^options.
```

Note: Some command line editing TSRs, such as CED, use the caret (^) as a delimiter. If you want to set **DOS16M** using the syntax above while one of these TSRs is resident, modify the TSR to use a different delimiter.

These are the options:

- 0x01** *check A20 line* -- This option forces *DOS/4GW* to wait until the A20 line is enabled before switching to protected mode. When *DOS/4GW* switches to real mode, this option suspends your program's execution until the A20 line is disabled, unless an XMS manager (such as HIMEM.SYS) is active. If an XMS manager is running, your program's execution is suspended until the A20 line is restored to the state it had when the CPU was last in real mode. Specify this option if you have a machine that runs *DOS/4GW* but is not truly AT-compatible. For more information on the A20 line, see the section entitled "Controlling Address Line 20" on page 34.
- 0x02** *prevent initialization of VCPI* -- By default, *DOS/4GW* searches for a VCPI server and, if one is present, forces it on. This option is useful if your application does not use EMS explicitly, is not a resident program, and may be used with 386-based EMS simulator software.
- 0x04** *directly pass down keyboard status calls* -- When this option is set, status requests are passed down immediately and unconditionally. When disabled, pass-downs are limited so the 8042 auxiliary processor does not become overloaded by keyboard polling loops.
- 0x10** *restore only changed interrupts* -- Normally, when a *DOS/4GW* program terminates, all interrupts are restored to the values they had at the time of program startup. When you use this option, only the interrupts changed by the *DOS/4GW* program are restored.
- 0x20** *set new memory to 00* -- When *DOS/4GW* allocates a new segment or increases the size of a segment, the memory is zeroed. This can help you find bugs having to do with uninitialized memory. You can also use it to provide a consistent working environment regardless of what programs were run earlier. This option only affects segment allocations or expansions that are made through the *DOS/4GW* kernel (with DOS function 48H or 4AH). This option does not affect memory allocated with a compiler's `malloc` function.
- 0x40** *set new memory to FF* -- When *DOS/4GW* allocates a new segment or increases the size of a segment, the memory is set to 0xFF bytes. This is helpful in making reproducible cases

of bugs caused by using uninitialized memory. This option only affects segment allocations or expansions that are made through the *DOS/4GW* kernel (with DOS function 48H or 4AH). This option does not affect memory allocated with a compiler's `malloc` function.

0x80 *new selector rotation* -- When *DOS/4GW* allocates a new selector, it usually looks for the first available (unused) selector in numerical order starting with the highest selector used when the program was loaded. When this option is set, the new selector search begins after the last selector that was allocated. This causes new selectors to rotate through the range. Use this option to find references to *stale* selectors, i.e., segments that have been cancelled or freed.

8.5 Controlling Address Line 20

This section explains how *DOS/4GW* uses address line 20 (A20) and describes the related **DOS16M** environment variable settings. It is unlikely that you will need to use these settings.

Because the 8086 and 8088 chips have 20-bit address spaces, their highest addressable memory location is one byte below 1MB. If you specify an address at 1MB or over, which would require a twenty-first bit to set, the address wraps back to zero. Some parts of DOS depend on this wrap, so on the 286 and 386, the twenty-first address bit is disabled. To address extended memory, *DOS/4GW* enables the twenty-first address bit (the A20 line). The A20 line must be enabled for the CPU to run in protected mode, but it may be either enabled or disabled in real mode.

By default, when *DOS/4GW* returns to real mode, it disables the A20 line. Some software depends on the line being enabled. *DOS/4GW* recognizes the most common software in this class, the XMS managers (such as HIMEM.SYS), and enables the A20 line when it returns to real mode if an XMS manager is present. For other software that requires the A20 line to be enabled, use the `A20` option. The `A20` option makes *DOS/4GW* restore the A20 line to the setting it had when *DOS/4GW* switched to protected mode. Set the environment variable as follows:

```
set DOS16M=A20
```

To specify more than one option on the command line, separate the options with spaces.

The **DOS16M** variable also lets you to specify the length of the delay between a *DOS/4GW* instruction to change the status of the A20 line and the next *DOS/4GW* operation. By default, this delay is 1 loop instruction when *DOS/4GW* is running on a 386 machine. In some cases, you may need to specify a longer delay for a machine that will run *DOS/4GW* but is not truly AT-compatible. To change the delay, set **DOS16M** to the desired number of loop instructions, preceded by a comma:

```
set DOS16M=,loops
```

9 VMM

The Virtual Memory Manager (VMM) uses a swap file on disk to augment RAM. With VMM you can use more memory than your machine actually has. When RAM is not sufficient, part of your program is swapped out to the disk file until it is needed again. The combination of the swap file and available RAM is the *virtual memory*.

Your program can use VMM if you set the DOS environment variable, **DOS4GVM**, as follows. To set the **DOS4GVM** environment variable, use the format shown below.

```
set DOS4GVM= [option[#value]] [option[#value]]
```

A "#" is used with options that take values since the DOS command shell will not accept "=".

If you set **DOS4GVM** equal to 1, the default parameters are used for all options.

Example:

```
C>set DOS4GVM=1
```

9.1 VMM Default Parameters

VMM parameters control the options listed below.

MINMEM	The minimum amount of RAM managed by VMM. The default is 512KB.
MAXMEM	The maximum amount of RAM managed by VMM. The default is 4MB.
SWAPMIN	The minimum or initial size of the swap file. If this option is not used, the size of the swap file is based on VIRTUALSIZE (see below).
SWAPINC	The size by which the swap file grows.
SWAPNAME	The swap file name. The default name is "DOS4GVM.SWP". By default the file is in the root directory of the current drive. Specify the complete path name if you want to keep the swap file somewhere else.
DELETESWAP	Whether the swap file is deleted when your program exits. By default the file is not deleted. Program startup is quicker if the file is not deleted.
VIRTUALSIZE	The size of the virtual memory space. The default is 16MB.

9.2 Changing the Defaults

You can change the defaults in two ways.

1. Specify different parameter values as arguments to the **DOS4GVM** environment variable, as shown in the example below.

```
set DOS4GVM=deleteswap maxmem#8192
```

2. Create a configuration file with the filetype extension ".VMC", and use that as an argument to the **DOS4GVM** environment variable, as shown below.

```
set DOS4GVM=@NEW4G.VMC
```

9.2.1 The .VMC File

A ".VMC" file contains VMM parameters and settings as shown in the example below. Comments are permitted. Comments on lines by themselves are preceded by an exclamation point (!). Comments that follow option settings are preceded by white space. Do not insert blank lines: processing stops at the first blank line.

```
!Sample .VMC file
!This file shows the default parameter values.
minmem = 512           At least 512K bytes of RAM is required.
maxmem = 4096          Uses no more than 4MB of RAM
virtualsize = 16384     Swap file plus allocated memory is 16MB
!To delete the swap file automatically when the program exits, add
!deleteswap
!To store the swap file in a directory called SWAPFILE, add
!swapname = c:\swapfile\dos4gvm.swp
```

10 Interrupt 21H Functions

When you call an Interrupt 21H function under *DOS/4GW*, the 32-bit registers in which you pass values are translated into the appropriate 16-bit registers, since DOS works only with 16 bits. However, you can use 32-bit values in your DOS calls. You can allocate blocks of memory larger than 64KB or use an address with a 32-bit offset, and *DOS/4GW* will translate the call appropriately, to use 16-bit registers. When the Interrupt 21H function returns, the value is widened - placed in a 32-bit register, with the high order bits zeroed.

DOS/4GW uses the following rules to manage registers:

- When you pass a parameter to an Interrupt 21H function that expects a 16-bit quantity in a general register (for example, AX), pass a 32-bit quantity in the corresponding extended register (for example, EAX). When a DOS function returns a 16-bit quantity in a general register, expect to receive it (with high-order zero bits) in the corresponding extended register.
- When an Interrupt 21H function expects to receive a 16:16 pointer in a segment:general register pair (for example, ES:BX), supply a 16:32 pointer using the same segment register and the corresponding extended general register (ES:EBX). *DOS/4GW* will copy data and translate pointers so that DOS ultimately receives a 16:16 real-mode pointer in the correct registers.
- When DOS returns a 16:16 real-mode pointer, *DOS/4GW* translates the segment value into an appropriate protected-mode selector and generates a 32-bit offset that results in a 16:32 pointer to the same location in the linear address space.
- Many DOS functions return an error code in AX if the function fails. *DOS/4GW* checks the status of the carry flag, and if it is set, indicating an error, zero-extends the code for EAX. It does not change any other registers.
- If the value is passed or returned in an 8-bit register (AL or AH, for example), *DOS/4GW* puts the value in the appropriate location and leaves the upper half of the 32-bit register untouched.

The table below lists all the Interrupt 21h functions. For each, it shows the registers that are widened or narrowed. Footnotes provide additional information about some of the interrupts that require special handling. Following the table is a section that provides a detailed explanation of interrupt handling under *DOS/4GW*.

Function	Purpose	Managed Registers
00H	Terminate Process	None
01H	Character Input with Echo	None
02H	Character Output	None
03H	Auxiliary Input	None
04H	Auxiliary Output	None
05H	Print Character	None
06H	Direct Console I/O	None
07H	Unfiltered Character Input Without Echo	None
08H	Character Input Without Echo	None
09H	Display String	EDX
0AH	Buffered Keyboard Input	EDX
0BH	Check Keyboard Status	None
0CH	Flush Buffer, Read Keyboard	EDX
0DH	Disk Reset	None
0EH	Select Disk	None
0FH	Open File with FCB	EDX
10H	Close File with FCB	EDX
11H	Find First File	EDX
12H	Find Next File	EDX
13H	Delete File	EDX
14H	Sequential Read	EDX
15H	Sequential Write	EDX
16H	Create File with FCB	EDX
17H	Rename File	EDX
19H	Get Current Disk	None
1AH	Set DTA Address	EDX
1BH	Get Default Drive Data	Returns in EBX, ECX, and EDX
1CH	Get Drive Data	Returns in EBX, ECX, and EDX
21H	Random Read	EDX
22H	Random Write	EDX
23H	Get File Size	EDX
24H	Set Relative Record	EDX
25H	Set Interrupt Vector	EDX
26H	Create New Program Segment Prefix	None
27H	Random Block Read	EDX, returns in ECX
28H	Random Block Write	EDX, returns in ECX
29H	Parse Filename	ESI, EDI, returns in EAX, ESI and EDI (1.)
2AH	Get Date	Returns in ECX
2BH	Set Date	None
2CH	Get Time	None
2DH	Set Time	None
2EH	Set/Reset Verify Flag	None
2FH	Get DTA Address	Returns in EBX
30H	Get MS-DOS Version Number	Returns in ECX
31H	Terminate and Stay Resident	None
33H	Get/Set Control-C Check Flag	None
34H	Return Address of InDOS Flag	Returns in EBX
35H	Get Interrupt Vector	Returns in EBX
36H	Get Disk Free Space	Returns in EAX, EBX, ECX, and EDX

38H	Get/Set Current Country	EDX, returns in EBX
39H	Create Directory	EDX
3AH	Remove Directory	EDX
3BH	Change Current Directory	EDX
3CH	Create File with Handle	EDX, returns in EAX
3DH	Open File with Handle	EDX, returns in EAX
3EH	Close File	None
3FH	Read File or Device	EBX, ECX, EDX, returns in EAX (2.)
40H	Write File or Device	EBX, ECX, EDX, returns in EAX (2.)
41H	Delete File	EDX
42H	Move File Pointer	Returns in EDX, EAX
43H	Get/Set File Attribute	EDX, returns in ECX
44H	IOCTL	(3.)
00H	Get Device Information	Returns in EDX
01H	SetDevice Information	None
02H	Read Control Data from CDD	EDX, returns in EAX
03H	Write Control Data to CDD	EDX, returns in EAX
04H	Read Control Data from BDD	EDX, returns in EAX
05H	Write Control Data to BDD	EDX, returns in EAX
06H	Check Input Status	None
07H	Check Output Status	None
08H	Check if Block Device is Removeable	Returns in EAX
09H	Check if Block Device is Remote	Returns in EDX
0AH	Check if Handle is Remote	Returns in EDX
0BH	Change Sharing Retry Count	None
0CH	Generic I/O Control for Character Devices	EDX
0DH	Generic I/O Control for Block Devices	EDX
0EH	Get Logical Drive Map	None
0FH	Set Logical Drive Map	None
45H	Duplicate File Handle	Returns in EAX
46H	Force Duplicate File Handle	None
47H	Get Current Directory	ESI
48H	Allocate Memory Block	Returns in EAX
49H	Free Memory Block	None
4AH	Resize Memory Block	None
4BH	Load and Execute Program (EXEC)	EBX, EDX (4.)
4CH	Terminate Process with Return Code	None
4DH	Get Return Code of Child Process	None
4EH	Find First File	EDX
4FH	Find Next File	None
52H	Get List of Lists	(not supported)
54H	Get Verify Flag	None
56H	Rename File	EDX, EDI
57H	Get/Set Date/Time of File	Returns in ECX, and EDX
58H	Get/Set Allocation Strategy	Returns in EAX
59H	Get Extended Error Information	Returns in EAX
5AH	Create Temporary File	EDX, returns in EAX and EDX
5BH	Create New File	EDX, returns in EAX
5CH	Lock/Unlock File Region	None
5EH	Network Machine Name/Printer Setup	
00H	Get Machine Name	EDX
02H	Set Printer Setup String	ESI

03H	Get Printer Setup String	EDI, returns in ECX
5FH	Get/Make Assign List Entry	
02H	Get Redirection List Entry	ESI, EDI, returns in ECX
03H	Redirect Device	ESI, EDI
04H	Cancel Device Redirection	ESI
62H	Get Program Segment Prefix Address	Returns in EBX
63H	Get Lead Byte Table (version 2.25 only)	Returns in ESI
65H	Get Extended Country Information	EDI
66H	Get or Set Code Page	None
67H	Set Handle Count	None

This list of functions is excerpted from *The MS-DOS Encyclopedia*, Copyright (c) 1988 by Microsoft Press. All Rights Reserved.

1. For Function 29H, DS:ESI and ES:EDI contain pointer values that are not changed by the call.
2. You can read and write quantities larger than 64KB with Functions 3FH and 40H. *DOS/4GW* breaks your request into chunks smaller than 64KB, and calls the DOS function once for each chunk.
3. You can't transfer more than 64KB using Function 44h, subfunctions 02H, 03H, 04H, or 05H. *DOS/4GW* does not break larger requests into DOS-sized chunks, as it does for Functions 3FH and 40H.
4. When you call Function 4B under *DOS/4GW*, you pass it a data structure that contains 16:32 bit pointers. *DOS/4GW* translates these into 16:16 bit pointers in the structure it passes to DOS.

10.1 Functions 25H and 35H: Interrupt Handling in Protected Mode

By default, interrupts that occur in protected mode are passed down: the entry in the IDT points to code in *DOS/4GW* that switches the CPU to real mode and resignals the interrupt. If you install an interrupt handler using Interrupt 21H, Function 25H, that handler will get control of any interrupts that occur while the processor is in protected mode. If the interrupt for which you installed the handler is in the *autopassup range*, your handler will also get control of interrupts signalled in real mode.

The *autopassup range* runs from 08H to 2EH inclusive, but excluding 21H. If the interrupt is in the *autopassup range*, the real-mode vector will be modified when you install the protected-mode handler to point to code in the *DOS/4GW* kernel. This code switches the processor to protected mode and resignals the interrupt-where your protected-mode handler will get control.

10.1.1 32-Bit Gates

The *DOS/4GW* kernel always assigns a 32-bit gate for the interrupt handlers it installs. It does not distinguish between 16-bit and 32-bit handlers for consistency with DPMI.

This 32-bit gate points into the *DOS/4GW* kernel. When *DOS/4GW* handles the interrupt, it switches to its own 16-bit stack, and from there it calls the interrupt handler (yours or the default). This translation is

transparent to the handler, with one exception: since the current stack is not the one on which the interrupt occurred, the handler cannot look up the stack for the address at which the interrupt occurred.

10.1.2 Chaining 16-bit and 32-bit Handlers

If your program hooks an interrupt, write a normal service routine that either handles the interrupt and IRETs or chains to the previous handler. As part of handling the interrupt, your handler can PUSHF/CALL to the previous handler. The handler *must* IRET (or IRETD) or chain.

For each protected-mode interrupt, *DOS/4GW* maintains separate chains of 16-bit and 32-bit handlers. If your 16-bit handler chains, the previous handler is a 16-bit program. If your 32-bit handler chains, the previous handler is a 32-bit program.

If a 16-bit program hooks a given interrupt before any 32-bit programs hook it, the 16-bit chain is executed first. If all the 16-bit handlers unhook later and a new 16-bit program hooks the interrupt while 32-bit handlers are still outstanding, the 32-bit handlers will be executed first.

If the first program to hook an interrupt is 32-bit, the 32-bit chain is executed first.

10.1.3 Getting the Address of the Interrupt Handler

When you signal Interrupt 21H, Function 35, it always returns a non-null address even if no other program of your bitness (i.e., 16-bit or 32-bit) has hooked the interrupt. The address points to a dummy handler that looks to you as though it does an IRET to end the chain. This means that you can't find an unused interrupt by looking for a NULL pointer. Since this technique is most frequently used by programs that are looking for an unclaimed *real-mode* interrupt on which to install a TSR, it shouldn't cause you problems.

11 Interrupt 31H DPMI Functions

When a *DOS/4GW* application runs under a DPMI host, such as Windows 3.1 in enhanced mode, an OS/2 virtual DOS machine, 386Max (with `DEBUG=DPMIXCOPY`), or QEMM/QDPMI (with `EXTCHKOFF`), the DPMI host provides the DPMI services, not *DOS/4GW*. The DPMI host also provides virtual memory, if any. Performance (speed and memory use) under different DPMI hosts varies greatly due to the quality of the DPMI implementation.

DPMI services are accessed using Interrupt 31H.

The following describes the services provided by *DOS/4GW* and *DOS/4GW Professional* in the absence of a DPMI host. *DOS/4GW* supports many of the common DPMI system services. Not all of the services described below are supported by other DPMI hosts.

Some of the information in this chapter was obtained from the the DOS Protected-Mode Interface (DPMI) specification.

The DPMI 1.0 specification can be obtained from the Open Watcom website <https://open-watcom.github.io/open-watcom-1.9/devel/dpmi10.pdf>. This file contains a PDF version of the DPMI 1.0 specification.

Alternatively the DPMI 1.0 specification can be obtained at <http://www.delorie.com/djgpp/doc/dpmi/dpmi-specs.tar.gz>. This file contains online HTML version of the DPMI 1.0 specification.

11.1 Using Interrupt 31H Function Calls

Interrupt 31H DPMI function calls can be used only by protected-mode programs.

The general ground rules for Interrupt 31H calls are as follows:

- All Interrupt 31H calls modify the AX register. Unsupported or unsuccessful calls return an error code in AX. Other registers are saved unless they contain specified return values.
- All Interrupt 31H calls modify flags: Unsupported or unsuccessful calls return with the carry flag set. Successful calls clear the carry flag. Only memory management and interrupt flag management calls modify the interrupt flag.
- Memory management calls can enable interrupts.
- All calls are reentrant.

The flag and register information for each call is listed in the following descriptions of supported Interrupt 31H function calls.

11.2 Int31H Function Calls

The Interrupt 31H subfunction calls supported by *DOS/4GW* are listed below by category:

- Local Descriptor Table (LDT) management services
- DOS memory management services
- Interrupt services
- Translation services
- DPMS version
- Memory management services
- Page locking services
- Demand paging performance tuning services
- Physical address mapping
- Virtual interrupt state functions
- Vendor specific extensions
- Coprocessor status

Only the most commonly used Interrupt 31H function calls are supported in this version.

11.2.1 Local Descriptor Table (LDT) Management Services

Function 0000H This function allocates a specified number of descriptors from the LDT and returns the base selector. Pass the following information:

AX = 0000H

CX = number of descriptors to be allocated

If the call succeeds, the carry flag is clear and the base selector is returned in AX. If the call fails, the carry flag is set.

An allocated descriptor is set to the present data type, with a base and limit of zero. The privilege level of an allocated descriptor is set to match the code segment privilege level of the application. To find out the privilege level of a descriptor, use the `lar` instruction.

Allocated descriptors must be filled in by the application. If more than one descriptor is allocated, the returned selector is the first of a contiguous array. Use Function 0003H to get the increment for the next selector in the array.

Function 0001H This function frees the descriptor specified. Pass the following information:

AX = 0001H

BX = the selector to free

Use the selector returned with function 0000h when the descriptor was allocated. To free an array of descriptors, call this function for each descriptor. Use Function 0003H to find out the increment for each descriptor in the array.

If the call succeeds, the carry flag is clear; if it fails, the carry flag is set.

You can use this function to free the descriptors allocated for the program's initial CS, DS, and SS segments, but you should not free other segments that were not allocated with Function 0000H or Function 000DH.

Function 0002H This function converts a real-mode segment to a descriptor that a protected-mode program can address. Pass the following information:

AX = 0002H

BX = real-mode segment address

If the call succeeds, it clears the carry flag and returns the selector mapped to the real-mode segment in AX. If the call fails, the carry flag is set.

If you call this function more than once with the same real-mode segment address, you get the same selector value each time. The descriptor limit is set to 64KB.

The purpose of this function is to give protected-mode programs easy access to commonly used real-mode segments. However, because you cannot modify or free descriptors created by this function, it should be used infrequently. Do not use this function to get descriptors for private data areas.

To examine real-mode addresses using the same selector, first allocate a descriptor, and then use Function 0007H to change the linear base address.

Function 0003H This function returns the increment value for the next selector. Use this function to get the value you add to the base address of an allocated array of descriptors to get the next selector address. Pass the following information:

AX = 0003H

This call always succeeds. The increment value is returned in AX. This value is always a power of two, but no other assumptions can be made.

Function 0006H This function gets the linear base address of a selector. Pass the following information:

AX = 0006H

BX = selector

If the call succeeds, the carry flag is clear and CX:DX contains the 32-bit linear base address of the segment. If the call fails, it sets the carry flag.

If the selector you specify in BX is invalid, the call fails.

Function 0007H This function changes the base address of a specified selector. Only descriptors allocated through Function 0000H should be modified. Pass the following information:

AX = 0007H

BX = selector

CX:DX = the new 32-bit linear base address for the segment

If the call succeeds, the carry flag is clear; if unsuccessful, the carry flag is set.

If the selector you specify in BX is invalid, the call fails.

Function 0008H This function sets the upper limit of a specified segment. Use this function to modify descriptors allocated with Function 0000H only. Pass the following information:

AX = 0008H

BX = selector

CX:DX = 32-bit segment limit

If the call succeeds, the carry flag is clear; if unsuccessful, the carry flag is set.

The call fails if the specified selector is invalid, or if the specified limit cannot be set.

Segment limits greater than 1MB must be page-aligned. This means that limits greater than 1MB must have the low 12 bits set.

To get the limit of a segment, use the 32-bit form of `lsl` for segment limits greater than 64KB.

Function 0009H This function sets the descriptor access rights. Use this function to modify descriptors allocated with Function 0000H only. To examine the access rights of a descriptor, use the `lar` instruction. Pass the following information:

AX = 0009H

BX = selector

CL = Access rights/type byte

CH = 386 extended access rights/type byte

If the call succeeds, the carry flag is clear; if unsuccessful, the carry flag is set. If the selector you specify in BX is invalid, the call fails. The call also fails if the access rights/type byte does not match the format and meet the requirements shown in the figures below.

The access rights/type byte passed in CL has the format shown in the figure below.




Figure 3. Access Rights/Type

The extended access rights/type byte passed in CH has the following format.




Figure 4. Extended Access Rights/Type

Function 000AH This function creates an alias to a code segment. This function creates a data descriptor that has the same base and limit as the specified code segment descriptor. Pass the following information:

AX = 000AH
BX = code segment selector

If the call succeeds, the carry flag is clear and the new data selector is returned in AX. If the call fails, the carry flag is set. The call fails if the selector passed in BX is not a valid code segment.

To deallocate an alias to a code segment, use Function 0001H.

After the alias is created, it does not change if the code segment descriptor changes. For example, if the base or limit of the code segment change later, the alias descriptor stays the same.

Function 000BH This function copies the descriptor table entry for a specified descriptor. The copy is written into an 8-byte buffer. Pass the following information:

AX = 000BH
BX = selector
ES:EDI = a pointer to the 8-byte buffer for the descriptor copy

If the call succeeds, the carry flag is clear and ES:EDI contains a pointer to the buffer that contains a copy of the descriptor. If the call fails, the carry flag is set. The call fails if the selector passed in BX is invalid or unallocated.

Function 000CH This function copies an 8-byte buffer into the LDT for a specified descriptor. The descriptor must first have been allocated with Function 0000H. Pass the following information:

AX = 000CH

BX = selector

ES:EDI = a pointer to the 8-byte buffer containing the descriptor

If the call succeeds, the carry flag is clear; if it fails, the carry flag is set. The call fails if the descriptor passed in BX is invalid.

The type byte, byte 5, has the same format and requirements as the access rights/type byte passed to Function 0009H in CL. The format is shown in the first figure presented with the description of Function 0009H.

The extended type byte, byte 6, has the same format and requirements as the extended access rights/type byte passed to Function 0009H in CH, except that the limit field can have any value, and the low order bits marked *reserved* are used to set the upper 4 bits of the descriptor limit. The format is shown in the second figure presented with the description of Function 0009H.

Function 000DH This function allocates a specific LDT descriptor. Pass the following information:

AX = 000DH

BX = selector

If the call succeeds, the carry flag is clear and the specified descriptor is allocated. If the call fails, the carry flag is set.

The call fails if the specified selector is already in use, or if it is not a valid LDT descriptor. The first 10h (16 decimal) descriptors are reserved for this function, and should not be used by the host. Some of these descriptors may be in use, however, if another client application is already loaded.

To free the descriptor, use Function 0001H.

11.2.2 DOS Memory Management Services

Function 0100H This function allocates memory from the DOS free memory pool. This function returns both the real-mode segment and one or more descriptors that can be used by protected-mode applications. Pass the following information:

AX = 0100H

BX = the number of paragraphs (16-byte blocks) requested

If the call succeeds, the carry flag is clear. AX contains the initial real-mode segment of the allocated block and DX contains the base selector for the allocated block.

If the call fails, the carry flag is set. AX contains the DOS error code. If memory is damaged, code 07H is returned. If there is not enough memory to satisfy the request, code 08H is returned. BX contains the number of paragraphs in the largest available block of DOS memory.

If you request a block larger than 64KB, contiguous descriptors are allocated. Use Function 0003H to find the value of the increment to the next descriptor. The limit of the first descriptor is set to the entire block. Subsequent descriptors have a limit of 64KB, except for the final descriptor, which has a limit of `blocksize MOD 64KB`.

You cannot modify or deallocate descriptors allocated with this function. Function 101H deallocates the descriptors automatically.

Function 0101H This function frees a DOS memory block allocated with function 0100H. Pass the following information:

AX = 0101H

DX = selector of the block to be freed

If the call succeeds, the carry flag is clear.

If the call fails, the carry flag is set and the DOS error code is returned in AX. If the incorrect segment was specified, code 09H is returned. If memory control blocks are damaged, code 07H is returned.

All descriptors allocated for the specified memory block are deallocated automatically and cannot be accessed correctly after the block is freed.

Function 0102H This function resizes a DOS memory block allocated with function 0100H. Pass the following information:

AX = 0102H

BX = the number of paragraphs (16-byte blocks) in the resized block

DX = selector of block to resize

If the call succeeds, the carry flag is clear.

If the call fails, the carry flag is set, the maximum number of paragraphs available is returned in BX, and the DOS error code is returned in AX. If memory code blocks are damaged, code 07H is returned. If there isn't enough memory to increase the size as requested, code 08H is returned. If the incorrect segment is specified, code 09h is returned.

Because of the difficulty of finding additional contiguous memory or descriptors, this function is not often used to increase the size of a memory block. Increasing the size of a memory block might well fail because other DOS allocations have used contiguous space. If the next descriptor in the LDT is not free, allocation also fails when the size of a block grows over the 64KB boundary.

If you shrink the size of a memory block, you may also free some descriptors allocated to the block. The initial selector remains unchanged, however; only the limits of subsequent selectors will change.

11.2.3 Interrupt Services

Function 0200H This function gets the value of the current task's real-mode interrupt vector for the specified interrupt. Pass the following information:

AX = 0200H
BL = interrupt number

This call always succeeds. All 100H (256 decimal) interrupt vectors are supported by the host. When the call returns, the carry flag is clear, and the `segment:offset` of the real-mode interrupt handler is returned in CX:DX.

Because the address returned in CX is a segment, and not a selector, you cannot put it into a protected-mode segment register. If you do, a general protection fault may occur.

Function 0201H This function sets the value of the current task's real-mode interrupt vector for the specified interrupt. Pass the following information:

AX = 0201H
BL = interrupt number
CX:DX = segment:offset of the real-mode interrupt handler

If the call succeeds, the carry flag is clear; if it fails, the carry flag is set.

The address passed in CX:DX should be a real-mode `segment:offset`, such as function 0200H returns. For this reason, the interrupt handler must reside in DOS addressable memory. You can use Function 0100H to allocate DOS memory. This version does not support the real-mode callback address function.

If you are hooking a hardware interrupt, you have to lock all segments involved. These segments include the segment in which the interrupt handler runs, and any segment it may touch at interrupt time.

Function 0202H This function gets the processor exception handler vector. This function returns the CS:EIP of the current protected-mode exception handler for the specified exception number. Pass the following information:

AX = 0202H
BL = exception/fault number (00h - 1Fh)

If the call succeeds, the carry flag is clear and the `selector:offset` of the protected-mode exception handler is returned in CX:EDX. If it fails, the carry flag is set.

The value returned in CX is a valid protected-mode selector, not a real-mode segment.

Function 0203H This function sets the processor exception handler vector. This function allows protected-mode applications to intercept processor exceptions that are not handled by the DPMI environment. Programs may wish to handle exceptions such as "not present segment faults" which would otherwise generate a fatal error. Pass the following information:

AX = 0203H

BL = exception/fault number (00h - 1Fh)

CX:EDX = selector:offset of the exception handler

If the call succeeds, the carry flag is clear. If it fails, the carry flag is set.

The address passed in CX must be a valid protected-mode selector, such as Function 204H returns, and not a real-mode segment. A 32-bit implementation must supply a 32-bit offset in the EDX register. If the handler chains to the next handler, it must use a 32-bit interrupt stack frame to do so.

The handler should return using a far return instruction. The original SS:ESP, CS:EIP and flags on the stack, including the interrupt flag, will be restored.

All fault stack frames have an error code. However the error code is only valid for exceptions 08h, 0Ah, 0Bh, 0Ch, 0Dh, and 0Eh.

The handler must preserve and restore all registers.

The exception handler will be called on a locked stack with interrupts disabled. The original SS, ESP, CS, and EIP will be pushed on the exception handler stack frame.

The handler must either return from the call by executing a far return or jump to the next handler in the chain (which will execute a far return or chain to the next handler).

The procedure can modify any of the values on the stack pertaining to the exception before returning. This can be used, for example, to jump to a procedure by modifying the CS:EIP on the stack. Note that the procedure must not modify the far return address on the stack — it must return to the original caller. The caller will then restore the flags, CS:EIP and SS:ESP from the stack frame.

If the DPMI client does not handle an exception, or jumps to the default exception handler, the host will reflect the exception as an interrupt for exceptions 0, 1, 2, 3, 4, 5 and 7. Exceptions 6 and 8 - 1Fh will be treated as fatal errors and the client will be terminated.

Exception handlers will only be called for exceptions that occur in protected mode.

Function 0204H This function gets the CS:EIP `selector:offset` of the current protected-mode interrupt handler for a specified interrupt number. Pass the following information:

AX = 0204H

BL = interrupt number

This call always succeeds. All 100H (256 decimal) interrupt vectors are supported by the host. When the call returns, the carry flag is clear and CX:EDX contains the protected-mode `selector:offset` of the exception handler.

A 32-bit offset is returned in the EDX register.

Function 0205H This function sets the address of the specified protected-mode interrupt vector. Pass the following information:

AX = 0205H

BL = interrupt number

CX:EDX = selector:offset of the exception handler

If the call succeeds, the carry flag is clear; if it fails, the carry flag is set.

The address passed in CX must be a valid protected-mode selector, such as Function 204H returns, and not a real-mode segment. A 32-bit implementation must supply a 32-bit offset in the EDX register. If the handler chains to the next handler, it must use a 32-bit interrupt stack frame to do so.

11.2.4 Translation Services

These services are provided so that protected-mode programs can call real-mode software that DPMI does not support directly. The protected-mode program must set up a data structure with the appropriate register values. This "real-mode call structure" is shown below.

Offset	Register
00H	EDI
04H	ESI
08H	EBP
0CH	Reserved by system
10H	EBX
14H	EDX
18H	ECX
1CH	EAX
20H	Flags
22H	ES
24H	DS
26H	FS
28H	GS
2AH	IP
2CH	CS
2EH	SP
30H	SS

After the call or interrupt is complete, all real-mode registers and flags except SS, SP, CS, and IP will be copied back to the real-mode call structure so that the caller can examine the real-mode return values.

The values in the segment registers should be real-mode segments, not protected-mode selectors.

The translation services will provide a real-mode stack if the SS:SP fields are zero. However, the stack provided is relatively small. If the real-mode procedure/interrupt routine uses more than 30 words of stack space then you should provide your own real-mode stack.

Function 0300H This function simulates a real-mode interrupt. This function simulates an interrupt in real mode. It will invoke the CS:IP specified by the real-mode interrupt vector and the handler must return by executing an `iret`. Pass the following information:

AX = 0300H

BL = interrupt number

BH = flags Bit 0 = 1 resets the interrupt controller and A20 line. Other flags are reserved and must be 0.

CX = number of words to copy from protected-mode stack to real-mode stack

ES:EDI = the selector:offset of real-mode call structure

If the call fails, the carry flag is set.

If the call succeeds, the carry flag is clear and ES:EDI contains the `selector:offset` of the modified real-mode call structure.

The CS:IP in the real-mode call structure is ignored by this service. The appropriate interrupt handler will be called based on the value passed in BL.

The flags specified in the real-mode call structure will be pushed on the real-mode stack `iret` frame. The interrupt handler will be called with the interrupt and trace flags clear.

It is up to the caller to remove any parameters that were pushed on the protected-mode stack.

The flag to reset the interrupt controller and the A20 line is ignored by DPMI implementations that run in Virtual 8086 mode. It causes DPMI implementations that return to real mode to set the interrupt controller and A20 address line hardware to its normal real-mode state.

Function 0301H (DOS/4GW Professional only) This function calls a real-mode procedure with a FAR return frame. The called procedure must execute a FAR return when it completes. Pass the following information:

AX = 0301H

BH = flags Bit 0 = 1 resets the interrupt controller and A20 line. Other flags reserved and must be 0.

CX = Number of words to copy from protected-mode to real-mode stack

ES:EDI = selector:offset of real-mode call structure

If the call succeeds, the carry flag is clear and ES:EDI contains the `selector:offset` of modified real-mode call structure.

If the call fails, the carry flag is set.

Notes:

1. The CS:IP in the real-mode call structure specifies the address of the real-mode procedure to call.
2. The real-mode procedure must execute a FAR return when it has completed.
3. If the SS:SP fields are zero then a real-mode stack will be provided by the DPMI host. Otherwise, the real-mode SS:SP will be set to the specified values before the procedure is called.
4. When the Int 31h returns, the real-mode call structure will contain the values that were returned by the real-mode procedure.

5. It is up to the caller to remove any parameters that were pushed on the protected-mode stack.
6. The flag to reset the interrupt controller and A20 line is ignored by DPMI implementations that run in Virtual 8086 mode. It causes DPMI implementations that return to real mode to set the interrupt controller and A20 address line hardware to its normal real-mode state.

Function 0302H (DOS/4GW Professional only) This function calls a real-mode procedure with an `iret` frame. The called procedure must execute an `iret` when it completes. Pass the following information:

AX = 0302H

BH = flags Bit 0 = 1 resets the interrupt controller and A20 line. Other flags reserved and must be 0.

CX = Number of words to copy from protected-mode to real-mode stack

ES:EDI = selector:offset of real-mode call structure

If the call succeeds, the carry flag is clear and ES:EDI contains the `selector:offset` of modified real-mode call structure.

If the call fails, the carry flag is set.

Notes:

1. The CS:IP in the real-mode call structure specifies the address of the real-mode procedure to call.
2. The real-mode procedure must execute an `iret` when it has completed.
3. If the SS:SP fields are zero then a real-mode stack will be provided by the DPMI host. Otherwise, the real-mode SS:SP will be set to the specified values before the procedure is called.
4. When the Int 31h returns, the real-mode call structure will contain the values that were returned by the real-mode procedure.
5. The flags specified in the real-mode call structure will be pushed the real-mode stack `iret` frame. The procedure will be called with the interrupt and trace flags clear.
6. It is up to the caller to remove any parameters that were pushed on the protected-mode stack.
7. The flag to reset the interrupt controller and A20 line is ignored by DPMI implementations that run in Virtual 8086 mode. It causes DPMI implementations that return to real mode to set the interrupt controller and A20 address line hardware to its normal real-mode state.

Function 0303H (DOS/4GW Professional only) This function allocates a real-mode callback address. This service is used to obtain a unique real-mode SEG:OFFSET that will transfer control from real mode to a protected-mode procedure.

At times it is necessary to hook a real-mode interrupt or device callback in a protected-mode driver. For example, many mouse drivers call an address whenever the mouse is moved. Software running in protected mode can use a real-mode callback to intercept the mouse driver calls. Pass the following information:

AX = 0303H

DS:ESI = selector:offset of procedure to call

ES:EDI = selector:offset of real-mode call structure

If the call succeeds, the carry flag is clear and CX:DX contains the `segment:offset` of real-mode callback address.

If the call fails, the carry flag is set.

Callback Procedure Parameters

Interrupts disabled
DS:ESI = selector:offset of real-mode SS:SP
ES:EDI = selector:offset of real-mode call structure
SS:ESP = Locked protected-mode API stack
All other registers undefined

Return from Callback Procedure

Execute an IRET to return
ES:EDI = selector:offset of real-mode call structure
to restore (see note)

Notes:

1. Since the real-mode call structure is static, you must be careful when writing code that may be reentered. The simplest method of avoiding reentrancy is to leave interrupts disabled throughout the entire call. However, if the amount of code executed by the callback is large then you will need to copy the real-mode call structure into another buffer. You can then return with ES:EDI pointing to the buffer you copied the data to — it does not have to point to the original real mode call structure.
2. The called procedure is responsible for modifying the real-mode CS:IP before returning. If the real-mode CS:IP is left unchanged then the real-mode callback will be executed immediately and your procedure will be called again. Normally you will want to pop a return address off of the real-mode stack and place it in the real-mode CS:IP. The example code in the next section demonstrates chaining to another interrupt handler and simulating a real-mode `iret`.
3. To return values to the real-mode caller, you must modify the real-mode call structure.
4. Remember that all segment values in the real-mode call structure will contain real-mode segments, not selectors. If you need to examine data pointed to by a real-mode `seg:offset` pointer, you should not use the segment to selector service to create a new selector. Instead, allocate a descriptor during initialization and change the descriptor's base to 16 times the real-mode segment's value. This is

important since selectors allocated through the segment to selector service can never be freed.

5. DPMI hosts should provide a minimum of 16 callback addresses per task.

The following code is a sample of a real-mode interrupt hook. It hooks the DOS Int 21h and returns an error for the delete file function (AH=41h). Other calls are passed through to DOS. This example is somewhat silly but it demonstrates the techniques used to hook a real mode interrupt. Note that since DOS calls are reflected from protected mode to real mode, the following code will intercept all DOS calls from both real mode and protected mode.

```
;*****
; This procedure gets the current Int 21h real-mode
; Seg:Offset, allocates a real-mode callback address,
; and sets the real-mode Int 21h vector to the call-
; back address.
;*****
Initialization_Code:
;
; Create a code segment alias to save data in
;
    mov     ax, 000Ah
    mov     bx, cs
    int     31h
    jc      ERROR
    mov     ds, ax
    ASSUMES DS,_TEXT
;
; Get current Int 21h real-mode SEG:OFFSET
;
    mov     ax, 0200h
    mov     bl, 21h
    int     31h
    jc      ERROR
    mov     [Orig_Real_Seg], cx
    mov     [Orig_Real_Offset], dx
;
; Allocate a real-mode callback
;
    mov     ax, 0303h
    push    ds
    mov     bx, cs
    mov     ds, bx
    mov     si, OFFSET My_Int_21_Hook
    pop     es
    mov     di, OFFSET My_Real_Mode_Call_Struct
    int     31h
    jc      ERROR
;
; Hook real-mode int 21h with the callback address
;
    mov     ax, 0201h
    mov     bl, 21h
    int     31h
    jc      ERROR

;*****
; This is the actual Int 21h hook code. It will return
; an "access denied" error for all calls made in real
; mode to delete a file. Other calls will be passed
; through to DOS.
;
; ENTRY:
;   DS:SI -> Real-mode SS:SP
;   ES:DI -> Real-mode call structure
;   Interrupts disabled
;
; EXIT:
;   ES:DI -> Real-mode call structure
;
;*****

My_Int_21_Hook:
    cmp     es:[di.RealMode_AH], 41h
    jne     Chain_To_DOS
;
; This is a delete file call (AH=41h). Simulate an
; iret on the real-mode stack, set the real-mode
; carry flag, and set the real-mode AX to 5 to indicate
; an access denied error.
;
```

```
        cld
        lodsw                ; Get real-mode ret IP
        mov     es:[di.RealMode_IP], ax
        lodsw                ; Get real-mode ret CS
        mov     es:[di.RealMode_CS], ax
        lodsw                ; Get real-mode flags
        or      ax, 1        ; Set carry flag
        mov     es:[di.RealMode_Flags], ax
        add     es:[di.RealMode_SP], 6
        mov     es:[di.RealMode_AX], 5
        jmp     My_Hook_Exit
;
; Chain to original Int 21h vector by replacing the
; real-mode CS:IP with the original Seg:Offset.
;
Chain_To_DOS:
        mov     ax, cs:[Orig_Real_Seg]
        mov     es:[di.RealMode_CS], ax
        mov     ax, cs:[Orig_Real_Offset]
        mov     es:[di.RealMode_IP], ax

My_Hook_Exit:
        iret
```

Function 0304H (DOS/4GW Professional only) This function frees a real-mode callback address that was allocated through the allocate real-mode callback address service. Pass the following information:

AX = 0304H

CX:DX = Real-mode callback address to free

If the call succeeds, the carry flag is clear; if it fails, the carry flag is set.

Notes:

1. Real-mode callbacks are a limited resource. Your code should free any break point that it is no longer using.

11.2.5 DPMI Version

Function 0400H This function returns the version of DPMI services supported. Note that this is not necessarily the version of any operating system that supports DPMI. It should be used by programs to determine what calls are legal in the current environment. Pass the following information:

AX = 0400H

The information returned is:

AH = Major version

AL = Minor version

BX = Flags Bit 0 = 1 if running under an 80386 DPMI implementation. Bit 1 = 1 if processor is returned to real mode for reflected interrupts (as opposed to Virtual 8086 mode). Bit 2 = 1 if virtual memory is supported. Bit 3 is reserved and undefined. All other bits are zero and reserved for later use.

CL = Processor type

02 = 80286
 03 = 80386
 04 = 80486
 05 = Pentium

DH = Current value of virtual master PIC base interrupt
DL = Current value of virtual slave PIC base interrupt
Carry flag clear (call cannot fail)

11.2.6 Memory Management Services

Function 0500H This function gets information about free memory. Pass the following information:

AX = 0500H

ES:EDI = the selector:offset of a 30H byte buffer.

If the call fails, the carry flag is set.

If the call succeeds, the carry flag is clear and ES:EDI contains the selector:offset of a buffer with the structure shown in the figure below.

Offset	Description
00H	Largest available block, in bytes
04H	Maximum unlocked page allocation
08H	Largest block of memory (in pages) that could be allocated and then locked
0CH	Total linear address space size, in pages, including already allocated pages
10H	Total number of free pages and pages currently unlocked and available for paging out
14H	Number of physical pages not in use
18H	Total number of physical pages managed by host
1CH	Free linear address space, in pages
20H	Size of paging/file partition, in pages
24H - 2FH	Reserved

Only the first field of the structure is guaranteed to contain a valid value. Any field that is not returned by *DOS/4GW* is set to -1 (0FFFFFFFH).

Function 0501H This function allocates and commits linear memory. Pass the following information:

AX = 0501H

BX:CX = size of memory to allocate, in bytes.

If the call succeeds, the carry flag is clear, BX:CX contains the linear address of the allocated memory, and SI:DI contains the memory block handle used to free or resize the block. If the call fails, the carry flag is set.

No selectors are allocated for the memory block. The caller must allocate and initialize selectors needed to access the memory.

If VMM is present, the memory is allocated as unlocked, page granular blocks. Because of the page granularity, memory should be allocated in multiples of 4KB.

Function 0502H This function frees a block of memory allocated through function 0501H. Pass the following information:

AX = 0502H

SI:DI = handle returned with function 0501H when memory was allocated

If the call succeeds, the carry flag is clear; if it fails, the carry flag is set. You must also free any selectors allocated to point to the freed memory block.

Function 0503H This function resizes a block of memory allocated through the 0501H function. If you resize a block of linear memory, it may have a new linear address and a new handle. Pass the following information:

AX = 0503H

BX:CX = new size of memory block, in bytes

SI:DI = handle returned with function 0501H when memory was allocated

If the call succeeds, the carry flag is clear, BX:CX contains the new linear address of the memory block, and SI:DI contains the new handle of the memory block. If the call fails, the carry flag is set.

If either the linear address or the handle has changed, update the selectors that point to the memory block. Use the new handle instead of the old one.

You cannot resize a memory block to zero bytes.

11.2.7 Page Locking Services

These services are only useful under DPMS implementations that support virtual memory. Although memory ranges are specified in bytes, the actual unit of memory that will be locked will be one or more pages. Page locks are maintained as a count. When the count is decremented to zero, the page is unlocked and can be swapped to disk. This means that if a region of memory is locked three times then it must be unlocked three times before the pages will be unlocked.

Function 0600H This function locks a specified linear address range. Pass the following information:

AX = 0600H

BX:CX = starting linear address of memory to lock

SI:DI = size of region to lock (in bytes)

If the call fails, the carry flag is set and none of the memory will be locked.

If the call succeeds, the carry flag is clear. If the specified region overlaps part of a page at the beginning or end of a region, the page(s) will be locked.

Function 0601H This function unlocks a specified linear address range that was previously locked using the "lock linear region" function (0600h). Pass the following information:

AX = 0601H

BX:CX = starting linear address of memory to unlock

SI:DI = size of region to unlock (in bytes)

If the call fails, the carry flag is set and none of the memory will be unlocked. An error will be returned if the memory was not previously locked or if the specified region is invalid.

If the call succeeds, the carry flag is clear. If the specified region overlaps part of a page at the beginning or end of a region, the page(s) will be unlocked. Even if the call succeeds, the memory will remain locked if the lock count is not decremented to zero.

Function 0604H This function gets the page size for Virtual Memory (VM) only. This function returns the size of a single memory page in bytes. Pass the following information:

AX = 0604H

If the call succeeds, the carry flag is clear and BX:CX = Page size in bytes.

If the call fails, the carry flag is set.

11.2.8 Demand Paging Performance Tuning Services

Some applications will discard memory objects or will not access objects for long periods of time. These services can be used to improve the performance of demand paging.

Although these functions are only relevant for DPMI implementations that support virtual memory, other implementations will ignore these functions (it will always return carry clear). Therefore your code can always call these functions regardless of the environment it is running under.

Since both of these functions are simply advisory functions, the operating system may choose to ignore them. In any case, your code should function properly even if the functions fail.

Function 0702H (DOS/4GW Professional only) This function marks a page as a demand paging candidate. This function is used to inform the operating system that a range of pages should be placed at the head of the page out candidate list. This will force these pages to be swapped to disk ahead of other pages even if the memory has been accessed recently. However, all memory contents will be preserved.

This is useful, for example, if a program knows that a given piece of data will not be accessed for a long period of time. That data is ideal for swapping to disk since the physical memory it now occupies can be used for other purposes. Pass the following information:

AX = 0702H

BX:CX = Starting linear address of pages to mark

SI:DI = Number of bytes to mark as paging candidates

If the call succeeds, the carry flag is clear; if it fails, the carry flag is set.

Notes:

1. This function does not force the pages to be swapped to disk immediately.
2. Partial pages will not be discarded.

Function 0703H (DOS/4GW Professional only) This function discards page contents. This function discards the entire contents of a given linear memory range. It is used after a memory object that occupied a given piece of memory has been discarded.

The contents of the region will be undefined the next time the memory is accessed. All values previously stored in this memory will be lost. Pass the following information:

AX = 0703H

BX:CX = Starting linear address of pages to discard

SI:DI = Number of bytes to discard

If the call succeeds, the carry flag is clear; if it fails, the carry flag is set.

Notes:

1. Partial pages will not be discarded.

11.2.9 Physical Address Mapping

Memory mapped devices such as network adapters and displays sometimes have memory mapped at physical addresses that lie outside of the normal 1Mb of memory that is addressable in real mode. Under many implementations of DPMI, all addresses are linear addresses since they use the paging mechanism of the 80386. This service can be used by device drivers to convert a physical address into a linear address. The linear address can then be used to access the device memory.

Function 0800H This function is used for Physical Address Mapping.

Some implementations of DPMI may not support this call because it could be used to circumvent system protection. This call should only be used by programs that absolutely require direct access to a memory mapped device.

Pass the following information:

AX = 0800H

BX:CX = Physical address of memory

SI:DI = Size of region to map in bytes

If the call succeeds, the carry flag is clear and BX:CX = Linear Address that can be used to access the physical memory.

If the call fails, the carry flag is set.

Notes:

1. Under DPMI implementations that do not use the 80386 paging mechanism, the call will always succeed and the address returned will be equal to the physical address parameter passed into this function.
2. It is up to the caller to build an appropriate selector to access the memory.
3. Do not use this service to access memory that is mapped in the first megabyte of address space (the real-mode addressable region).

Function 0801H This function is used to free Physical Address Mapping. Pass the following information:

AX = 0801H

BX:CX = Linear address returned by Function 0800H.

If the call succeeds, the carry flag is clear; if it fails, the carry flag is set.

Notes:

1. The client should call this function when it is finished using a device previously mapped to linear addresses with the Physical Address Mapping function (Function 0800H).

11.2.10 Virtual Interrupt State Functions

Under many implementations of DPMI, the interrupt flag in protected mode will always be set (interrupts enabled). This is because the program is running under a protected operating system that cannot allow programs to disable physical hardware interrupts. However, the operating system will maintain a "virtual" interrupt state for protected-mode programs. When the program executes a CLI instruction, the program's virtual interrupt state will be disabled, and the program will not receive any hardware interrupts until it executes an STI to reenable interrupts (or calls service 0901h).

When a protected-mode program executes a PUSHF instruction, the real processor flags will be pushed onto the stack. Thus, examining the flags pushed on the stack is not sufficient to determine the state of the program's virtual interrupt flag. These services enable programs to get and modify the state of their virtual interrupt flag.

The following sample code enters an interrupt critical section and then restores the virtual interrupt state to its previous state.

```
;
; Disable interrupts and get previous interrupt state
;
        mov     ax, 0900h
        int     31h
;
; At this point AX = 0900h or 0901h
;
        .
        .
        .
;
; Restore previous state (assumes AX unchanged)
;
        int     31h
```

Function 0900H This function gets and disables Virtual Interrupt State. This function will disable the virtual interrupt flag and return the previous state of the virtual interrupt flag. Pass the following information:

AX = 0900H

After the call, the carry flag is clear (this function always succeeds) and virtual interrupts are disabled.

AL = 0 if virtual interrupts were previously disabled.

AL = 1 if virtual interrupts were previously enabled.

Notes:

1. AH will not be changed by this procedure. Therefore, to restore the previous state, simply execute an Int 31h.

Function 0901H This function gets and enables the Virtual Interrupt State. This function will enable the virtual interrupt flag and return the previous state of the virtual interrupt flag. Pass the following information:

AX = 0901H

After the call, the carry flag is clear (this function always succeeds) and virtual interrupts are enabled.

AL = 0 if virtual interrupts were previously disabled.

AL = 1 if virtual interrupts were previously enabled.

Notes:

1. AH will not be changed by this procedure. Therefore, to restore the previous state, simply execute an Int 31h.

Function 0902H This function gets the Virtual Interrupt State. This function will return the current state of the virtual interrupt flag. Pass the following information:

AX = 0902H

After the call, the carry flag is clear (this function always succeeds).

AL = 0 if virtual interrupts are disabled.

AL = 1 if virtual interrupts are enabled.

11.2.11 Vendor Specific Extensions

Some DOS extenders provide extensions to the standard set of DPMI calls. This call is used to obtain an address which must be called to use the extensions. The caller points DS:ESI to a null terminated string that specifies the vendor name or some other unique identifier to obtain the specific extension entry point.

Function 0A00H This function gets Tenberry Software's API Entry Point. Pass the following information:

AX = 0A00H

DS:ESI = Pointer to null terminated string "RATIONAL DOS/4G"

If the call succeeds, the carry flag is clear and ES:EDI = Extended API entry point. DS, FS, GS, EAX, EBX, ECX, EDX, ESI, and EBP may be modified.

If the call fails, the carry flag is set.

Notes:

1. Execute a far call to call the API entry point.
2. All extended API parameters are specified by the vendor.
3. The string comparison used to return the API entry point is case sensitive.

11.2.12 Coprocessor Status

Function 0E00H This function gets the coprocessor status. Pass the following information:

AX = 0E00H

If the call succeeds, the carry flag is clear and AX contains the coprocessor status.

Bit	Significance
0	MPv (MP bit in the virtual MSW/CR0). 0 = Numeric coprocessor is disabled for this client. 1 = Numeric coprocessor is disabled for this client.
1	EMv (EM bit in the virtual MSW/CR0). 0 = Client is not emulating coprocessor instructions. 1 = Client is emulating coprocessor instructions.
2	MPr (MP bit from the actual MSW/CR0). 0 = Numeric coprocessor is not present. 1 = Numeric coprocessor is present.

1	EMr (EM bit from the actual MSW/CR0). 0 = Host is not emulating coprocessor instructions. 1 = Host is emulating coprocessor instructions.
4-7	Coprocessor type. 00H = no coprocessor. 02H = 80287 03H = 80387 04H = 80486 with numeric coprocessor 05H = Pentium
8-15	Not applicable.

If the call fails, the carry flag is set.

Notes:

1. If the real EM (EMr) bit is set, the host is supplying or is capable of supplying floating-point emulation.
2. If the MPv bit is not set, the host may not need to save the coprocessor state for this virtual machine to improve system performance.
3. The MPr bit setting should be consistent with the setting of the coprocessor type information. Ignore MPr bit information if it is in conflict with the coprocessor type information.
4. If the virtual EM (EMv) bit is set, the host delivers all coprocessor exceptions to the client, and the client is performing its own floating-point emulation (whether or not a coprocessor is present or the host also has a floating-point emulator). In other words, if the EMv bit is set, the host sets the EM bit in the real CR0 while the virtual machine is active, and reflects coprocessor not present faults (int 7) to the virtual machine.
5. A client can determine the CPU type with int 31H Function 0400H, but a client should not draw any conclusions about the presence or absence of a coprocessor based on the CPU type alone.

Function 0E01H This function sets coprocessor emulation. Pass the following information:

AX = 0E01H
BX = coprocessor bits

<i>Bit</i>	<i>Significance</i>
0	New value of MPv bit for client's virtual CR0. 0 = Disable numeric coprocessor for this client. 1 = Enable numeric coprocessor for this client.
1	New value of EMv bit for client's virtual CR0. 0 = client will not supply coprocessor emulation. 1 = client will supply coprocessor emulation.
2-15	Not applicable.

If the call succeeds, the carry flag is clear; if it fails, the carry flag is set.

12 Utilities

This chapter describes the Tenberry Software *DOS/4GW* utility programs provided with the Open Watcom F77 package. Each program is described using the following format:

Purpose: This is a brief statement of what the utility program does. More specific information is provided under "Notes".

Syntax: This shows the syntax of the program. The fixed portion of each command is in a typewriter font, while variable parts of the command are in *italics*. Optional parts are enclosed in [brackets].

Notes: These are explanatory remarks noting major features and possible pitfalls. We explain anything special that you might need to know about the program.

See Also: This is a cross-reference to any information that is related to the program.

Example: You'll find one or more sample uses of the utility program with an explanation of what the program is doing.

Some of the utilities are *DOS/4GW*-based, protected-mode programs. To determine which programs run in protected mode and which in real, run the program. If you see the *DOS/4GW* banner, the program runs in protected mode.

12.1 DOS4GW

Purpose: Loads and executes linear executables.

Syntax: *linear_executable*

Notes: The stub program at the beginning of the linear executable invokes this program, which loads the linear executable and starts up the DOS extender. The stub program must be able to find DOS4GW: make sure it is in the path.

12.2 PMINFO

Purpose: Measures the performance of protected/real-mode switching and extended memory.

Syntax: PMINFO.EXE

Notes: We encourage you to distribute this program to your users.

The time-based measurements made by PMINFO may vary slightly from run to run.

Example: The following example shows the output of the PMINFO program on a 386 AT-compatible machine.

```
C>pminfo
Protected Mode and Extended Memory Performance Measurement -- 5.00
Copyright (c) Tenberry Software, Inc. 1987 - 1993

DOS memory   Extended memory   CPU performance equivalent to 67.0 MHz 80486
-----
      736           8012   K bytes configured (according to BIOS).
      640          15360   K bytes physically present (SETUP).
      651           7887   K bytes available for DOS/16M programs.
22.0 (3.0)      18.9 (4.0) MB/sec word transfer rate (wait states).
42.9 (3.0)      37.0 (4.0) MB/sec 32-bit transfer rate (wait states).

Overall cpu and memory performance (non-floating point) for typical
DOS programs is 10.36 ± 1.04 times an 8MHz IBM PC/AT.

Protected/Real switch rate = 36156/sec (27 usec/switch, 15 up + 11 down),
DOS/16M switch mode 11 (VCPI).
```

The top information line shows that the CPU performance is equivalent to a 67.0 MHz 80486. Below are the configuration and timings for both the DOS memory and extended memory. If the computer is not equipped with extended memory, or none is available for *DOS/4GW*, the extended memory measurements may be omitted ("--").

The line "according to BIOS" shows the information provided by the BIOS (interrupts 12h and 15h function 88h). The line "SETUP", if displayed, is the configuration obtained directly from the CMOS RAM as set by the computer's setup program. It is displayed only if the numbers are different from those in the BIOS line. They will be different for computers where the BIOS has reserved memory for itself or if another program has allocated some memory and is intercepting the BIOS configuration requests to report less memory available than is physically configured. The "DOS/16M memory range", if displayed, shows the low and high addresses available to *DOS/4GW* in extended memory.

Below the configuration information is information on the memory speed (*transfer rate*). PMINFO tries to determine the memory architecture. Some architectures will perform well under some circumstances and poorly under others; PMINFO will show both the best and worst cases. The architectures detected are cache, interleaved, page-mode (or static column), and direct. Measurements are made using 32-bit accesses and reported as the number of megabytes per second that can be transferred. The number of wait states is reported in parentheses. The wait states can be a fractional number, like 0.5, if there is a wait state on writes but not on reads. Memory bandwidth (i.e., how fast the CPU can access memory) accounts for 60% to 70% of the performance for typical programs (that are not heavily dependent on floating-point math).

A performance metric developed by Tenberry Software is displayed, showing the expected throughput for the computer relative to a standard 8MHz IBM PC/AT (disk accesses and floating point are excluded). Finally, the speed with which the computer can switch between real and protected mode is displayed, both as the maximum number of round-trip switches that can occur per second, and the time for a single round-trip switch, broken out into the real-to-protected (up) and protected-to-real (down) components.

12.3 PRIVATXM

Purpose: Creates a private pool of memory for *DOS/4GW* programs.

Syntax: PRIVATXM [-r]

Notes: This program may be distributed to your users.

Without PRIVATXM, a *DOS/4GW* program that starts up while another *DOS/4GW* program is active uses the pool of memory built by the first program. The new program cannot change the parameters of this memory pool, so setting **DOS16M** to increase the size of the pool has no effect. To specify that the two programs use different pools of memory, use PRIVATXM.

PRIVATXM marks the active *DOS/4GW* programs as private, preventing subsequent *DOS/4GW* programs from using the same memory pool. The first *DOS/4GW* program to start after PRIVATXM sets up a new pool of memory for itself and any subsequent *DOS/4GW* programs. To release the memory used by the private programs, use the PRIVATXM -r option.

PRIVATXM is a TSR that requires less than 500 bytes of memory. It is not supported under DPML.

Example: The following example creates a 512KB memory pool that is shared by two *DOS/4GW* TSRs. Subsequent *DOS/4GW* programs use a different memory pool.

C>set DOS16M=:512	Specifies the size of the memory pool.
C>TSR1	Sets up the memory pool at startup.
C>TSR2	This TSR shares the pool built by TSR1.
C>PRIVATXM	Makes subsequent <i>DOS/4GW</i> programs use a new memory pool.
C>set DOS16M=	Specifies an unlimited size for the new pool.
C>PROGRAM3	This program uses the new memory pool.
C>PRIVATXM -R	Releases the 512KB memory pool used by the TSRs. (If the TSRs shut down, their memory is not released unless PRIVATXM is released.)

12.4 RMINFO

Purpose: Supplies configuration information and the basis for real/protected-mode switching in your machine.

Syntax: RMINFO.EXE

Notes: This program may be distributed to your users.

RMINFO starts up *DOS/4GW*, but stops your machine just short of switching from real mode to protected mode and displays configuration information about your computer. The information shown by RMINFO can help determine why *DOS/4GW* applications won't run on a particular machine. Run RMINFO if PMINFO does not run to completion.

Example: The following example shows the output of the RMINFO program on an 386 AT-compatible machine.

```
C>rminfo

DOS/16M Real Mode Information Program 5.00
Copyright (C) Tenberry Software, Inc. 1987 - 1993

Machine and Environment:
  Processor:           i386, coprocessor present
  Machine type:        10 (AT-compatible)
  A20 now:             enabled
  A20 switch rigor:    disabled
  DPMS host found
Switching Functions:
  To PM switch:        DPMS
  To RM switch:        DPMS
  Nominal switch mode: 0
  Switch control flags: 0000
Memory Interfaces:
  DPMS may provide:    16384K returnable
  Contiguous DOS memory: 463K
```

The information provided by RMINFO includes:

Machine and Environment:

Processor: processor type, coprocessor present/not present

Machine type:

(NEC 9801)
 (PS/2-compatible)
 (AT-compatible)
 (FM R)
 (AT&T 6300+)
 (AT-compatible)
 (C&T 230 chipset)
 (AT-compatible)
 (AT-compatible)
 (Acer)
 (Zenith)
 (Hitachi)
 (Okidata)
 (PS/55)

A20 now: Current state of Address line 20.

A20 switch rigor: Whether DOS4GW rigorously controls enabling and disabling of Address line 20 when switching modes.

PS feature flag

XMS host found Whether your system has any software using extended memory under the XMS discipline.

VCPI host found Whether your system has any software using extended memory under the VCPI discipline.

page table 0 at: x000h

DPMI host found

DOS/16M resident with private/public memory

Switching Functions:

A20 switching:

To PM switch: reset catch:
 pre-PM prep:
 post-PM-switch:

To RM switch:
 pre-RM prep:
 reset method:
 post-reset:
 reset uncatch:

Nominal switch mode: *x*

Switch control flags: *xxxxh*

Memory Interfaces:

(VCPI remapping in effect)

DPMI may provide: xxxxxK returnable

VCPI may provide: xxxxxK returnable

Top-down

Other16M

Forced

Contiguous DOS memory:

13 Error Messages

The following lists DOS/4G error messages, with descriptions of the circumstances in which the error is most likely to occur, and suggestions for remedying the problem. Some error messages pertaining to features — like DLLs — that are not supported in *DOS/4GW* will not arise with that product. In the following descriptions, references to DOS/4G, DOS4G, or DOS4G.EXE may be replaced by DOS/4GW, DOS4GW, or DOS4GW.EXE should the error message arise when using *DOS/4GW*.

13.1 Kernel Error Messages

This section describes error messages from the DOS/16M kernel embedded in DOS/4G. Kernel error messages may occur because of severe resource shortages, corruption of DOS4GW.EXE, corruption of memory, operating system incompatibilities, or internal errors in DOS/4GW. All of these messages are quite rare.

0. involuntary switch to real mode

The computer was in protected mode but switched to real mode without going through DOS/16M. This error most often occurs because of an unrecoverable stack segment exception (stack overflow), but can also occur if the Global Descriptor Table or Interrupt Descriptor Table is corrupted. Increase the stack size, recompile your program with stack overflow checking, or look into ways that the descriptor tables may have been overwritten.

1. not enough extended memory

2. not a DOS/16M executable <filename>

DOS4G.EXE, or a bound DOS/4G application, has probably been corrupted in some way. Rebuild or recopy the file.

3. no DOS memory for transparent segment

4. cannot make transparent segment

5. too many transparent segments

6. not enough memory to load program

There is not enough memory to load DOS/4G. Make more memory available and try again.

7. no relocation segment

8. cannot open file <filename>

The DOS/16M loader cannot load DOS/4G, probably because DOS has run out of file units. Set a larger FILES= entry in CONFIG.SYS, reboot, and try again.

9. cannot allocate tstack

There is not enough memory to load DOS/4G. Make more memory available and try again.

10. cannot allocate memory for GDT

There is not enough memory to load DOS/4G. Make more memory available and try again.

11. no passup stack selectors -- GDT too small

This error indicates an internal error in DOS/4G or an incompatibility with other software.

12. no control program selectors -- GDT too small

This error indicates an internal error in DOS/4G or an incompatibility with other software.

13. cannot allocate transfer buffer

There is not enough memory to load DOS/4G. Make more memory available and try again.

14. premature EOF

DOS4G.EXE, or a bound DOS/4G application, has probably been corrupted in some way. Rebuild or recopy the file.

15. protected mode available only with 386 or 486

DOS/4G requires an 80386 (or later) CPU. It cannot run on an 80286 or earlier CPU.

16. cannot run under OS/2

17. system software does not follow VCPI or DPML specifications

Some memory resident program has put your 386 or 486 CPU into Virtual 8086 mode. This is done to provide special memory services to DOS programs, such as EMS simulation (EMS interface without EMS hardware) or high memory. In this mode, it is not possible to switch into protected mode unless the resident software follows a standard that DOS/16M supports (DPML, VCPI, and XMS are the most common). Contact the vendor of your memory management software.

18. you must specify an extended memory range (SET DOS16M=)

On some Japanese machines that are not IBM AT-compatible, and have no protocol for managing extended memory, you must set the DOS16M environment variable to specify the range of available extended memory.

19. computer must be AT- or PS/2- compatible

20. unsupported DOS16M switchmode choice

21. requires DOS 3.0 or later

22. cannot free memory

This error probably indicates that memory was corrupted during execution of your program.

23. *no memory for VCPI page table*

There is not enough memory to load DOS/4G. Make more memory available and try again.

24. *VCPI page table address incorrect*

This is an internal error.

25. *cannot initialize VCPI*

This error indicates an incompatibility with other software. DOS/16M has detected that VCPI is present, but VCPI returns an error when DOS/16M tries to initialize the interface.

26. *8042 timeout***27. *extended memory is configured but it cannot be allocated*****28. *memory error, avail loop***

This error probably indicates that memory was corrupted during execution of your program. Using an invalid or stale alias selector may cause this error. Incorrect manipulation of segment descriptors may also cause it.

29. *memory error, out of range*

This error probably indicates that memory was corrupted during execution of your program. Writing through an invalid or stale alias selector may cause this error.

30. *program must be built -AUTO for DPMI***31. *protected mode already in use in this DPMI virtual machine*****32. *DPMI host error (possibly insufficient memory)*****33. *DPMI host error (need 64K XMS)*****34. *DPMI host error (cannot lock stack)***

Any of these errors (32, 33, 34) probably indicate insufficient memory under DPMI. Under Windows, you might try making more physical memory available by eliminating or reducing any RAM drives or disk caches. You might also try editing DEFAULT.PIF so that at least 64KB of XMS memory is available to non-Windows programs. Under OS/2, you want to increase the DPMI_MEMORY_LIMIT in the DOS box settings.

35. *General Protection Fault*

This message probably indicates an internal error in DOS/4G. Faults generated by your program should cause error 2001 instead.

36. *The DOS16M.386 virtual device driver was never loaded***37. *Unable to reserve selectors for DOS16M.386 Windows driver***

38. Cannot use extended memory: HIMEM.SYS not version 2

This error indicates an incompatibility with an old version of HIMEM.SYS.

39. An obsolete version of DOS16M.386 was loaded

40. not enough available extended memory (XMIN)

This message probably indicates an incompatibility with your memory manager or its configuration. Try configuring the memory manager to provide more extended memory, or change memory managers.

13.2 DOS/4G Errors

1000 "can't hook interrupts"

A DPMI host has prevented DOS/4G from loading. Please contact Tenberry Technical Support.

1001 "error in interrupt chain"

DOS/4G internal error. Please contact Tenberry Technical Support.

1003 "can't lock extender kernel in memory"

DOS/4G couldn't lock the kernel in physical memory, probably because of a memory shortage.

1004 "syntax is DOS4G <executable.xxx>"

You must specify a program name.

1005 "not enough memory for dispatcher data"

There is not enough memory for DOS/4G to manage user-installed interrupt handlers properly. Free some memory for the DOS/4G application.

1007 "can't find file <program> to load"

DOS/4G could not open the specified program. Probably the file didn't exist. It is possible that DOS ran out of file handles, or that a network or similar utility has prohibited read access to the program. Make sure that the file name was spelled correctly.

1008 "can't load executable format for file <filename> [<error code>]"

DOS/4G did not recognize the specified file as a valid executable file. DOS/4G can load linear executables (LE and LX) and EXPs (BW). The error code is for Tenberry Software's use.

1009 "program <filename> is not bound"

This message does not occur in DOS/4G, only DOS/4GW Professional; the latter requires that the DOS extender be bound to the program file. The error signals an attempt to load

1010 "can't initialize loader <loader> [<error code>]"

DOS/4G could not initialize the named loader, probably because of a resource shortage. Try making more memory available. If that doesn't work, please contact Tenberry Technical Support. The error code is for Tenberry Software' use.

1011 "VMM initialization error [<error code>]"

DOS/4G could not initialize the Virtual Memory Manager, probably because of a resource shortage. Try making more memory available. If that doesn't work, please contact Tenberry Technical Support. The error code is for Tenberry Software' use.

1012 "<filename> is not a WATCOM program"

This message does not occur in DOS/4G, only DOS/4GW and DOS/4GW Professional. Those extenders only support WATCOM 32-bit compilers.

1013 "int 31h initialization error"

DOS/4G was unable to initialize the code that handles Interrupt 31h, probably because of an internal error. Please call Tenberry Technical Support.

1100 "assertion \"<statement>\" failed (<file>:<line>)"

DOS/4G internal error. Please contact Tenberry Technical Support.

1200 "invalid EXP executable format"

DOS/4G tried to load an EXP, but couldn't. The executable file is probably corrupted.

1201 "program must be built -AUTO for DPMI"

Under DPMI, DOS/4G can only load EXPs that have been linked with the GLU -AUTO or -DPMI switch.

1202 "can't allocate memory for GDT"

There is not enough memory available for DOS/4G to build a Global Descriptor Table. Make more memory available.

1203 "premature EOF"

DOS/4G tried to load an EXP but couldn't. The file is probably corrupted.

1204 "not enough memory to load program"

There is not enough memory available for DOS/4G to load your program. Make more memory available.

1301 "invalid linear executable format"

DOS/4G cannot recognize the program file as a LINEXE format. Make sure that you specified the correct file name.

1304 "file I/O seek error"

DOS/4G was unable to seek to a file location that should exist. This usually indicates truncated program files or problems with the storage device from which your program loads. Run CHKDSK or a similar utility to begin determining possible causes.

1305 "file I/O read error"

DOS/4G was unable to read a file location that should contain program data. This usually indicates truncated program files or problems with the storage device from which your program loads. Run CHKDSK or a similar utility to begin determining possible causes.

1307 "not enough memory"

As it attempted to load your program, DOS/4G ran out of memory. Make more memory available, or enable VMM.

1308 "can't load requested program"

1309 "can't load requested program"

1311 "can't load requested program"

1312 "can't load requested program"

DOS/4G cannot load your program for some reason. Contact Tenberry Technical Support.

1313 "can't resolve external references"

DOS/4G was unable to resolve all references to DLLs for the requested program, or the program contained unsupported fixup types. Use EXEHDR or a similar LINEXE dump utility to see what references your program makes and what special fixup records might be present.

1314 "not enough lockable memory"

As it attempted to load your program, DOS/4G encountered a refusal to lock a virtual memory region. Some memory must be locked in order to handle demand-load page faults. Make more physical memory available.

1315 "can't load requested program"

1316 "can't load requested program"

DOS/4G cannot load your program for some reason. Contact Tenberry Technical Support.

1317 "program has no stack"

DOS/4G reports this error when you try to run a program with no stack. Rebuild your program, building in a stack.

2000 "deinitializing twice"

DOS/4G internal error. Please contact Tenberry Technical Support.

2001 "exception <exception_number> (<exception_description>) at <selector:offset>"

Your program has generated an exception. For information about interpreting this message, see the file COMMON.DOC.

2002 "transfer stack overflow at <selector:offset>"

Your program has overflowed the DOS/4G transfer stack. For information about interpreting this message, see the file COMMON.DOC.

2300 "can't find <DLL>.<ordinal> - referenced from <module>"

DOS/4G could not find the ordinal listed in the specified DLL, or it could not find the DLL at all. Correct or remove the reference, and make sure that DOS/4G can find the DLL.

DOS/4G looks for DLLs in the following directories:

- The directory specified by the Libpath32 configuration option (which defaults to the directory of the main application file).
- The directory or directories specified by the LIBPATH32 environment variable.
- Directories specified in the PATH.

2301 "can't find <DLL>.<name> - referenced from <module>"

DOS/4G could not find the entry point named in the specified module. Correct or remove the reference, and make sure that DOS/4G can find the DLL.

2302 "DLL modules not supported"

This DOS/4GW Professional error message arises when an application references or tries to explicitly load a DLL. DOS/4GW Professional does not support DLLs.

2303 "internal LINEXE object limit reached"

DOS/4G currently handles a maximum of 128 LINEXE objects, including all .DLL and .EXE files. Most .EXE or .DLL files use only three or four objects. If possible, reduce the number of objects, or contact Tenberry Technical Support.

2500 "can't connect to extender kernel"

DOS/4G internal error. Please contact Tenberry Technical Support.

2503 "not enough disk space for swapping - <count> bytes required"

VMM was unable to create a swap file of the required size. Increase the amount of disk space available.

2504 "can't create swap file \<filename>\\""

VMM was unable to create the swap file. This could be because the swap file is specified for a nonexistent drive or on a drive that is read-only. Set the SWAPNAME parameter to change the location of the swap file.

2505 "not enough memory for <table>"

VMM was unable to get sufficient extended memory for internal tables. Make more memory available. If <table> is page buffer, make more DOS memory available.

2506 "not enough physical memory (minmem)"

There is less physical memory available than the amount specified by the MINMEM parameter. Make more memory available.

2511 "swap out error [<error code>]"

Unknown disk error. The error code is for Tenberry Software' use.

2512 "swap in error [<error code>]"

Unknown disk error. The error code is for Tenberry Software' use.

2514 "can't open trace file"

VMM could not open the VMM.TRC file in the current directory for writing. If the directory already has a VMM.TRC file, delete it. If not, there may not be enough memory on the drive for the trace file, or DOS may not have any more file handles.

2520 "can't hook int 31h"

DOS/4G internal error. Please contact Tenberry Technical Support.

2523 "page fault on non-present mapped page"

Your program references memory that has been mapped to a nonexistent physical device, using DPMI function 508h. Make sure the device is present, or remove the reference.

2524 "page fault on uncommitted page"

Your program references memory reserved with a call to DPMI function

504h, but never committed (using a DPMI 507h or 508h call). Commit the memory before you reference it.

3301 "unhandled EMPTYFWD, GATE16, or unknown relocation"

3302 "unhandled ALIAS16 reference to unaliased object"

3304 "unhandled or unknown relocation"

If your program was built for another platform that supports the LINEXE format, it may contain a construct that DOS/4G does not currently support, such as a call gate. This message may also occur if your program has a problem mixing 16- and 32-bit code. A linker error is another likely cause.

14 *DOS/4GW Commonly Asked Questions*

The following information has been provided by Tenberry Software, Inc. for their DOS/4GW and DOS/4GW Professional product. The content of this chapter has been edited by Open Watcom. In most cases, the information is applicable to both products.

This chapter covers the following topics:

- Differences within the DOS/4G product line
- Addressing
- Interrupt and exception handling
- Memory management
- DOS, BIOS, and mouse services
- Virtual memory
- Debugging
- Compatibility

1a. References.

The DOS/4GW documentation from Open Watcom is the primary reference for DOS/4GW Professional as well. Another useful reference is the DPMI specification.

The DPMI 1.0 specification can be obtained from the Open Watcom website <https://open-watcom.github.io/open-watcom-1.9/devel/dpmi10.pdf>. This file contains a PDF version of the DPMI 1.0 specification.

Alternatively the DPMI 1.0 specification can be obtained at <http://www.delorie.com/djgpp/doc/dpmi/dpmi-specs.tar.gz>. This file contains online HTML version of the DPMI 1.0 specification.

14.1 Differences Within the DOS/4G Product Line

2a. DOS/4GW Professional versus DOS/4GW

DOS/4GW Professional was designed to be a higher-performance version of DOS/4GW suitable for commercial applications. Here is a summary of the advantages of DOS/4GW Professional with respect to DOS/4GW:

- Extender binds to the application program file
- Extender startup time has been reduced
- Support for Open Watcom floating-point emulator has been optimized
- Virtual memory manager performance has been greatly improved
- Under VMM, programs are demand loaded
- Virtual address space is 4 GB instead of 32 MB
- Extender memory requirements have been reduced by more than 50K
- Extender disk space requirements have been reduced by 40K
- Can omit virtual memory manager to save 50K more disk space
- Support for INT 31h functions 301h-304h and 702h-703h

DOS/4GW Professional is intended to be fully compatible with programs written for DOS/4GW 1.9 and up. The only functional difference is that the extender is bound to your program instead of residing in a separate file. Not only does this help reduce startup time, but it eliminates version-control problems when someone has both DOS/4GW and DOS/4GW Professional applications present on one machine.

2b. DOS/4GW Professional versus DOS/4G.

DOS/4GW Professional is not intended to provide any other new DOS extender functionality. Tenberry Software's top-of-the-line 32-bit extender, DOS/4G, is not sold on a retail basis but is of special interest to developers who require more flexibility (such as OEMs). DOS/4G offers these additional features beyond DOS/4GW and DOS/4GW Professional:

- Complete documentation
- DLL support
- TSR support
- Support for INT 31h functions 301h-306h, 504h-50Ah, 702h-703h
- A C language API that offers more control over interrupt handling and program loading, as well as making it easier to use the extender
- An optional (more protected) nonzero-based flat memory model
- Remappable error messages
- More configuration options
- The D32 debugger, GLU linker, and other tools
- Support for other compilers besides Open Watcom
- A higher level of technical support

- Custom work is available (e.g., support for additional executable formats, operating system API emulations, mixed 16-bit and 32-bit code)

Please contact Tenberry Software if you have questions about other products (present or future) in the DOS/4G line.

2c. DPMI functions supported by DOS/4GW.

Note that when a DOS/4GW application runs under a DPMI host, such as Windows 3.1 in enhanced mode, an OS/2 virtual DOS machine, 386Max (with DEBUG=DPMIXCOPY), or QDPMI (with EXTCHKOFF), the DPMI host provides the DPMI services, not DOS/4GW. The DPMI host also provides virtual memory, if any. Performance (speed and memory use) under different DPMI hosts varies greatly due to the quality of the DPMI implementation.

These are the services provided by DOS/4GW and DOS/4GW Professional in the absence of a DPMI host.

0000	Allocate LDT Descriptors
0001	Free LDT Descriptor
0002	Map Real-Mode Segment to Descriptor
0003	Get Selector Increment Value
0006	Get Segment Base Address
0007	Set Segment Base Address
0008	Set Segment Limit
0009	Set Descriptor Access Rights
000A	Create Alias Descriptor
000B	Get Descriptor
000C	Set Descriptor
000D	Allocate Specific LDT Descriptor
0100	Allocate DOS Memory Block
0101	Free DOS Memory Block
0102	Resize DOS Memory Block
0200	Get Real-Mode Interrupt Vector
0201	Set Real-Mode Interrupt Vector
0202	Get Processor Exception Handler
0203	Set Processor Exception Handler
0204	Get Protected-Mode Interrupt Vector
0205	Set Protected-Mode Interrupt Vector
0300	Simulate Real-Mode Interrupt
0301	Call Real-Mode Procedure with Far Return Frame (DOS/4GW Professional only)
0302	Call Real-Mode Procedure with IRET Frame (DOS/4GW Professional only)
0303	Allocate Real-Mode Callback Address (DOS/4GW Professional only)
0304	Free Real-Mode Callback Address (DOS/4GW Professional only)
0400	Get DPMI Version
0500	Get Free Memory Information
0501	Allocate Memory Block
0502	Free Memory Block

0503	Resize Memory Block
0600	Lock Linear Region
0601	Unlock Linear Region
0604	Get Page Size (VM only)
0702	Mark Page as Demand Paging Candidate (DOS/4GW Professional only)
0703	Discard Page Contents (DOS/4GW Professional only)
0800	Physical Address Mapping
0801	Free Physical Address Mapping
0900	Get and Disable Virtual Interrupt State
0901	Get and Enable Virtual Interrupt State
0902	Get Virtual Interrupt State
0A00	Get Tenberry Software API Entry Point
0E00	Get Coprocessor Status
0E01	Set Coprocessor Emulation

14.2 Addressing

3a. Converting between pointers and linear addresses.

Because DOS/4GW uses a zero-based flat memory model, converting between pointers and linear addresses is trivial. A pointer value is always relative to the current segment (the value in CS for a code pointer, or in DS or SS for a data pointer). The segment bases for the default DS, SS, and CS are all zero. Hence a near pointer is exactly the same thing as a linear address: a null pointer points to linear address 0, and a pointer with value 0x10000 points to linear address 0x10000.

3b. Converting between code and data pointers.

Because DS and CS have the same base address, they are natural aliases for each other. To create a data alias for a code pointer, merely create a data pointer and set it equal to the code pointer. It's not necessary for you to create your own alias descriptor. Similarly, to create a code alias for a data pointer, merely create a code pointer and set it equal to the data pointer.

3c. Converting between pointers and low memory addresses.

Linear addresses under 1 MB map directly to physical memory. Hence the real-mode interrupt vector table is at address 0, the BIOS data segment is at address 0x400, the monochrome video memory is at address 0xB0000, and the color video memory is at address 0xB8000. To read and write any of these, you can just use a pointer set to the proper address. You don't need to create a far pointer, using some magic segment value.

3d. Converting between linear and physical addresses.

Linear addresses at or above 1 MB do not map directly to physical memory, so you can not in general read or write extended memory directly, nor can you tell how a particular block of extended memory has been used.

DOS/4GW supports the DPMI call INT 31h/800h, which maps physical addresses to linear addresses. In other words, if you have a peripheral device in your machine that has memory at a physical address of 256 MB, you can issue this call to create a linear address that points to that physical memory. The linear address is the same thing as a near pointer to the memory and can be manipulated as such.

There is no way in a DPMI environment to determine the physical address corresponding to a given linear address. This is part of the design of DPMI. You must design your application accordingly.

3e. Null pointer checking.

DOS/4GW will trap references to the first sixteen bytes of physical memory if you set the environment variable DOS4G=NULLP. This is currently the only null-pointer check facility provided by DOS/4GW.

As of release 1.95, DOS/4GW traps both reads and writes. Prior to this, it only trapped writes.

You may experience problems if you set DOS4G=NULLP and use some versions of the Open Watcom Debugger with a 1.95 or later extender. These problems have been corrected in later versions of the Open Watcom Debugger.

14.3 Interrupt and Exception Handling

4a. Handling asynchronous interrupts.

Under DOS/4GW, there is a convenient way to handle asynchronous interrupts and an efficient way to handle them.

Because your CPU may be in either protected mode (when 32-bit code is executing) or real mode (a DOS or BIOS call) when a hardware interrupt comes in, you have to be prepared to handle interrupts in either mode. Otherwise, you may miss interrupts.

You can handle both real-mode and protected-mode interrupts with a single handler, if 1) the interrupt is in the auto-passup range, 8 to 2Eh; and 2) you install a handler with INT 21h/25h or `_dos_setvect()`; 3) you do not install a handler for the same interrupt using any other mechanism. DOS/4GW will route both protected-mode interrupts and real-mode interrupts to your protected-mode handler. This is the convenient way.

The efficient way is to install separate real-mode and protected-mode handlers for your interrupt, so your CPU won't need to do unnecessary mode switches. Writing a real-mode handler is tricky; all you can reasonably expect to do is save data in a buffer and IRET. Your protected-mode code can periodically check the buffer and process any queued data. (Remember, protected-mode code can access data and execute code in low memory, but real-mode code can't access data or execute code in extended memory.)

For performance, it doesn't matter how you install the real-mode handler, but we recommend the DPMI function INT 31h/201h for portability.

It does matter how you install the protected-mode handler. You can't install it directly into the IDT, because a DPML provider must distinguish between interrupts and exceptions and maintain separate handler chains. Installing with INT 31h/205h is the recommended way to install your protected-mode handler for both performance and portability.

If you install a protected-mode handler with INT 21h/25h, both interrupts and exceptions will be funneled to your handler, to mimic DOS. Since DPML exception handlers and interrupt handlers are called with different stack frames, DOS/4GW executes a layer of code to cover these differences up; the same layer is used to support the DOS/4G API (not part of DOS/4GW). This layer is the reason that hooking with INT 21h/25h is less efficient than hooking with INT 31h/205h.

4b. Handling asynchronous interrupts in the second IRQ range.

Because the second IRQ range (normally INTs 70h-77h) is outside the DOS/4GW auto-passup range (8-2Eh, excluding 21h) you may not handle these interrupts with a single handler, as described above (the "convenient" method). You must install separate real-mode and protected-mode handlers (the "efficient" method).

DOS/4G does allow you to specify additional passup interrupts, however.

4c. Asynchronous interrupt handlers and DPML.

The DPML specification requires that all code and data referenced by a hardware interrupt handler **MUST** be locked at interrupt time. A DPML virtual memory manager can use the DOS file system to swap pages of memory to and from the disk; because DOS is not reentrant, a DPML host is not required to be able to handle page faults during asynchronous interrupts. Use INT 31h/600h (Lock Linear Region) to lock an address range in memory.

If you fail to lock all of your code and data, your program may run under DOS/4GW, but fail under the DOS/4GW Virtual Memory Manager or under another DPML host such as Windows or OS/2.

You should also lock the code and data of a mouse callback function.

4d. Open Watcom signal() function and Ctrl-Break.

In earlier versions of the Watcom C/C++ library, there was a bug that caused signal(SIGBREAK) not to work. Calling signal(SIGBREAK) did not actually install an interrupt handler for Ctrl-Break (INT 1Bh), so Ctrl-Break would terminate the application rather than invoking the signal handler.

With these earlier versions of the library, you could work around this problem by hooking INT 1Bh directly. With Watcom release 10.0, this problem has been fixed.

4e. More tips on writing hardware interrupt handlers.

- It's more like handling interrupts in real mode than not.

The same problems arise when writing hardware interrupt handlers for protected mode as arise for real mode. We assume you know how to write real-mode handlers; if our suggestions don't seem clear, you might want to brush up on real-mode interrupt programming.

- Minimize the amount of time spent in your interrupt handlers.

When your interrupt handlers are called, interrupts are disabled. This means that no other system tasks can be performed until you enable interrupts (an STI instruction) or until your handler returns. In general, it's a good idea to handle interrupts as quickly as possible.

- Minimize the amount of time spent in the DOS extender by installing separate real-mode and protected-mode handlers.

If you use a passup interrupt handler, so that interrupts received in real mode are resigalled in protected mode by the extender, your application has to switch from real mode to protected mode to real mode once per interrupt. Mode switching is a time-consuming process, and interrupts are disabled during a mode switch. Therefore, if you're concerned about performance, you should install separate handlers for real-mode and protected-mode interrupts, eliminating the mode switch.

- If you can't just set a flag and return, enable interrupts (STI).

Handlers that do more than just set a flag or store data in a buffer should re-enable interrupts as soon as it's safe to do so. In other words, save your registers on the stack, establish your addressing conventions, switch stacks if you're going to — and then enable interrupts (STI), to give priority to other hardware interrupts.

- If you enable interrupts (STI), you should disable interrupts (CLI).

Because some DPMI hosts virtualize the interrupt flag, if you do an STI in your handler, you should be sure to do a CLI before you return. (CLI, then switch back to the original stack if you switched away, then restore registers, then IRET.) If you don't do this, the IRET will not necessarily restore the previous interrupt flag state, and your program may crash. This is a difference from real-mode programming, and it tends to show up as a problem when you try running your program in a Windows or OS/2 DOS box for the first time (but not before).

- Add a reentrancy check.

If your handler doesn't complete its work by the time the next interrupt is signalled, then interrupts can quickly nest to the point of overflowing the transfer stack. This is a design flaw in your program, not in the DOS extender; a real-mode DOS program can have exactly the same behavior. If you can conceive of a situation where your interrupt handler can be called again before the first instance returns, you need to code in a reentrancy check of some sort (before you switch stacks and enable interrupts (STI), obviously).

Remember that interrupts can take different amounts of time to execute on different machines; the CPU manufacturer, CPU speed, speed of memory accesses, and CMOS settings (e.g. "system BIOS shadowing") can all affect performance in subtle ways. We recommend you program defensively and always check for unexpected reentry, to avoid transfer stack overflows.

- Switch to your own stack.

Interrupt handlers are called on a stack that typically has only a small amount of stack available (512 bytes or less). If you need to use more stack than this, you have to switch to your own stack on entry into the handler, and switch back before returning.

If you want to use C run-time library functions, which are compiled for flat memory model (SS == DS, and the base of CS == the base of DS), you need to switch back to a stack in the flat data segment first.

Note that switching stacks by itself won't prevent transfer stack overflows of the kind described above.

14.4 Memory Management

5a. Using the realloc() function.

In versions of Watcom C/C++ prior to 9.5b, there was a bug in the library implementation of realloc() under DOS/4GW and DOS/4GW Professional. This bug was corrected by Watcom in the 9.5b release.

5b. Using all of physical memory.

DOS/4GW Professional is currently limited to 64 MB of physical memory. We do not expect to be able to fix this problem for at least six months. If you need more than 64 MB of memory, you must use virtual memory.

14.5 DOS, BIOS, and Mouse Services

6a. Speeding up file I/O.

The best way to speed up DOS file I/O in DOS/4GW is to write large blocks (up to 65535 bytes, or the largest number that will fit in a 16-bit int) at a time from a buffer in low memory. In this case, DOS/4GW has to copy the least amount of data and make the fewest number of DOS calls in order to process the I/O request.

Low memory is allocated through INT 31h/0100h, Allocate DOS Memory Block. You can convert the real-mode segment address returned by INT 31h/0100h to a pointer (suitable for passing to setvbuf()) by shifting it left four bits.

6b. Spawning.

It is possible to spawn one DOS/4GW application from another. However, two copies of the DOS extender will be loaded into memory. DOS/4G supports loading of multiple programs atop a single extender, as well as DLLs.

6c. Mouse callbacks.

DOS/4GW Professional now supports the INT 31h interface for managing real-mode callbacks. However, you don't need to bother with them for their single most important application: mouse callback functions. Just register your protected-mode mouse callback function as you would in real mode, by issuing INT 33h/0Ch with the event mask in CX and the function address in ES:EDX, and your function will work as expected.

Because a mouse callback function is called asynchronously, the same locking requirement exists for a mouse callback function as for a hardware interrupt handler. See (4c) above.

6d. VESA support.

While DOS/4GW automatically handles most INT 10h functions so that you can issue them from protected mode, it does not translate the INT 10h VESA extensions. The workaround is to use INT 31h/300h (Simulate Real-Mode Interrupt).

14.6 Virtual Memory

7a. Testing for the presence of VMM.

INT 31h/400h returns a value (BX, bit 2) that tells if virtual memory is available. Under a DPMI host such as Windows 3.1, this will be the host's virtual memory manager, not DOS/4GW's.

You can test for the presence of a DOS/4G-family DOS extender with INT 31h/0A00h, with a pointer to the null-terminated string "RATIONAL DOS/4G" in DS:ESI. If the function returns with carry clear, a DOS/4G-family extender is running.

7b. Reserving memory for a spawned application.

If you spawn one DOS/4GW application from another, you should set the DELETESWAP configuration option (i.e., SET DOS4GVM=deleteswap) so that the two applications don't try to use the same swap file. You should also set the MAXMEM option low enough so that the parent application doesn't take all available physical memory; memory that's been reserved by the parent application is not available to the child application.

7c. Instability under VMM.

A program that hooks hardware interrupts, and works fine without VMM but crashes sporadically with it, probably needs to lock the code and data for its hardware interrupt handlers down in memory. DOS/4GW does not support page faults during hardware interrupts, because DOS services may not be available at that time. See (4c) and (6c) above.

Memory can be locked down with INT 31h/600h (Lock Linear Region).

7d. Running out of memory with a huge virtual address space.

Because DOS/4GW has to create page tables to describe your virtual address space, we recommend that you set your VIRTUALSIZE parameter just large enough to accommodate your program. If you set your VIRTUALSIZE to 4 GB, the physical memory occupied by the page tables will be 4 MB, and that memory will not be available to DOS/4GW.

7e. Reducing the size of the swap file.

DOS/4GW will normally create a swap file equal to your VIRTUALSIZE setting, for efficiency. However, if you set the SWAPMIN parameter to a size (in KB), DOS/4GW will start with a swap file of that size, and will grow the swap file when it has to. The SWAPINC value (default 64 KB) controls the incremental size by which the swap file will grow.

7f. Deleting the swap file.

The DELETESWAP option has two effects: telling DOS/4GW to delete the swap file when it exits, and causing DOS/4GW to provide a unique swap file name if an explicit SWAPNAME setting was not given.

DELETESWAP is required if one DOS/4GW application is to spawn another; see (7b) above.

7g. Improving demand-load performance of large static arrays.

DOS/4GW demand-loading feature normally cuts the load time of a large program drastically. However, if your program has large amounts of global, zero-initialized data (storage class BSS), the Watcom startup code will explicitly zero it (version 9.5a or earlier). Because the zeroing operation touches every page of the data, the benefits of demand-loading are lost.

Demand loading can be made fast again by taking advantage of the fact that DOS/4GW automatically zeroes pages of BSS data as they are loaded. You can make this change yourself by inserting a few lines into the startup routine, assembling it (MASM 6.0 will work), and listing the modified object module first when you link your program.

Here are the changes for \watcom\src\startup\386\cstart3r.asm (startup module from the Watcom C/C++ 9.5 compiler, library using register calling conventions). You can modify the workaround easily for other Open Watcom compilers:

```
...                                ; cstart3r.asm, circa line 332
                                ; end of _BSS segment (start of STACK)
mov     ecx,offset DGROUP:_end
                                ; start of _BSS segment
mov     edi,offset DGROUP:_edata
;-----; RSI OPTIMIZATION
mov     eax,edi                 ; minimize _BSS initialization loop
or      eax,0FFFh               ; compute address of first page after
inc     eax                     ; start of _BSS
cmp     eax,ecx                 ; if _BSS extends onto that page,
jae     allzero                 ; then we can rely on the loader
mov     ecx,eax                 ; zeroing the remaining pages
allzero:
;-----; END RSI OPTIMIZATION
sub     ecx,edi                 ; calc # of bytes in _BSS segment
mov     dl,cl                   ; save bottom 2 bits of count in edx
shr     ecx,2                   ; calc # of dwords
sub     eax,eax                 ; zero the _BSS segment
rep     stosd                   ; ...
mov     cl,dl                   ; get bottom 2 bits of count
and     cl,3                    ; ...
rep     stosb                   ; ...
...
```

Note that the Watcom 9.5b and later versions of the Open Watcom C library already contain this enhancement.

7h. How should I configure VM for best performance?

Here are some recommendations for setting up the DOS/4GW virtual memory manager.

- VIRTUALSIZE** Set to no more than twice the total amount of memory (virtual and otherwise) your program requires. If your program has 16 MB of code and data, set to 32 MB. (There is only a small penalty for setting this value larger than you will need, but your program won't run if you set it too low.) See (7d) above.
- MINMEM** Set to the minimum hardware requirement for running your application. (If you require a 2 MB machine, set to 2048).
- MAXMEM** Set to the maximum amount of memory you want your application to use. If you don't spawn any other applications, set this large (e.g., 32000) to make sure you can use all available physical memory. If you do spawn, see (7b) above.
- SWAPMIN** Don't use this if you want the best possible VM performance. The trade-off is that DOS/4GW will create a swap file as big as your VIRTUALSIZE.
- SWAPINC** Don't use this if you want the best possible VM performance.
- DELETESWAP** DOS/4GW's VM will start up slightly slower if it has to create the swap file afresh each time. However, unless your swap file is very large, DELETESWAP is a reasonable choice; it may be required if you spawn another DOS/4GW program at the same time. See (7b) above.

14.7 Debugging

8a. Attempting to debug a bound application.

You can't debug a bound application. The 4GWBIND utility (included with DOS/4GW Professional) will allow you to take apart a bound application so that you can debug it:

```
4GWBIND -U <boundapp.exe> <yourapp.exe>
```

8b. Debugging with an old version of the Watcom debugger.

DOS/4GW supports versions 8.5 and up of the Open Watcom C, C++ and FORTRAN compilers. However, in order to debug your unbound application with a Open Watcom debugger, you must have version 9.5a or later of the debugger.

If you have an older version of the debugger, we strongly recommend that you contact Open Watcom to upgrade your compiler and tools. The only way to debug a DOS/4GW Professional application with an old version of the debugger is to rename 4GWPRO.EXE to DOS4GW.EXE and make sure that it's either in the current directory or the first DOS4GW.EXE on the DOS PATH.

Tenberry will not provide technical support for this configuration; it's up to you to keep track of which DOS extender is which.

8c. Meaning of "unexpected interrupt" message/error 2001.

In version 1.95 of DOS/4GW, we revised the "unexpected interrupt" message to make it easier to understand.

For example, the message:

```
Unexpected interrupt 0E (code 0) at 168:10421034
```

is now printed:

```
error (2001): exception 0Eh (page fault) at 168:10421034
```

followed by a register dump, as before.

This message indicates that the processor detected some form of programming error and signaled an exception, which DOS/4GW trapped and reported. Exceptions which can be trapped include:

00h	divide by zero
01h	debug exception OR null pointer used
03h	breakpoint
04h	overflow
05h	bounds
06h	invalid opcode
07h	device not available
08h	double fault
09h	overrun
0Ah	invalid TSS
0Bh	segment not present
0Ch	stack fault
0Dh	general protection fault
0Eh	page fault

When you receive this message, this is the recommended course of action:

1. Record all of the information from the register dump.
2. Determine the circumstances under which your program fails.
3. Consult your debugger manual, or an Intel 386, 486 or Pentium Programmer's Reference Manual, to determine the circumstances under which the processor will generate the reported exception.
4. Get the program to fail under your debugger, which should stop the program as soon as the exception occurs.
5. Determine from the exception context why the processor generated an exception in this particular instance.

8d. Meaning of "transfer stack overflow" message/error 2002.

In version 1.95 of DOS/4GW, we added more information to the "transfer stack overflow" message. The message (which is now followed by a register dump) is printed:

```
error (2002): transfer stack overflow
on interrupt <number> at <address>
```

This message means DOS/4GW detected an overflow on its interrupt handling stack. It usually indicates either a recursive fault, or a hardware interrupt handler that can't keep up with the rate at which interrupts are occurring. The best way to understand the problem is to use the VERBOSE option in DOS/4GW to dump the interrupt history on the transfer stack; see (8e) below.

8e. Making the most of a DOS/4GW register dump.

If you can't understand your problem by running it under a debugger, the DOS/4GW register dump is your best debugging tool. To maximize the information available for postmortem debugging, set the environment variable DOS4G to VERBOSE, then reproduce the crash and record the output.

Here's a typical register dump with VERBOSE turned on, with annotations.

```
1 DOS/4GW error (2001): exception 0Eh (page fault)
                                     at 170:0042C1B2
2 TSF32: prev_tsf32 67D8
3 SS      178 DS      178 ES      178 FS      0 GS      20
  EAX 1F000000 EBX      0 ECX  43201C EDX      E
  ESI      E EDI      0 EBP  431410 ESP  4313FC
  CS:IP 170:0042C1B2 ID 0E COD      0 FLG  10246
4 CS= 170, USE32, page granular, limit FFFFFFFF, base      0, acc CF9B
  SS= 178, USE32, page granular, limit FFFFFFFF, base      0, acc CF93
  DS= 178, USE32, page granular, limit FFFFFFFF, base      0, acc CF93
  ES= 178, USE32, page granular, limit FFFFFFFF, base      0, acc CF93
  FS=  0, USE16, byte granular, limit      0, base      15, acc  0
  GS= 20, USE16, byte granular, limit FFFF, base  6AA0, acc 93
5 CR0: PG:1 ET:1 TS:0 EM:0 MP:0 PE:1 CR2: 1F000000 CR3: 9067
6 Crash address (unrelocated) = 1:000001B2
7 Opcode stream: 8A 18 31 D2 88 DA EB 0E 50 68 39 00 43 00 E8 1D
  Stack:
8 0178:004313FC 000E 0000 0000 0000 C2D5 0042 C057 0042 0170 0000 0000 0000
  0178:00431414 0450 0043 0452 0043 0000 0000 1430 0043 CBEF 0042 011C 0000
  0178:0043142C C568 0042 0000 0000 0000 0000 0000 0000 F248 0042 F5F8 0042
  0178:00431444 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000
  0178:0043145C 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000
  0178:00431474 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000
9 Last 4 ints: 21 @ 170:42CF48/21 @ 170:42CF48/21 @ 170:42CF48/E @ 170:42C1B2/
```

1. The error message includes a synopsis of the problem. In this case, the processor signaled a page fault exception while executing at address 170:0042C1B2.
2. The prev_tsf32 field is not usually of interest.
3. These are the register values at the time of the exception. The interrupt number and error code (pushed on the stack by the processor for certain exceptions) are also printed.
4. The descriptors referenced by each segment register are described for your convenience. USE32 segments in general belong to your program; USE16 segments generally belong to the DOS extender. Here, CS points to your program's code segment, and SS, DS, and ES point to your data segment. FS is NULL and GS points to a DOS extender segment.

5. The control register information is not of any general interest, except on a page fault, when CR2 contains the address value that caused the fault. Since EAX in this case contains the same value, an attempt to dereference EAX could have caused this particular fault.
6. If the crash address (unrelocated) appears, it tells you where the crash occurred relative to your program's link map. You can therefore tell where a crash occurred even if you can't reproduce the crash in a debugger.
7. The opcode stream, if it appears, shows the next 16 bytes from the code segment at the point of the exception. If you disassemble these instructions, you can tell what instructions caused the crash, even without using a debugger. In this case, 8A 18 is the instruction `mov bl, [eax]`.
8. 72 words from the top of the stack, at the point of the exception, may be listed next. You may be able to recognize function calls or data from your program on the stack.
9. The four interrupts least to most recently handled by DOS/4GW in protected mode are listed next. In this example, the last interrupt issued before the page fault occurred was an INT 21h (DOS call) at address 170:42CF48. Sometimes, this information provides helpful context.

Here's an abridged register dump from a stack overflow.

```
DOS/4GW error (2002): transfer stack overflow
                        on interrupt 70h at 170:0042C002
TSF32: prev_tsf32 48C8
SS      C8 DS      170 ES      28 FS      0 GS      20
EAX AAAAAAAAA EBX BBBBBBBB ECX CCCCCCCC EDX DDDDDDDD
ESI 51515151 EDI D1D1D1D1 EBP B1B1B1B1 ESP 4884
1 CS:IP 170:0042C002 ID 70 COD 0 FLG 2
...
2 Previous TSF:
TSF32: prev_tsf32 498C
SS      C8 DS      170 ES      28 FS      0 GS      20
EAX AAAAAAAAA EBX BBBBBBBB ECX CCCCCCCC EDX DDDDDDDD
ESI 51515151 EDI D1D1D1D1 EBP B1B1B1B1 ESP 4960
3 CS:IP 170:0042C002 ID 70 COD 0 FLG 2
...
Previous TSF:
TSF32: prev_tsf32 67E4
SS      178 DS      170 ES      28 FS      0 GS      20
EAX AAAAAAAAA EBX BBBBBBBB ECX CCCCCCCC EDX DDDDDDDD
ESI 51515151 EDI D1D1D1D1 EBP B1B1B1B1 ESP 42FFE0
4 CS:IP 170:0042C039 ID 70 COD 0 FLG 202
5 Opcode stream: CF 66 B8 62 25 66 8C CB 66 8E DB BA 00 C0 42 00
Last 4 ints: 70 @ 170:42C002/70 @ 170:42C002/70 @ 170:42C002/70 @ 170:42C002/
```

1. We overflowed the transfer stack while trying to process an interrupt 70h at 170:0042C002.
2. The entire interrupt history from the transfer stack is printed next. The prev_tsf32 numbers increase as we progress from most recent to least recent interrupt. All of these interrupts are still pending, which is why we ran out of stack space.
3. Before we overflowed the stack, we got the same interrupt at the same address. For a recursive interrupt situation, this is typical.
4. The oldest frame on the transfer stack shows the recursion was touched off at a slightly different address. Looking at this address may help you understand the recursion.
5. The opcode stream and last four interrupt information comes from the newest transfer stack frame, not the oldest.

14.8 Compatibility

9a. Running DOS/4GW applications from inside Lotus 1-2-3.

In order to run DOS/4GW applications while "shelled out" from Lotus 1-2-3, you must use the PRIVATXM program included with your Open Watcom compiler. Otherwise, 1-2-3 will take all of the memory on your machine and prevent DOS/4GW from using it.

Before starting 1-2-3, you must set the DOS16M environment variable to limit Lotus' memory use (see your Open Watcom manual). After shelling out, you must run PRIVATXM, then clear the DOS16M environment variable before running your application.

9b. EMM386.EXE provided with DOS 6.0.

We know of at least three serious bugs in the EMM386.EXE distributed with MS-DOS 6.0, one involving mis-counting the amount of available memory, one involving mapping too little of the High Memory Area (HMA) into its page tables, and one involving allocation of EMS memory. Version 1.95 of DOS/4GW contains workarounds for some of these problems.

If you are having problems with DOS/4GW and you are using an EMM386.EXE dated 3-10-93 at 6:00:00, or later, you may wish to try the following workarounds, in sequence, until the problem goes away.

- Configure EMM386 with both the NOEMS and NOVCPi options.
- Convert the DEVICEHIGH statements in your CONFIG.SYS to DEVICE statements, and remove the LH (Load High) commands from your AUTOEXEC.BAT.
- Run in a Windows DOS box.
- Replace EMM386 with another memory manager, such as QEMM-386, 386Max, or an older version of EMM386.
- Run with HIMEM.SYS alone.

You may also wish to contact Microsoft Corporation to inquire about the availability of a fix.

9c. Spawning under OS/2 2.1.

We know of a bug in OS/2 2.1 that prevents one DOS/4GW application from spawning another over and over again. The actual number of repeated spawns that are possible under OS/2 varies from machine to machine, but is generally about 30.

This bug also affects programs running under other DOS extenders, and we have not yet found a workaround, other than linking your two programs together as a single program.

9d. "DPMI host error: cannot lock stack".

This error message almost always indicates insufficient memory, rather than a real incompatibility. If you see it under an OS/2 DOS box, you probably need to edit your DOS Session settings and make DPMI_MEMORY_LIMIT larger.

9e. Bug in Novell TCPIP driver.

Some versions of a program from Novell called TCPIP.EXE, a real-mode program, will cause the high words of EAX and EDX to be altered during a hardware interrupt. This bug breaks protected-mode software (and other real-mode software that uses the 80386 registers). Novell has released a newer version of TCPIP that fixes the problem; contact Novell to obtain the fix.

9f. Bugs in Windows NT.

The initial release of Windows NT includes a DPMI host, DOSX.EXE, with several serious bugs, some of which apparently cannot be worked around. We cannot warranty operation of DOS/4GW under Windows NT at this time, but we are continuing to exercise our best efforts to work around these problems.

You may wish to contact Microsoft Corporation to inquire about the availability of a new version of DOSX.EXE.

Windows 3.x Programming Guide

15 Creating 16-bit Windows 3.x Applications

This chapter describes how to compile and link 16-bit Windows 3.x applications simply and quickly. In this chapter, we look at applications written to exploit the Windows 3.x Application Programming Interface (API).

We will illustrate the steps to creating 16-bit Windows 3.x applications by taking a small sample application and showing you how to compile, link, run and debug it.

15.1 The Sample GUI Application

To demonstrate the creation of 16-bit Windows 3.x applications, we introduce a simple sample program. For our example, we are going to use the "sieve" program.

```
*$include winapi.fi
*$noreference
    INTEGER*2 FUNCTION FWINMAIN( hInstance, hPrevInstance,
    &                               lpszCmdLine, nCmdShow )
*$reference
    IMPLICIT NONE
    INTEGER*2 hInstance, hPrevInstance, nCmdShow
    INTEGER*4 lpszCmdLine
    include 'windows.fi'

* This program computes the prime numbers between 1 and 10,000
* using the Sieve of Eratosthenes algorithm.

    INTEGER UBOUND
    PARAMETER (UBOUND=10000)
    INTEGER I, K, PRIMES
    LOGICAL*1 NUMBERS(2:UBOUND)
    CHARACTER*11 FORM
    CHARACTER*60 BUFFER
    PARAMETER (FORM=' (A,I5,A,I5) ')
    DO I = 2, UBOUND
        NUMBERS(I) = .TRUE.
    ENDDO
    PRIMES = 0
    DO I = 2, UBOUND
        IF ( NUMBERS(I) ) THEN
            PRIMES = PRIMES + 1
            DO K = I + I, UBOUND, I
                NUMBERS(K) = .FALSE.
            ENDDO
        ENDIF
    ENDDO
    WRITE(BUFFER, FORM) 'The Number of Primes between 1 and ',
    &    UBOUND, ' are: ', PRIMES
    CALL MessageBox( 0, BUFFER,
    &    'Sieve of Eratosthenes',
    &    MB_OK .OR. MB_TASKMODAL )
    END
```

The goal of this program is to count the prime numbers between 1 and 10,000. It uses the famous *Sieve of Eratosthenes* algorithm to accomplish this task. We will take you through the steps necessary to produce this result.

15.2 Building and Running the GUI Application

To compile and link our example program which is stored in the file `sieve.for`, enter the following command:

```
C>wfl -l=windows -win sieve.for
```

The typical messages that appear on the screen are shown in the following illustration.

```
C>wfl -l=windows -win sieve.for
Open Watcom F77/16 Compile and Link Utility
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1990-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
      wfc sieve.for -win
Open Watcom FORTRAN 77/16 Optimizing Compiler
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1984-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
sieve.for: 4305 statements, 356 bytes, 1524 extensions, 0 warnings, 0 errors

Open Watcom Linker
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1985-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
loading object files
searching libraries
creating a Windows 16-bit executable
```

If you examine the current directory, you will find that two files have been created. These are `sieve.obj` (the result of compiling `sieve.for`) and `sieve.exe` (the result of linking `sieve.obj` with the appropriate Open Watcom FORTRAN 77 libraries).

The resultant 16-bit Windows 3.x application `SIEVE.EXE` can now be run under Windows 3.x.

15.3 Debugging the GUI Application

Let us assume that you wish to debug your application in order to locate an error in programming. In the previous section, the "sieve" program was compiled with default compile and link options. When debugging an application, it is useful to refer to the symbolic names of routines and variables. It is also convenient to debug at the source line level rather than the machine language level. To do this, we must direct both the compiler and linker to include additional debugging information in the object and executable files. Using the **WFL** command, this is fairly straightforward. **WFL** recognizes the Open Watcom F77 compiler "debug" options and will create the appropriate debug directives for the Open Watcom Linker.

For example, to compile and link the "sieve" program with debugging information, the following command may be issued.

```
C>wfl -l=windows -win -d2 sieve.for
```

The typical messages that appear on the screen are shown in the following illustration.

```
C>wfl -l=windows -win -d2 sieve.for
Open Watcom F77/16 Compile and Link Utility
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1990-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
    wfc sieve.for -win -d2
Open Watcom FORTRAN 77/16 Optimizing Compiler
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1984-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
sieve.for: 4305 statements, 467 bytes, 1524 extensions, 0 warnings, 0 errors

Open Watcom Linker
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1985-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
loading object files
searching libraries
creating a Windows 16-bit executable
```

The "d2" option requests the maximum amount of debugging information that can be provided by the Open Watcom F77 compiler. **WFL** will make sure that this debugging information is included in the executable file that is produced by the linker.

The "bytes" value is larger than in the previous example since selection of the "d2" option results in fewer code optimizations by default. You can request more optimization by specifying the appropriate options. However, you do so at the risk of making it more difficult for yourself to determine the relationship between the object code and the original source language code.

To request the Open Watcom Debugger to assist in debugging the application, select the Open Watcom Debugger icon. It would be too ambitious to describe the debugger in this introductory chapter so we refer you to the book entitled *Open Watcom Debugger User's Guide*.

There are more extensive examples of Windows applications written in FORTRAN 77 in the `\WATCOM\SAMPLES\FORTRAN\WIN` directory. The example programs are `ELLIPSE.FOR` and `FWCOPY.FOR`.

16 Porting Non-GUI Applications to 16-bit Windows 3.x

Generally, an application that is to run in a windowed environment must be written in such a way as to exploit the Windows Application Programming Interface (API). To take an existing character-based (i.e., non-graphical) application that ran under a system such as DOS and adapt it to run under Windows can require some considerable effort. There is a steep learning curve associated with the API function libraries.

This chapter describes how to create a Windows application quickly and simply from an application that does not use the Windows API. The application will make use of Open Watcom's default windowing support.

Suppose you have a set of FORTRAN 77 applications that previously ran under a system like DOS and you now wish to run them under Windows 3.x. To achieve this, you can simply recompile your application with the appropriate options and link with the appropriate libraries. We provide a default windowing system that turns your character-mode application into a simple Windows 3.x Graphical User Interface (GUI) application.

Normally, a Windows 3.x GUI application makes use of user-interface tools such as menus, icons, scroll bars, etc. However, an application that was not designed as a windowed application (such as a DOS application) can run as a GUI application. This is achieved by our default windowing system. The following sections describe the default windowing system.

16.1 Console Device in a Windowed Environment

In a FORTRAN 77 application that runs under DOS, *unit 5* and *unit 6* are connected to the standard input and standard output devices respectively. It is not a recommended practice to read directly from the standard input device or write to the standard output device when running in a windowed environment. For this reason, a default windowing environment is created for FORTRAN 77 applications that read from *unit 5* or write to *unit 6*. When your application is started, a window is created in which output to *unit 6* is displayed and input from *unit 5* is requested.

In addition to the standard I/O device, it is also possible to perform I/O to the console by explicitly opening a file whose name is "CON". When this occurs, another window is created and displayed. This window is different from the one created for standard input and standard output. In fact, every time you open the console device a different window is created. This provides a simple multi-windowing system for multiple streams of data to and from the console device.

16.2 The Sample Non-GUI Application

To demonstrate the creation of 16-bit Windows 3.x applications, we introduce a simple sample program. For our example, we are going to use the "sieve" program.

```
* This program computes the prime numbers between 1 and 10,000
* using the Sieve of Eratosthenes algorithm.

      IMPLICIT NONE
      INTEGER UPBOUND
      PARAMETER (UPBOUND=10000)
      INTEGER I, K, PRIMES
      LOGICAL*1 NUMBERS(2:UPBOUND)
      CHARACTER*11 FORM
      PARAMETER (FORM='(A,I5,A,I5)')
      DO I = 2, UPBOUND
         NUMBERS(I) = .TRUE.
      ENDDO
      PRIMES = 0
      DO I = 2, UPBOUND
         IF( NUMBERS(I) ) THEN
            PRIMES = PRIMES + 1
            DO K = I + I, UPBOUND, I
               NUMBERS(K) = .FALSE.
            ENDDO
         ENDIF
      ENDDO
      PRINT FORM, 'The Number of Primes between 1 and ', UPBOUND,
1      ' are: ', PRIMES
      END
```

The goal of this program is to count the prime numbers between 1 and 10,000. It uses the famous *Sieve of Eratosthenes* algorithm to accomplish this task. We will take you through the steps necessary to produce this result.

16.3 Building and Running the Non-GUI Application

Very little effort is required to port an existing FORTRAN 77 application to Windows 3.x.

You must compile and link the file `sieve.for` specifying the "bw" option.

```
C>wfl -l=windows -bw -win sieve.for
```

The typical messages that appear on the screen are shown in the following illustration.


```
C>wfl -l=windows -bw -win sieve.for
Open Watcom F77/16 Compile and Link Utility
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1990-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
    wfc sieve.for -bw -win
Open Watcom FORTRAN 77/16 Optimizing Compiler
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1984-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
sieve.for: 21 statements, 311 bytes, 6 extensions, 0 warnings, 0 errors

Open Watcom Linker
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1985-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
loading object files
searching libraries
creating a Windows 16-bit executable
```

If you examine the current directory, you will find that two files have been created. These are `sieve.obj` (the result of compiling `sieve.for`) and `sieve.exe` (the result of linking `sieve.obj` with the appropriate Open Watcom FORTRAN 77 libraries).

The resultant 16-bit Windows 3.x application `SIEVE.EXE` can now be run under Windows 3.x as a Windows GUI application.

16.4 Debugging the Non-GUI Application

Let us assume that you wish to debug your application in order to locate an error in programming. In the previous section, the "sieve" program was compiled with default compile and link options. When debugging an application, it is useful to refer to the symbolic names of routines and variables. It is also convenient to debug at the source line level rather than the machine language level. To do this, we must direct both the compiler and linker to include additional debugging information in the object and executable files. Using the **WFL** command, this is fairly straightforward. **WFL** recognizes the Open Watcom F77 compiler "debug" options and will create the appropriate debug directives for the Open Watcom Linker.

For example, to compile and link the "sieve" program with debugging information, the following command may be issued.

```
C>wfl -l=windows -bw -win -d2 sieve.for
```

The typical messages that appear on the screen are shown in the following illustration.

```
C>wfl -l=windows -bw -win -d2 sieve.for
Open Watcom F77/16 Compile and Link Utility
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1990-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
    wfc sieve.for -bw -win -d2
Open Watcom FORTRAN 77/16 Optimizing Compiler
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1984-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
sieve.for: 21 statements, 392 bytes, 6 extensions, 0 warnings, 0 errors

Open Watcom Linker
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1985-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
loading object files
searching libraries
creating a Windows 16-bit executable
```

The "d2" option requests the maximum amount of debugging information that can be provided by the Open Watcom F77 compiler. **WFL** will make sure that this debugging information is included in the executable file that is produced by the linker.

The "bytes" value is larger than in the previous example since selection of the "d2" option results in fewer code optimizations by default. You can request more optimization by specifying the appropriate options. However, you do so at the risk of making it more difficult for yourself to determine the relationship between the object code and the original source language code.

To request the Open Watcom Debugger to assist in debugging the application, select the Open Watcom Debugger icon. It would be too ambitious to describe the debugger in this introductory chapter so we refer you to the book entitled *Open Watcom Debugger User's Guide*.

16.5 Default Windowing Library Functions

A few library functions have been written to enable some simple customization of the default windowing system's behaviour. The following functions are supplied:

dwfDeleteOnClose

```
integer function dwfDeleteOnClose( unit )
integer unit
```

This function tells the console window that it should close itself when the file is closed. You must pass to it the unit number associated with the opened console.

dwfSetAboutDlg

```
integer function dwfSetAboutDlg( title, text )
character*(*) title
character*(*) text
```

This function sets the about dialog box of the default windowing system. The "title" points to the string that will replace the current title. If title is CHAR(0) then the title will not be replaced. The "text" points to a string which will be placed in the about box. To get multiple lines, embed a new

line after each logical line in the string. If "text" is CHAR(0), then the current text in the about box will not be replaced.

dwfSetTitle

```
integer function dwfSetTitle( title )  
character*(*) title
```

This function sets the main window's title.

dwfSetConTitle

```
integer function dwfSetConTitle( unit, title )  
integer unit  
character*(*) title
```

This function sets the console window's title which corresponds to the unit number passed to it.

dwfShutDown

```
integer function dwfShutDown()
```

This function shuts down the default windowing I/O system. The application will continue to execute but no windows will be available for output.

dwfYield

```
integer function dwfYield()
```

This function yields control back to the operating system, thereby giving other processes a chance to run.

These functions are described more fully in the *Open Watcom FORTRAN 77 User's Guide*.

17 Creating 32-bit Windows 3.x Applications

This chapter describes how to compile and link 32-bit Windows 3.x applications simply and quickly. In this chapter, we look at applications written to exploit the Windows 3.x Application Programming Interface (API).

We will illustrate the steps to creating 32-bit Windows 3.x applications by taking a small sample application and showing you how to compile, link, run and debug it.

17.1 The Sample GUI Application

To demonstrate the creation of 32-bit Windows 3.x applications, we introduce a simple sample program. For our example, we are going to use the "sieve" program.

```
*$include winapi.fi
*$noreference
    INTEGER*2 FUNCTION FWINMAIN( hInstance, hPrevInstance,
    &                               lpszCmdLine, nCmdShow )
*$reference
    IMPLICIT NONE
    INTEGER*2 hInstance, hPrevInstance, nCmdShow
    INTEGER*4 lpszCmdLine
    include 'windows.fi'

* This program computes the prime numbers between 1 and 10,000
* using the Sieve of Eratosthenes algorithm.

    INTEGER UPBOUND
    PARAMETER (UPBOUND=10000)
    INTEGER I, K, PRIMES
    LOGICAL*1 NUMBERS(2:UPBOUND)
    CHARACTER*11 FORM
    CHARACTER*60 BUFFER
    PARAMETER (FORM=' (A,I5,A,I5) ')
    DO I = 2, UPBOUND
        NUMBERS(I) = .TRUE.
    ENDDO
    PRIMES = 0
    DO I = 2, UPBOUND
        IF ( NUMBERS(I) ) THEN
            PRIMES = PRIMES + 1
            DO K = I + I, UPBOUND, I
                NUMBERS(K) = .FALSE.
            ENDDO
        ENDIF
    ENDDO
    WRITE(BUFFER, FORM) 'The Number of Primes between 1 and ',
    & UPBOUND, ' are: ', PRIMES
    CALL MessageBox( 0, BUFFER,
    & 'Sieve of Eratosthenes',
    & MB_OK .OR. MB_TASKMODAL )
END
```

The goal of this program is to count the prime numbers between 1 and 10,000. It uses the famous *Sieve of Eratosthenes* algorithm to accomplish this task. We will take you through the steps necessary to produce this result.

17.2 Building and Running the GUI Application

To compile and link our example program which is stored in the file `sieve.for`, enter the following command:

```
C>wfl386 -l=win386 sieve.for
```

The typical messages that appear on the screen are shown in the following illustration.

```
C>wfl386 -l=win386 sieve.for
Open Watcom F77/32 Compile and Link Utility
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1990-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
wfc386 sieve.for
Open Watcom FORTRAN 77/32 Optimizing Compiler
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1984-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
sieve.for: 4390 statements, 207 bytes, 1585 extensions, 0 warnings, 0 errors

Open Watcom Linker
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1985-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
loading object files
searching libraries
creating a Windows 32-bit executable
```

If you examine the current directory, you will find that two files have been created. These are `sieve.obj` (the result of compiling `sieve.for`) and `sieve.rex` (the result of linking `sieve.obj` with the appropriate Open Watcom FORTRAN 77 libraries). The ".rex" file must now be combined with Open Watcom's 32-bit Windows supervisor `WIN386.EXE` using the Open Watcom Bind utility. `WBIND.EXE` combines your 32-bit application code and data (".rex" file) with the 32-bit Windows supervisor. The process involves the following steps:

1. `WBIND` copies `WIN386.EXE` into the current directory.
2. `WBIND.EXE` optionally runs the resource compiler on the 32-bit Windows supervisor so that the 32-bit executable can have access to the applications resources.
3. `WBIND.EXE` concatenates `WIN386.EXE` and the ".rex" file, and creates a ".exe" file with the same name as the ".rex" file.

The following describes the syntax of the `WBIND` command.

WBIND file_spec [-d] [-n] [-q] [-s supervisor] [-R rc_options]

The square brackets [] denote items which are optional.

WBIND	is the name of the Open Watcom Bind utility.
file_spec	is the name of the 32-bit Windows EXE to bind.
-d	requests that a 32-bit DLL be built.
-n	indicates that the resource compiler is NOT to be invoked.
-q	requests that WBIND run in quiet mode (no informational messages are displayed).
-s supervisor	specifies the path and name of the Windows supervisor to be bound with the application. If not specified, a search of the paths listed in the PATH environment variable is performed. If this search is not successful and the WATCOM environment variable is defined, the %WATCOM%\BINW directory is searched.
-R rc_options	all options after -R are passed to the resource compiler.

To bind our example program, the following command may be issued:

```
C>wbind sieve -n
```

If the "s" option is specified, it must identify the location of the WIN386.EXE file or the W386DLL.EXE file (if you are building a DLL).

Example:

```
C>wbind sieve -n -s c:\watcom\binw\win386.exe
```

If the "s" option is not specified, then the **WATCOM** environment variable must be defined or the "BINW" directory must be listed in your **PATH** environment variable.

Example:

```
C>set watcom=c:\watcom
or
C>path c:\watcom\binw;c:\dos;c:\windows
```

The resultant 32-bit Windows 3.x application SIEVE.EXE can now be run under Windows 3.x.

17.3 Debugging the GUI Application

Let us assume that you wish to debug your application in order to locate an error in programming. In the previous section, the "sieve" program was compiled with default compile and link options. When debugging an application, it is useful to refer to the symbolic names of routines and variables. It is also convenient to debug at the source line level rather than the machine language level. To do this, we must direct both the compiler and linker to include additional debugging information in the object and executable files. Using the **WFL386** command, this is fairly straightforward. **WFL386** recognizes the Open Watcom F77 compiler "debug" options and will create the appropriate debug directives for the Open Watcom Linker.

For example, to compile and link the "sieve" program with debugging information, the following command may be issued.

```
C>wfl386 -l=win386 -d2 sieve.for
```

The typical messages that appear on the screen are shown in the following illustration.

```
C>wfl386 -l=win386 -d2 sieve.for
Open Watcom F77/32 Compile and Link Utility
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1990-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
      wfc386 sieve.for -d2
Open Watcom FORTRAN 77/32 Optimizing Compiler
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1984-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
sieve.for: 4390 statements, 293 bytes, 1585 extensions, 0 warnings, 0 errors

Open Watcom Linker
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1985-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
loading object files
searching libraries
creating a Windows 32-bit executable
```

The "d2" option requests the maximum amount of debugging information that can be provided by the Open Watcom F77 compiler. **WFL386** will make sure that this debugging information is included in the executable file that is produced by the linker.

The "bytes" value is larger than in the previous example since selection of the "d2" option results in fewer code optimizations by default. You can request more optimization by specifying the appropriate options. However, you do so at the risk of making it more difficult for yourself to determine the relationship between the object code and the original source language code.

Once again, the ".rex" file must be combined with Open Watcom's 32-bit Windows supervisor WIN386.EXT using the Open Watcom Bind utility. This step is described in the previous section.

To request the Open Watcom Debugger to assist in debugging the application, select the Open Watcom Debugger icon. It would be too ambitious to describe the debugger in this introductory chapter so we refer you to the book entitled *Open Watcom Debugger User's Guide*.

There are more extensive examples of Windows applications written in FORTRAN 77 in the \WATCOM\SAMPLES\FORTRAN\WIN directory. The example programs are ELLIPSE.FOR and FWCOPY.FOR.

18 Porting Non-GUI Applications to 32-bit Windows 3.x

Generally, an application that is to run in a windowed environment must be written in such a way as to exploit the Windows Application Programming Interface (API). To take an existing character-based (i.e., non-graphical) application that ran under a system such as DOS and adapt it to run under Windows can require some considerable effort. There is a steep learning curve associated with the API function libraries.

This chapter describes how to create a Windows application quickly and simply from an application that does not use the Windows API. The application will make use of Open Watcom's default windowing support.

Suppose you have a set of FORTRAN 77 applications that previously ran under a system like DOS and you now wish to run them under Windows 3.x. To achieve this, you can simply recompile your application with the appropriate options and link with the appropriate libraries. We provide a default windowing system that turns your character-mode application into a simple Windows 3.x Graphical User Interface (GUI) application.

Normally, a Windows 3.x GUI application makes use of user-interface tools such as menus, icons, scroll bars, etc. However, an application that was not designed as a windowed application (such as a DOS application) can run as a GUI application. This is achieved by our default windowing system. The following sections describe the default windowing system.

18.1 Console Device in a Windowed Environment

In a FORTRAN 77 application that runs under DOS, *unit 5* and *unit 6* are connected to the standard input and standard output devices respectively. It is not a recommended practice to read directly from the standard input device or write to the standard output device when running in a windowed environment. For this reason, a default windowing environment is created for FORTRAN 77 applications that read from *unit 5* or write to *unit 6*. When your application is started, a window is created in which output to *unit 6* is displayed and input from *unit 5* is requested.

In addition to the standard I/O device, it is also possible to perform I/O to the console by explicitly opening a file whose name is "CON". When this occurs, another window is created and displayed. This window is different from the one created for standard input and standard output. In fact, every time you open the console device a different window is created. This provides a simple multi-windowing system for multiple streams of data to and from the console device.

18.2 The Sample Non-GUI Application

To demonstrate the creation of 32-bit Windows 3.x applications, we introduce a simple sample program. For our example, we are going to use the "sieve" program.

* This program computes the prime numbers between 1 and 10,000
* using the Sieve of Eratosthenes algorithm.

```
IMPLICIT NONE
INTEGER UPBOUND
PARAMETER (UPBOUND=10000)
INTEGER I, K, PRIMES
LOGICAL*1 NUMBERS(2:UPBOUND)
CHARACTER*11 FORM
PARAMETER (FORM='(A,I5,A,I5)')
DO I = 2, UPBOUND
    NUMBERS(I) = .TRUE.
ENDDO
PRIMES = 0
DO I = 2, UPBOUND
    IF( NUMBERS(I) ) THEN
        PRIMES = PRIMES + 1
        DO K = I + I, UPBOUND, I
            NUMBERS(K) = .FALSE.
        ENDDO
    ENDIF
ENDDO
PRINT FORM, 'The Number of Primes between 1 and ', UPBOUND,
1      ' are: ', PRIMES
END
```

The goal of this program is to count the prime numbers between 1 and 10,000. It uses the famous *Sieve of Eratosthenes* algorithm to accomplish this task. We will take you through the steps necessary to produce this result.

18.3 Building and Running the Non-GUI Application

Very little effort is required to port an existing FORTRAN 77 application to Windows 3.x.

You must compile and link the file `sieve.for` specifying the "bw" option.

```
C>wfl386 -l=win386 -bw sieve.for
```

The typical messages that appear on the screen are shown in the following illustration.

```
C>wfl386 -l=win386 sieve.for
Open Watcom F77/32 Compile and Link Utility
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1990-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
    wfc386 sieve.for -bw
Open Watcom FORTRAN 77/32 Optimizing Compiler
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1984-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
sieve.for: 21 statements, 172 bytes, 6 extensions, 0 warnings, 0 errors

Open Watcom Linker
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1985-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
loading object files
searching libraries
creating a Windows 32-bit executable
```

If you examine the current directory, you will find that two files have been created. These are `sieve.obj` (the result of compiling `sieve.for`) and `sieve.rex` (the result of linking `sieve.obj` with the appropriate Open Watcom FORTRAN 77 libraries). The ".rex" file must now be combined with Open Watcom's 32-bit Windows supervisor `WIN386.EXT` using the Open Watcom Bind utility. `WBIND.EXE` combines your 32-bit application code and data (".rex" file) with the 32-bit Windows supervisor. The process involves the following steps:

1. `WBIND` copies `WIN386.EXT` into the current directory.
2. `WBIND.EXE` optionally runs the resource compiler on the 32-bit Windows supervisor so that the 32-bit executable can have access to the applications resources.
3. `WBIND.EXE` concatenates `WIN386.EXT` and the ".rex" file, and creates a ".exe" file with the same name as the ".rex" file.

The following describes the syntax of the `WBIND` command.

WBIND file_spec [-d] [-n] [-q] [-s supervisor] [-R rc_options]

The square brackets [] denote items which are optional.

<i>WBIND</i>	is the name of the Open Watcom Bind utility.
<i>file_spec</i>	is the name of the 32-bit Windows EXE to bind.
<i>-d</i>	requests that a 32-bit DLL be built.
<i>-n</i>	indicates that the resource compiler is NOT to be invoked.
<i>-q</i>	requests that <code>WBIND</code> run in quiet mode (no informational messages are displayed).
<i>-s supervisor</i>	specifies the path and name of the Windows supervisor to be bound with the application. If not specified, a search of the paths listed in the PATH

environment variable is performed. If this search is not successful and the **WATCOM** environment variable is defined, the %WATCOM%\BINW directory is searched.

-R rc_options all options after -R are passed to the resource compiler.

To bind our example program, the following command may be issued:

```
C>wbind sieve -n
```

If the "s" option is specified, it must identify the location of the WIN386.EXE file or the W386DLL.EXE file (if you are building a DLL).

Example:

```
C>wbind sieve -n -s c:\watcom\binw\win386.ext
```

If the "s" option is not specified, then the **WATCOM** environment variable must be defined or the "BINW" directory must be listed in your **PATH** environment variable.

Example:

```
C>set watcom=c:\watcom
or
C>path c:\watcom\binw;c:\dos;c:\windows
```

The resultant 32-bit Windows 3.x application SIEVE.EXE can now be run under Windows 3.x as a Windows GUI application.

18.4 Debugging the Non-GUI Application

Let us assume that you wish to debug your application in order to locate an error in programming. In the previous section, the "sieve" program was compiled with default compile and link options. When debugging an application, it is useful to refer to the symbolic names of routines and variables. It is also convenient to debug at the source line level rather than the machine language level. To do this, we must direct both the compiler and linker to include additional debugging information in the object and executable files. Using the **WFL386** command, this is fairly straightforward. **WFL386** recognizes the Open Watcom F77 compiler "debug" options and will create the appropriate debug directives for the Open Watcom Linker.

For example, to compile and link the "sieve" program with debugging information, the following command may be issued.

```
C>wfl386 -l=win386 -bw -d2 sieve.for
```

The typical messages that appear on the screen are shown in the following illustration.

```
C>wfl386 -l=win386 -d2 sieve.for
Open Watcom F77/32 Compile and Link Utility
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1990-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
    wfc386 sieve.for -bw -d2
Open Watcom FORTRAN 77/32 Optimizing Compiler
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1984-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
sieve.for: 21 statements, 237 bytes, 6 extensions, 0 warnings, 0 errors

Open Watcom Linker
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1985-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
loading object files
searching libraries
creating a Windows 32-bit executable
```

The "d2" option requests the maximum amount of debugging information that can be provided by the Open Watcom F77 compiler. **WFL386** will make sure that this debugging information is included in the executable file that is produced by the linker.

The "bytes" value is larger than in the previous example since selection of the "d2" option results in fewer code optimizations by default. You can request more optimization by specifying the appropriate options. However, you do so at the risk of making it more difficult for yourself to determine the relationship between the object code and the original source language code.

Once again, the ".rex" file must be combined with Open Watcom's 32-bit Windows supervisor WIN386.EXT using the Open Watcom Bind utility. This step is described in the previous section.

To request the Open Watcom Debugger to assist in debugging the application, select the Open Watcom Debugger icon. It would be too ambitious to describe the debugger in this introductory chapter so we refer you to the book entitled *Open Watcom Debugger User's Guide*.

18.5 Default Windowing Library Functions

A few library functions have been written to enable some simple customization of the default windowing system's behaviour. The following functions are supplied:

dwfDeleteOnClose

```
integer function dwfDeleteOnClose( unit )
integer unit
```

This function tells the console window that it should close itself when the file is closed. You must pass to it the unit number associated with the opened console.

dwfSetAboutDlg

```
integer function dwfSetAboutDlg( title, text )
character*(*) title
character*(*) text
```

This function sets the about dialog box of the default windowing system. The "title" points to the string that will replace the current title. If title is CHAR(0) then the title will not be replaced. The "text" points to a string which will be placed in the about box. To get multiple lines, embed a new line after each logical line in the string. If "text" is CHAR(0), then the current text in the about box will not be replaced.

dwfSetAppTitle

```
integer function dwfSetAppTitle( title )  
character*(*) title
```

This function sets the main window's title.

dwfSetConTitle

```
integer function dwfSetConTitle( unit, title )  
integer unit  
character*(*) title
```

This function sets the console window's title which corresponds to the unit number passed to it.

dwfShutDown

```
integer function dwfShutDown()
```

This function shuts down the default windowing I/O system. The application will continue to execute but no windows will be available for output.

dwfYield

```
integer function dwfYield()
```

This function yields control back to the operating system, thereby giving other processes a chance to run.

These functions are described more fully in the *Open Watcom FORTRAN 77 User's Guide*.

19 The Open Watcom 32-bit Windows 3.x Extender

Open Watcom FORTRAN 77 contains the necessary tools and libraries to create 32-bit applications for Windows 3.x. Using Open Watcom FORTRAN 77 gives the programmer the benefits of a 32-bit flat memory model and access to the full Windows API (along with the usual FORTRAN 77 and C library functions).

The general model of the environment is as follows: The 32-bit flat memory model program is linked against a special 32-bit Windows library. This library contains the necessary information to invoke special 16-bit functions, which lie in the supervisor (`WIN386.EXT`). The 32-bit program is then bound (using `WBIND.EXE`) with the supervisor to create a Windows executable. At the same time as the 32-bit program is being bound, the resource compiler is run on the supervisor, and all the resources for the application are placed there. When the application is started, the supervisor obtains the 32-bit memory, relocates the 32-bit application into the memory, and invokes the 32-bit application.

All Windows functions are invoked from the supervisor, and all callback routines lie within the supervisor. The local heap resides within the supervisor as well.

If you are starting from a 16-bit Windows application, most of the code will not change when you port it to the 32-bit Windows environment. However, because of the nature of the Windows API and its implicit dependencies on a 16-bit environment, some source changes are necessary. These source changes are minimal, and are backwards compatible with the 16-bit environment.

19.1 Pointers

Throughout this document, there will be references to both *near* and *far*, and *16-bit* and *32-bit* pointers. Since this can rapidly become confusing, some initial explanations will be given here.

A *far pointer* is a pointer that is composed of both a selector and an offset. A selector determines a specific region of memory, and the offset is relative to the start of this region. A *near pointer* is a pointer that has an offset only, the selector is automatically assumed by the CPU.

The problem with far pointers is the selector overhead. Using a far pointer is much more expensive than using a near pointer. This is the advantage of the 32-bit flat memory model - all pointers within the program are near, and yet you can address up to 4 gigabytes of memory.

A *16-bit near pointer* occupies 2 bytes of memory (i.e., the offset is 16 bits long). This pointer can reference up to 64K of data.

A *16-bit far pointer* occupies 4 bytes of memory. There is a 16-bit selector and a 16-bit offset. This pointer can reference up to 64K of data.

A *32-bit near pointer* occupies 4 bytes of memory (i.e., the offset is 32 bits long). This pointer can reference up to 4 gigabytes of data.

A *32-bit far pointer* occupies 6 bytes of memory. There is a 16-bit selector and a 32-bit offset. This pointer can reference up to 4 gigabytes of data.

Windows, in general, uses 16-bit far pointers to pass information around. These 16-bit far pointers can also be used by a 32-bit Windows application. The conversion from a 16-bit pointer to a 32-bit pointer will occur automatically when you map a dynamically allocatable array to the memory pointed to by the 16-bit pointer using the *LOCATION=* specifier of the *ALLOCATE* statement. You must also declare the allocatable array as *far* using the *array* pragma. The syntax for the *array* pragma is:

```
$*pragma array ARRAY_NAME far
```

where *ARRAY_NAME* is the array name.

19.2 Implementation Overview

This section provides an overview of the issues that require consideration when creating a 32-bit Windows application for a 16-bit Windows environment.

First, all modules have to be recompiled for the 32-bit flat memory model with a compiler capable of generating 32-bit instructions.

Pointers to data passed to Windows are all far pointers. We will be passing pointers to data in our 32-bit flat address space, and these pointers are near pointers. As well, notice that these 32-bit near pointers are the same size as as their 16-bit far pointer counterparts (4 bytes). This is good, since all data structures containing pointers will remain the same size.

Windows cannot be called from 32-bit code on a 32-bit stack. This means that in order to call the API functions, it is necessary to write a set of cover functions that will accept the parameters, switch into a 16-bit environment, and then call Windows. There is another issue, though. Windows only understands 16-bit pointers, so before calling Windows, all pointers being passed to Windows must be converted to 16-bit far pointers.

It turns out that Windows can also call back to your application. Windows can only call 16-bit code, though, so there is a need for a bridge from the 16-bit side to the 32-bit side. It is necessary to allocate 16-bit call back routines that can be passed to Windows. These call back routines will then switch into the 32-bit environment and call whatever 32-bit function is required. The 32-bit call back has to be declared as a far function, since it is necessary to issue a far call to enter it from the 16-bit side. If it is a far function, then the compiler will generate the appropriate code for it to return from the far call.

Once Windows calls you back, it can hand you 16-bit far pointers in a long (4 byte) parameter. This pointer can only be used in the 32-bit environment if it is a 32-bit far pointer, not a 16-bit far pointer. The conversion is simple: the 16-bit offset is extended to a 32-bit offset (the high word is zeroed out). Any far pointer that Windows hands to you must be converted in this way. This conversion is performed automatically when a dynamically allocatable array is mapped to a 16-bit far pointer using the *LOCATION* specifier of the Open Watcom FORTRAN 77 *ALLOCATE* statement and the *array* pragma. The syntax for the *array* pragma is:

```
$*pragma array ARRAY_NAME far
```

where *ARRAY_NAME* is the array name.

Example:

```

subroutine DLLSUB( arg_list )

structure /argtypes/
  integer    w1
  integer    w2
  integer    w3
  integer    sum
end structure
record /argtypes/ args(:)
*$pragma array args far
integer*4 arg_list

allocate( args(1), location=arg_list )
    
```

In the preceding example, `arg_list` is a 16-bit far pointer to a structure with the elements described by the `argtypes` structure. The allocatable array `args` is described as *far* using the *array* pragma.

Sometimes, a Windows application wants to call a procedure in a DLL. The procedure address is a 16-bit far pointer. It is not possible to issue an indirect call to this address from the 32-bit environment, so some sort of interface is needed. This interface would switch into the 16-bit environment, and then call the 16-bit function.

These issues, along with other minor items, are handled by Open Watcom FORTRAN 77, and are discussed in more technical detail in later sections.

19.3 System Structure




Figure 5. WIN386 Structure




Figure 6. 32-bit Application Structure

19.4 System Overview

- `WIN386.EXT` is the key component of a 32-bit Windows application. It is a 16-bit Windows application which contains:

- All application resources.
- A 16-bit local heap.
- A 16-bit stack.

- `W386DLL.EXT` is similar to `WIN386.EXT`, only it provides a DLL interface.

`WIN386.EXT` is bound to your 32-bit application to create a 32-bit application that will run under Windows 3.x. `WIN386.EXT` provides the following functionality:

- supervisor to bring the 32-bit application into memory and start it running.
- "glue" functions to connect to Windows for both API and DOS functionality. This interface is designed to transparently set up the calling functions' pointers and parameters to their 16-bit counterparts.
- "glue-back" functions to allow Windows to call back 32-bit routines.
- special code to allow debugging of 32-bit applications.
- A number of files with file extension `.fi` are located in the `\WATCOM\SRC\FORTRAN\WIN` directory. The file `WINAPI.FI` describes the calling convention of each Windows API function. Other files define Windows constants and data structures.
- `WIN386.LIB` contains all the necessary library functions to connect to the 32-bit supervisor `WIN386.EXT`. All Windows API calls and Open Watcom FORTRAN 77 library DOS calls are found here.
- The standard FORTRAN 77 and C library functions, specially modified to run in the 32-bit environment, are located in the `\WATCOM\LIB386\WIN` directory.
- `WBIND.EXE` merges your 32-bit executable and the appropriate Supervisor into a single executable.

19.5 Steps to Obtaining a 32-bit Application

The following is an overview of the procedure for creating a 32-bit Windows Application:

1. If you are starting with a 16-bit Windows application, you must adapt your source code to the 32-bit environment.
2. You must compile the application using a 32-bit compiler.
3. You must link the application with the 32-bit libraries.
4. You must bind the 32-bit application with the 32-bit supervisor.
5. You can then run and/or debug the application.

20 Windows 3.x 32-bit Programming Overview

This chapter includes the following topics:

- WINAPI.FI and WINDOWS.FI
- Environment Notes
- Floating-point Emulation
- Multiple Instances
- Pointer Handling
- When To Convert Incoming Pointers
- When To Convert Outgoing Pointers
- SendMessage and SendDlgItemMessage
- GlobalAlloc and LocalAlloc
- Callback Function Pointers
- Window Sub-classing
- Calling 16-bit DLLs
- _16 Functions

20.1 WINAPI.FI

When developing programs, make sure `WINAPI.FI` is included at the start of all source files and the necessary include files (particularly `WINDOWS.FI`) are included in each function or subroutine.

It is especially important to get the correct function and argument typing information for Windows API functions. Due to the default typing rules of FORTRAN, many Windows API functions have a default result type of `REAL` when they may in fact return an `INTEGER` or `INTEGER*2` result. By including the appropriate include files, you ensure that this never happens. For example, the function `CreateDialog` is described in `WINDLG.FI` as a function returning an `INTEGER*2` result.

Example:

```
external CreateDialog
integer*2 CreateDialog
```

Failure to specify the correct type of a function will result in code that looks correct but does not execute correctly. Similarly, you should make sure that all symbolic constants are properly defined by including the appropriate include files. For example, the constant `DEFAULT_QUALITY` is described in `WINFONT.FI` as an `INTEGER` constant whose value is 0.

Example:

```
integer DEFAULT_QUALITY
parameter ( DEFAULT_QUALITY = 0 )
```

Without this information, `DEFAULT_QUALITY` would be assumed to be a `REAL` variable and would not have any assigned value.

The "EXPLICIT" compiler option is useful in this regard. It requires that all symbols be explicitly typed.

20.2 Environment Notes

- The 32-bit Windows Supervisor uses the first 256 bytes of the 32-bit application's stack to save state information. If this is corrupted, your application will abnormally terminate.
- The 32-bit Windows Supervisor provides resources for up to 512 callback routines. Note that this restriction is only on the maximum number of active callbacks.

20.3 Floating-point Emulation

The file `WEMU387.386` is included to support floating-point emulation for 32-bit applications running under Windows. This file is installed in the `[386Enh]` section of your `SYSTEM.INI` file. By using the floating-point emulator, your application can be compiled with the "fpi87" option to use inline floating-point instructions, and it will run on a machine without a numeric coprocessor.

Only one of `WEMU387.386` and `WDEBUG.386` may be installed in your `[386Enh]` section. `WEMU387.386` may be distributed with your application.

20.4 Multiple Instances

Since the 32-bit application resides in a flat memory space, it is NOT possible to share code with other instances. This means that you must register new window classes with callbacks into the new instance's code space. A simple way of accomplishing this is as follows:

```
integer*2 function FWINMAIN( hInstance,
&                               hPrevInstance,
&                               lpszCmdLine,
&                               nCmdShow )

integer*2 hInstance
integer*2 hPrevInstance
integer*2 nCmdShow
integer*4 lpszCmdLine
```

```

include 'win386.fi'
include 'wincreat.fi'
include 'wincurs.fi'
include 'windefn.fi'
include 'windisp.fi'
include 'winmsg.fi'
include 'winmsgs.fi'
include 'windtool.fi'
include 'winutil.fi'

external WndProc

integer*2                hWnd
record /MSG/              msg
record /WNDCLASS/         wndclass
character*14              class

wndclass.style = CS_HREDRAW .or. CS_VREDRAW
wndclass.lpfnWndProc = loc( WndProc )
wndclass.cbClsExtra = 0
wndclass.cbWndExtra = 0
wndclass.hInstance = hInstance
wndclass.hIcon = NULL_HANDLE
wndclass.hCursor = LoadCursor( NULL_HANDLE, IDC_ARROW )
wndclass.hbrBackground = GetStockObject( WHITE_BRUSH )
wndclass.lpszMenuName = NULL
write( class, '( 'Ellipses' ',i5.5,a)' ) hInstance, char(0)
wndclass.lpszClassName = Loc( class )
call RegisterClass( wndclass )

hWnd = CreateWindow( class,
&                'Application'c,
&                WS_OVERLAPPEDWINDOW,
&                CW_USEDEFAULT,
&                0,
&                CW_USEDEFAULT,
&                0,
&                NULL_HANDLE,
&                NULL_HANDLE,
&                hInstance,
&                NULL )

```

The variable *class* contains a unique name based on the instance of the application.

20.5 Pointer Handling

Windows 3.x is a 16-bit operating system. Function pointers that Windows deals with are 16-bit far pointers, and any data you communicate to Windows with are 16-bit far pointers. 16-bit far pointers occupy 4 bytes of data, and are capable of addressing up to 64K. For data objects larger than 64K, huge pointers are used (a sequence of far pointers that map out consecutive 64K segments for the data object). 16-bit far pointers are expensive to use due to the overhead of selector loads (each time you use the pointer, a segment register must have a value put in it). 16-bit huge pointers are even more expensive: not only is there the overhead of selector loads, but a run-time call is necessary to perform any pointer arithmetic.

In a 32-bit flat memory model, such as that of the Open Watcom F77 for Windows environment, all pointers are 32-bit near pointers (occupying 4 bytes of data as well). However, these pointers may access objects of up to 4 gigabytes in size, and there is no selector load overhead.

For a 32-bit environment to communicate with Windows 3.x, there are some considerations. All pointers sent to Windows must be converted from 32-bit near pointers to 16-bit far pointers. These conversions are handled by the Supervisor.

It is important to remember that all API functions which accept pointers (with the exception of functions that accept function pointers) accept 32-bit near pointers in this 32-bit model. If you attempt to pass a 32-bit far pointer, the conversion will not take place correctly.

16-bit far pointers to data may be passed into the API functions, and the Supervisor will not do any conversion.

Incoming pointers must be converted from 16-bit far pointers to 32-bit far pointers. This conversion is a trivial one: the offset portion of the 16-bit far pointer is extended to 32-bits. The pointer conversion will occur automatically when you map a dynamically allocatable array to the memory pointed to by the 16-bit pointer using the *LOCATION*= specifier of the *ALLOCATE* statement. You must also declare the array as *far* using the *array* pragma. The syntax for the *array* pragma is:

```
$*pragma array ARRAY_NAME far
```

where *ARRAY_NAME* is the array name. Pointers from Windows are by their nature far (that is, the data is pointed to by its own selector), and must be used as far in the 32-bit environment. Of course, conversions are only required if you actually need to reference the pointer.

Function pointers (i.e., pointers to callback routines) used by Windows are not converted from 32-bit to 16-bit. Rather, a 16-bit thunking layer that transfers control from the 16-bit environment to the 32-bit environment must be used. This thunking layer is provided by the Supervisor.

20.5.1 When To Convert Incoming Pointers

Whenever you wish to use a pointer passed to you by Windows, you must map a dynamically allocatable array to the memory pointed to by the pointer using the *LOCATION* specifier of the *ALLOCATE* statement. You must also declare the array as *far* using the *array* pragma. The pointer conversion will occur automatically.

Some places where pointer conversion may be required are:

- LocalLock
- GlobalLock
- the *lParam* in a window callback routine (if it is a pointer)

20.5.2 When To Convert Outgoing Pointers

Typically, there is no need to do any kind of conversions on your pointers when passing them to Windows. The Supervisor handles all 32-bit to 16-bit translations for you, in the case of the regular Windows API functions. However, if you are passing a 32-bit pointer to some other 16-bit application in the Windows environment, then pointer conversions must be done. There are two types of "outgoing" pointers: data pointers and function pointers.

Function pointers (to callback routines) must have a thunking layer provided, using the *GetProc16* function (this is explained in detail in a later section).

Data pointers can be translated from 32-bit to 16-bit using the *AllocAlias16* and *AllocHugeAlias16* functions. These functions create 16-bit far pointers that have the same linear address as the 32-bit near pointer that was converted.

It is important to remember that when passing a pointer to a data structure in this fashion, any pointers in the data structure must also be converted.

The Supervisor will convert any pointers that it knows about; but there are some complications created by the fact that Windows allows you to pass pointers in functions that are prototyped to take a long integer.

The Windows API functions *SendMessage* and *SendDlgItemMessage* rely on other fields determining the nature of the long data item that they accept; this is discussed in detail in the next section.

20.5.2.1 SendMessage and SendDlgItemMessage

SendMessage and *SendDlgItemMessage* have special cover functions that determine when the 32-bit integer is really a pointer and needs to be converted. These cover functions are used automatically, unless the macro *NOCOVERSENDS* is defined before including `WINAPI.FI` as in the following example.

```
*$define NOCOVERSENDS
*$include winapi.fi
```

SendMessage and *SendDlgItemMessage* will do pointer conversions automatically using *AllocAlias16* and *FreeAlias16* (unless *NOCOVERSENDS* is defined) for the following message types:

- combo boxes (CB_ messages)
- edit controls (EM_ messages)
- list boxes (LB_ messages)
- certain windows messages (WM_ messages);

The messages that are intercepted by the cover functions for *SendMessage* and *SendDlgItemMessage* are:

CB_ADDSTRING	CB_DIR	CB_FINDSTRING
CB_FINDSTRINGEXACT	CB_GETLBTEXT	CB_INSERTSTRING
CB_SELECTSTRING		
EM_GETLINE	EM_GETRECT	EM_REPLACESEL
EM_SETRECT	EM_SETRECTNP	EM_SETTABSTOPS
LB_ADDSTRING	LB_DIR	LB_FINDSTRING
LB_FINDSTRINGEXACT	LB_GETITEMRECT	LB_GETSELITEMS
LB_GETTEXT	LB_INSERTSTRING	LB_SELECTSTRING
LB_SETTABSTOPS		
WM_MDICREATE	WM_NCCALCSIZE	

Note that for *SendMessage* and *SendDlgItemMessage*, some of the messages may NOT require pointer conversion:

- CB_ADDSTRING, CB_FINDSTRING, CB_FINDSTRINGEXACT, CB_INSERTSTRING will not need a conversion if the combo box was created as owner-draw style without CBS_HASSTRINGS style.
- LB_ADDSTRING, LB_FINDSTRING, LB_FINDSTRINGEXACT, LB_INSERTSTRING will not need a conversion if the list box was created as owner-draw style without LBS_HASSTRINGS style.

The macro *NOCOVERSENDS* should be defined in modules where messages like these are being sent. With these messages, the lParam data item does not contain a pointer, and the automatic pointer conversion would be incorrect. By doing

```
*$define NOCOVERSENDS
*$include winapi.fi
```

modules that send messages like the above will not have the pointer conversion performed.

20.5.3 GlobalAlloc and LocalAlloc

The functions *GlobalAlloc* and *LocalAlloc* are the typical way of allocating memory in the 16-bit Windows environment. In the 32-bit environment, there is no need to use these functions. The only time *GlobalAlloc* is needed is when allocating shared memory, i.e., *GMEM_DDESHARE*.

The *ALLOCATE* and *DEALLOCATE* statements can be used to allocate memory from your 32-bit near heap. By allocating memory in this way, you may create data objects as large as the enhanced mode Windows memory manager will permit.

20.5.4 Callback Function Pointers

To access a callback function, an instance of it must be created using *MakeProcInstance*. This creates a "thunk" (a special piece of code) that automatically puts the application's data segment into the AX register, and then calls the specified callback function.

In Windows 3.x, it is not possible to do a *MakeProcInstance* directly on a 32-bit callback routine, since Windows 3.x does not understand 32-bit applications. Therefore, it is necessary to use a 16-bit callback routine that passes control to the 32-bit callback routine. This 16-bit callback routine is automatically created by the Supervisor when using any of the standard Windows API functions that accept a callback routine.

The 16-bit callback routine for a 32-bit application is a special layer that transfers the parameters from a 16-bit stack to the 32-bit stack, and then passes control to 32-bit code. These 16-bit callback routines are found in the Supervisor. The function *GetProc16* provides pointers to these 16-bit callback routines.

However, it is not often necessary to use the *GetProc16* function to obtain a 16-bit/32-bit callback interface function.

In the general case, one would have to write code as follows:

```
integer*4 pCb, fpProc

pCb = GetProc16( A_Function, GETPROC_callbacktype )
fpProc = MakeProcInstance( pCb, hInstance )

*   do stuff

    call Do_it( ..., fpProc, ... )

*   do more stuff

    call FreeProcInstance( fpProc )
    call ReleaseProc16( pCb )
```

It is not necessary to use this general code in the case of the regular Windows API functions. The following functions will automatically allocate the correct 16-bit/32-bit callback interface functions:

- ChooseColor
- ChooseFont
- CreateDialog
- CreateDialogIndirect
- CreateDialogIndirectParam
- CreateDialogParam
- DialogBox
- DialogBoxIndirect
- DialogBoxIndirectParam
- DialogBoxParam
- EnumChildWindows
- EnumFonts
- EnumMetaFile
- EnumObjects
- EnumProps
- EnumTaskWindows
- EnumWindows
- Escape (SETABORTPROC option)
- FindText
- GetOpenFileName
- GetSaveFileName
- GlobalNotify
- GrayString
- LineDDA
- PrintDlg
- RegisterClass
- ReplaceText
- SetClassLong (GCL_WNDPROC option)
- SetResourceHandler
- SetTimer
- SetWindowLong (GWL_WNDPROC option)
- SetWindowsHook

As well, the following functions are covered to provide support for automatic creation of 16-bit callback routines:

- FreeProcInstance
- MakeProcInstance
- UnhookWindowsHook

If you need to get a callback that is not used by one of the above functions, then you must code the general case. Typically, this is required when a DLL needs a callback routine. In modules where this is necessary, you define the macro **NOAUTOPROCS** before you include `WINAPI.FI` as in the following example.

```
*$define NOAUTOPROCS
*$include winapi.fi
```

Be careful of the following when using **NOAUTOPROCS**.

1. The call to **MakeProcInstance** and **FreeProcInstance** for the callback function occurs in a module with **NOAUTOPROCS** defined.
2. No Windows API functions (listed above) are used in the module with **NOAUTOPROCS** defined. If they are, you must code the general case to use them.

Note that *NOAUTOPROCS* is in effect on a module-to-module basis only.

RegisterClass automatically does a GetProc16 for the callback function, unless the macro NOCOVERRC is specified before including WINAPI.FI as in the following example.

```
*$define NOCOVERRC
*$include winapi.fi
```

20.5.4.1 Window Sub-classing

Sub-classing a Windows control in the 32-bit environment is straightforward. In fact, the code is identical to the code used in the 16-bit environment. A simple example is:

```
*$include winapi.fi
*$pragma aux (callback) SubClassProc parm( value, value, value, value )

integer*4 function SubClassProc( hWnd, msg, wp, lp )
integer*2 hWnd
integer*2 msg
integer*2 wp
integer*4 lp

include 'windows.fi'

common fpOldProc
integer*4 fpOldProc

! code for sub-classing here

SubClassProc = CallWindowProc( fpOldProc, hWnd, msg, wp, lp )

end

program SubClassDemo
integer*2 hControl
common fpOldProc
integer*4 fpOldProc
integer*4 fp;

include 'windows.fi'

integer*4 SubClassProc
external SubClassProc
integer*4 ProgramInstance
external ProgramInstance

! assume hControl gets created in here

fpOldProc = GetWindowLong( hControl, GWL_WNDPROC )
fp = MakeProcInstance( SubClassProc, ProgramInstance )
call SetWindowLong( hControl, GWL_WNDPROC, fp )

! set it back
call SetWindowLong( hControl, GWL_WNDPROC, fpOldProc )
call FreeProcInstance( fp )

end
```

Note that *SetWindowLong* is covered to recognize *GWL_WNDPROC* and automatically creates a 16-bit callback for the 32-bit callback. When replacing the callback routine with the original 16-bit routine, the covered version of *SetWindowLong* recognizes that the function is not a 32-bit callback, and so passes the pointer right through to Windows unchanged.

20.6 Calling 16-bit DLLs

A 16-bit function in a DLL can be called using the `_Call16` function. The first argument to `_Call16` is the address of the 16-bit function. This address is usually obtained by calling `GetProcAddress` with the name of the desired function. The second argument to `_Call16` is a string identifying the types of the parameters to be passed to the 16-bit function.

Character Parameter Type

<i>c</i>	call a 'cdecl' function as opposed to a 'pascal' function (if specified, it must be listed first)
<i>b</i>	unsigned BYTE
<i>w</i>	16-bit WORD
<i>d</i>	32-bit DWORD
<i>f</i>	double precision floating-point
<i>p</i>	32-bit flat pointer (converted to 16:16 far pointer)

The 16-bit function must use either the **PASCAL** or **CDECL** calling convention. **PASCAL** calling convention is the default. If the function uses the **CDECL** calling convention, then you must specify the letter "c" as the first character of the argument type string.

Pointer types will automatically be converted from 32-bit near pointers to 16-bit far pointers before the function is invoked. Note that this pointer is only valid over the period of the call; after control returns to the 32-bit application, the 16-bit pointer created by the Supervisor is no longer valid.

The return value from `_Call16` is a **DWORD**.

```
*$include winapi.fi

integer*2 function FWinMain( hInstance,
&                           hPrevInstance,
&                           lpszCmdLine,
&                           nCmdShow )
integer*2 hInstance
integer*2 hPrevInstance
integer*4 lpszCmdLine
integer*2 nCmdShow

include 'windows.fi'

integer*2 hDrv, hWnd
integer*4 lpfn, cb

hDrv = LoadLibrary( 'your.dll'c )
if( hDrv .lt. 32 )then
    return
end if
lpfn = GetProcAddress( hDrv, 'ExtDeviceMode'c )
if( lpfn .eq. 0 )then
    return
end if
! Invoke the function.
cb = _Call16( lpfn, 'wwdppddw'c,
&           hWnd, hDrv, NULL,
&           'POSTSCRIPT PRINTER'c,
&           'LPT1'c,
&           NULL, NULL, 0 )
```

20.7 _16 Functions

Every Windows API function that accepts a pointer has a corresponding _16 function. The _16 version of the function will not convert any of the pointers that it accepts; it will assume that all pointers are 16-bit far pointers already. This applies to both data and function pointers.

21 Windows 32-Bit Dynamic Link Libraries

21.1 Introduction to 32-Bit DLLs

Open Watcom FORTRAN 77 allows the creation of 32-bit Dynamic Link Libraries (DLL). In fact, 32-bit DLLs are simpler to write than 16-bit DLLs. A 16-bit DLL runs on the caller's stack, and thus $DS \neq SS$. This creates difficulties in the small and medium memory models because near pointers to local variables are different from near pointers to global variables. The 32-bit DLL runs on its own stack, in the usual flat memory space, which eliminates these concerns.

There is a special version of the supervisor, `W386DLL.EXT` that performs a similar job to `WIN386.EXT`. However, the 32-bit DLL supervisor is a 16-bit Windows DLL, rather than a 16-bit Windows application. On the first use of the 32-bit DLL, the DLL supervisor loads the 32-bit DLL and invokes the 32-bit initialization routine (the DLL's `FWinMain` routine). The initialization routine declares all entry points (via `DefineDLEntry`) and performs any other necessary initialization. An index number in the range 1 to 128 is used to identify all external 32-bit DLL routines. `DefineDLEntry` is used to assign an index number to each routine, as well as to identify the arguments.

The DLL supervisor contains a general entry point for Windows applications to call into called `Win386LibEntry`. It also contains 128 specific entry points called `DLL1` to `DLL128` which correspond to the entry points established via `DefineDLEntry` (the first argument to `DefineDLEntry` is an index number in the range 1 to 128). All applications call into the 32-bit DLL via these entry points. They build the necessary stack frame and switch to the 32-bit DLL's data space.

If you call via `Win386LibEntry` then you pass the DLL entry point number or index (1 to 128) as the last argument. `Win386LibEntry` uses this index number to call the appropriate 32-bit DLL routine. From a pseudo-code point of view, the 16-bit supervisor might look like the following:

```
DLL1::  set index=1
        invoke 32bitDLLindirect

DLL2::  set index=2
        invoke 32bitDLLindirect
        .
        .
        .
DLL128:: set index=128
        invoke 32bitDLLindirect

Win386LibEntry::
        set index from index_argument
        invoke 32bitDLLindirect

32bitDLLindirect:
        set up stack frame
        switch to 32-bit data space
        call indirect registration_list[ index ]
        .
        .
        .
```

When you are creating a 32-bit DLL, keep in mind that the entry points you define may be invoked by a 16-bit application as well as a 32-bit application. It is for this reason that all far pointers passed to a 32-bit DLL are 16-bit far pointers. Hence, whenever a pointer is passed as an argument to a 32-bit DLL entry point and you wish to access the data it points to, you must convert the pointer appropriately. To do this, you must map a dynamically allocatable array to the memory pointed to by the 16-bit far pointer.

21.2 A Sample 32-bit DLL

Let us begin our discussion of DLLs by showing the code for a simple DLL. The source code for these examples is provided in the \WATCOM\SAMPLES\FORTRAN\WIN\DLL directory. We describe how to compile and link the examples in the section entitled "Creating and Debugging Dynamic Link Libraries" on page 147.

```
*$include winapi.fi

* WINDLLV.FOR

* Setup:          set finclude=\WATCOM\src\fortran\win
* Compile and Link: wfl386 windllv -explicit -d2 -bd -l=win386
* Bind:           wbind windllv -d -n

*$pragma aux (dll_function) Add3

    integer function Add3( w1, w2, w3 )
    integer*4 w1, w2, w3

    include 'windows.fi'

    character*128 str

    write( str, '(16hDLL 1 arguments:, 3i10, a)' ) w1, w2, w3,
    &                                           char(0)
    call MessageBox( NULL, str, 'DLL Function 1'c, MB_OK )
    Add3 = w1 + w2 + w3

    end

*$pragma aux (dll_function) Add2

    integer function Add2( w1, w2 )
    integer*4 w1, w2

    include 'windows.fi'

    character*128 str

    write( str, '(16hDLL 2 arguments:, 2i10, a)' ) w1, w2, char(0)
    call MessageBox( NULL, str, 'DLL Function 2'c, MB_OK )
    Add2 = w1 + w2

    end

    integer*2 function FWinMain( hInstance,
    &                             hPrevInstance,
    &                             lpszCmdLine,
    &                             nCmdShow )
    integer*2 hInstance
    integer*2 hPrevInstance
    integer*4 lpszCmdLine
    integer*2 nCmdShow

    include 'windows.fi'

    external Add3, Add2
    integer rc
```



```

call BreakPoint
rc = DefineDLEntry( 1, Add3, DLL_DWORD, DLL_DWORD, DLL_DWORD,
& DLL_ENDLIST )
if( rc .ne. 0 )then
    FWinMain = 0
    return
end if
rc = DefineDLEntry( 2, Add2, DLL_DWORD, DLL_DWORD,
& DLL_ENDLIST )
if( rc .ne. 0 )then
    FWinMain = 0
    return
end if
call MessageBox( NULL, '32-bit DLL started'c,
& 'WINDLLV'c, MB_OK )
FWinMain = 1

end

```

There are two entry points defined, Add3 (index number 1) and Add2 (index number 2). Add3 has three INTEGER*4 arguments and Add2 has two INTEGER*4 arguments. The argument lists are described by calling DefineDLEntry. All arguments are passed by value. As previously mentioned, all pointers passed to 32-bit DLLs are 16-bit far pointers. Since, by default, FORTRAN 77 passes arguments by reference (a pointer to the data is passed instead of the actual data), a level of complexity is introduced since some pointer conversions must take place when accessing the data pointed to by a 16-bit far pointer in a 32-bit environment. We will deal with this problem in a following example. First, let us deal with passing arguments by value to 32-bit DLLs from 16 and 32-bit Windows applications.

Note that each entry name must be given the `dll_function` attribute using an auxiliary pragma. This alias name is defined in the file `WINAPI.FI`.

FWinMain returns zero to notify Windows that the DLL initialization failed, and returns a one if initialization succeeds.

FWinMain accepts the same arguments as the FWinMain procedure of a regular Windows program, however, only two arguments are used. `hInstance` is the DLL handle and `lpszCmdLine` is the command line passed to the DLL.

21.3 Calling Functions in a 32-bit DLL from a 16-bit Application

The following is a 16-bit Windows program that demonstrates how to call the two routines defined in our DLL example.

```

*$include winapi.fi

* GEN16V.FOR

* Setup:          set finclude=\WATCOM\src\fortran\win
* Compile and Link: wfl gen16v -explicit -d2 -windows -l=windows
*                  -"op desc '16-bit DLL Test'"

*$pragma aux (dll32_call) indirect_1 \
*             parm( value*4, value*4, value*4 )
*$pragma aux (dll32_call) indirect_2 \
*             parm( value*4, value*4 )

```

```
integer*2 function FWinMain( hInstance,
&                               hPrevInstance,
&                               lpszCmdLine,
&                               nCmdShow )
integer*2 hInstance
integer*2 hPrevInstance
integer*4 lpszCmdLine
integer*2 nCmdShow

include 'windows.fi'

integer*2 hlib
integer*4 indirect_1, indirect_2
integer*4 dll_1, dll_2, cb
character*128 str

hlib = LoadLibrary( 'windllv.dll'c )
if( hlib .lt. 32 ) then
    call MessageBox( NULL, 'Can't load WINDLLV'c,
&                    'Gen16V'c, MB_OK )
    stop
endif

dll_1 = GetProcAddress( hlib, 'DLL1'c )
dll_2 = GetProcAddress( hlib, 'DLL2'c )

cb = indirect_1( 111, 22222, 3333, dll_1 )
write( str, '(15hDLL 1 returned , i10, a)' ) cb, char(0)
call MessageBox( NULL, str, 'Gen16V Test 1'c, MB_OK )

cb = indirect_2( 4444, 55, dll_2 )
write( str, '(15hDLL 2 returned , i10, a)' ) cb, char(0)
call MessageBox( NULL, str, 'Gen16V Test 2'c, MB_OK )

FWinMain = 0

end
```

The addresses of the routines DLL1 and DLL2 in the 32-bit DLL are obtained and stored in the variables `dll_1` and `dll_2`. Since the FORTRAN 77 language does not support indirect function calls, we need a mechanism to call these functions indirectly. We do this using the two indirect functions called `indirect_1` and `indirect_2`. These two functions are given the `dll32_call` attribute using an auxiliary pragma which is defined in the file `WINAPI.FI`. Note that the last argument of the calls to `indirect_1` or `indirect_2` is the actual address of the DLL routine.

What you should realize is that the `indirect_1` and `indirect_2` functions do not really exist. The code that is generated for statements like the following is really an indirect call to the function whose address is represented in the last argument.

```
cb = indirect_1( 111, 22222, 3333, dll_1 )
cb = indirect_2( 4444, 55, dll_2 )
```

This is a result of using the `dll32_call` auxiliary pragma attribute to describe both `indirect_1` and `indirect_2`. You can verify this by disassembling the object file that is generated when this code is compiled.

21.4 Calling Functions in a 32-bit DLL from a 32-bit Application

The following is a 32-bit Windows program that demonstrates how to call the two routines defined in our 32-bit DLL example. Since this is a 32-bit Windows program, we will use the `_Call16` function to call functions in our 32-bit DLL. Note that we get to the 32-bit DLL functions by going indirectly through the 16-bit supervisor that forms the "front end" for our 32-bit DLL.

```
*$include winapi.fi

* GEN32V.FOR

* Setup:          set finclude=\WATCOM\src\fortran\win
* Compile and Link: wfl386 gen32v -explicit -d2 -l=win386
* Bind:           wbind gen32v -n -D "32-bit DLL Test"

integer*2 function FWinMain( hInstance,
&                             hPrevInstance,
&                             lpszCmdLine,
&                             nCmdShow )
integer*2 hInstance
integer*2 hPrevInstance
integer*4 lpszCmdLine
integer*2 nCmdShow

include 'windows.fi'

integer*2 hlib
integer*4 dll_1, dll_2, cb
character*128 str

hlib = LoadLibrary( 'windllv.dll'c )
if( hlib .lt. 32 ) then
    call MessageBox( NULL, 'Can't load WINDLLV'c,
&                    'Gen32V'c, MB_OK )
    stop
endif

dll_1 = GetProcAddress( hlib, 'DLL1'c )
dll_2 = GetProcAddress( hlib, 'DLL2'c )

cb = _Call16( dll_1, 'ddd'c, 111, 22222, 3333 )
write( str, '(15hDLL 1 returned , i10, a)' ) cb, char(0)
call MessageBox( NULL, str, 'Gen32V Test 1'c, MB_OK )

cb = _Call16( dll_2, 'dd'c, 4444, 55 )
write( str, '(15hDLL 2 returned , i10, a)' ) cb, char(0)
call MessageBox( NULL, str, 'Gen32V Test 2'c, MB_OK )

FWinMain = 0

end
```

Note that the first argument of a call to `_Call16` is the DLL function address returned by `GetProcAddress` and must be a 32-bit argument. The second argument of a call to `_Call16` is a string describing the types of arguments that will be passed to the DLL function.

21.5 A Sample 32-bit DLL Using a Structure

As previously mentioned, passing pointers from a 16 or 32-bit Windows application to a 32-bit DLL poses a problem since all pointers are passed as 16-bit far pointers. The pointer must be converted from a 16-bit far pointer to a 32-bit far pointer. This is achieved by mapping a dynamically allocatable array to each argument that is passed by reference using the `LOCATION` specifier of the `ALLOCATE` statement. Furthermore, you must specify the `far` attribute for each such array using the `array` pragma. Since this is cumbersome if you wish to pass many arguments, it is recommended that a single argument be passed that is actually a pointer to a structure that contains the actual arguments. Furthermore, since each call to a DLL routine is made indirectly through one of `Win386LibEntry` or `DLL1` through `DLL128`, you should also return any values in the same structure since the return value from any of these functions is only 32-bits wide.

The following example is a 32-bit DLL that receives its arguments and returns values using a structure. The source code for these examples is provided in the `\WATCOM\SAMPLES\FORTRAN\WIN\DLL` directory. We describe how to compile and link the examples in the section entitled "Creating and Debugging Dynamic Link Libraries" on page 147.

```
*$include winapi.fi

* WINDLL.FOR

* Setup:          set finclude=\WATCOM\src\fortran\win
* Compile and Link: wfl386 windll -explicit -d2 -bd -l=win386
* Bind:           wbind windll -d -n

*$pragma aux (dll_function) Add3

      subroutine Add3( arg_list )
      integer*4 arg_list
      structure /argtypes/
        integer    w1
        integer    w2
        integer    w3
        integer    sum
      end structure
      record /argtypes/ args(:)
*$pragma array args far

      include 'windows.fi'

      character*128 str

      allocate( args(1), location=arg_list )
      write( str, '(16hDLL 1 arguments:, 3i10, a)' ) args(1).w1,
&                                                    args(1).w2,
&                                                    args(1).w3,
&                                                    char(0)
      call MessageBox( NULL, str, 'DLL Function 1'c, MB_OK )
      args(1).sum = args(1).w1 + args(1).w2 + args(1).w3
      deallocate( args )

      end
```

```

*$pragma aux (dll_function) Add2

      subroutine Add2( arg_list )
      integer*4 arg_list
      structure /argtypes/
        real      w1
        real      w2
        real      sum
      end structure
      record /argtypes/ args(:)
*$pragma array args far

      include 'windows.fi'

      character*128 str

      allocate( args(1), location=arg_list )
      write( str, '(16hDLL 2 arguments:, 2f10.2, a)' ) args(1).w1,
&                                                    args(1).w2,
&                                                    char(0)
      call MessageBox( NULL, str, 'DLL Function 2'c, MB_OK )
      args(1).sum = args(1).w1 + args(1).w2
      deallocate( args )

      end

      integer*2 function FWinMain( hInstance,
&                                hPrevInstance,
&                                lpszCmdLine,
&                                nCmdShow )
      integer*2 hInstance
      integer*2 hPrevInstance
      integer*4 lpszCmdLine
      integer*2 nCmdShow

      include 'windows.fi'

      external Add3, Add2
      integer rc

      call BreakPoint
      rc = DefineDLEntry( 1, Add3, DLL_PTR, DLL_ENDLIST )
      if( rc .ne. 0 )then
        FWinMain = 0
        return
      end if
      rc = DefineDLEntry( 2, Add2, DLL_PTR, DLL_ENDLIST )
      if( rc .ne. 0 )then
        FWinMain = 0
        return
      end if
      call MessageBox( NULL, '32-bit DLL started'c,
&                    'WINDLL'c, MB_OK )
      FWinMain = 1

      end

```

The following example is a 16-bit Windows application that passes arguments to a 32-bit DLL using a structure.

```

*$include winapi.fi

* GEN16.FOR

* Setup:          set finclude=\WATCOM\src\fortran\win
* Compile and Link: wfl gen16 -explicit -d2 -windows -l=windows
*                  -"op desc '16-bit DLL Test'"

*$pragma aux (dll32_call) indirect_1 parm( reference, value*4 )
*$pragma aux (dll32_call) indirect_2 parm( reference, value*4 )

```

```
integer*2 function FWinMain( hInstance,
&                               hPrevInstance,
&                               lpszCmdLine,
&                               nCmdShow )
integer*2 hInstance
integer*2 hPrevInstance
integer*4 lpszCmdLine
integer*2 nCmdShow

include 'windows.fi'

integer*2 hlib
integer*4 dll_1, dll_2
character*128 str

structure /args_1/
integer    w1
integer    w2
integer    w3
integer    sum
end structure

structure /args_2/
real       w1
real       w2
real       sum
end structure

record /args_1/ args_1/111, 22222, 3333, 0/
record /args_2/ args_2/714.3, 35.7, 0.0/

hlib = LoadLibrary( 'windll.dll'c )
if( hlib .lt. 32 ) then
    call MessageBox( NULL, 'Can't load WINDLL'c,
&                    'Gen16'c, MB_OK )
    stop
endif

dll_1 = GetProcAddress( hlib, 'DLL1'c )
dll_2 = GetProcAddress( hlib, 'DLL2'c )

call indirect_1( args_1, dll_1 )
write( str, '(15hDLL 1 returned , i10, a)' ) args_1.sum,
&                                           char(0)
call MessageBox( NULL, str, 'Gen16 Test 1'c, MB_OK )

call indirect_2( args_2, dll_2 )
write( str, '(15hDLL 2 returned , f10.2, a)' ) args_2.sum,
&                                           char(0)
call MessageBox( NULL, str, 'Gen16 Test 2'c, MB_OK )

FWinMain = 0

end
```

The following example is a 32-bit Windows application that passes arguments to a 32-bit DLL using a structure.

```
*$include winapi.fi

* GEN32.FOR

* Setup:          set finclude=\WATCOM\src\fortran\win
* Compile and Link: wfl386 gen32 -explicit -d2 -l=win386
* Bind:           wbind gen32 -n -D "32-bit DLL Test"
```

```

integer*2 function FWinMain( hInstance,
&                               hPrevInstance,
&                               lpszCmdLine,
&                               nCmdShow )
integer*2 hInstance
integer*2 hPrevInstance
integer*4 lpszCmdLine
integer*2 nCmdShow

include 'windows.fi'

integer*2 hlib
integer*4 dll_1, dll_2, cb
character*128 str

structure /args_1/
integer    w1
integer    w2
integer    w3
integer    sum
end structure

structure /args_2/
real       w1
real       w2
real       sum
end structure

record /args_1/ args_1/111, 22222, 3333, 0/
record /args_2/ args_2/714.3, 35.7, 0.0/

hlib = LoadLibrary( 'windll.dll'c )
if( hlib .lt. 32 ) then
    call MessageBox( NULL, 'Can't load WINDLL'c,
&                    'Gen32'c, MB_OK )
    stop
endif

dll_1 = GetProcAddress( hlib, 'DLL1'c )
dll_2 = GetProcAddress( hlib, 'DLL2'c )

cb = _Call16( dll_1, 'p'c, loc(args_1) )
write( str, '(15hDLL 1 returned , i10, a)' ) args_1.sum,
&                                           char(0)
call MessageBox( NULL, str, 'Gen32 Test 1'c, MB_OK )

cb = _Call16( dll_2, 'p'c, loc(args_2) )
write( str, '(15hDLL 2 returned , f10.2, a)' ) args_2.sum,
&                                           char(0)
call MessageBox( NULL, str, 'Gen32 Test 2'c, MB_OK )

FWinMain = 0

end

```

21.6 Creating and Debugging Dynamic Link Libraries

In the following sections, we will take you through the steps of compiling, linking, and debugging 32-bit Dynamic Link Libraries (DLLs).

We will use example programs that are provided in source-code form in the Open Watcom F77 package. The files described in this chapter are located in the directory \WATCOM\SAMPLES\FORTRAN\WIN\DLL. The following files are provided:

WINDLLV.FOR	is the source code for a simple 32-bit DLL containing two library routines that use integer arguments to pass information.
GEN16V.FOR	is the source code for a generic 16-bit Windows application that calls functions in the "WINDLLV" 32-bit Windows DLL.
GEN32V.FOR	is the source code for a generic 32-bit Windows application that calls functions in the "WINDLLV" 32-bit Windows DLL.
WINDLL.FOR	is the source code for a simple 32-bit DLL containing two library routines that use structures to pass information.
GEN16.FOR	is the source code for a generic 16-bit Windows application that calls functions in the "WINDLL" 32-bit Windows DLL.
GEN32.FOR	is the source code for a generic 32-bit Windows application that calls functions in the "WINDLL" 32-bit Windows DLL.
MAKEFILE	is a makefile for compiling and linking the programs described above.

21.6.1 Building the Applications

To create the DLLs and test applications, we will use the WATCOM Open Watcom Make utility and the supplied makefile.

Example:

```
C>wmake -f makefile
```

21.6.2 Installing the Examples under Windows

Start up Microsoft Windows 3.x if you have not already done so. Add the GEN16V.EXE and GEN32V.EXE files to one of your Window groups using the Microsoft Program Manager.

1. Select the "New..." entry from the "File" menu of the Microsoft Windows Program Manager.
2. Select "Program Item" from the "New Program Object" window and press the "OK" button.
3. Enter "16-bit DLL Test" as a description for the GEN16V program. Enter the full path to the GEN16V program as a command line.

Example:

```
Description:    16-bit DLL Test
Command Line:   c:\work\dll\gen16v.exe
```

4. Enter "32-bit DLL Test" as a description for the GEN32V program. Enter the full path to the GEN32V program as a command line.

Example:

Description: 32-bit DLL Test
Command Line: c:\work\dll\gen32v.exe

Use a similar procedure to install the GEN16.EXE and GEN32.EXE programs.

21.6.3 Running the Examples

Start the 16-bit application by double clicking on its icon. A number of message boxes are presented. You may wish to compare the output in each message box with the source code of the program to determine if the correct results are being obtained. Click on the "OK" button as each of them are displayed.

Similarly, start the 32-bit application by double-clicking on its icon and observe the results.

21.6.4 Debugging a 32-bit DLL

The Open Watcom Debugger can be used to debug a DLL. To debug a 32-bit DLL, a "breakpoint" instruction must be inserted into the source code for the DLL at the "FWinMain" entry point. This is done using the "pragma" compiler directive. We have already added the breakpoint to the source code for the 32-bit DLL.

```

        .
        .
        .
integer*2 function FWinMain( hInstance,
&                               hPrevInstance,
&                               lpszCmdLine,
&                               nCmdShow )
integer*2 hInstance
integer*2 hPrevInstance
integer*4 lpszCmdLine
integer*2 nCmdShow

include 'windows.fi'

external Add3, Add2
integer rc

call BreakPoint
rc = DefineDLEntry( 1, Add3, DLL_DWORD, DLL_DWORD, DLL_DWORD,
&                               DLL_ENDLIST )
if( rc .ne. 0 )then
    FWinMain = 0
    return
end if
        .
        .
        .
    
```

The pragma for "BreakPoint" is defined in the "WINAPI.FI" file.

Start up Microsoft Windows 3.x if you have not already done so. Start the debugger by double-clicking on the Open Watcom Debugger icon. At the prompt, enter the path specification for the application. When the debugger has successfully loaded GEN32v, start execution of the program. When the breakpoint is encountered in the 32-bit DLL, the debugger is re-entered. The debugger will automatically skip past the breakpoint.

From this point on, you can symbolically debug the 32-bit DLL. You might, for example, set breakpoints at the start of each DLL routine to debug each of them as they are called.

21.6.5 Summary

Note that the "FWinMain" entry point is only called once, at the start of any application requesting it. After this, the "FWinMain" entry point is no longer called. You may have to restart Windows to debug this section of code a second or third time.

22 Interfacing Visual Basic and Open Watcom FORTRAN 77 DLLs

22.1 Introduction to Visual Basic and DLLs

This chapter describes how to interface Microsoft Visual Basic 3.0 applications and 32-bit Dynamic Link Libraries (DLLs) created by Open Watcom FORTRAN 77. It describes how to write functions for a 32-bit DLL, how to compile and link them, and how to call these functions from Visual Basic. One of the proposed techniques involves the use of a set of cover functions in a 16-bit DLL so, indirectly, this chapter also describes interfacing to 16-bit DLLs.

It is possible to invoke the `Win386LibEntry` function (Open Watcom's 32-bit function entry point, described below) directly from Visual Basic. However, this technique limits the arguments that can be passed to a 32-bit DLL. The procedure and problems are explained below.

To work around the problem, a 16-bit DLL can be created, that covers the 32-bit DLL. Within the 16-bit DLL, we will place cover functions that will call the corresponding 32-bit function in the 32-bit DLL. We illustrate the creation of the 16-bit DLL using the 16-bit C compiler in Open Watcom C/C++.

Before we begin our example, there are some important technical issues to consider.

The discussion in this chapter assumes that you, the developer, have a working knowledge of Visual Basic, including how to bring up the general declarations screen, how to create command buttons, and how to associate code with command buttons. You must use Visual Basic 3.0 or later. Visual Basic Version 2.x will not work because of a deficiency in this product regarding the calling of functions in DLLs.

For the purposes of the following discussion, you should have installed the 32-bit version of Open Watcom FORTRAN 77, as well as version 3.0 or later of Visual Basic. If you also have the 16-bit Open Watcom C compiler, you can use this to create a 16-bit DLL containing the 16-bit cover functions. Ensure that the **PATH** and **FINCLUDE** environment variables are defined to include at least the directories indicated. We have assumed that Open Watcom FORTRAN 77 is installed in the `c:\watcom` directory, and Visual Basic is in the `c:\vb` directory:

```
set path=c:\watcom\binw;c:\vb;c:\dos;c:\windows
set finclude=c:\watcom\src\fortran\win
```

Open Watcom's 32-bit DLL supervisor contains a general entry point for Windows applications to call into called `Win386LibEntry`. It also contains 128 specific entry points called `DLL1` to `DLL128` which correspond to the entry points established via `DefinedLLEntry` (the first argument to `DefinedLLEntry` is an index number in the range 1 to 128). All applications call into the 32-bit DLL via these entry points. They build the necessary stack frame and switch to the 32-bit DLL's data space.

If you call via `Win386LibEntry` then you pass the DLL entry point number or index (1 to 128) as the last argument. `Win386LibEntry` uses this index number to call the appropriate 32-bit DLL routine.

In many languages and programs (such as C and Microsoft Excel), function calls are very flexible. In other words, a function can be called with different argument types each time. This is generally necessary for

calling `Win386LibEntry` in a 32-bit extended DLL function. The reason is that this function takes the same arguments as the function being called, as well as the index number of the called function. After the 32-bit flat model has been set up, `Win386LibEntry` then calls this function. In Visual Basic, once a function is declared as having certain arguments, it cannot be redeclared. For example, suppose we have a declaration as follows:

Example:

```
Declare Function Win386LibEntry Lib "c:\path\vbdll132.dll"  
=> (ByVal v1 As Long, ByVal v2 As Long, ByVal  
=> v3 As Long, ByVal I As Integer) As Long
```

(Note: the `=>` means to continue the statement on the same line.) In this example, we could only call a function in any 32-bit extended DLL with three 32-bit integers as arguments. There are three ways to work around this deficiency in Visual Basic:

1. Use the Visual Basic "Alias" attribute to declare `Win386LibEntry` differently for each DLL routine. Reference the different DLL routines using these aliases.
2. Use the specific entry point, one of `DLL1` through `DLL128`, corresponding to the DLL routine that you want to call. Each entry point can be described to take different arguments. We can still use the "Alias" attribute to make the link between the name we use in the Visual Basic function and the name in the 32-bit extended DLL. This is the method that we will use in the "Direct Call" technique discussed below. It is simpler to use since it requires one less argument (you don't require the index number).
3. Use a method which involves calling functions in a 16-bit "cover" DLL written in a flexible-argument language, which then calls the functions in the 32-bit DLL. This is the "Indirect Call" method discussed below.

22.2 A Working Example

The best way to demonstrate these techniques is through an example. This example consists of a Visual Basic application with 3 push buttons. The first push button invokes a direct call to a 32-bit DLL which will display a message window with its arguments, the second push button invokes an indirect call to the same function through a 16-bit DLL, and the third button exits the Visual Basic application.

To create a Visual Basic application:

- (1) **Start up a new project folder** from the "File" menu.
- (2) **Select "View Form"** from the "Project" window.
- (3) **Draw three command buttons** on the form by selecting command buttons from the "Toolbox" window.
- (4) **Change the caption on each button.** To do this, highlight the first button. Then, open the "Properties" window. Double click on the "Caption window", and change the caption to "Direct call". Highlight the second button, and change its caption to "Indirect call". Highlight the third, changing the caption to "Exit".

Now, your Visual Basic application should have three push buttons, "Direct call", "Indirect call", and "Exit".

(5) Double click on the "Direct Call" button.

An edit window will pop up. Enter the following code:

```
Sub Command1_Click ()
    Dim var1, var2, var3, worked As Long

    var1 = 230
    var2 = 215
    var3 = 32
    worked = Add3(var1, var2, var3)
    Print worked
    worked = Add2(var2, var3)
    Print worked
End Sub
```

(6) Double click on the "Indirect Call" button.

Another edit window will pop up. Enter the following code:

```
Sub Command2_Click ()
    Dim var1, var2, var3, worked As Long

    var1 = 230
    var2 = 215
    var3 = 32
    worked = Function1(var1, var2, var3)
    Print worked
    worked = Function2(var2, var3)
    Print worked
End Sub
```

(7) Double click on the "Exit" command button and enter the following code in the pop-up window:

```
Sub Command3_Click ()
    End
End Sub
```

(8) Select "View Code" from the "Project" window. To interface these Visual Basic functions to the DLLs, the following code is needed in the

```
Object: [general] Proc: [declarations]
```

section of the code. This code assumes that VBDLL32.DLL and COVER16.DLL are in the c:\path directory. Modify the pathnames appropriately if this is not the case. (Note: the => means to continue the statement on the same line.)

```
Declare Function Function1 Lib "c:\path\cover16.dll"  
=> (ByVal v1 As Long, ByVal v2 As Long, ByVal v3 As Long)  
=> As Long  
  
Declare Function Function2 Lib "c:\path\cover16.dll"  
=> (ByVal v1 As Long, ByVal v2 As Long) As Long  
  
Declare Function Add3 Lib "c:\path\vbdll32.dll"  
=> Alias "DLL1"  
=> (ByVal v1 As Long, ByVal v2 As Long, ByVal v3 As Long)  
=> As Long  
  
Declare Function Add2 Lib "c:\path\vbdll32.dll"  
=> Alias "DLL2"  
=> (ByVal v1 As Long, ByVal v2 As Long) As Long
```

Now, when all of the code below is compiled correctly, and the Visual Basic program is run, the "Direct call" button will call the DLL1 and DLL2 functions directly, aliased as the functions Add3 and Add2 respectively. The "Indirect call" button will call the 16-bit DLL, which will then call the 32-bit DLL, for both Function1 and Function2. To run the Visual Basic program, select "Start" from the "Run" menu.

22.3 Sample Visual Basic DLL Programs

The sample programs provided below are for a 32-bit DLL, and a 16-bit cover DLL, which will call the two functions contained in the 32-bit DLL.

22.3.1 Source Code for VBDLL32.DLL

```
*$include winapi.fi  
  
* VBDLL32.FOR  
  
* Setup:          set finclude=\WATCOM\src\fortran\win  
* Compile and Link: wfl386 vbdll32 -explicit -d2 -bd -l=win386  
* Bind:           wbind vbdll32 -d -n  
  
*$pragma aux (dll_function) Add3  
  
    integer function Add3( w1, w2, w3 )  
    integer w1, w2, w3  
  
    include 'windows.fi'  
  
    character*128 str  
  
    write( str, '(16hDLL 1 arguments:, 3i10, a)' ) w1, w2, w3,  
    &                                           char(0)  
    call MessageBox( NULL, str, 'F77 VBDLL32'c, MB_OK )  
    Add3 = w1 + w2 + w3  
  
end
```

```
*$pragma aux (dll_function) Add2

integer function Add2( w1, w2 )
integer w1, w2

include 'windows.fi'

character*128 str

write( str, '(16hDLL 2 arguments:, 2i10, a)' ) w1, w2, char(0)
call MessageBox( NULL, str, 'F77 VBDLL32'c, MB_OK )
Add2 = w1 + w2

end

integer*2 function FWinMain( hInstance,
&                           hPrevInstance,
&                           lpszCmdLine,
&                           nCmdShow )
integer*2 hInstance
integer*2 hPrevInstance
integer*4 lpszCmdLine
integer*2 nCmdShow

include 'windows.fi'

external Add3, Add2
integer rc

rc = DefineDLEntry( 1, Add3, DLL_DWORD, DLL_DWORD, DLL_DWORD,
&                  DLL_ENDLIST )
if( rc .ne. 0 )then
    FWinMain = 0
    return
end if
rc = DefineDLEntry( 2, Add2, DLL_DWORD, DLL_DWORD,
&                  DLL_ENDLIST )
if( rc .ne. 0 )then
    FWinMain = 0
    return
end if
call MessageBox( NULL, '32-bit DLL started'c,
&                'F77 VBDLL32'c, MB_OK )
FWinMain = 1

end
```

22.3.2 Source code for COVER16.DLL

The functions in this 16-bit DLL will call the functions in the 32-bit DLL, VBDLL32.DLL, shown above, with the appropriate Win386LibEntry call for each function.

```
/*
 * COVER16.C
 */

#include <stdio.h>
#include <windows.h> /* required for all Windows applications */

typedef long (FAR PASCAL *FPROC)();

FPROC DLL_1;
FPROC DLL_2;
```

```
long FAR PASCAL __export Function1( long var1,
                                   long var2,
                                   long var3 )
{
    return( (long) DLL_1( var1, var2, var3 ) );
}

long FAR PASCAL __export Function2( long var1, long var2 )
{
    return( (long) DLL_2( var1, var2 ) );
}

#pragma off (unreferenced);
BOOL FAR PASCAL LibMain( HANDLE hInstance, WORD wDataSegment,
                        WORD wHeapSize, LPSTR lpszCmdLine )
#pragma on (unreferenced);
{
    HANDLE hlib;

    /* Do our DLL initialization */
    hlib = LoadLibrary( "vbdll32.dll" );
    if( hlib < 32 ) {
        MessageBox( NULL,
                    "Make sure your PATH contains VBDLL32.DLL",
                    "COVER16", MB_OK | MB_ICONEXCLAMATION );
        return( FALSE );
    }
    DLL_1 = (FPROC) GetProcAddress( hlib, "DLL1" );
    DLL_2 = (FPROC) GetProcAddress( hlib, "DLL2" );
    return( TRUE );
}
```

22.4 Compiling and Linking the Examples

To create the 32-bit DLL `VBDLL32.DLL`, type the following at the command line (make sure that `VBDLL32.for` is in your current directory):

```
set finclude=c:\watcom\src\fortran\win
wfl386 vbdll32 -explicit -bd -d2 -l=win386
wbind vbdll32 -d -n
```

To create the 16-bit DLL `COVER16.DLL`, type the following at the command line (make sure that `COVER16.C` are in your current directory):

```
wcl cover16 -mc -bt=windows -bd -zu -d2 -l=windows_dll
```

Notes:

1. An object file is provided for `COVER16.C` if you do not have access to the 16-bit Open Watcom C compiler. In this case, the DLL can be generated from the object file using the following command:

```
wfl cover16.obj -d2 -l=windows_dll
```

2. The "mc" option selects the compact memory model (small code, big data). The code for 16-bit DLLs must be compiled with one of the big data models.
3. The "bd" option indicates that a DLL will be created from the object files.
4. The "bt" option selects the "windows" target. This option causes the C or C++ compiler to generate Windows prologue/epilogue code sequences which are required for Microsoft Windows

applications. It also causes the compiler to use the **WINDOWS_INCLUDE** environment variable for header file searches. It also causes the compiler to define the macro `__WINDOWS__` and, for the 32-bit C or C++ compiler only, the macro `__WINDOWS_386__`.

5. The "zu" option is used when compiling 16-bit code that is to be placed in a Dynamic Link Library (DLL) since the SS register points to the stack segment of the calling application upon entry to the function.
6. The "d2" option is used to disable optimizations and include debugging information in the object file and DLL. The techniques for debugging DLLs are described in the chapter entitled "Windows 32-Bit Dynamic Link Libraries" on page 139.

You are now ready to run the Visual Basic application.

23 *WIN386 Library Subprograms*

Each special Windows subprogram in the Open Watcom F77 library is described in this chapter. Each description consists of a number of subsections:

Synopsis: This subsection gives the include files that should be included within a source file that references the subprogram. It also shows an appropriate declaration for the function or for a function that could be substituted for a macro. This declaration is not included in your program; only the include file(s) should be included.

Description: This subsection is a description of the subprogram.

Returns: This subsection describes the return value (if any) for the subprogram.

See Also: This optional subsection provides a list of related subprograms.

Example: This optional subsection consists of one or more examples of the use of the subprogram. The examples are often just fragments of code (not complete programs) for illustration purposes.

Classification: This subsection provides an indication of where the subprogram is commonly found. The subprograms in this section are all classified as "WIN386" (i.e., they pertain to 32-bit Windows programming).

Synopsis: c\$include 'winapi.fi'
 integer*4 function AllocAlias16(ptr)
 integer*4 ptr

Description: The AllocAlias16 function obtains a 16-bit far pointer equivalent of a 32-bit near pointer. These pointers are used when passing data pointers to Windows through functions that have INTEGER*4 arguments, and for any pointers within data structures passed this way.

Returns: The AllocAlias16 function returns a 16-bit far pointer (as an INTEGER*4) usable by Windows, or returns 0 if the alias cannot be allocated.

See Also: FreeAlias16

Example: integer*4 mcs_16
 record /MDICREATESTRUCT/ mcs
 mcs.szTitle = AllocAlias16(loc('Title'c))
 mcs.szClass = AllocAlias16(loc('mdichild'c))
 mcs.hOwner = hInst
 mcs.x = mcs.cx = CW_USEDEFAULT
 mcs.y = mcs.cy = CW_USEDEFAULT
 mcs.style = 0
 ! Send a message to an MDI client to create a window.
 ! Since the pointer to the structure is passed in an
 ! argument that may not be a pointer (depending on the
 ! type of message), there is no implicit 32 to 16-bit
 ! conversion done so the conversion must be done by the
 ! programmer.
 mcs_16 = AllocAlias16(loc(mcs))
 hwnd = SendMessage(hwndMDI, WM_MDICREATE, 0, mcs_16)
 FreeAlias16(mcs_16)
 FreeAlias16(mcs.szClass)
 FreeAlias16(mcs.szTitle)

Classification: WIN386

Synopsis: c\$include 'winapi.fi'
 integer*4 function AllocHugeAlias16(ptr, size)
 integer*4 ptr, size

Description: The AllocHugeAlias16 function obtains a 16-bit far pointer to a 32-bit memory object that is *size* bytes in size. This is similar to the function AllocAlias16, except that AllocAlias16 will only give 16-bit far pointers to 32-bit memory objects of up to 64K in size. To get 16-bit far pointers to 32-bit memory objects larger than 64K, AllocHugeAlias16 should be used.

Returns: The AllocHugeAlias16 function returns a 16-bit far pointer (as an INTEGER*4) usable by Windows, or returns 0 if the alias cannot be allocated.

See Also: AllocAlias16, FreeAlias16, FreeHugeAlias16

Example: integer ierr, SIZE
 integer*4 alias
 parameter (SIZE=300000)
 integer*1 tmp(SIZE)

 allocate(tmp(SIZE), stat=ierr)
 if(ierr .ne. 0)then
 alias = AllocHugeAlias16(loc(tmp), SIZE)

 ! Windows calls using the alias ...

 call FreeHugeAlias16(alias, SIZE)
 endif

Classification: WIN386

Synopsis:

```
c$include 'winapi.fi'
integer*4 function _Call16( lpFunc, fmt, ... )
integer*4 lpFunc
character*(*) fmt
```

Description: The `_Call16` function performs the same function as `GetIndirectFunctionHandle`, `InvokeIndirectFunctionHandle` and `FreeIndirectFunctionHandle` but is much easier to use. The first argument *lpFunc* is the address of the 16-bit function to be called. This address is usually obtained by calling `GetProcAddress` with the name of the desired function. The second argument *fmt* is a string identifying the types of the parameters to be passed to the 16-bit function.

<i>Character</i>	<i>Parameter Type</i>
------------------	-----------------------

<i>c</i>	call a 'cdecl' function as opposed to a 'pascal' function (if specified, it must be listed first)
<i>b</i>	unsigned BYTE
<i>w</i>	16-bit WORD (INTEGER*2)
<i>d</i>	32-bit DWORD (INTEGER*4, REAL*4)
<i>f</i>	double precision floating-point (DOUBLE PRECISION, REAL*8)
<i>p</i>	32-bit flat pointer (converted to 16:16 far pointer) (LOC(x))

The 16-bit function must use either the PASCAL or CDECL calling convention. PASCAL calling convention is the default. If the function uses the CDECL calling convention, then you must specify the letter "c" as the first character of the argument type string.

Pointer types will automatically be converted from 32-bit near pointers to 16-bit far pointers before the function is invoked. Note that this pointer is only valid over the period of the call; after control returns to the 32-bit application, the 16-bit pointer created by the Supervisor is no longer valid.

Returns: The `_Call16` function returns a 32-bit DWORD (as an INTEGER*4) which represents the return value from the 16-bit function that was called.

See Also: `GetIndirectFunctionHandle`, `FreeIndirectFunctionHandle`

Example:

```
c$include winapi.fi

include 'windows.fi'
integer*2 hlib
integer*4 dll_1, cb
character*128 str

hlib = LoadLibrary( 'windllv.dll'c )
dll_1 = GetProcAddress( hlib, 'DLL1'c )

cb = _Call16( dll_1, 'ddd'c, 111, 22222, 3333 )
```

Classification: WIN386

Synopsis:

```
c$include 'winapi.fi'
      integer*4 function DefineDLEntry(index, routine, ...)
      integer*4 index
      external routine
```

Description: The DefineDLEntry function defines an *index* number for the 32-bit DLL procedure *routine*. The parameter *index* defines the index number that must be used in order to invoke the 32-bit FAR procedure *routine*. The variable argument list defines the types of parameters that will be received by the 32-bit DLL *routine*. Valid parameter types are:

DLL_PTR	16-bit far pointer
DLL_DWORD	32-bits
DLL_WORD	16-bits
DLL_CHAR	8-bits
DLL_ENDLIST	Marks the end of the variable argument list.

In the above synopsis, "..." in the argument list is not valid FORTRAN 77 syntax; it is used only to indicate a variable number of arguments.

Note that all pointers are received as 16-bit far pointers. To access the data from the 32-bit DLL, a dynamically allocatable array must be mapped to the memory pointed to by the 16-bit far pointer using the LOCATION specifier of the ALLOCATE statement and assigning the FAR attribute to the array using the array pragma.

Returns: The DefineDLEntry function returns zero if successful, and a non-zero value otherwise.

Example:

```
c$include winapi.fi

c$pragma aux (dll_function) DLL_1

      integer function DLL_1( w1, w2, w3 )
      integer*4 w1, w2, w3
      include 'win386.fi'
      include 'windefn.fi'
      include 'winerror.fi'
      character*128 str
      write( str, ' (16hDLL 1 arguments:, 3i10, a)' ) w1,
&                                                    w2, w3, char(0)
      call MessageBox( NULL, str,
&                      'DLL Function 1'c, MB_OK )
      DLL_1 = w1 + w2 + w3
      end

      integer*2 function FWINMAIN( hInstance,
&                                hPrevInstance,
&                                lpszCmdLine,
&                                nCmdShow )
      integer*2 hInstance, hPrevInstance, nCmdShow
      integer*4 lpszCmdLine
      include 'win386.fi'
      include 'windefn.fi'
      include 'winerror.fi'
      external DLL_1
      integer rc
      rc = DefineDLEntry( 1, DLL_1, DLL_DWORD,
```



```
&                                DLL_DWORD, DLL_DWORD,  
&                                DLL_ENDLIST )  
    if( rc .ne. 0 )then  
        FWinMain = 0  
        return  
    end if  
    call MessageBox( NULL,  
&                    '32-bit DLL started'c,  
&                    '32-bit DLL'c, MB_OK )  
    FWinMain = 1  
end
```

Classification: WIN386

Synopsis:

```
*$include 'winapi.fi'
      integer*4 function DefineUserProc16(typ, routine, ...)
      integer*4 typ
      external routine
```

Description: The DefineUserProc16 function defines the arguments accepted by the user defined callback procedure *routine*. There may be up to 32 user defined callbacks. The parameter *typ* indicates which one of GETPROC_USERDEFINED_1 through GETPROC_USERDEFINED_32 is being defined (see GetProc16). The callback routine must be declared as FAR PASCAL, or as FAR cdecl. The variable argument list defines the types of parameters that will be received by the user defined callback procedure *routine*. Valid parameter types are:

UDP16_PTR	16-bit far pointer
UDP16_DWORD	32-bits
UDP16_WORD	16-bits
UDP16_CHAR	8-bits
UDP16_CDECL	callback routine will be declared as type cdecl rather than as type PASCAL. This keyword may be placed anywhere before the UDP16_ENDLIST keyword.
UDP16_ENDLIST	Marks the end of the variable argument list.

Once the DefineUserProc16 function has been used to declare the user callback routine, then GetProc16 may be used to get a 16-bit function pointer that may be used by Windows.

In the above synopsis, "..." in the argument list is not valid FORTRAN 77 syntax; it is used only to indicate a variable number of arguments.

Returns: The DefineUserProc16 function returns zero if it succeeds; and non-zero if it fails.

See Also: GetProc16

Example:

```
c$include winapi.fi

c$pragma aux TestProc parm( value )

    subroutine TestProc( i )
    integer i
    character*128 str
    write( str, '(2hi=, i10, a)' ) i, char(0)
    call MessageBox( NULL, str, 'TEST'c, MB_OK )
    end

    integer function DefineTest()
    integer*4 cb
    external TestProc
    call DefineUserProc16( GETPROC_USERDEFINED_1,
&                          TestProc,
&                          UDP16_DWORD,
&                          UDP16_ENDLIST )
    cb = GetProc16( TestProc, GETPROC_USERDEFINED_1 )
    ! cb may then be used whenever a pointer to the
    ! callback is required by 16-bit Windows
    end
```

Classification: WIN386

Synopsis: c\$include 'winapi.fi'
 subroutine FreeAlias16(fp16)
 integer*4 fp16

Description: FreeAlias16 frees a 16-bit far pointer alias for a 32-bit near pointer that was allocated with AllocAlias16. This is important to do when there is no further use for the pointer since there are a limited number of 16-bit aliases available (due to limited space in the local descriptor table).

Returns: FreeAlias16 is a subroutine.

See Also: AllocAlias16

Example: integer*4 mcs_16
 record /MDICREATESTRUCT/ mcs
 mcs.szTitle = AllocAlias16(loc('Title'c))
 mcs.szClass = AllocAlias16(loc('mdichild'c))
 mcs.hOwner = hInst
 mcs.x = mcs.cx = CW_USEDEFAULT
 mcs.y = mcs.cy = CW_USEDEFAULT
 mcs.style = 0
 ! Send a message to an MDI client to create a window.
 ! Since the pointer to the structure is passed in an
 ! argument that may not be a pointer (depending on the
 ! type of message), there is no implicit 32 to 16-bit
 ! conversion done so the conversion must be done by the
 ! programmer.
 mcs_16 = AllocAlias16(loc(mcs))
 hwnd = SendMessage(hwndMDI, WM_MDICREATE, 0, mcs_16)
 FreeAlias16(mcs_16)
 FreeAlias16(mcs.szClass)
 FreeAlias16(mcs.szTitle)

Classification: WIN386

Synopsis: c\$include 'winapi.fi'
 subroutine FreeHugeAlias16(fp16, size)
 integer*4 fp16, size

Description: FreeHugeAlias16 frees a 16-bit far pointer alias that was allocated with AllocHugeAlias16. The size of the original 32-bit memory object must be specified. It is important to use FreeHugeAlias16 when there is no further use for the pointer, since there are a limited number of 16-bit aliases available (due to limited space in the local descriptor table).

Returns: FreeHugeAlias16 is a subroutine.

See Also: AllocHugeAlias16, AllocAlias16, FreeAlias16

Example: integer ierr, SIZE
 integer*4 alias
 parameter (SIZE=300000)
 integer*1 tmp(SIZE)

 allocate(tmp(SIZE), stat=ierr)
 if(ierr .ne. 0)then
 alias = AllocHugeAlias16(loc(tmp), SIZE)

 ! Windows calls using the alias ...

 call FreeHugeAlias16(alias, SIZE)
 endif

Classification: WIN386

FreeIndirectFunctionHandle

Synopsis: c\$include 'winapi.fi'
 subroutine FreeIndirectFunctionHandle(handle)
 integer*4 handle

Description: FreeIndirectFunctionHandle frees a handle that was obtained using GetIndirectFunctionHandle. This is important to do when there is no further use for the pointer since there are a limited number of 16-bit aliases available (due to limited space in the local descriptor table).

Returns: FreeIndirectFunctionHandle is a subroutine.

See Also: _Call16, GetIndirectFunctionHandle, InvokeIndirectFunction

Example: c\$include winapi.fi

```

integer*2 hDrv
integer*4 lpfn

hDrv = LoadLibrary( 'your.lib'c )
if( hDrv .lt. 32 ) return
lpfn = GetProcAddress( hDrv,
&                        'ExtDeviceMode'c )
if( lpfn .eq 0 ) return

hIndir = GetIndirectFunctionHandle(
&      lpfn,
&      INDIR_WORD,
&      INDIR_WORD,
&      INDIR_DWORD,
&      INDIR_PTR,
&      INDIR_PTR,
&      INDIR_DWORD,
&      INDIR_DWORD,
&      INDIR_WORD,
&      INDIR_ENDLIST )

cb = InvokeIndirectFunction(
&      hIndir,
&      hwnd,
&      hDrv,
&      NULL,
&      'POSTSCRIPT PRINTER'c,
&      'LPT1'c,
&      NULL,
&      NULL,
&      0 )
call FreeIndirectFunctionHandle( hIndir )
```

Classification: WIN386

Description: The GetIndirectFunctionHandle function gets a handle for a 16-bit procedure that is to be invoked indirectly. The procedure is assumed to have PASCAL calling convention, unless the `INDIR_CDECL` parameter is used, to indicate that Microsoft C calling convention is to be used. The 16-bit far pointer *prc* is supplied to GetIndirectFunctionHandle, and a list of the type of each parameter (in the order that they will be passed to the 16-bit function). The parameter types are:

In the above synopsis, "..." in the argument list is not valid FORTRAN 77 syntax; it is used only to indicate a variable number of arguments.

This handle is a data structure that was created using the `malloc` function. To free the handle, just use one of the `FreeIndirectFunctionHandle` or `free` functions.

You may find it easier to use `_Call16` rather than `GetIndirectFunctionHandle` followed by a call to `InvokeIndirectFunction`.

See Also: `_Call16`, `FreeIndirectFunctionHandle`, `InvokeIndirectFunction`

172 WIN386 Library Subprograms


```
hIndir = GetIndirectFunctionHandle(  
&      lpfn,  
&      INDIR_WORD,  
&      INDIR_WORD,  
&      INDIR_DWORD,  
&      INDIR_PTR,  
&      INDIR_PTR,  
&      INDIR_DWORD,  
&      INDIR_DWORD,  
&      INDIR_WORD,  
&      INDIR_ENDLIST )  
  
cb = InvokeIndirectFunction(  
&      hIndir,  
&      hwnd,  
&      hDrv,  
&      NULL,  
&      'POSTSCRIPT PRINTER'c,  
&      'LPT1'c,  
&      NULL,  
&      NULL,  
&      0 )  
call FreeIndirectFunctionHandle( hIndir )
```

Classification: WIN386

Synopsis:

```
c$include 'winapi.fi'
integer*4 function GetProcAddress( fcn, type )
integer*4 fcn, type
```

Description: The `GetProcAddress` function returns a 16-bit far function pointer suitable for use as a Windows callback function. This callback function will invoke the 32-bit far procedure specified by *fcn*. The types of callback functions that may be allocated are:

GETPROC_CALLBACK This is the most common form of callback; suitable as the callback routine for a window.

GETPROC_ABORTPROC This is the callback type used for trapping abort requests when printing.

GETPROC_ENUMCHILDWINDOWS This callback is used with the `EnumChildWindows` Windows function.

GETPROC_ENUMFONTS This callback type is used with the `EnumFonts` Windows function.

GETPROC_ENUMMETAFILE This callback is used with the `EnumMetaFile` Windows function.

GETPROC_ENUMOBJECTS This callback is used with the `EnumObjects` Windows function.

GETPROC_ENUMPROPS_FIXED_DS This callback is used with the `EnumProps` Windows function, when the fixed data segments callback is needed.

GETPROC_ENUMPROPS_MOVEABLE_DS This callback is used with the `EnumProps` Windows function, when the moveable data segments callback is needed.

GETPROC_ENUMTASKWINDOWS This callback is used with the `EnumTaskWindows` Windows function.

GETPROC_ENUMWINDOWS This callback is used with the `EnumWindows` Windows function.

GETPROC_GLOBALNOTIFY This callback is used with the `GlobalNotify` Windows function.

GETPROC_GRAYSTRING This callback is used with the `GrayString` Windows function.

GETPROC_LINEDDA This callback is used with the `LinedDA` Windows function.

GETPROC_SETRESOURCEHANDLER This callback is used with the `SetResourceHandler` Windows function.

GETPROC_SETTIMER This callback is used with the `SetTimer` Windows function.

GETPROC_SETWINDOWSHOOK This callback is used with the `SetWindowsHook` Windows function.

GETPROC_USERDEFINED_x This callback is used in conjunction with `DefineUserProc16` function to create a callback routine with an arbitrary set of parameters. Up to 32 user defined callbacks are allowed, they are identified by using `GETPROC_USERDEFINED_1` through `GETPROC_USERDEFINED_32`. The user defined callback must be declared as a FAR PASCAL function, or as a FAR cdecl function.

Returns: The GetProc16 function returns a 16-bit far pointer to a callback procedure. This pointer may then be fed to any Windows function that requires a pointer to a function within the 32-bit program. Note that the callback function within the 32-bit program must be declared as FAR.

See Also: ReleaseProc16

Example: `c$include winapi.fi`

```
integer*4 cbp
integer*4 lpProcAbout

! get a 16-bit callback routine to point at
! our About dialogue procedure, then create
! the dialogue.
cbp = GetProc16( About, GETPROC_CALLBACK )
lpProcAbout = MakeProcInstance( cbp, hInst )
call DialogBox( hInst,
&                'AboutBox'c,
                hWnd,
                lpProcAbout )
call FreeProcInstance( lpProcAbout )
call ReleaseProc16( cbp )
```

Classification: WIN386

InvokeIndirectFunction

Synopsis: c\$include 'winapi.fi'
 integer*4 function InvokeIndirectFunction(handle, ...)
 integer*4 handle

Description: The InvokeIndirectFunction function invokes the 16-bit function pointed to by the specified handle. The handle must have been previously allocated using the GetIndirectFunctionHandle function. The handle is followed by the list of parameters to be passed to the 16-bit function.

If you specified INDIR_PTR as a parameter when allocating the handle, then a 16-bit pointer is allocated for a 32-bit pointer that you pass. However, this pointer is freed when the 16-bit function being invoked returns.

In the above synopsis, "..." in the argument list is not valid FORTRAN 77 syntax; it is used only to indicate a variable number of arguments.

Returns: The InvokeIndirectFunction function returns the value which the 16-bit function returned.

See Also: _Call16, FreeIndirectFunctionHandle, GetIndirectFunctionHandle

Example: c\$include winapi.fi

```

      integer*2 hDrv
      integer*4 lpfn

      hDrv = LoadLibrary( 'your.lib'c )
      if( hDrv .lt. 32 ) return
      lpfn = GetProcAddress( hDrv,
&                               'ExtDeviceMode'c )
      if( lpfn .eq 0 ) return

      hIndir = GetIndirectFunctionHandle(
&          lpfn,
&          INDIR_WORD,
&          INDIR_WORD,
&          INDIR_DWORD,
&          INDIR_PTR,
&          INDIR_PTR,
&          INDIR_DWORD,
&          INDIR_DWORD,
&          INDIR_WORD,
&          INDIR_ENDLIST )

      cb = InvokeIndirectFunction(
&          hIndir,
&          hwnd,
&          hDrv,
&          NULL,
&          'POSTSCRIPT PRINTER'c,
&          'LPT1'c,
&          NULL,
&          NULL,
&          0 )
      call FreeIndirectFunctionHandle( hIndir )
```

Classification: WIN386

MapAliasToFlat

Synopsis: c\$include 'winapi.fi'
 integer*4 function MapAliasToFlat(alias)
 integer*4 alias

Description: The MapAliasToFlat function returns a 32-bit near pointer equivalent of a pointer allocated previously with AllocAlias16 or AllocHugeAlias16. This is useful if you are communicating with a 16-bit application that is returning pointers that you previously gave it.

Returns: The MapAliasToFlat function returns a 32-bit near pointer (as an INTEGER*4) usable by the 32-bit application.

See Also: AllocAlias16, AllocHugeAlias16

Example: c\$include winapi.fi

 integer alias
 integer ptr

 alias = AllocAlias16(loc(alias))
 alias += 5
 ptr = MapAliasToFlat(alias)
 if(ptr .eq. loc(alias) + 5)then
 call MessageBox(NULL, 'It Worked'c, ''c, MB_OK)
 else
 call MessageBox(NULL, 'It Failed'c, ''c, MB_OK)
 end if

Classification: WIN386

Synopsis: `c$include 'winapi.fi'`
 `integer*4 function PASS_WORD_AS_POINTER( dw )`
 `integer*4 dw`

Description: Some Windows API functions have pointer parameters that do not always take pointers. Sometimes these parameters are pure data. In order to stop the supervisor from trying to convert the data into a 16-bit far pointer, the `PASS_WORD_AS_POINTER` function is used.

Returns: The `PASS_WORD_AS_POINTER` returns a 32-bit "near" pointer, that is really the parameter *dw*.

Example: `c$include winapi.fi`

 `call Func( PASS_WORD_AS_POINTER(1) )`

Classification: WIN386

ReleaseProc16

Synopsis: c\$include 'winapi.fi'
 subroutine ReleaseProc16(cbp)
 integer*4 cbp

Description: ReleaseProc16 releases the callback function allocated by GetProc16. Since the callback routines are a limited resource, it is important to release the routines when they are no longer required.

Returns: ReleaseProc16 is a subroutine.

See Also: GetProc16

Example: c\$include winapi.fi

```

                integer*4 cbp
                integer*4 lpProcAbout
                ! get a 16-bit callback routine to point at
                ! our About dialogue procedure, then create
                ! the dialogue.
                cbp = GetProc16( About, GETPROC_CALLBACK )
                lpProcAbout = MakeProcInstance( cbp, hInst )
                call DialogBox( hInst, 'AboutBox'c,
&                    hWnd, lpProcAbout )
                call FreeProcInstance( lpProcAbout )
                call ReleaseProc16( cbp )
```

Classification: WIN386

24 32-bit Extended Windows Application Development

The purpose of this chapter is to anticipate some common questions about 32-bit Windows application development.

The following topics are discussed in this chapter:

- Can you call 16-bit code from a 32-bit code?
- How do I add my Windows resources?
- What size of function pointers passed to Windows?
- Why are 32-bit callback routines FAR?
- Why use the `_16` API functions?

24.1 Can you call 16-bit code from a 32-bit code?

A 32-bit Windows application can make a call to 16-bit code through the use of the Open Watcom `_Call16` or `InvokeIndirectFunction` procedures. These functions ensure that the Open Watcom Windows Supervisor prepares the stack for the 16-bit call and return to the 32-bit code. The 32-bit application uses `LoadLibrary` function to bring the 16-bit DLL into memory and then calls the 16-bit procedures. To invoke 16-bit procedures, use `GetProcAddress` to get the 16-bit far pointer to the function. Use the `_Call16` procedure to call the 16-bit function since it is simpler to use than the `GetIndirectFunctionHandle`, `InvokeIndirectFunction`, and `FreeIndirectFunctionHandle` sequence. An example of this process is provided under the `_Call16` Windows library function description.

This method can be used to call any 16-bit Dynamic Link Library (DLL) procedure or any 32-bit extended DLL procedure from within a 32-bit application, including DLLs that are available as products through Independent Software Vendors (ISVs).

24.2 How do I add my Windows resources?

The `WBIND` utility automatically runs the resource compiler to add the resources to the 32-bit Windows supervisor (since the supervisor is a 16-bit Windows application). Note that resource compiler options may be specified by using the "R" option of `WBIND`.

24.3 What size of function pointers passed to Windows?

All function pointers passed to Windows must be 16-bit far pointers since no translation is applied to any function pointers passed to Windows. Translation is often not possible, since any functions that Windows is to call back must be exported, and only 16-bit functions can be exported.

A 16-bit far pointer to a function is obtained in one of two ways: either Windows gives it to you (via `GetProcAddress`, for example), or you obtain a pointer from the supervisor, via `GetProcAddress`.

Function pointers obtained from Windows may either be fed into other Windows functions requiring function pointers, or called indirectly by using `_Call16` or by using the `GetIndirectFunctionHandle`, `InvokeIndirectFunction`, and `FreeIndirectFunctionHandle` sequence.

The function `GetProcAddress` returns a 16-bit far pointer to a callback function that Windows can use. This callback function will direct control into the desired 32-bit routine.

24.4 Why are 32-bit callback routines FAR?

The callback routines are declared as FAR so that the compiler will generate a far return from the procedure. This is necessary since the 32-bit callback routine is "far" called from the supervisor.

The callback routine is still "near" in the sense that it lies within the 32-bit flat address space of the application. This means that `GetProcAddress` only needs the offset of the 32-bit callback function in order to set up the 16-bit procedure to call back correctly. Thus, `GetProcAddress` accepts type `PROC_PTR` which is in fact only 4 bytes long. The compiler will provide the offset only, which is, as already stated, all that is needed.

24.5 Why use the _16 API functions?

The regular Windows API functions used in Open Watcom F77 automatically convert any pointers to 16-bit far pointers for use by Windows. Sometimes, you may have a set of pointers that are 16-bit far pointers already (e.g., obtained from `GlobalLock`), and do not need any conversion. The "_16..." API functions do not convert pointers, they simply pass them on directly to Windows. See the appendix entitled "Special Windows API Functions" on page 183 for a list of the "_16..." API functions.

25 *Special Windows API Functions*

On rare occasions, you want to use 16-bit far pointers directly in a Windows function. Since all Windows functions in the 32-bit environment are expecting 32-bit near pointers, you cannot simply use the 16-bit far pointer directly in the function.

The following functions are special versions of Windows API functions that do NOT convert any of the pointers from 32-bit to 16-bit. There are `_16` versions of all Windows API functions that accept data pointers.

- `_16AddAtom`
- `_16AddFontResource`
- `_16AdjustWindowRect`
- `_16AdjustWindowRectEx`
- `_16AnimatePalette`
- `_16AnsiLower`
- `_16AnsiLowerBuff`
- `_16AnsiToOem`
- `_16AnsiToOemBuff`
- `_16AnsiUpper`
- `_16AnsiUpperBuff`
- `_16BuildCommDCB`
- `_16CallMsgFilter`
- `_16ChangeMenu`
- `_16ClientToScreen`
- `_16ClipCursor`
- `_16CopyMetaFile`
- `_16CopyRect`
- `_16CreateBitmap`
- `_16CreateBitmapIndirect`
- `_16CreateBrushIndirect`
- `_16CreateCursor`
- `_16CreateDC`
- `_16CreateDialog`
- `_16CreateDialogIndirect`
- `_16CreateDialogIndirectParam`
- `_16CreateDialogParam`
- `_16CreateDIBitmap`
- `_16CreateEllipticRgnIndirect`
- `_16CreateFont`
- `_16CreateFontIndirect`
- `_16CreateIC`
- `_16CreateIcon`
- `_16CreateMetaFile`
- `_16CreatePalette`
- `_16CreatePenIndirect`
- `_16CreatePolygonRgn`
- `_16CreatePolyPolygonRgn`
- `_16CreateRectRgnIndirect`

- _16CreateWindow
- _16CreateWindowEx
- _16DialogBox
- _16DialogBoxIndirect
- _16DialogBoxIndirectParam
- _16DialogBoxParam
- _16DispatchMessage
- _16DlgDirList
- _16DlgDirListComboBox
- _16DlgDirSelect
- _16DlgDirSelectComboBox
- _16DPtoLP
- _16DrawFocusRect
- _16DrawText
- _16EndPaint
- _16EnumChildWindows
- _16EnumFonts
- _16EnumMetaFile
- _16EnumObjects
- _16EnumProps
- _16EnumTaskWindows
- _16EnumWindows
- _16EqualRect
- _16Escape
- _16ExtTextOut
- _16FillRect
- _16FindAtom
- _16FindResource
- _16FindWindow
- _16FrameRect
- _16FreeProcInstance
- _16GetAtomName
- _16GetBitmapBits
- _16GetCaretPos
- _16GetCharWidth
- _16GetClassInfo
- _16GetClassName
- _16GetClientRect
- _16GetClipboardFormatName
- _16GetClipBox
- _16GetCodeInfo
- _16GetCommError
- _16GetCommState
- _16GetCursorPos
- _16GetDIBits
- _16GetDlgItemInt
- _16GetDlgItemText
- _16GetEnvironment
- _16GetKeyboardState
- _16GetKeyNameText
- _16GetMenuString
- _16GetMetaFile
- _16GetModuleFileName
- _16GetModuleHandle

- _16GetObject
- _16GetPaletteEntries
- _16GetPriorityClipboardFormat
- _16GetPrivateProfileInt
- _16GetPrivateProfileString
- _16GetProcAddress
- _16GetProfileInt
- _16GetProfileString
- _16GetProp
- _16GetRgnBox
- _16GetScrollRange
- _16GetSystemDirectory
- _16GetSystemPaletteEntries
- _16GetTabbedTextExtent
- _16GetTempFileName
- _16GetTextExtent
- _16GetTextFace
- _16GetTextMetrics
- _16GetUpdateRect
- _16GetWindowRect
- _16GetWindowsDirectory
- _16GetWindowText
- _16GlobalAddAtom
- _16GlobalFindAtom
- _16GlobalGetAtomName
- _16GlobalNotify
- _16GrayString
- _16InflateRect
- _16IntersectRect
- _16InvalidateRect
- _16InvertRect
- _16IsDialogMessage
- _16IsRectEmpty
- _16LineDDA
- _16LoadAccelerators
- _16LoadBitmap
- _16LoadCursor
- _16LoadIcon
- _16LoadLibrary
- _16LoadMenu
- _16LoadMenuIndirect
- _16LoadModule
- _16LoadString
- _16LPtoDP
- _16MakeProcInstance
- _16MapDialogRect
- _16MessageBox
- _16OemToAnsi
- _16OemToAnsiBuff
- _16OffsetRect
- _16OpenComm
- _16OpenFile
- _16OutputDebugString
- _16PlayMetaFileRecord

- _16Polygon
- _16Polyline
- _16PolyPolygon
- _16PtInRect
- _16ReadComm
- _16RectInRegion
- _16RectVisible
- _16RegisterClipboardFormat
- _16RegisterWindowMessage
- _16RemoveFontResource
- _16RemoveProp
- _16ScreenToClient
- _16ScrollDC
- _16ScrollWindow
- _16SetBitmapBits
- _16SetCommState
- _16SetDIBits
- _16SetDIBitsToDevice
- _16SetDlgItemText
- _16SetEnvironment
- _16SetKeyboardState
- _16SetPaletteEntries
- _16SetProp
- _16SetRect
- _16SetRectEmpty
- _16SetResourceHandler
- _16SetSysColors
- _16SetTimer
- _16SetWindowsHook
- _16SetWindowText
- _16StretchDIBits
- _16TabbedTextOut
- _16TextOut
- _16ToAscii
- _16TrackPopupMenu
- _16TranslateAccelerator
- _16TranslateMDISysAccel
- _16TranslateMessage
- _16UnhookWindowsHook
- _16UnionRect
- _16UnregisterClass
- _16ValidateRect
- _16WinExec
- _16WinHelp
- _16WriteComm
- _16WritePrivateProfileString
- _16WriteProfileString
- _16_lread
- _16_lwrite

Windows NT Programming Guide

26 Windows NT Programming Overview

Windows NT supports both non-windowed character-mode applications and windowed Graphical User Interface (GUI) applications. In addition, Windows NT supports Dynamic Link Libraries and applications with multiple threads of execution.

We have supplied all the necessary tools for native development on Windows NT. You can also cross develop for Windows NT using either the DOS-hosted compilers and tools, the Windows 95-hosted compilers and tools, or the OS/2-hosted compilers and tools. Testing and debugging of your Windows NT application must be done on Windows NT or Windows 95.

If you are creating a character-mode application, you may also be interested in a special DOS extender from Phar Lap (TNT) that can run your Windows NT character-mode application under DOS.

26.1 Windows NT Character-mode Versus GUI

Basically, there are two classes of FORTRAN 77 applications that can run in a windowed environment like Windows NT.

The first are those FORTRAN 77 applications that do not use any of the Win32 API functions; they are strictly FORTRAN 77 applications that do not rely on the features of a particular operating system.

- This Application must be created as Windows NT Character-mode Application.

The second class of FORTRAN 77 applications are those that actually call Win32 API functions directly. These are applications that have been tailored for the Win32 operating environment.

- Open Watcom FORTRAN 77 does not provide direct support for these types of applications. While we do provide include files that map out 16-bit Windows structures and the interface to 16-bit Windows API calls, we do not provide this for Win32 API. The Win32 application developer must create these as required.
- An alternate solution, for those so-inclined, is to develop the GUI part of the interface in C and call these functions from FORTRAN code.

A subsequent chapters deal with the creation of different application types for Windows NT target.

27 Creating Windows NT Character-mode Applications

This chapter describes how to compile and link Windows NT Character-mode applications simply and quickly. In this chapter, we look at applications written to exploit the Windows NT Application Programming Interface (API).

We will illustrate the steps to creating Windows NT Character-mode applications by taking a small sample application and showing you how to compile, link, run and debug it.

27.1 The Sample Character-mode Application

To demonstrate the creation of Windows NT Character-mode applications, we introduce a simple sample program. For our example, we are going to use the "sieve" program.

```
* This program computes the prime numbers between 1 and 10,000
* using the Sieve of Eratosthenes algorithm.

      IMPLICIT NONE
      INTEGER UPBOUND
      PARAMETER (UPBOUND=10000)
      INTEGER I, K, PRIMES
      LOGICAL*1 NUMBERS(2:UPBOUND)
      CHARACTER*11 FORM
      CHARACTER*60 BUFFER
      PARAMETER (FORM='(A,I5,A,I5)')
      DO I = 2, UPBOUND
         NUMBERS(I) = .TRUE.
      ENDDO
      PRIMES = 0
      DO I = 2, UPBOUND
         IF( NUMBERS(I) )THEN
            PRIMES = PRIMES + 1
            DO K = I + I, UPBOUND, I
               NUMBERS(K) = .FALSE.
            ENDDO
         ENDIF
      ENDDO
      PRINT FORM, 'The Number of Primes between 1 and ', UPBOUND,
1      ' are: ', PRIMES
      END
```

The goal of this program is to count the prime numbers between 1 and 10,000. It uses the famous *Sieve of Eratosthenes* algorithm to accomplish this task. We will take you through the steps necessary to produce this result.

27.2 Building and Running the Character-mode Application

To compile and link our example program which is stored in the file `sieve.for`, enter the following command:

```
C>wfl386 -l=nt sieve.for
```

The typical messages that appear on the screen are shown in the following illustration.

```
C>wfl386 -l=nt sieve.for
Open Watcom F77/32 Compile and Link Utility
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1990-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
wfc386 sieve.for
Open Watcom FORTRAN 77/32 Optimizing Compiler
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1984-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
sieve.for: 4390 statements, 207 bytes, 1585 extensions, 0 warnings, 0 errors

Open Watcom Linker
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1985-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
loading object files
searching libraries
creating a Windows NT Character-mode executable
```

If you examine the current directory, you will find that two files have been created. These are `sieve.obj` (the result of compiling `sieve.for`) and `sieve.exe` (the result of linking `sieve.obj` with the appropriate Open Watcom FORTRAN 77 libraries).

The resultant Windows NT Character-mode application `SIEVE.EXE` can now be run under Windows NT.

27.3 Debugging the Character-mode Application

Let us assume that you wish to debug your application in order to locate an error in programming. In the previous section, the "sieve" program was compiled with default compile and link options. When debugging an application, it is useful to refer to the symbolic names of routines and variables. It is also convenient to debug at the source line level rather than the machine language level. To do this, we must direct both the compiler and linker to include additional debugging information in the object and executable files. Using the `WFL386` command, this is fairly straightforward. `WFL386` recognizes the Open Watcom F77 compiler "debug" options and will create the appropriate debug directives for the Open Watcom Linker.

For example, to compile and link the "sieve" program with debugging information, the following command may be issued.

```
C>wfl386 -l=nt -d2 sieve.for
```

The typical messages that appear on the screen are shown in the following illustration.

```
C>wfl386 -l=nt -d2 sieve.for
Open Watcom F77/32 Compile and Link Utility
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1990-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
    wfc386 sieve.for -d2
Open Watcom FORTRAN 77/32 Optimizing Compiler
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1984-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
sieve.for: 4390 statements, 293 bytes, 1585 extensions, 0 warnings, 0 errors

Open Watcom Linker
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1985-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
loading object files
searching libraries
creating a Windows NT Character-mode executable
```

The "d2" option requests the maximum amount of debugging information that can be provided by the Open Watcom F77 compiler. WFL386 will make sure that this debugging information is included in the executable file that is produced by the linker.

The "bytes" value is larger than in the previous example since selection of the "d2" option results in fewer code optimizations by default. You can request more optimization by specifying the appropriate options. However, you do so at the risk of making it more difficult for yourself to determine the relationship between the object code and the original source language code.

To request the Open Watcom Debugger to assist in debugging the application, select the Open Watcom Debugger icon. It would be too ambitious to describe the debugger in this introductory chapter so we refer you to the book entitled *Open Watcom Debugger User's Guide*.

There are more extensive examples of Windows applications written in FORTRAN 77 in the \WATCOM\SAMPLES\FORTRAN\WIN directory. The example programs are ELLIPSE.FOR and FWCOPY.FOR.

28 Windows NT Multi-threaded Applications

This chapter describes how to create multi-threaded applications. A multi-threaded application is one whose tasks are divided among several threads of execution. A process is an executing application and the resources it uses. A thread is the smallest unit of execution within a process. Each thread has its own stack and a set of machine registers and shares all resources with its parent process. The path of execution of one thread does not affect that of another; each thread is an independent entity.

Typically, an application has a single thread of execution. In this type of application, all tasks, once initiated, are completed before the next task begins. In contrast, tasks in a multi-threaded application can be performed concurrently since more than one thread is executing at once. For example, each thread may be designed to perform a separate task.

28.1 Programming Considerations

Since a multi-threaded application consists of many threads of execution, there are a number of issues that you must consider.

Since threads share the resources of its parent, it may be necessary to serialize access to these resources. For example, if your application contains more than one thread of execution and each thread uses the `PRINT` statement to display output to the console, it would be necessary for the I/O support routines to allow only one thread to use the `PRINT` facility at any time. That is, once a thread issues a `PRINT` request, the I/O support routines should ensure that no other thread displays information until all information for the initial thread has been displayed.

28.2 Creating Threads

Each application initially contains a single thread. The run-time libraries contain two functions that create and terminate threads of execution. The function `beginthread` creates a thread of execution and the function `endthread` ends a thread of execution. The function `threadid` can be used to determine the current thread identifier.

WARNING! If any thread uses an I/O statement or calls an intrinsic function, you must use the `beginthread` function to create the thread. Do not use the `CreateThread` API function.

28.2.1 Creating a New Thread

The `beginthread` function creates a new thread. It is defined as follows.

```
integer function beginthread( start_address,
                             stack_size )
integer stack_size
end
```

<i>where</i>	<i>description</i>
--------------	--------------------

<i>start_address</i>	is the address of the subroutine that will be called when the newly created thread is executed. When the thread returns from that subroutine, the thread will be terminated. Note that a call to the <code>endthread</code> subroutine will also terminate the thread.
----------------------	--

<i>stack_size</i>	specifies the size of the stack to be allocated by the operating system for the new thread. The stack size should be a multiple of 4K.
-------------------	--

If a new thread is successfully created, the thread identifier of the new thread is returned. Otherwise, a value of -1 is returned.

The include file `thread.fi` contains the definition of the `beginthread` function.

Another thread related function for Windows NT is `_beginthreadex`. See the *Open Watcom C Library Reference* for more information.

28.2.2 Terminating the Current Thread

The `endthread` subroutine terminates the current thread. It is defined as follows.

```
subroutine endthread()  
end
```

The include file `thread.fi` contains the definition of the `endthread` function.

28.2.3 Getting the Current Thread Identifier

The `threadid` function can be used to determine the current thread identifier. It is defined as follows.

```
integer function threadid()  
end
```

The include file `thread.fi` contains the definition of the `threadid` function.

28.3 A Multi-threaded Example

Let us create a simple multi-threaded application. The source code for this example can be found in `\watcom\samples\fortran\win32`.


```
* MTHREAD.FOR

*$pragma aux (__stdcall) Sleep parm( value )
*$pragma aux (__stdcall) InitializeCriticalSection parm( reference )
*$pragma aux (__stdcall) DeleteCriticalSection parm( reference )
*$pragma aux (__stdcall) EnterCriticalSection parm( reference )
*$pragma aux (__stdcall) LeaveCriticalSection parm( reference )

        structure /RTL_CRITICAL_SECTION/
            integer*4 DebugInfo
            integer*4 LockCount
            integer*4 RecursionCount
            integer*4 OwningThread
            integer*4 LockSemaphore
            integer*4 Reserved
        end structure

        integer NumThreads
        logical HoldThreads
        volatile HoldThreads, NumThreads
        record /RTL_CRITICAL_SECTION/ CriticalSection
        common NumThreads, HoldThreads, CriticalSection

        integer STACK_SIZE
        parameter (STACK_SIZE=8192)
        integer NUM_THREADS
        parameter (NUM_THREADS=5)

        integer i, threadid, beginthread
        external a_thread

        print '(''main thread id = '',i4)', threadid()
        NumThreads = 0
        HoldThreads = .true.
        ! main thread counts as 1
        call InitializeCriticalSection( CriticalSection )
        do i = 2, NUM_THREADS
            if( beginthread( a_thread, STACK_SIZE ) .eq. -1 )then
                print '(''creation of thread'',i4,'failed'')', i
            else
                NumThreads = NumThreads + 1
            end if
        end do
        HoldThreads = .false.
        while( NumThreads .ne. 0 ) do
            call Sleep( 1 )
        end while
        call DeleteCriticalSection( CriticalSection )
        end

        subroutine a_thread()

        structure /RTL_CRITICAL_SECTION/
            integer*4 DebugInfo
            integer*4 LockCount
            integer*4 RecursionCount
            integer*4 OwningThread
            integer*4 LockSemaphore
            integer*4 Reserved
        end structure

        integer NumThreads
        logical HoldThreads
        volatile HoldThreads
        record /RTL_CRITICAL_SECTION/ CriticalSection
        common NumThreads, HoldThreads, CriticalSection

        integer threadid
```

```
while( HoldThreads ) do
    call Sleep( 1 )
end while
print '('Hi from thread ', i4)', threadid()
call EnterCriticalSection( CriticalSection )
NumThreads = NumThreads - 1
call LeaveCriticalSection( CriticalSection )
call endthread()
end
```

Note:

1. In the subroutine `a_thread`, `EnterCriticalSection` and `LeaveCriticalSection` are called when we modify the variable `NumThreads`. This ensures that the action of extracting the value of `NumThreads` from memory, incrementing the value, and storing the new result into memory, occurs without interruption. If these functions were not called, it would be possible for two threads to extract the value of `NumThreads` from memory before an update occurred.

Let us assume that the file `mthread.for` contains the above example. Before compiling the file, make sure that the **WATCOM** environment variable is set to the directory in which you installed Open Watcom FORTRAN 77. Also, the **FINCLUDE** environment variable must contain the `\watcom\src\fortran` directory where "**WATCOM**" is the name of the directory in which you installed Open Watcom FORTRAN 77.

We can now compile and link the application by issuing the following command.

```
C:\>wfl386 -bm -l=nt mthread
```

The "`bm`" option must be specified since we are creating a multi-threaded application. If your multi-threaded application contains more than one module, each module must be compiled using the "`bm`" switch.

The "`l`" option specifies the target system for which the application is to be linked. The system name `nt` is defined in the file `wlssystem.lnk` which is located in the "`BINW`" subdirectory of the directory in which you installed Open Watcom FORTRAN 77.

The multi-threaded application is now ready to be run.

29 Windows NT Dynamic Link Libraries

A dynamic link library, like a standard library, is a library of functions. When an application uses functions from a standard library, the library functions referenced by the application become part of the executable module. This form of linking is called static linking. When an application uses functions from a dynamic link library, the library functions referenced by the application are not included in the executable module. Instead, the executable module contains references to these functions which are resolved when the application is loaded. This form of linking is called dynamic linking.

Let us consider some of the advantages of using dynamic link libraries over standard libraries.

1. Functions in dynamic link libraries are not linked into your program. Only references to the functions in dynamic link libraries are placed in the program module. These references are called import definitions. As a result, the linking time is reduced and disk space is saved. If many applications reference the same dynamic link library, the saving in disk space can be significant.
2. Since program modules only reference dynamic link libraries and do not contain the actual executable code, a dynamic link library can be updated without re-linking your application. When your application is executed, it will use the updated version of the dynamic link library.
3. Dynamic link libraries also allow sharing of code and data between the applications that use them. If many applications that use the same dynamic link library are executing concurrently, the sharing of code and data segments improves memory utilization.

29.1 Creating Dynamic Link Libraries

Once you have developed the source for a library of functions, a number of steps are required to create a dynamic link library containing those functions.

First, you must compile your source using the "bd" compiler option. This option tells the compiler that the module you are compiling is part of a dynamic link library. Once you have successfully compiled your source, you must create a linker directive file that describes the attributes of your dynamic link library. The following lists the most common linker directives required to create a dynamic link library.

1. The "SYSTEM" directive is used to specify that a dynamic link library is to be created.
2. The "EXPORT" directive is used to specify which functions in the dynamic link library are to be exported.
3. The "OPTION" directive is used to specify attributes such as the name of the dynamic link library and how to allocate the automatic data segment when the dynamic link library is referenced.
4. The "SEGMENT" directive is used to specify attributes of segments. For example, a segment may be read-only or read-write.

Once the dynamic link library is created, you must allow access to the dynamic link library to client applications that wish to use it. This can be done by creating an import library for the dynamic link library

or creating a linker directive file that contains "IMPORT" directives for each of the entry points in the dynamic link library.

29.2 Creating a Sample Dynamic Link Library

Let us now create a dynamic link library using the following example. The source code for this example can be found in `\watcom\samples\fortran\win\dll`. Unlike applications developed in the C or C++ language, the FORTRAN 77 developer must not provide a `LibMain` entry point. This entry point is already defined in the Open Watcom FORTRAN 77 run-time libraries. The run-time system's `LibMain` provides for the proper initialization of the FORTRAN 77 run-time system and includes hooks to call developer-written process attach/detach and thread attach/detach routines. These routines are optional but we show skeleton versions in the following example so that you can develop your own if required.

```
* DLLSAMP.FOR

      integer function __fdll_initialize_()
*      Called from LibMain during "DLL PROCESS ATTACH"

*      do process initialization
      print *, 'Hi from process attach'

*      returning 0 indicates failure
      __fdll_initialize_ = 1
      return
      end

      integer function __fthrd_initialize_()
*      Called from LibMain during "DLL THREAD ATTACH"

*      do thread initialization
      print *, 'Hi from thread attach'

*      returning 0 indicates failure
      __fthrd_initialize_ = 1
      return
      end

      integer function __fthrd_terminate_()
*      Called from LibMain during "DLL THREAD DETACH"

*      do thread cleanup
      print *, 'Hi from thread detach'

*      returning 0 indicates failure
      __fthrd_terminate_ = 1
      return
      end

      integer function __fdll_terminate_()
*      Called from LibMain during "DLL PROCESS DETACH"

*      do process cleanup
      print *, 'Hi from process detach'

*      returning 0 indicates failure
      __fdll_terminate_ = 1
      return
      end

      subroutine dll_entry_1()
      print *, 'Hi from dll entry #1'
      end

      subroutine dll_entry_2()
      print *, 'Hi from dll entry #2'
      end
```

Here are some explanatory notes on this example.

Function	Description
-----------------	--------------------

__FDLL_INITIALIZE_ This function is called when the DLL is attaching to the address space of the current process as a result of the process starting up or as a result of a call to `LoadLibrary`. A DLL can use this opportunity to initialize any instance data.

During initial process startup or after a call to `LoadLibrary`, the operating system scans the list of loaded DLLs for the process. For each DLL that has not already been called with the `DLL_PROCESS_ATTACH` value, the system calls the DLL's `LibMain` entry-point (in the Open Watcom FORTRAN 77 run-time system). This call is made in the context of the thread that caused the process address space to change, such as the primary thread of the process or the thread that called `LoadLibrary`.

__FTHRD_INITIALIZE_ This function is called when the current process is creating a new thread. When this occurs, the system calls the `LibMain` entry-point (in the Open Watcom FORTRAN 77 run-time system) of all DLLs currently attached to the process. The call is made in the context of the new thread. DLLs can use this opportunity to initialize thread specific data. A thread calling the DLL's `LibMain` with the `DLL_PROCESS_ATTACH` value does not call `LibMain` with the `DLL_THREAD_ATTACH` value. Note that `LibMain` is called with this value only by threads created after the DLL is attached to the process. When a DLL is attached by `LoadLibrary`, existing threads do not call the `LibMain` entry-point of the newly loaded DLL.

__FTHRD_TERMINATE_ This function is called when a thread is exiting normally. The DLL uses this opportunity to do thread specific termination such as closing files that were opened by the thread. The operating system calls the `LibMain` entry-point (in the Open Watcom FORTRAN 77 run-time system) of all currently loaded DLLs with this value. The call is made in the context of the exiting thread. There are cases in which `LibMain` is called for a terminating thread even if the DLL never attached to the thread. For example, `LibMain` is never called with the `DLL_THREAD_ATTACH` value in the context of the thread in either of these two situations:

- The thread was the initial thread in the process, so the system called `LibMain` with the `DLL_PROCESS_ATTACH` value.
- The thread was already running when a call to the `LoadLibrary` function was made, so the system never called `LibMain` for it.

__FDLL_TERMINATE_ This function is called when the DLL is detaching from the address space of the calling process as a result of either a normal termination or of a call to `FreeLibrary`. When a DLL detaches from a process as a result of process termination or as a result of a call to `FreeLibrary`, the operating system does not call the DLL's `LibMain` with the `DLL_THREAD_DETACH` value for the individual threads of the process. The DLL is only given `DLL_PROCESS_DETACH` notification. DLLs can take this opportunity to clean up all resources for all threads attached and known to the DLL.

Note: These functions return 1 if initialization succeeds or 0 if initialization fails. Subsequently, this value will be returned by the run-time system's `LibMain` function.

If the return value is 0 when `LibMain` is called because the process uses the `LoadLibrary` function, `LoadLibrary` returns NULL.

If the return value is 0 when `LibMain` is called during process initialization, the process terminates with an error.

DLL_ENTRY_1, DLL_ENTRY_2 These are sample DLL entry points that we will call from our simple test program.

Some further explanation and an example are provided in a later section.

Assume the above example is contained in the file `dllsamp.for`. We can compile the file using the following command. Note that we must specify the "bd" compiler option.

```
C:\>wfc386 -bd dllsamp
```

Before we can link our example, we must create a linker directive file that describes the attributes and entry points of our dynamic link library. The following is a linker directive file, called `dllsamp.lnk`, that can be used to create the dynamic link library.

```
system nt_dll initinstance terminstance
export DLL_ENTRY_1
export DLL_ENTRY_2
file dllsamp
```

Notes:

1. The "SYSTEM" directive specifies that we are creating a Windows NT dynamic link library.
2. When a dynamic link library uses the Open Watcom FORTRAN 77 run-time libraries, an automatic data segment is created each time a new process accesses the dynamic link library. For this reason, initialization code must be executed when a process accesses the dynamic link library for the first time. To achieve this, "INITINSTANCE" must be specified in the "SYSTEM" directive. Similarly, "TERMINSTANCE" must be specified so that the termination code is executed when a process has completed its access to the dynamic link library. If the Open Watcom FORTRAN 77 run-time libraries are not used, these options are not required.
3. The "EXPORT" directive specifies the entry points into the dynamic link library. Note that in Open Watcom FORTRAN 77, names of all symbols are uppercased. Regardless of the case used in source files, linker directives must use uppercased symbol names. The linker is case sensitive by default, although the "OP NOCASEEXACT" directive may be used to override this.

We can now create our dynamic link library by issuing the following command.

```
C:\>wlink @dllsamp
```

A file called `dllsamp.dll` will be created.

29.3 Using Dynamic Link Libraries

Once we have created a dynamic link library, we must allow other applications to access the functions available in the dynamic link library. There are two ways to achieve this.

The first method is to create a linker directive file which contains an "IMPORT" directive for all entry points in the dynamic link library. The "IMPORT" directive provides the name of the entry point and the name of the dynamic link library. When creating an application that references a function in the dynamic

link library, this linker directive file would be included as part of the linking process that created the application.

The second method is to use import libraries. An import library is a standard library that is created from a dynamic link library by using the Open Watcom Library Manager. It contains object modules that describe the entry points in a dynamic link library. The resulting import library can then be specified in a "LIBRARY" directive in the same way one would specify a standard library.

Using an import library is the preferred method of providing references to functions in dynamic link libraries. When a dynamic link library is modified, typically the import library corresponding to the modified dynamic link library is updated to reflect the changes. Hence, any directive file that specifies the import library in a "LIBRARY" directive need not be modified. However, if you are using "IMPORT" directives, you may have to modify the "IMPORT" directives to reflect the changes in the dynamic link library.

Let us create an import library for our sample dynamic link library we created in the previous section. We do this by issuing the following command.

```
C:\>wlib dllsamp +dllsamp.dll
```

A standard library called `dllsamp.lib` will be created.

Suppose the following sample program, contained in the file `dlltest.for`, calls the functions from our sample dynamic link library.

```
* DLLTEST.FOR
      call dll_entry_1()
      call dll_entry_2()
      end
```

We can compile and link our sample application by issuing the following command.

```
C:\>wfl386 -l=nt dlltest dllsamp.lib
```

If we had created a linker directive file of "IMPORT" directives instead of an import library for the dynamic link library, the linker directive file, say `dllimps.lnk`, would be as follows.

```
import DLL_ENTRY_1 dllsamp
import DLL_ENTRY_2 dllsamp
```

To compile and link our sample application, we would issue the following command.

```
C:\>wfl386 -l=nt dlltest -"@dllimps"
```

29.4 The Dynamic Link Library Data Area

The Open Watcom FORTRAN 77 32-bit run-time library does not support the general case operation of DLLs in an execution environment where there is only one instance of the DATA segment (DGROUP) for that DLL.

There are two cases that can lead to a DLL executing with only one instance of the DGROUP.

1. DLLs linked for 32-bit OS/2 without the MANYAUTODATA option.

2. DLLs linked for the Win32 API and executing under Win32s.

In these cases the run-time library startup code detects that there is only one instance of the DGROUP when a second process attempts to attach to the DLL. At that point, it issues a diagnostic for the user and then notifies the operating system that the second process cannot attach to the DLL.

Developers who require DLLs to operate when there is only one instance of the DGROUP can suppress the function which issues the diagnostic and notifies the operating system that the second process cannot attach to the DLL.

Doing so requires good behaviour on the part of processes attaching to the DLL. This good behaviour consists primarily of ensuring that the first process to attach to the DLL is also the last process to detach from the DLL thereby ensuring that the DATA segment is not released back to the free memory pool.

To suppress the function which issues the diagnostic and notifies the operating system that the second process cannot attach to the DLL, the developer must provide a replacement entry point with the following prototype:

```
int __disallow_single_dgroup( int );
```

This function should return zero to indicate that the detected single copy of the DATA segment is allowed.

29.5 Dynamic Link Library Initialization/Termination

Each dynamic link library (DLL) has an initialization and termination routine associated with it. The initialization routine can either be called the first time any process accesses the DLL ("INITGLOBAL" is specified at link time) or each time a process accesses the DLL ("INITINSTANCE" is specified at link time). Similarly, the termination routine can either be called when all processes have completed their access of the DLL ("TERMGLOBAL" is specified at link time) or each time a process completes its access of the DLL ("TERMINSTANCE" is specified at link time).

For a DLL that uses the FORTRAN 77 run-time libraries, initialization and termination of the FORTRAN 77 run-time environment is performed automatically. It is also possible for a DLL to do its own special initialization and termination process.

The FORTRAN 77 run-time environment provides a method for calling user-written DLL initialization and termination code. The `__fdll_initialize_` routine is called for DLL process initialization. The `__fthrd_initialize_` routine is called for DLL thread initialization. The `__fthrd_terminate_` routine is called for DLL thread termination. The `__fdll_terminate_` routine is called for DLL process termination. Default stub versions of these routines are included in the run-time library. If you wish to perform additional initialization/termination processing that is specific to your dynamic link library, you may write your own versions of these routines.

When a process first attaches to the DLL, the FORTRAN 77 run-time environment is initialized and then the routine `__fdll_initialize_` is called. When a thread is started, the routine `__fthrd_initialize_` is called. When a thread is terminated, the routine `__fthrd_terminate_` is called. When the main process relinquishes the DLL, the routine `__fdll_terminate_` is called and then the FORTRAN 77 run-time environment is terminated.

The initialization and termination routines return an integer. A value of 0 indicates failure; a value of 1 indicates success. The following example illustrates sample initialization/termination routines.


```

* DLLINIT.FOR

integer function __fdll_initialize_()
integer __fthrd_initialize_, __fthrd_terminate_
integer __fdll_terminate_, dll_entry

integer WORKING_SIZE
parameter ( WORKING_SIZE = 16*1024 )
integer ierr, WorkingStorage
dimension WorkingStorage(:)

allocate( WorkingStorage(WORKING_SIZE), stat=ierr )
if( ierr .eq. 0 )then
    __fdll_initialize_ = 1
else
    __fdll_initialize_ = 0
endif
return

entry __fthrd_initialize_()
__fthrd_initialize_ = 1
return

entry __fthrd_terminate_()
__fthrd_terminate_ = 1
return

entry __fdll_terminate_()
* Note: no run-time calls allowed under OS/2 Warp
deallocate( WorkingStorage )
__fdll_terminate_ = 1
return

entry dll_entry()
! use WorkingStorage
return

end

```

In the above example, the process initialization routine allocates storage that the dynamic link library needs, the routine `dll_entry` uses the storage, and the process termination routine frees the storage allocated in the initialization routine.

OS/2 Programming Guide

30 Creating 16-bit OS/2 1.x Applications

An OS/2 application can be one of the following; a fullscreen application, a PM-compatible application, or a Presentation Manager application. A fullscreen application runs in its own screen group. A PM-compatible application will run in an OS/2 fullscreen environment or in a window in the Presentation Manager screen group but does not take direct advantage of menus, mouse or other features available in the Presentation Manager. A Presentation Manager application has full access to the complete set of user-interface tools such as menus, icons, scroll bars, etc.

This chapter deals with the creation of OS/2 fullscreen applications. For information on creating Presentation Manager applications, refer to the section entitled "Programming for OS/2 Presentation Manager" on page 227.

We will illustrate the steps to creating 16-bit OS/2 1.x applications by taking a small sample application and showing you how to compile, link, run and debug it.

30.1 The Sample Application

To demonstrate the creation of 16-bit OS/2 1.x applications using command-line oriented tools, we introduce a simple sample program. For our example, we are going to use the "sieve" program.

- * This program computes the prime numbers between 1 and 10,000
- * using the Sieve of Eratosthenes algorithm.

```
IMPLICIT NONE
INTEGER UBOUND
PARAMETER (UBOUND=10000)
INTEGER I, K, PRIMES
LOGICAL*1 NUMBERS(2:UBOUND)
CHARACTER*11 FORM
PARAMETER (FORM='(A,I5,A,I5)')
DO I = 2, UBOUND
    NUMBERS(I) = .TRUE.
ENDDO
PRIMES = 0
DO I = 2, UBOUND
    IF( NUMBERS(I) )THEN
        PRIMES = PRIMES + 1
        DO K = I + I, UBOUND, I
            NUMBERS(K) = .FALSE.
        ENDDO
    ENDIF
ENDDO
PRINT FORM, 'The Number of Primes between 1 and ', UBOUND,
1      ' are: ', PRIMES
END
```

The goal of this program is to count the prime numbers between 1 and 10,000. It uses the famous *Sieve of Eratosthenes* algorithm to accomplish this task. We will take you through the steps necessary to produce this result.

30.2 Building and Running the Sample OS/2 1.x Application

To compile and link our example program which is stored in the file `sieve.for`, enter the following command:

```
[C:\]wfl -l=os2 sieve.for
```

The typical messages that appear on the screen are shown in the following illustration.

```
[C:\]wfl -l=os2 sieve.for
Open Watcom F77/16 Compile and Link Utility
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1990-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
      wfc sieve.for
Open Watcom FORTRAN 77/16 Optimizing Compiler
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1984-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
sieve.for: 21 statements, 311 bytes, 6 extensions, 0 warnings, 0 errors

Open Watcom Linker
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1985-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
loading object files
searching libraries
creating an OS/2 16-bit executable
```

Provided that no errors were encountered during the compile or link phases, the "sieve" program may now be run.

```
[C:\]sieve
The Number of Primes between 1 and 10000 are: 1229
```

If you examine the current directory, you will find that two files have been created. These are `sieve.obj` (the result of compiling `sieve.for`) and `sieve.exe` (the result of linking `sieve.obj` with the appropriate Open Watcom FORTRAN 77 libraries). It is `sieve.exe` that is run by OS/2 when you enter the "sieve" command.

30.3 Debugging the Sample OS/2 1.x Application

Let us assume that you wish to debug your application in order to locate an error in programming. In the previous section, the "sieve" program was compiled with default compile and link options. When debugging an application, it is useful to refer to the symbolic names of routines and variables. It is also convenient to debug at the source line level rather than the machine language level. To do this, we must direct both the compiler and linker to include additional debugging information in the object and executable files. Using the `WFL` command, this is fairly straightforward. `WFL` recognizes the Open Watcom F77 compiler "debug" options and will create the appropriate debug directives for the Open Watcom Linker.

For example, to compile and link the "sieve" program with debugging information, the following command may be issued.

```
[C:\]wfl -l=os2 -d2 sieve.for
```

The typical messages that appear on the screen are shown in the following illustration.

```
[C:\]wfl -l=os2 -d2 sieve.for
Open Watcom F77/16 Compile and Link Utility
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1990-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
      wfc sieve.for -d2
Open Watcom FORTRAN 77/16 Optimizing Compiler
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1984-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
sieve.for: 21 statements, 392 bytes, 6 extensions, 0 warnings, 0 errors

Open Watcom Linker
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1985-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
loading object files
searching libraries
creating an OS/2 16-bit executable
```

The "d2" option requests the maximum amount of debugging information that can be provided by the Open Watcom F77 compiler. WFL will make sure that this debugging information is included in the executable file that is produced by the linker.

The "bytes" value is larger than in the previous example since selection of the "d2" option results in fewer code optimizations by default. You can request more optimization by specifying the appropriate options. However, you do so at the risk of making it more difficult for yourself to determine the relationship between the object code and the original source language code.

For OS/2, you should also include the BINP\DLL directory in the "LIBPATH" directive of the system configuration file CONFIG.SYS. It contains the Open Watcom Debugger Dynamic Link Libraries (DLLs).

Example:

```
libpath=c:\watcom\binp\dll
```

To request the Open Watcom Debugger to assist in debugging the application, the following command may be issued.

```
[C:\]wd sieve
```

It would be too ambitious to describe the debugger in this introductory chapter so we refer you to the book entitled *Open Watcom Debugger User's Guide*.

31 Creating 32-bit OS/2 Applications

An OS/2 application can be one of the following; a fullscreen application, a PM-compatible application, or a Presentation Manager application. A fullscreen application runs in its own screen group. A PM-compatible application will run in an OS/2 fullscreen environment or in a window in the Presentation Manager screen group but does not take direct advantage of menus, mouse or other features available in the Presentation Manager. A Presentation Manager application has full access to the complete set of user-interface tools such as menus, icons, scroll bars, etc.

This chapter deals with the creation of OS/2 fullscreen applications. For information on creating Presentation Manager applications, refer to the section entitled "Programming for OS/2 Presentation Manager" on page 227.

We will illustrate the steps to creating 32-bit OS/2 applications by taking a small sample application and showing you how to compile, link, run and debug it.

31.1 The Sample Application

To demonstrate the creation of 32-bit OS/2 applications using command-line oriented tools, we introduce a simple sample program. For our example, we are going to use the "sieve" program.

* This program computes the prime numbers between 1 and 10,000
* using the Sieve of Eratosthenes algorithm.

```
IMPLICIT NONE
INTEGER UBOUND
PARAMETER (UBOUND=10000)
INTEGER I, K, PRIMES
LOGICAL*1 NUMBERS(2:UBOUND)
CHARACTER*11 FORM
PARAMETER (FORM='(A,I5,A,I5)')
DO I = 2, UBOUND
    NUMBERS(I) = .TRUE.
ENDDO
PRIMES = 0
DO I = 2, UBOUND
    IF( NUMBERS(I) )THEN
        PRIMES = PRIMES + 1
        DO K = I + I, UBOUND, I
            NUMBERS(K) = .FALSE.
        ENDDO
    ENDIF
ENDDO
PRINT FORM, 'The Number of Primes between 1 and ', UBOUND,
1      ' are: ', PRIMES
END
```

The goal of this program is to count the prime numbers between 1 and 10,000. It uses the famous *Sieve of Eratosthenes* algorithm to accomplish this task. We will take you through the steps necessary to produce this result.

31.2 Building and Running the Sample OS/2 Application

To compile and link our example program which is stored in the file `sieve.for`, enter the following command:

```
[C:\]wfl386 -l=os2v2 sieve.for
```

The typical messages that appear on the screen are shown in the following illustration.

```
[C:\]wfl386 -l=os2v2 sieve.for
Open Watcom F77/32 Compile and Link Utility
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1990-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
wfc386 sieve.for
Open Watcom FORTRAN 77/32 Optimizing Compiler
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1984-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
sieve.for: 21 statements, 172 bytes, 6 extensions, 0 warnings, 0 errors

Open Watcom Linker
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1985-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
loading object files
searching libraries
creating an OS/2 32-bit executable
```

Provided that no errors were encountered during the compile or link phases, the "sieve" program may now be run.

```
[C:\]sieve
The Number of Primes between 1 and 10000 are: 1229
```

If you examine the current directory, you will find that two files have been created. These are `sieve.obj` (the result of compiling `sieve.for`) and `sieve.exe` (the result of linking `sieve.obj` with the appropriate Open Watcom FORTRAN 77 libraries). It is `sieve.exe` that is run by OS/2 when you enter the "sieve" command.

31.3 Debugging the Sample OS/2 Application

Let us assume that you wish to debug your application in order to locate an error in programming. In the previous section, the "sieve" program was compiled with default compile and link options. When debugging an application, it is useful to refer to the symbolic names of routines and variables. It is also convenient to debug at the source line level rather than the machine language level. To do this, we must direct both the compiler and linker to include additional debugging information in the object and executable files. Using the `WFL386` command, this is fairly straightforward. `WFL386` recognizes the Open Watcom F77 compiler "debug" options and will create the appropriate debug directives for the Open Watcom Linker.

For example, to compile and link the "sieve" program with debugging information, the following command may be issued.

```
[C:\]wfl386 -l=os2v2 -d2 sieve.for
```

The typical messages that appear on the screen are shown in the following illustration.

```
[C:\]wfl386 -l=os2v2 -d2 sieve.for
Open Watcom F77/32 Compile and Link Utility
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1990-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
    wfc386 sieve.for -d2
Open Watcom FORTRAN 77/32 Optimizing Compiler
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1984-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
sieve.for: 21 statements, 237 bytes, 6 extensions, 0 warnings, 0 errors

Open Watcom Linker
Copyright (c) 2002-2026 the Open Watcom Contributors. All Rights Reserved.
Portions Copyright (c) 1985-2002 Sybase, Inc. All Rights Reserved.
Source code is available under the Sybase Open Watcom Public License.
See https://github.com/open-watcom/open-watcom-v2 for details.
loading object files
searching libraries
creating an OS/2 32-bit executable
```

The "d2" option requests the maximum amount of debugging information that can be provided by the Open Watcom F77 compiler. WFL386 will make sure that this debugging information is included in the executable file that is produced by the linker.

The "bytes" value is larger than in the previous example since selection of the "d2" option results in fewer code optimizations by default. You can request more optimization by specifying the appropriate options. However, you do so at the risk of making it more difficult for yourself to determine the relationship between the object code and the original source language code.

To request the Open Watcom Debugger to assist in debugging the application, the following command may be issued.

```
[C:\]wd sieve
```

It would be too ambitious to describe the debugger in this introductory chapter so we refer you to the book entitled *Open Watcom Debugger User's Guide*.

32 OS/2 2.x Multi-threaded Applications

This chapter describes how to create multi-threaded applications. A multi-threaded application is one whose tasks are divided among several threads of execution. A process is an executing application and the resources it uses. A thread is the smallest unit of execution within a process. Each thread has its own stack and a set of machine registers and shares all resources with its parent process. The path of execution of one thread does not affect that of another; each thread is an independent entity.

Typically, an application has a single thread of execution. In this type of application, all tasks, once initiated, are completed before the next task begins. In contrast, tasks in a multi-threaded application can be performed concurrently since more than one thread is executing at once. For example, each thread may be designed to perform a separate task.

32.1 Programming Considerations

Since a multi-threaded application consists of many threads of execution, there are a number of issues that you must consider.

Since threads share the resources of its parent, it may be necessary to serialize access to these resources. For example, if your application contains more than one thread of execution and each thread uses the `PRINT` statement to display output to the console, it would be necessary for the I/O support routines to allow only one thread to use the `PRINT` facility at any time. That is, once a thread issues a `PRINT` request, the I/O support routines should ensure that no other thread displays information until all information for the initial thread has been displayed.

32.2 Creating Threads

Each application initially contains a single thread. The run-time libraries contain two functions that create and terminate threads of execution. The function `beginthread` creates a thread of execution and the function `endthread` ends a thread of execution. The function `threadid` can be used to determine the current thread identifier.

WARNING! If any thread uses an I/O statement or calls an intrinsic function, you must use the `beginthread` function to create the thread. Do not use the `DosCreateThread` API function.

32.2.1 Creating a New Thread

The `beginthread` function creates a new thread. It is defined as follows.

```
integer function beginthread( start_address,
                             stack_size )
integer stack_size
end
```

<i>where</i>	<i>description</i>
<i>start_address</i>	is the address of the subroutine that will be called when the newly created thread is executed. When the thread returns from that subroutine, the thread will be terminated. Note that a call to the <code>endthread</code> subroutine will also terminate the thread.
<i>stack_size</i>	specifies the size of the stack to be allocated by the operating system for the new thread. The stack size should be a multiple of 4K.

If a new thread is successfully created, the thread identifier of the new thread is returned. Otherwise, a value of -1 is returned.

The include file `thread.fi` contains the definition of the `beginthread` function.

32.2.2 Terminating the Current Thread

The `endthread` subroutine terminates the current thread. It is defined as follows.

```
subroutine endthread()  
end
```

The include file `thread.fi` contains the definition of the `endthread` function.

32.2.3 Getting the Current Thread Identifier

The `threadid` function can be used to determine the current thread identifier. It is defined as follows.

```
integer function threadid()  
end
```

The include file `thread.fi` contains the definition of the `threadid` function.

32.3 A Multi-threaded Example

Let us create a simple multi-threaded application. The source code for this example can be found in `\watcom\samples\fortran\os2.`

```
* MTHREAD.FOR  
  
*$pragma aux DosSleep parm( value ) [] caller  
  
integer NumThreads  
logical HoldThreads  
common NumThreads, HoldThreads  
  
integer STACK_SIZE  
parameter (STACK_SIZE=32768)  
integer NUM_THREADS  
parameter (NUM_THREADS=5)  
  
integer i, threadid, beginthread  
external a_thread
```

```

print ' (''main thread id = '', i4)', threadid()
NumThreads = 0
HoldThreads = .true.
! main thread counts as 1
do i = 2, NUM_THREADS
  if( beginthread( a_thread, STACK_SIZE ) .eq. -1 )then
    print ' (''creation of thread'', i4, ''failed'')', i
  else
    NumThreads = NumThreads + 1
  end if
end do
HoldThreads = .false.
while( NumThreads .ne. 0 )do
  call DosSleep( 1 )
end while
end

subroutine a_thread()
integer NumThreads
logical HoldThreads
common NumThreads, HoldThreads
integer threadid
while( HoldThreads )do
  call DosSleep( 1 )
end while
call DosEnterCritSec()
print ' (''Hi from thread '', i4)', threadid()
NumThreads = NumThreads - 1
call DosExitCritSec()
call endthread()
end

```

Note:

1. In the subroutine `a_thread`, `DosEnterCritSec` and `DosExitCritSec` are called when we modify the variable `NumThreads`. This ensures that the action of extracting the value of `NumThreads` from memory, incrementing the value, and storing the new result into memory, occurs without interruption. If these functions were not called, it would be possible for two threads to extract the value of `NumThreads` from memory before an update occurred.

Let us assume that the file `mthread.for` contains the above example. Before compiling the file, make sure that the **WATCOM** environment variable is set to the directory in which you installed Open Watcom FORTRAN 77. Also, the **FINCLUDE** environment variable must contain the `\watcom\src\fortran\os2` directory where "**WATCOM**" is the name of the directory in which you installed Open Watcom FORTRAN 77.

We can now compile and link the application by issuing the following command.

```
[C:\>]wfl386 -bm -l=os2v2 mthread
```

The "bm" option must be specified since we are creating a multi-threaded application. If your multi-threaded application contains more than one module, each module must be compiled using the "bm" switch.

The "l" option specifies the target system for which the application is to be linked. The system name `os2v2` is defined in the file `wlssystem.lnk` which is located in the "BINW" subdirectory of the directory in which you installed Open Watcom FORTRAN 77.

The multi-threaded application is now ready to be run.

32.4 Thread Limits

There is a limit to the number of threads an application can create under 16-bit OS/2. The default limit is 32. This limit can be adjusted by defining the integer function `__getmaxthreads` which returns the new thread limit.

Under 32-bit OS/2, there is no limit to the number of threads an application can create. However, due to the way in which multiple threads are supported in the Open Watcom libraries, there is a small performance penalty once the number of threads exceeds the default limit of 32 (this number includes the initial thread). If you are creating more than 32 threads and wish to avoid this performance penalty, you can redefine the threshold value of 32. You can statically initialize the global variable `__MaxThreads`.

This limit can be adjusted by defining the integer function `__getmaxthreads` which returns the new thread limit. By defining `__getmaxthreads` as follows, the new threshold value will be set to 48.

```
integer function __getmaxthreads()  
__getmaxthreads = 48  
end
```

This version of `__getmaxthreads` will replace the default function that is included in the run-time library. The default function simply returns the current value of the internal variable `__MaxThreads`. Your version of this function will return a new value for this variable. Internally, the run-time system executes code similar to the following:

```
.  
. .  
. .  
__MaxThreads = __getmaxthreads()  
. .  
. .  
. .
```

Thus, the default `__getmaxthreads` function does not alter the value of `__MaxThreads` but your version will.

33 OS/2 2.x Dynamic Link Libraries

A dynamic link library, like a standard library, is a library of functions. When an application uses functions from a standard library, the library functions referenced by the application become part of the executable module. This form of linking is called static linking. When an application uses functions from a dynamic link library, the library functions referenced by the application are not included in the executable module. Instead, the executable module contains references to these functions which are resolved when the application is loaded. This form of linking is called dynamic linking.

Let us consider some of the advantages of using dynamic link libraries over standard libraries.

1. Functions in dynamic link libraries are not linked into your program. Only references to the functions in dynamic link libraries are placed in the program module. These references are called import definitions. As a result, the linking time is reduced and disk space is saved. If many applications reference the same dynamic link library, the saving in disk space can be significant.
2. Since program modules only reference dynamic link libraries and do not contain the actual executable code, a dynamic link library can be updated without re-linking your application. When your application is executed, it will use the updated version of the dynamic link library.
3. Dynamic link libraries also allow sharing of code and data between the applications that use them. If many applications that use the same dynamic link library are executing concurrently, the sharing of code and data segments improves memory utilization.

33.1 Creating Dynamic Link Libraries

Once you have developed the source for a library of functions, a number of steps are required to create a dynamic link library containing those functions.

First, you must compile your source using the "bd" compiler option. This option tells the compiler that the module you are compiling is part of a dynamic link library. Once you have successfully compiled your source, you must create a linker directive file that describes the attributes of your dynamic link library. The following lists the most common linker directives required to create a dynamic link library.

1. The "SYSTEM" directive is used to specify that a dynamic link library is to be created.
2. The "EXPORT" directive is used to specify which functions in the dynamic link library are to be exported.
3. The "OPTION" directive is used to specify attributes such as the name of the dynamic link library and how to allocate the automatic data segment when the dynamic link library is referenced.
4. The "SEGMENT" directive is used to specify attributes of segments. For example, a segment may be read-only or read-write.

Once the dynamic link library is created, you must allow access to the dynamic link library to client applications that wish to use it. This can be done by creating an import library for the dynamic link library

or creating a linker directive file that contains "IMPORT" directives for each of the entry points in the dynamic link library.

33.2 Creating a Sample Dynamic Link Library

Let us now create a dynamic link library using the following example. The source code for this example can be found in `\watcom\samples\fortran\os2\dll`. Unlike applications developed in the C or C++ language, the FORTRAN 77 developer must not provide a `LibMain` entry point. This entry point is already defined in the Open Watcom FORTRAN 77 run-time libraries. The run-time system's `LibMain` provides for the proper initialization of the FORTRAN 77 run-time system and includes hooks to call developer-written process attach/detach and thread attach/detach routines. These routines are optional but we show skeleton versions in the following example so that you can develop your own if required.

```
* DLLSAMP.FOR

      integer function __fdll_initialize_()
*      Called from LibMain during "DLL PROCESS ATTACH"

*      do process initialization
      print *, 'Hi from process attach'

*      returning 0 indicates failure
      __fdll_initialize_ = 1
      return
      end

      integer function __fdll_terminate_()
*      Called from LibMain during "DLL PROCESS DETACH"

*      do process cleanup
      print *, 'Hi from process detach'

*      returning 0 indicates failure
      __fdll_terminate_ = 1
      return
      end

      subroutine dll_entry_1()
      print *, 'Hi from dll entry #1'
      end

      subroutine dll_entry_2()
      print *, 'Hi from dll entry #2'
      end
```

Assume the above example is contained in the file `dllsamp.for`. We can compile the file using the following command. Note that we must specify the "bd" compiler option.

```
[C:\]wfc386 -bd dllsamp
```

Before we can link our example, we must create a linker directive file that describes the attributes and entry points of our dynamic link library. The following is a linker directive file, called `dllsamp.lnk`, that can be used to create the dynamic link library.

```
system os2v2 dll initinstance terminstance
option manyautodata
export DLL_ENTRY_1
export DLL_ENTRY_2
file dllsamp
```

Notes:

1. The "SYSTEM" directive specifies that we are creating a 32-bit OS/2 dynamic link library.
2. The "MANYAUTODATA" option specifies that the automatic data segment is allocated for every instance of the dynamic link library. This option must be specified only for a dynamic link library that uses the Open Watcom FORTRAN 77 run-time libraries. If the Open Watcom FORTRAN 77 run-time libraries are not used, this option is not required. Our example does use the Open Watcom FORTRAN 77 run-time libraries so we must specify the "MANYAUTODATA" option.

As was just mentioned, when a dynamic link library uses the Open Watcom FORTRAN 77 run-time libraries, an automatic data segment is created each time a process accesses the dynamic link library. For this reason, initialization code must be executed when a process accesses the dynamic link library for the first time. To achieve this, "INITINSTANCE" must be specified in the "SYSTEM" directive. Similarly, "TERMINSTANCE" must be specified so that the termination code is executed when a process has completed its access to the dynamic link library. If the Open Watcom FORTRAN 77 run-time libraries are not used, these options are not required.

3. The "EXPORT" directive specifies the entry points into the dynamic link library. Note that in Open Watcom FORTRAN 77, names of all symbols are uppercased. Regardless of the case used in source files, linker directives must use uppercased symbol names. The linker is case sensitive by default, although the "OP NOCASEEXACT" directive may be used to override this.

We can now create our dynamic link library by issuing the following command.

```
[C:\]wlink @dllsamp
```

A file called `dllsamp.dll` will be created.

33.3 Using Dynamic Link Libraries

Once we have created a dynamic link library, we must allow other applications to access the functions available in the dynamic link library. There are two ways to achieve this.

The first method is to create a linker directive file which contains an "IMPORT" directive for all entry points in the dynamic link library. The "IMPORT" directive provides the name of the entry point and the name of the dynamic link library. When creating an application that references a function in the dynamic link library, this linker directive file would be included as part of the linking process that created the application.

The second method is to use import libraries. An import library is a standard library that is created from a dynamic link library by using the Open Watcom Library Manager. It contains object modules that describe the entry points in a dynamic link library. The resulting import library can then be specified in a "LIBRARY" directive in the same way one would specify a standard library.

Using an import library is the preferred method of providing references to functions in dynamic link libraries. When a dynamic link library is modified, typically the import library corresponding to the modified dynamic link library is updated to reflect the changes. Hence, any directive file that specifies the import library in a "LIBRARY" directive need not be modified. However, if you are using "IMPORT"

directives, you may have to modify the "IMPORT" directives to reflect the changes in the dynamic link library.

Let us create an import library for our sample dynamic link library we created in the previous section. We do this by issuing the following command.

```
[C:\]wlib dllsamp +dllsamp.dll
```

A standard library called `dllsamp.lib` will be created.

Suppose the following sample program, contained in the file `dlltest.for`, calls the functions from our sample dynamic link library.

```
* DLLTEST.FOR  
  
    call dll_entry_1()  
    call dll_entry_2()  
end
```

We can compile and link our sample application by issuing the following command.

```
[C:\]wfl386 -l=os2v2 dlltest dllsamp.lib
```

If we had created a linker directive file of "IMPORT" directives instead of an import library for the dynamic link library, the linker directive file, say `dllimps.lnk`, would be as follows.

```
import DLL_ENTRY_1 dllsamp  
import DLL_ENTRY_2 dllsamp
```

To compile and link our sample application, we would issue the following command.

```
[C:\]wfl386 -l=os2v2 dlltest -"@dllimps"
```

33.4 The Dynamic Link Library Data Area

The Open Watcom FORTRAN 77 32-bit run-time library does not support the general case operation of DLLs in an execution environment where there is only one instance of the DATA segment (DGROUP) for that DLL.

There are two cases that can lead to a DLL executing with only one instance of the DGROUP.

1. DLLs linked for 32-bit OS/2 without the MANYAUTODATA option.
2. DLLs linked for the Win32 API and executing under Win32s.

In these cases the run-time library startup code detects that there is only one instance of the DGROUP when a second process attempts to attach to the DLL. At that point, it issues a diagnostic for the user and then notifies the operating system that the second process cannot attach to the DLL.

Developers who require DLLs to operate when there is only one instance of the DGROUP can suppress the function which issues the diagnostic and notifies the operating system that the second process cannot attach to the DLL.

Doing so requires good behaviour on the part of processes attaching to the DLL. This good behaviour consists primarily of ensuring that the first process to attach to the DLL is also the last process to detach from the DLL thereby ensuring that the DATA segment is not released back to the free memory pool.

To suppress the function which issues the diagnostic and notifies the operating system that the second process cannot attach to the DLL, the developer must provide a replacement entry point with the following prototype:

```
int __disallow_single_dgroup( int );
```

This function should return zero to indicate that the detected single copy of the DATA segment is allowed.

33.5 Dynamic Link Library Initialization/Termination

Each dynamic link library (DLL) has an initialization and termination routine associated with it. The initialization routine can either be called the first time any process accesses the DLL ("INITGLOBAL" is specified at link time) or each time a process accesses the DLL ("INITINSTANCE" is specified at link time). Similarly, the termination routine can either be called when all processes have completed their access of the DLL ("TERMGLOBAL" is specified at link time) or each time a process completes its access of the DLL ("TERMINSTANCE" is specified at link time).

For a DLL that uses the FORTRAN 77 run-time libraries, initialization and termination of the FORTRAN 77 run-time environment is performed automatically. It is also possible for a DLL to do its own special initialization and termination process.

The FORTRAN 77 run-time environment provides a method for calling user-written DLL initialization and termination code. The `__fdll_initialize_` routine is called for DLL process initialization. The `__fdll_terminate_` routine is called for DLL process termination. Default stub versions of these routines are included in the run-time library. If you wish to perform additional initialization/termination processing that is specific to your dynamic link library, you may write your own versions of these routines.

Once the FORTRAN 77 run-time environment is initialized, the routine `__fdll_initialize_` is called. After the FORTRAN 77 run-time environment is terminated, the routine `__fdll_terminate_` is called. This last point is important since it means that you cannot do any run-time calls in the termination routine.

The initialization and termination routines return an integer. A value of 0 indicates failure; a value of 1 indicates success. The following example illustrates sample initialization/termination routines.

```
* DLLINIT.FOR

integer function __fdll_initialize_()
integer __fdll_terminate_, dll_entry

integer WORKING_SIZE
parameter ( WORKING_SIZE = 16*1024 )
integer ierr, WorkingStorage
dimension WorkingStorage(:)

allocate( WorkingStorage(WORKING_SIZE), stat=ierr )
if( ierr .eq. 0 )then
    __fdll_initialize_ = 1
else
    __fdll_initialize_ = 0
endif
return
```

```
        entry __fdll_terminate_  
* Note: no run-time calls allowed under OS/2 Warp  
        deallocate( WorkingStorage )  
        __fdll_terminate_ = 1  
        return  
  
        entry dll_entry()  
        ! use WorkingStorage  
        return  
  
    end
```

In the above example, the process initialization routine allocates storage that the dynamic link library needs, the routine `dll_entry` uses the storage, and the process termination routine frees the storage allocated in the initialization routine.

34 Programming for OS/2 Presentation Manager

Basically, there are two classes of FORTRAN 77 applications that can run in a windowed environment.

The first are those FORTRAN 77 applications that do not use any of the Presentation Manager API functions; they are strictly FORTRAN 77 applications that do not rely on the features of a particular operating system.

The second class of FORTRAN 77 applications are those that actually call Presentation Manager API functions directly. These are applications that have been tailored for the Presentation Manager operating environment.

It is assumed that the reader is familiar with the concepts of Presentation Manager programming.

34.1 Porting Existing FORTRAN 77 Applications

Suppose you have a set of FORTRAN 77 applications that previously ran under DOS and you now wish to run them under OS/2. To achieve this, simply recompile your application and link with the appropriate libraries. Depending on the method with which you linked your application, it can run in an OS/2 fullscreen environment, a PM-compatible window, or as a Presentation Manager application. An OS/2 fullscreen application runs in its own screen group. A PM-compatible application will run in an OS/2 fullscreen environment or in a window in the Presentation Manager screen group but does not take direct advantage of menus, mouse or other features available in the Presentation Manager. A Presentation Manager application has full access to the complete set of user-interface tools such as menus, icons, scroll bars, etc. However, porting a console oriented application to Presentation Manager often requires significant effort and a substantial redesign of the application.

34.1.1 An Example

Very little effort is required to port an existing FORTRAN 77 application to OS/2. Let us try to run the following sample program (contained in the file `hello.for`).

```
print *, 'Hello world!'
end
```

First we must compile the file `hello.for` by issuing the following command.

```
[C:\]wfc386 hello
```

Once we have successfully compiled the file, we can link it by issuing the following command.

```
[C:\]wlink sys os2v2 file hello
```

It is also possible to compile and link in one step, by issuing the following command.

```
[C:\]wfl386 -l=os2v2 hello
```

This will create a PM-compatible application. If you wish to create a fullscreen application, link with the following command.

```
[C:\]wlink sys os2v2 fullscreen file hello
```

34.2 Calling Presentation Manager API Functions

It is also possible for a FORTRAN 77 application to create its own windowing environment. This is achieved by calling PM API functions directly from your FORTRAN 77 program. The techniques for developing these applications can be found in the *OS/2 Technical Library*.

A number of FORTRAN 77 include files (files with extension `.fi` or `.fap`) are provided which define Presentation Manager data structures and constants. They are located in the `\watcom\src\fortran\os2` directory. These include files are equivalent to the C header files that are available with the IBM OS/2 Developer's Toolkit.

A sample FORTRAN 77 Presentation Manager application is also located in the `\watcom\samples\fortran\os2` directory. It is contained in the files `fshapes.for` and `fshapes.fi`. The file `fshapes.for` contains the following.

```
c$define INCL_WINFRAMEMGR
c$define INCL_WINMESSAGEMGR
c$define INCL_WINWINDOWMGR
c$define INCL_WINTIMER
c$define INCL_GPIPRIMITIVES
c$include os2.fap

      program fshapes

      integer          style
      record /QMSG/    qmsg

      character*7      watcom
      parameter        (watcom='WATCOM'c)

      include 'fshapes.fi'

      AnchorBlock = WinInitialize( 0 )
      if( AnchorBlock .eq. 0 ) stop
      hMessageQueue = WinCreateMsgQueue( AnchorBlock, 0 )
      if( hMessageQueue .eq. 0 ) stop
      if( WinRegisterClass( AnchorBlock, watcom, MainDriver,
+                          CS_SIZEREDRAW, 0 ) .eq. 0 ) stop
      style = FCF_TITLEBAR .or. FCF_SYSMENU .or. FCF_SIZEBORDER .or.
+           FCF_MINMAX .or. FCF_SHELLPOSITION .or. FCF_TASKLIST
      FrameHandle = WinCreateStdWindow( HWND_DESKTOP, WS_VISIBLE,
+                                     style, watcom,
+                                     char(0), 0, NULL,
+                                     0, WinHandle )
      if( FrameHandle .eq. 0 ) stop

      while( WinGetMsg( AnchorBlock, qmsg, NULL, 0, 0 ) ) do
        call WinDispatchMsg( AnchorBlock, qmsg )
      end while

      call WinDestroyWindow( FrameHandle )
      call WinDestroyMsgQueue( hMessageQueue )
      call WinTerminate( AnchorBlock )

      end
```



```
function MainDriver( hwnd, msg, mp1, mp2 )

integer hwnd
integer msg
integer mp1
integer mp2

include 'fshapes.fi'

integer      ps
record /RECTL/  rcl

select case ( msg )
case ( WM_CREATE )
    WinHandle = hwnd
    call WinStartTimer( AnchorBlock, WinHandle, 1, 150 )
case ( WM_TIMER )
    call DrawEllipse()
    MainDriver = 0
    return
case ( WM_SIZE )
    SizeX = SHORT1FROMMP( mp2 )
    SizeY = SHORT2FROMMP( mp2 )
    MainDriver = 0
    return
case ( WM_PAINT )
    ps = WinBeginPaint( WinHandle, NULL, NULL_POINTER )
    call WinQueryWindowRect( WinHandle, rcl )
    call WinFillRect( ps, rcl, CLR_WHITE )
    call WinEndPaint( ps )
    MainDriver = 0
    return
end select

MainDriver = WinDefWindowProc( WinHandle, msg, mp1, mp2 )
return

end

subroutine DrawEllipse

record /POINTL/      ptl
integer              ps
integer              Odd /0/
integer              parm1
integer              parm2

include 'fshapes.fi'

ps = WinGetPS( WinHandle )
ptl.x = Random( SizeX )
ptl.y = Random( SizeY )
call GpiMove( ps, ptl )
ptl.x = Random( SizeX )
ptl.y = Random( SizeY )
parm1 = Random( 32767 )
parm2 = Random( 32767 )
if( Random( 10 ) .ge. 5 ) then
    execute NewColor
    call GpiBox( ps, DRO_FILL, ptl, 0, 0 )
    execute NewColor
    call GpiBox( ps, DRO_OUTLINE, ptl, 0, 0 )
else
    execute NewColor
    call GpiBox( ps, DRO_FILL, ptl, parm1, parm2 )
    execute NewColor
    call GpiBox( ps, DRO_OUTLINE, ptl, parm1, parm2 )
end if
```

```
Odd = Odd + 1
Odd = Odd .and. 1
call WinReleasePS( ps )

remote block NewColor
call GpiSetColor( ps, Random( 15 ) + 1 )
end block

end

integer function Random( high )
integer          high

external          urand
real              urand
integer           seed /75347/
Random = urand( seed ) * high

end
```

The include file `fshapes.fi` contains the following.

```
include 'os2.fi'

integer          SizeX
integer          SizeY
integer          FrameHandle
integer          WinHandle
integer          hMessageQueue
integer          AnchorBlock

common /globals/
+ SizeX,
+ SizeY,
+ FrameHandle,
+ WinHandle,
+ hMessageQueue,
+ AnchorBlock

external          Random
integer           Random

external          MainDriver
integer           MainDriver
c$pragma aux (FNWP) MainDriver
```

Notes:

1. Include files with extension `.fap` define the calling conventions for each of the OS/2 API functions. These files must be included at the top of each FORTRAN 77 source module.
2. Include files with extension `.fi` define the data structures and constants used by the OS/2 API functions. These files must be included in each subprogram that requires them.
3. Each call-back function (i.e. window procedure) must be defined using the following pragma.

```
c$pragma aux (FNWP) WindowProc
```

4. The include file `os2.fap` is included at the beginning of the source file and `os2.fi` is included in each subprogram. Also note that a number of macros were defined at the top of the file. By defining these macros, only those components of the OS/2 API required by the module will be compiled.

You can compile, link and run this demonstration by issuing the following commands.

```
[C:\]set finclude=\watcom\src\fortran\os2  
[C:\]wfl386 -l=os2v2_pm fshapes  
[C:\]fshapes
```


Novell NLM Programming Guide

35 Creating NetWare 386 NLM Applications

Open Watcom FORTRAN 77 supports version 4.0 of the Netware 386 API. We include the following components:

header files Header files for the Netware 4.0 API are located in the \WATCOM\NOVH directory.

import libraries Import libraries for the Netware 4.0 API are located in the \WATCOM\NOVI directory.

libraries The FORTRAN 77 libraries for Netware 4.0 is located in the \WATCOM\LIB386 and \WATCOM\LIB386\NETWARE directories.

debug servers Servers for remote debugging of Netware 4.0 NLMs are located in the \WATCOM\NLM directory. The same directory also contains the Open Watcom Execution Sampler for NLMs.

Applications built for version 4.0 will run on 4.1. We do not include support for any API specific to version 4.1. Netware developers must use the support included with Open Watcom FORTRAN 77 version 10.0 or greater since the version supplied by Novell only works with Open Watcom FORTRAN 77 version 9.5. Netware 4.1 support requires modification to the header files supplied by Novell. Contact Novell for more information.

The following special notes apply to developing applications for NetWare.

1. You must compile your source files with the small memory model option ("ms").
2. You must compile your source files with the stack-based calling convention option ("sc").
3. You must specify

system NETWARE

when linking an NLM. This is automatic if you are using WFL386 and the "-l=NETWARE" option.

Mixed Language Programming

36 *Inter-Language calls: C and FORTRAN*

The purpose of this chapter is to anticipate common questions about mixed-language development using Open Watcom C/C++ and Open Watcom FORTRAN 77.

The following topics are discussed in this chapter:

- Symbol Naming Convention
- Argument Passing Convention
- Memory Model Compatibility
- Integer Type Compatibility
- How do I pass integers from C to a FORTRAN function?
- How do I pass integers from FORTRAN to a C function?
- How do I pass a string from a C function to FORTRAN?
- How do I pass a string from FORTRAN to a C function?
- How do I access a FORTRAN common block from within C?
- How do I call a C function that accepts a variable number of arguments?

36.1 *Symbol Naming Convention*

The symbol naming convention describes how a symbol in source form is mapped to its object form. Because of this mapping, the name generated in the object file may differ from its original source form.

Default symbol naming conventions vary between compilers. Open Watcom C/C++ prefixes an underscore character to the beginning of variable names and appends an underscore to the end of function names during the compilation process. Open Watcom FORTRAN 77 converts symbols to upper case. Auxiliary pragmas can be used to resolve this inconsistency.

Pragmas are compiler directives which can provide several capabilities, one of which is to provide information used for code generation. When calling a FORTRAN subprogram from C, we want to instruct the compiler NOT to append the underscore at the end of the function name and to convert the name to upper case. This is achieved by using the following C auxiliary pragma:

```
#pragma aux ftname "^";
```

The "^" character tells the compiler to convert the symbol name "ftname" to upper case; no underscore character will be appended. This solves potential linker problems with "ftname" since (by C convention) the linker would attempt to resolve a reference to "ftname_".

When calling C functions from FORTRAN, we need to instruct the compiler to add the underscore at the end of the function name, and to convert the name to lower case. Since the FORTRAN compiler automatically converts identifiers to uppercase, it is necessary to force the compiler to emit an equivalent lowercase name. Both of these things can be done with the following FORTRAN auxiliary pragma:

```
*$pragma aux CNAME "!"_"
```

There is another less convenient way to do this as shown in the following:

```
*$pragma aux CNAME "cname_"
```

In the latter example, the case of the name in quotation marks is preserved.

Use of these pragmas resolves the naming differences, however, the issue of argument passing must still be resolved.

36.2 Argument Passing Convention

In general, C uses call-by-value (passes argument values) while FORTRAN uses call-by-reference (passes pointers to argument values). This implies that to pass arguments to a FORTRAN subprogram we must pass the addresses of arguments rather than their values. C uses the "&" character to signify "address of".

Example:

```
result = ftnname( &arg );
```

When calling a C function from FORTRAN, the pragma used to correct the naming conventions must also instruct the compiler that the C function is expecting values, not addresses.

```
*$pragma aux CNAME "!"_" parm (value)
```

The "parm (value)" addition instructs the FORTRAN compiler to pass values, instead of addresses.

Character data (strings) are an exception to the general case when used as arguments. In C, strings are not thought of as a whole entity, but rather as an "array of characters". Since strings are not considered scalar arguments, they are referenced differently in both C and FORTRAN. This is described in more detail in a following section.

36.3 Memory Model Compatibility

While it is really not an issue with the 32-bit compilers (both use the default "flat" memory model), it is important to know that the default memory model used in Open Watcom FORTRAN 77 applications is the "large" memory model ("ml") with "medium" and "huge" memory models as options. Since the 16-bit Open Watcom C/C++ default is the "small" memory model, you must specify the correct memory model when compiling your C/C++ code with the 16-bit C or C++ compiler.

36.4 Linking Considerations

When both C/C++ and FORTRAN object files are combined into an executable program or dynamic link library, it is important that you list a least one of the FORTRAN object files first in the Open Watcom Linker (WLINK) "FILES" directive to guarantee the proper search order of the FORTRAN and C run-time libraries. If you place a C/C++ object file first, you may inadvertently cause the wrong version of run-time initialization routines to be loaded by the linker.

36.5 Integer Type Compatibility

In general, the number of bytes used to store an integer type is implementation dependent. In FORTRAN, the default size of an integer type is always 4 bytes, while in C/C++, the size is architecture dependent. The size of an "int" is 2 bytes for the 16-bit Open Watcom C/C++ compilers and 4 bytes for the 32-bit compilers while the size of a "long" is 4 bytes regardless of architecture. It is safest to prototype the function in C, specifying exactly what size integers are being used. The byte sizes are as follows:

1. LONG - 4 bytes
2. SHORT - 2 bytes

Since FORTRAN uses a default of 4 bytes, we should specify the "long" keyword in C for integer types.

Example:

```
long int ftname( long int *, long int *, long int * );
```

In this case, "ftname" takes three "pointers to long ints" as arguments, and returns a "long int". By specifying that the arguments are pointers, and not values, and by specifying "long int" for the return type, this prototype has solved the problems of argument passing and integer type compatibility.

36.6 How do I pass integers from C to a FORTRAN function?

The following Open Watcom C/C++ routine passes three integers to a FORTRAN function that returns an integer value.

```
/* MIX1C.C - This C program calls a FORTRAN function to
 *           compute the max of three numbers.
 *
 * Compile/Link: wcl /ml mix1c mix1f.obj /fe=mix1
 *               wcl386 mix1c mix1f.obj /fe=mix1
 */

#include <stdio.h>

#pragma aux tmax3 "^";
long int tmax3( long int *, long int *, long int * );
```

```
void main()
{
    long int  result;
    long int  i, j, k;

    i = -1;
    j = 12;
    k = 5;
    result = tmax3( &i, &j, &k );
    printf( "Maximum is %ld\n", result );
}
```

The FORTRAN function:

```
* MIX1F.FOR - This FORTRAN function accepts three integer
*              arguments and returns their maximum.
```

```
* Compile: wfc[386] mix1f.for
```

```
integer function tmax3( arga, argb, argc )
integer arga, argb, argc

tmax3 = arga
if ( argb .gt. tmax3 ) tmax3 = argb
if ( argc .gt. tmax3 ) tmax3 = argc
end
```

36.7 How do I pass integers from FORTRAN to a C function?

The following Open Watcom FORTRAN 77 routine passes three integers to a Open Watcom C/C++ function that returns an integer value.

```
* MIX2F.FOR - This FORTRAN program calls a C function to
*              compute the max of three numbers.
```

```
* Compile/Link: wfl[386] mix2f mix2c.obj /fe=mix2
```

```
*$pragma aux tmax3 "!" parm (value)
```

```
program mix2f

integer*4  tmax3
integer*4  result
integer*4  i, j, k

i = -1
j = 12
k = 5
result = tmax3( i, j, k )
print *, 'Maximum is ', result
end
```

The C function "tmax3" is shown below.

```

/* MIX2C.C - This C function accepts 3 integer arguments
 *           and returns their maximum.
 *
 * Compile: wcc /ml mix2c
 *           wcc386  mix2c
 */

long int tmax3( long int arga,
                long int argb,
                long int argc )
{
    long int  result;
    result = arga;
    if( argb > result ) result = argb;
    if( argc > result ) result = argc;
    return( result );
}

```

36.8 How do I pass a string from a C function to FORTRAN?

Character strings are referenced differently in C and FORTRAN. The C language terminates its strings with a null character as an End-Of-String (EOS) marker. In this case, C need not store the length of the string in memory. FORTRAN, however, does not use any EOS marker; hence it must store each string's length in memory.

The structure FORTRAN uses to keep track of character data is called a "string descriptor" which consists of a pointer to the character data (2, 4, or 6 bytes, depending on the data model) followed by an unsigned integer length (2 bytes or 4 bytes, depending on the data model).

system	option	size of pointer	size of length
16-bit	/MM	16 bits	16 bits
16-bit	/ML	32 bits	16 bits
32-bit	/MF	32 bits	32 bits
32-bit	/ML	48 bits	32 bits

In order to access character data, FORTRAN needs to have access to the data's string descriptor. Hence, FORTRAN expects a pointer to a string descriptor to be passed as an argument for character data.

Passing string arguments between C and FORTRAN is a simple task of describing a struct type in C containing the two fields described above. The first field must contain the pointer to the character data, and the second field must contain the length of the string being passed. A pointer to this structure can then be passed to FORTRAN.

```

* MIX3F.FOR - This FORTRAN program calls a function written
*             in C that passes back a string.
*
* Compile/Link: wfl[386] mix3f mix3c.obj /fe=mix3

      program mix3f

      character*80 sendstr
      character*80 cstring

```

```
cstring = sendstr()
print *, cstring(1:lentrim(cstring))
end
```

The C function "sendstr" is shown below.

```
/* MIX3C.C - This C function passes a string back to its
 *           calling FORTRAN program.
 *
 * Compile:  wcc /ml mix3c
 *           wcc386 mix3c
 */

#include <string.h>

#pragma aux sendstr "^";

typedef struct descriptor {
    char          *addr;
    unsigned      len;
} descriptor;

void sendstr( descriptor *ftn_str_desc )
{
    ftn_str_desc->addr = "This is a C string";
    ftn_str_desc->len  = strlen( ftn_str_desc->addr );
}
```

36.9 How do I pass a string from FORTRAN to a C function?

By default, FORTRAN passes the address of the string descriptor when passing strings. If the C function knows it is being passed a string descriptor address, then it is very similar to the above example. If the C function is expecting normal C-type strings, then a FORTRAN pragma can be used to pass the string correctly. When the Open Watcom FORTRAN 77 compiler pragma to pass by value is used for strings, then just a pointer to the string is passed.

Example:

```
*$pragma aux cname "!"_ parm (value)
```

The following example FORTRAN mainline defines a string, and passes it to a C function that prints it out.

```
* MIX4F.FOR - This FORTRAN program calls a function written
 *           in C and passes it a string.
 *
 * Compile/Link: wfl[386] mix4f mix4c.obj /fe=mix4

*$pragma aux cstr "!"_ parm (value)

program mix4f

character*80 forstring

forstring = 'This is a FORTRAN string'//char(0)
call cstr( forstring )
end
```


The C function:

```
/* MIX4C.C - This C function prints a string passed from
 *          FORTRAN.
 *
 * Compile: wcc /ml mix4c
 *          wcc386  mix4c
 */

#include <stdio.h>

void cstr( char *instring )
{
    printf( "%s\n", instring );
}
```

36.10 How do I access a FORTRAN common block from within C?

The following code demonstrates a technique for accessing a FORTRAN common block in a C routine. The C routine defines an extern struct to correspond to the FORTRAN common block.

```
* MIX5F.FOR - This program shows how a FORTRAN common
*             block can be accessed from C.
*
* Compile/Link: wfl[386] mix5f mix5c.obj /fe=mix5

      program mix5f
      external put
      common/cblk/i,j

      i=12
      j=10
      call put
      print *, 'i = ', i
      print *, 'j = ', j
      end
```

The C function:

```
/* MIX5C.C - This code shows how to access a FORTRAN
 *          common block from C.
 *
 * Compile: wcc /ml mix5c
 *          wcc386  mix5c
 */

#include <stdio.h>

#pragma aux put    "^";
#pragma aux cblk  "^";
```

```
#ifdef __386__
#define FAR
#else
#define FAR far
#endif

extern struct cb {
    long int i,j;
} FAR cblk;

void put( void )
{
    printf( "i = %ld\n", cblk.i );
    printf( "j = %ld\n", cblk.j );
    cblk.i++;
    cblk.j++;
}
```

For the 16-bit C compiler, the common block "cblk" is described as `far` to force a load of the segment portion of the address. Otherwise, since the object is smaller than 32K (the default data threshold), it is assumed to be located in the DGROUP group which is accessed through the SS segment register.

36.11 How do I call a C function that accepts a variable number of arguments?

One capability that C possesses is the ability to define functions that accept variable number of arguments. This feature is not present, however, in the definition of the FORTRAN 77 language. As a result, a special pragma is required to call these kinds of functions.

```
*$pragma aux printf "!_" parm (value) caller []
```

The "caller" specifies that the caller will pop the arguments from the stack. The "[]" indicates that there are no arguments passed in registers because the `printf` function takes a variable number of arguments passed on the stack. The following example is a FORTRAN function that uses this pragma. It calls the `printf` function to print the value 47 on the screen.

```
* MIX6.FOR - This FORTRAN program calls the C
*           printf function.

* Compile/Link: wfl[386] mix6

*$pragma aux printf "!_" parm (value) caller []
```

```
program mix6

character cr/z0d/, nullchar/z00/

call printf( 'Value is %ld.'//cr//nullchar, 47 )
end
```

For more information on the pragmas that are used extensively during inter-language programming, please refer to the chapter entitled "Pragmas" in both the *Open Watcom C/C++ User's Guide* and the *Open Watcom FORTRAN 77 User's Guide*.

Common Problems

37 Commonly Asked Questions and Answers

As with any sophisticated piece of software, there are topics that are not directly addressed by the descriptive portions of the manuals. The purpose of this chapter is to anticipate common questions concerning Open Watcom F77. It is difficult to predict what topics will prove to be useful but with that in mind, we hope that this chapter will help our customers make full use of Open Watcom F77.

A number of example programs are presented throughout. The source text for these files can be found in the `\WATCOM\SAMPLES\GOODIES` directory.

The purpose of this chapter is to present some of the more commonly asked questions from our users and the answers to these questions. The following topics are discussed:

- How do I convert to Open Watcom F77?
- What should I know about optimization?
- How do I read a stream of binary data from a file?
- How do I redefine math error handling with Open Watcom F77?
- Why can't the compiler find my include files?
- Why does the linker report a "stack segment not found" error?
- How do I resolve an "Undefined Reference" linker error?
- Why aren't local variable values maintained between subprogram calls?
- What does "Stack Overflow!" mean?
- What are the probable causes of a General Protection Fault in 32-bit applications?
- Which floating-point compiler option should I use for my application?
- How can I open more than 20 files at a time?
- How can I see my source files in the debugger?
- What is the difference between the "d1" and "d2" compiler options?
- What is the difference between the "debug" and "d2" compiler options?

37.1 Converting to Open Watcom F77

Applications written in ANSI standard FORTRAN 77 code usually only need to be recompiled with the Open Watcom F77 compiler. In addition to the ANSI standard, many compilers support specific extensions. If you are porting code from a UNIX platform or other DOS compilers, check *Appendix A - Extensions to Standard FORTRAN 77* of the *Open Watcom FORTRAN 77 Language Reference*, to determine which FORTRAN 77 extensions are supported.

By default, most FORTRAN 77 compilers preserve the values of local variables in a subprogram between calls. With Open Watcom F77, local variables are kept on the stack and their values are popped from the stack when exiting a subprogram. To preserve local variables, use the FORTRAN 77 SAVE statement for variables that you wish to preserve, or the "save" compiler option to preserve all local variables. Note that the use of the "save" compiler option causes an overall performance degradation.

Open Watcom F77 uses register-based parameter passing as a default, however, the compiler is flexible enough to use different calling conventions on a per function basis. Auxiliary pragmas can be used to specify the calling convention that is to be used to interface with assembler code. This enables you to explicitly state how parameters are to be passed to the assembler code. This topic is described in the

"Pragmas" chapter of the *Open Watcom F77 User's Guide* under "Describing Argument Information". See also the chapter entitled "Inter-Language calls: C and FORTRAN" on page 239.

37.2 What you should know about optimization

The *Open Watcom F77 User's Guide* contains a detailed description for each of the optimization options supported by the compiler. These options allow you to customize the type of code optimizations that are performed. For instance, the "OS" option can be used to reduce the size of your code, but this may affect the execution speed. To ensure that the speed of the code is optimized, possibly at the cost of code size, use the "OT" option. The "OX" option, intended for the maximum number of optimizations, generates code that is a combination of "OM" (inline math functions), "OL" (loop), "OT" (time) and the "OR" (instruction scheduling) optimization options. Note that when you are using the "OM" option to generate inline math functions no argument validation will be done for the intrinsic math functions such as "sin" or "cos". Consider the needs of your application and select the optimization options that best meet your requirements.

Hint: The definitive reference on compiler design is the "dragon" book "Compilers - Principles, Techniques, and Tools", Alfred V. Aho, Ravi Sethi, and Jeffrey D. Ullman, published by Addison-Wesley, Reading, Massachusetts, 1986. The authors of the "dragon" book advocate a conservative approach to code generation where optimizations must preserve the semantics of the original program. The conservative approach is used throughout the Open Watcom F77 compiler to ensure that programmers can use the compiler without worrying about the semantics of their program being changed.

There are certain pieces of information which the compiler cannot derive from the source code. The "\$pragma" compiler directive is used to provide extra information to the compiler. It is necessary to have a complete understanding of both FORTRAN 77 and the machine architecture (i.e., 80x86) before using the powerful pragma compiler directives. See the "Pragmas" chapter in the *Open Watcom F77 User's Guide* for more details.

Debugging optimized programs is difficult because variables can be assigned to different locations (i.e., memory or registers) in different parts of the function. The "d2" compiler option will restrict the amount of optimization so that variables occupy one location and can be easily displayed. It follows that the "d2" option is useful for initial development but production programs should be compiled with only the "d1" option for the best code quality. Before you distribute your application to others, you may wish to use the Open Watcom Strip Utility (WSTRIP) to remove debugging information from the executable image on disk thereby reducing disk space requirements.

Hint: The "d2" compiler option will generate symbolic information (for every local variable) and line number information for the source file. The "d1" compiler option will only generate line number information for the source file. The use of these options determines what kind of information will be available for the particular module during the debugging session.

37.3 Reading a stream of binary data from a file

The Open Watcom F77 allows for three types of binary data file formats:

- Variable length, unformatted sequential access binary records,
- Fixed length, unformatted direct access binary records, and
- Unformatted, sequential, binary data with a fixed record type.

Variable length binary records are preceded by a four byte descriptor that indicates the length of the record in bytes. The end of the binary record is also marked by another descriptor tag specifying the length. Binary records that are of a fixed length are kept in a direct access, unformatted file. Refer to the *Open Watcom F77 User's Guide* section on File Handling for more information on file formats.

Binary data files that have no structure or record length information may be read if you open the file as a sequential, unformatted file with a fixed record type. This allows you to read files that contain a stream of binary data without any inherent record format. If you know the type of data that is contained in the binary file, you may then read the binary data directly into variables. The following program provides an example of reading binary stream data.

- * BINDATA.FOR – This program demonstrates how to read a
- * binary data file that does not have any defined records.

```
program bindata

integer BinArray(20)
integer i

open( unit=1, file='bindata.fil',
+     access='sequential',
+     form='unformatted',
+     recordtype='fixed' )

* Read 20 integers from the binary data file
do i = 1, 20
    read( 1 ) BinArray( i )
end do

* Write the extracted values to standard output
do i = 1, 20
    write( *, * ) BinArray( i )
end do
end
```

37.4 Redefining math error handling with Open Watcom F77

If you wish to customize math error handling for your application, you can create your own math error handling procedure. The following illustrates the procedures for trapping errors by way of an example. See the *Open Watcom F77 User's Guide* for a description of the `fsignal` subroutine and math library error handling.

The main program example "MATHDEMO" is a FORTRAN program that contains a floating-point divide by zero error, a floating-point overflow error, a floating-point underflow error, and an invalid argument to a math library function.

```
program mathdemo

* MATHDEMO.FOR - This program forms part of a collection of FORTRAN
*                code that demonstrates how to take over control of
*                math error handling from the run-time system.

* Compile: wfl[386] mathdemo cw87 _matherr

* Notes:
* (1) We call "cw87" to enable underflow exceptions which are
*     masked (ignored) by default.
* (2) The signal handler must be re-installed after each signal
*     (it can also be re-installed even when there is no signal).
* (3) To prevent compile-time constant folding in expressions,
*     we add log(1.0) which is 0. We do this for the sake of
*     demonstrating exception handling.

implicit none

double precision x, y, z

call cw87          ! init 80x87 control word

call resetFPE      ! install signal handler
print *, ' '
print *, 'Divide by zero will be attempted'
x = 1.0d0 + DLOG( 1.0d0 )
y = 0.0d0
z = x / y
call chkFPE        ! check for exception
print *, z

call resetFPE      ! install signal handler
print *, ' '
print *, 'Overflow will be attempted'
x = 1.2d300 + DLOG( 1.0d0 )
y = 1.2d300
z = x * y
call chkFPE        ! check for exception
print *, z

call resetFPE      ! install signal handler
print *, ' '
print *, 'Underflow will be attempted'
x = 1.14d-300 + DLOG( 1.0d0 )
y = 2.24d-308
z = x * y
call chkFPE        ! check for exception
print *, z

call resetFPE      ! install signal handler
print *, ' '
print *, 'Math error will be attempted'
x = -12.0
! an exception will not be generated since the intrinsic function
! will validate the argument - if you compile with -om, the "fsqrt"
! 80x87 instruction will be generated in-line and an exception
! will occur
y = SQRT( x )
call chkFPE        ! check for exception
print *, x, y
end
```



```

subroutine resetFPE
include 'fsignal.fi'
external fpe_handler
logical fpe_flag
integer fpe_sig, fpe_fpe
common fpe_flag, fpe_sig, fpe_fpe
fpe_flag = .false.
call fsignal( SIGFPE, fpe_handler )
end

*$pragma aux fpe_handler parm( value )

subroutine fpe_handler( sig, fpe )
integer sig, fpe
logical fpe_flag
integer fpe_sig, fpe_fpe
common fpe_flag, fpe_sig, fpe_fpe
fpe_flag = .true.
fpe_sig = sig
fpe_fpe = fpe
end

*$pragma aux fwait = "fwait"

subroutine chkFPE
include 'fsignal.fi'
logical fpe_flag
integer fpe_sig, fpe_fpe
common fpe_flag, fpe_sig, fpe_fpe
* Notes:
* (1) An fwait is required to make sure that the last
* floating-point instruction has completed.
* (2) "volatile" is not needed here but would be
* needed in main program if it references "fpe_flag"
call fwait()
if( volatile( fpe_flag ) ) then
  print *, '*ERROR* exception occurred',
    & fpe_sig, fpe_fpe
  if( fpe_fpe .eq. FPE_INVALID )then
    print *, 'Invalid'
  else if( fpe_fpe .eq. FPE_DENORMAL )then
    print *, 'Denormalized operand error'
  else if( fpe_fpe .eq. FPE_ZERODIVIDE )then
    print *, 'Divide by zero error'
  else if( fpe_fpe .eq. FPE_OVERFLOW )then
    print *, 'Overflow error'
  else if( fpe_fpe .eq. FPE_UNDERFLOW )then
    print *, 'Underflow error'
  else if( fpe_fpe .eq. FPE_INEXACT )then
    print *, 'Inexact result (precision)then error'
  else if( fpe_fpe .eq. FPE_UNEMULATED )then
    print *, 'Unemulated instruction error'
  else if( fpe_fpe .eq. FPE_SQRTNEG )then
    print *, 'Square root of a negative number error'
  else if( fpe_fpe .eq. FPE_STACKOVERFLOW )then
    print *, 'NDP stack overflow error'
  else if( fpe_fpe .eq. FPE_STACKUNDERFLOW )then
    print *, 'NDP stack underflow error'
  else if( fpe_fpe .eq. FPE_EXPLICITGEN )then
    print *, 'SIGFPE signal raised (software)'
  else if( fpe_fpe .eq. FPE_IOVERFLOW )then
    print *, 'Integer overflow error'
  endif
else
  print *, '*OK* no exception occurred'
endif
end

```

The following subroutine illustrates how to enable or disable particular types of floating-point exceptions.

```
subroutine cw87

* CW87.FOR
* This subroutine uses the C Library routine "_control87"
* to modify the math coprocessor exception mask.

* Compile: wfc[386] cw87

include 'fsignal.fi'

character*9 status(0:1)/' disabled',' enabled'/
integer fp_cw, fp_mask, bits, i

* Enable floating-point underflow since default is disabled.
* The mask defines which bits we want to change (1 means change,
* 0 means not change). The corresponding bit in the control
* word (fp_cw) is set to 0 to enable the exception or 1 to disable
* the exception. In this example, we change only the underflow
* bit and leave the others unchanged.

fp_mask = EM_UNDERFLOW      ! mask for the bits to set/reset
fp_cw = '0000'x             ! new bit settings (0=enable/1=disable)
fp_cw = _control87( fp_cw, fp_mask )

* Now get up-to-date setting

fp_cw = _control87( 0, 0 )

bits = IAND( fp_cw, MCW_EM )
print ' (a,lx,z4)', 'Interrupt Exception Mask', bits
i = 0
if( IAND(fp_cw, EM_INVALID) .eq. 0 ) i = 1
print *, ' Invalid Operation exception', status(i)
i = 0
if( IAND(fp_cw, EM_DENORMAL) .eq. 0 ) i = 1
print *, ' Denormalized exception', status(i)
i = 0
if( IAND(fp_cw, EM_ZERODIVIDE) .eq. 0 ) i = 1
print *, ' Divide-By-Zero exception', status(i)
i = 0
if( IAND(fp_cw, EM_OVERFLOW) .eq. 0 ) i = 1
print *, ' Overflow exception', status(i)
i = 0
if( IAND(fp_cw, EM_UNDERFLOW) .eq. 0 ) i = 1
print *, ' Underflow exception', status(i)
i = 0
if( IAND(fp_cw, EM_PRECISION) .eq. 0 ) i = 1
print *, ' Precision exception', status(i)
end
```

The following subroutine illustrates how to replace the run-time system's math error handler. Source code similar to this example is provided with the software (look for the file `_matherr.for`).

```
*
* _MATHERR.FOR : math error handler
*
* Compile: wfc[386] _matherr

*$pragma aux __imath2err "*" parm( value, reference, reference )
*$pragma aux __amath1err "*" parm( value, reference )
*$pragma aux __amath2err "*" parm( value, reference, reference )
*$pragma aux __math1err "*" parm( value, reference )
*$pragma aux __math2err "*" parm( value, reference, reference )
*$pragma aux __zmath2err "*" parm( value, reference, reference )
*$pragma aux __qmath2err "*" parm( value, reference, reference )
```

```
integer function __imath2err( err_info, arg1, arg2 )
integer err_info
integer arg1, arg2
include 'mathcode.fi'
arg1 = arg1      ! to avoid unreferenced warning message
arg2 = arg2      ! to avoid unreferenced warning message
if( ( err_info .and. M_DOMAIN ) .ne. 0 )then
  select( err_info .and. FUNC_MASK )
  case( FUNC_POW )
    print *, 'arg2 cannot be <= 0'
  case( FUNC_MOD )
    print *, 'arg2 cannot be 0'
  end select
end if
__imath2err = 0
end

real function __amathlerr( err_info, arg1 )
integer err_info
real arg1
include 'mathcode.fi'
arg1 = arg1      ! to avoid unreferenced warning message
if( ( err_info .and. M_DOMAIN ) .ne. 0 )then
  select( err_info .and. FUNC_MASK )
  case( FUNC_COTAN )
    print *, 'overflow'
  end select
end if
__amathlerr = 0.0
end

real function __amath2err( err_info, arg1, arg2 )
integer err_info
real arg1, arg2
include 'mathcode.fi'
arg1 = arg1      ! to avoid unreferenced warning message
arg2 = arg2      ! to avoid unreferenced warning message
if( ( err_info .and. M_DOMAIN ) .ne. 0 )then
  select( err_info .and. FUNC_MASK )
  case( FUNC_MOD )
    print *, 'arg2 cannot be 0'
  end select
end if
__amath2err = 0.0
end

double precision function __mathlerr( err_info, arg1 )
integer err_info
double precision arg1, __math2err
__mathlerr = __math2err( err_info, arg1, arg1 )
end

double precision function __math2err( err_info, arg1, arg2 )
integer err_info
double precision arg1, arg2
include 'mathcode.fi'
arg1 = arg1      ! to avoid unreferenced warning message
arg2 = arg2      ! to avoid unreferenced warning message
```

```
if( ( err_info .and. M_DOMAIN ) .ne. 0 )then
  select( err_info .and. FUNC_MASK )
  case( FUNC_SQRT )
    print *, 'argument cannot be negative'
  case( FUNC_ASIN, FUNC_ACOS )
    print *, 'argument must be less than or equal to one'
  case( FUNC_ATAN2 )
    print *, 'both arguments must not be zero'
  case( FUNC_POW )
    if( arg1 .eq. 0.0 )then
      print *, 'a zero base cannot be raised to a ',
&        'negative power'
    else ! base < 0 and non-integer power
      print *, 'a negative base cannot be raised to a ',
&        'non-integral power'
    endif
  case( FUNC_LOG, FUNC_LOG10 )
    print *, 'argument must not be negative'
  end select
else if( ( err_info .and. M_SING ) .ne. 0 )then
  if( ( ( err_info .and. FUNC_MASK ) .eq. FUNC_LOG ) .or.
&    ( ( err_info .and. FUNC_MASK ) .eq. FUNC_LOG10 ) )then
    print *, 'argument must not be zero'
  endif
else if( ( err_info .and. M_OVERFLOW ) .ne. 0 )then
  print *, 'value of argument will cause overflow condition'
else if( ( err_info .and. M_UNDERFLOW ) .ne. 0 )then
  print *, 'value of argument will cause underflow ',
&    'condition - return zero'
end if
__math2err = 0
end

complex function __zmath2err( err_info, arg1, arg2 )
integer err_info
complex arg1, arg2
include 'mathcode.fi'
arg1 = arg1      ! to avoid unreferenced warning message
arg2 = arg2      ! to avoid unreferenced warning message
if( ( err_info .and. M_DOMAIN ) .ne. 0 )then
  select( err_info .and. FUNC_MASK )
  case( FUNC_POW )
    ! arg1 is (0,0)
    if( imag( arg2 ) .ne. 0 )then
      print *, 'a zero base cannot be raised to a ',
&        'complex power with non-zero imaginary part'
    else
      print *, 'a zero base cannot be raised to a ',
&        'complex power with non-positive real part'
    endif
  end select
end if
__zmath2err = (0,0)
end
```

```
double complex function __qmath2err( err_info, arg1, arg2 )
integer err_info
double complex arg1, arg2
include 'mathcode.fi'
arg1 = arg1      ! to avoid unreferenced warning message
arg2 = arg2      ! to avoid unreferenced warning message
if( ( err_info .and. M_DOMAIN ) .ne. 0 )then
  select( err_info .and. FUNC_MASK )
  case( FUNC_POW )
    ! arg1 is (0,0)
    if( imag( arg2 ) .ne. 0 )then
      print *, 'a zero base cannot be raised to a',
&      ' complex power with non-zero imaginary part'
    else
      print *, 'a zero base cannot be raised to a',
&      ' complex power with non-positive real part'
    endif
  end select
end if
__qmath2err = (0,0)
end
```

37.5 The compiler cannot find my include files

In order to locate your INCLUDE files, the compiler first searches your current directory, then each directory listed in the **FINCLUDE** environment variable (in the order that they are specified). If the compiler reports that it is unable to find one of your include files, change the **FINCLUDE** environment variable by adding the directory path to your include files. For more information on setting the environment variable, refer to the "Compiling an Application" chapter of the *Open Watcom F77 User's Guide*.

37.6 The linker reports a "stack segment not found" error

The linker usually reports the error "1014: stack segment not found" when it is unable to find the run-time libraries required to link your application. To ensure you are linking with the correct run-time libraries, check to see that your link command contains the correct "SYSTEM" directive. As well, the **WATCOM** environment variable should be pointing to the directory containing the Open Watcom F77 software. For a Open Watcom F77 application, if this error is followed by the error "1023: no starting address found", it may indicate that you are attempting to link code that does not have a main program procedure. Ensure that you include your main program object module with your linker directives.

With FORTRAN 77, "STACK" is a reserved word. If you use "STACK" as the name of a common block, this may also result in the "Stack Segment Not Found" error. Check the names of your common blocks and rename them if necessary.

37.7 Resolving an "Undefined Reference" linker error

The Open Watcom Linker builds an executable file by a process of resolving references to functions or data items that are declared in other source files. Certain conditions arise that cause the linker to generate an "Undefined Reference" error message. An "Undefined Reference" error message will be displayed by the linker when it cannot find a function or data item that was referenced in the program. Verify that you have included all the required object modules in the linker command and that you are linking with the correct libraries.

The "SYSTEM" linker directive should be used to indicate the target environment for the executable. This directive specifies the format of the executable and the libraries for the target environment. Verify that the **WATCOM** environment variable is set to the directory that Open Watcom F77 was installed in since it is used to complete the library path in the "SYSTEM" directive. You may also explicitly include a library using the "LIBRARY" linker directive.

If the linker reports an unresolved reference for "_cstart_", this indicates that the linker could not find the FORTRAN 77 run-time libraries. In 16-bit applications, the FORTRAN 77 run-time libraries for the medium memory model (-mm) and the floating-point calls floating-point model (-fpc) would be `flibm.lib`. In 32-bit applications, the FORTRAN 77 run-time libraries for the flat memory model would be `flib.lib`. Verify that the "LIB" environment variable has been set to point to the correct **WATCOM** library directories and that the library corresponds to the memory and floating-point model that you selected.

37.8 Why local variable values are not maintained between subprogram calls

By default, the local variables for a subprogram are stored on the stack and are not initialized. When the subprogram returns, the variables are popped off the stack and their values are lost. If you want to preserve the value of a local variable, after the execution of a RETURN or END statement in a subprogram, the FORTRAN 77 SAVE statement or the "save" compiler option can be used.

Using the FORTRAN 77 SAVE statement in your program allows you to explicitly select which values you wish to preserve. The SAVE statement ensures that space is allocated for a local variable from static memory and not the stack. Include a SAVE statement in your FORTRAN 77 code for each local variable that you wish to preserve.

To automatically preserve all local variables, you can use the "save" compiler option. This option adds code to initialize and allocate space for each local variable in the program. This is equivalent to specifying a SAVE statement. The "save" option makes it easier to ensure that all the variables are preserved during program execution, but it increases the size of the code that is generated. You may wish to use this option during debugging to help diagnose bugs caused by corrupted local values. Usually, it is more efficient to use SAVE statements rather than the general "save" compiler option. You should selectively use the SAVE statement for each subprogram variable that you want to preserve until the next call. This leads to smaller code than the "save" option and avoids the overhead of allocating space and initializing values unnecessarily.

37.9 What "Stack Overflow!" means

The memory used for local variables is allocated from the function call stack although the Open Watcom compilers will often use registers for local variables. The size of the function call stack is limited at link-time and it is possible to exceed the amount of stack space during execution.

There are various ways of protecting against stack overflow errors. First, one should minimize the number of recursive functions used in an application program. This can be done by recoding recursive functions to use loops.

The user may also optionally force the compiler to use static storage for all local variables (Open Watcom F77 "save" option). This would eliminate most stack problems for FORTRAN programs. These techniques will reduce the amount of stack space required but there still may be times where the default

amount of stack space is insufficient. The Open Watcom Linker (WLINK) allows the user to set the amount of stack space at link-time through the directive "OPTION STACK=size" where size may be specified in bytes with an optional "k" suffix for kilobytes (1024 bytes).

Example:

```
option stack=9k
```

Note that with the Open Watcom F77 run-time system, the I/O routines require 4k of stack space. If your application requires 5K of stack space, set aside 9K to allow for 4K of I/O stack space in addition to the stack space required by the application.

Debugging a program that reports a stack overflow error can be accomplished with the following sequence.

1. Load your application into the debugger
2. Set a breakpoint at `__STKOVERFLOW`
3. Run the application until the breakpoint at `__STKOVERFLOW` is triggered
4. Issue the debugger "show calls" command. This will display a stack traceback giving you the path of calls that led up to the stack overflow situation.

The solution to the stack overflow problem at this point depends on the programmer.

37.10 What are the probable causes of a General Protection Fault in 32-bit applications?

If you are running a 32-bit application using DOS/4GW, a program crash may report an "Unexpected Interrupt 0D" general protection fault error. The Phar Lap DOS extender would report an "Abnormal Program Termination" general protection fault error. This often indicates that something in your program has tried to access an invalid memory location. In a Open Watcom F77 application, the most likely causes of a general protection fault are:

- Attempting to access an array out of bounds.
- Running out of stack space.
- Passing incorrect parameter types to a function.

To help locate the cause of the protection fault, compile your program with the "debug" and "stack" options. With these options, code will be added to your application to help identify these problems and generate run-time error messages when they are encountered. In addition, the "stack" option checks for stack overflow conditions by including code at the beginning of each subprogram.

If you still encounter general protection faults after compiling with "debug" and "stack", then debug the program using the debugger. This will help to identify the location of the crash and the state of your parameters and variables at the time of the crash.

37.11 Which floating-point compiler option should I use for my application?

The answer to this question depends on the expected target machines for your application. If you know that a co-processor will be available, use the "fpi87" compiler option to optimize the run-time performance of the application.

When you are running a FORTRAN 77 application on a machine with or without a co-processor and you want to favour the use of emulation libraries over code size, use the "fpc" option. The "fpc" option will also take advantage of an 80x87 co-processor if it is available. If your application needs to be flexible enough to run either with or without a co-processor, the "fpc" option is recommended.

The "fpi" option is the default floating-point option and can be used with or without a co-processor. On machines that do not have a co-processor, you may notice that programs compiled using "fpc" run faster than those compiled with "fpi". This occurs because the "fpc" option uses the floating-point libraries directly whereas the "fpi" option interfaces with an emulator interrupt handler. Although the "fpi" option is slower than "fpc" without a co-processor, the code that it generates is smaller.

When you are running an application that has been compiled with "fpi", the startup code checks to determine whether a math co-processor is present. If it is not present, the emulator hook is installed at the INT 7h interrupt to manage the co-processor requests and convert them to the emulation library calls. Each time a floating-point operation is requested, the processor issues an INT 7h.

For 16-bit applications, the interrupt handler overhead accounts for the performance discrepancy between the "fpc" and "fpi" options.

For 32-bit applications, the manner in which this interrupt is handled depends on the DOS extender. Depending on the DOS extender, there are two methods of managing floating-point instructions through the interrupt handler. The DOS extender will either pass the interrupt directly to the INT 7h handler or it will perform some intermediary steps. Similarly, there is a delay after the interrupt as control is passed back through the DOS extender. Passing floating-point handling from the DOS extender to the interrupt handler results in the performance degradation. This performance degradation may vary across DOS extenders. It is the overhead of transferring the call through an interrupt that leads to the speed difference between "fpi" and "fpc". If you need to run an application on machines without math co-processors, and you want to ensure that your performance is optimal, build your application using the "fpc" option rather than "fpi".

In a Windows environment, both the "fpi87" and the "fpi" options will use floating-point emulation if a co-processor is not available. Windows floating-point emulation is provided through Open Watcom's "WEMU387.386". "WEMU387.386" is royalty free and may be redistributed with your application. For machines that do not have a math co-processor, install "WEMU387.386" as a device in the [386Enh] section of the Windows SYSTEM.INI file to handle the floating-point operations. Note that the speed of code using "WEMU387.386" on machines without a co-processor will be much slower than code compiled with the "fpc" option that always uses floating-point libraries.

37.12 How more than 20 files at a time can be opened

The number of file handles allowed by Open Watcom F77 is initialized to 20 in `stdio.h`, but this can be changed by the application developer. To change the number of file handles allowed with Open Watcom F77, follow the steps outlined below.

1. Let `n` represent the number of files the application developer wishes to have open. Ensure that the `stdin`, `stdout`, `stderr`, `stdaux`, and `stdprn` files are included in the count.
2. Change the `CONFIG.SYS` file to include "files=`n`" where "`n`" is the number of file handles required by the application plus an additional 5 handles for the standard files (this applies to DOS 5.0). The number "`n`" may vary depending on your operating system and version. If you are running a network such as Novell's NetWare, this will also affect the number of available file handles. In this case, you may have to increase the number specified in the "files=`n`" statement.
3. Add a call to `GROWHANDLES` in your application.

The following example illustrates the use of `GROWHANDLES`.

Example:

```
* FHANDLES.FOR
*
* This FORTRAN program grows the number of file handles so
* more than 16 files can be opened. This program
* illustrates the interaction between GROWHANDLES and
* the DOS 5.0 file system. If you are running a network
* such as Novell's NetWare, this will also affect the
* number of available file handles. In the actual trial,
* FILES=40 was specified in CONFIG.SYS.

* Compile/Link: set finclude=\watcom\src\fortran
*               wfl[386] fhandles

* Get proper typing information from include file
* include 'fsublib.fi'

integer i, j, maxh, maxo
integer tmpfile
integer units(7:57)

do i = 25, 40
*   Count 5 for stdin, stdout, stderr, stdaux,
*   and stdprn
  print 100, 5 + i
  maxh = growhandles( 5 + i )
  print *, 'Growhandles=', maxh
  maxo = 0
```

```

        do j = 7, 7 + maxh
            print *, 'Attempting file', j
            units(j) = tmpfile(j)
            if( units(j) .eq. 0 )goto 10
            maxo = maxo + 1
        enddo

10      print 101, maxo, maxh
        do j = 7, 7 + maxo
            close( units(j) )
        enddo
    enddo

100     format('Trying for ',I2,' handles... ', $)
101     format(I2,'/',I2,' temp files opened')
    end

integer function tmpfile( un )
integer un, ios
open( unit=un, status='SCRATCH', iostat=ios )
if( ios .eq. 0 )then
    write( unit=un, fmt='(I2)', err=20 ) un
    tmpfile = un
    return
endif
20     tmpfile = 0
end
```

37.13 How source files can be seen in the debugger

The selection and use of debugging information is important for getting the most out of the Open Watcom Debugger. If you are not able to see your source code in the Open Watcom Debugger source window, there are three areas where things may have gone wrong, namely:

1. using the correct option for the Open Watcom F77.
2. using the correct directives for the Open Watcom Linker.
3. using the right commands in the Open Watcom Debugger.

The Open Watcom F77 compiler takes FORTRAN 77 source and creates an object file containing the generated code. By default, no debugging information is included in the object file. The compiler will output debugging information into the object file if you specify a debugging option during the compile. There are two levels of debugging information that the compiler can generate:

1. Line numbers and local variables ("d2" option)
2. Line numbers ("d1" option)

The options are used to determine how much debugging information will be visible when you are debugging a particular module. If you use the "d2" option, you will be able to see your source file and display your local variables. The "d1" option will display the source but will not give you access to local variable information.

The Open Watcom Linker (WLINK) is the tool that puts together a complete program and sets up the debugging information for all the modules in the executable file. There is a linker directive that indicates to

the linker when it should include debugging information from the modules. There are five levels of debugging information that can be collected during the link. These are:

1. global names (DEBUG)
2. global names, line numbers (DEBUG LINE)
3. global names, types (DEBUG TYPES)
4. global names, local variables (DEBUG LOCALS)
5. all of the above (DEBUG ALL)

Notice that global names will always be included in any request for debugging information. The debugging options can be combined

```
DEBUG LINE, TYPES
```

with the above directive resulting in full line number and typing information being available during debugging. The directives are position dependent so you must precede any object files and libraries with the debugging directive. For instance, if the file `mylink.lnk` contained:

```
#
# invoke with: wlink @mylink
#
file main
debug line
file input, output
debug all
file process
```

then the modules `input` and `output` will have global names and source line information available during debugging. All debugging information in the module `process` will be available during debugging.

Hint: A subtle point to debugging information is that all the modules will have global names available if any debugging directive is used. In the above example, the module `main` will have global name information even though it does not have a `DEBUG` directive preceding it.

It is preferable to have one `DEBUG` directive before any `FILE` and `LIBRARY` directives. You might wonder if this increases the size of the executable file so that it will occupy too much memory during debugging. The debugging information is loaded "on demand" by the debugger during the debugging session. A small amount of memory (40k default, selectable with the Open Watcom Debugger "dynamic" command line option) is used to hold the most recently used module debugging information. In practice, this approach saves a lot of memory because most debugging information is never used. The overhead of accessing the disk for debugging information is negligible compared to accessing the source file information. In other words, you can have as much debugging information as you want included in the executable file without sacrificing memory required by the program. See the section entitled "The `DEBUG` Directive" in the *Open Watcom Linker User's Guide* for more details.

If the previous steps have been followed, you should be well on your way to debugging your programs with source line information. There are instances where the Open Watcom Debugger cannot find the appropriate source file even though it knows all the line numbers. The problem that has surfaced involves how the source file is associated with the debugging information of the module. The original location of the source file is included in the debugging information for a module. The name that is included in the debugging information is the original name that was on the Open Watcom F77 command line. If the original filename is no longer valid (i.e., you have moved the executable to another directory), the Open Watcom Debugger must be told where to find the source files. The Open Watcom Debugger "Source Path"

menu item (under "File") can be used to supply new directories to search for source files. If your source files are located in two directories, the following paths can be added in the Open Watcom Debugger:

```
c:\program\fortran\*.for  
c:\program\new\fortran\*.for
```

The "*" character indicates where the module name will be inserted while the Open Watcom Debugger is searching for the source file. See the description of the "Source Path" menu item in the *Open Watcom Debugger User's Guide* for more details.

37.14 The difference between the "d1" and "d2" compiler options

The reason that there are two levels of debugging information available is that the code optimizer can perform many more optimizations and still maintain "d1" (line) information. The "d2" option forces the code optimizer to ensure that any local variable can be displayed at any time in the function.

The "d2" option will always generate code and debugging information so that you can print the value of any variable during the execution of the function. In order to get the best code possible and still see your source file while debugging, the "d1" option only generates line number information into the object file. With line number information, much better code can be generated. The debugging of programs that have undergone extensive optimization can be difficult, but with the source line information it is much easier. To summarize, use the "d2" compiler option if you are developing a module and you would like to be able to display each local variable. The "d1" compiler option will give you line number information and the best generated code possible. There is absolutely no reason not to specify the "d1" option because the code quality will be identical to code generated without the "d1" option.

37.15 The difference between the "debug" and "d2" compiler options

The "d2" (and "d1") compiler options are used to add debugging information to your executable. The "d2" option makes line numbering, local symbol and typing information available to the debugger whereas "d1" only provides line number debugging information to the debugger. This information is used during a debugging session to examine the state of variables and to provide the source code display.

The "debug" option provides run-time error messages that are independent of the Open Watcom Debugger. The "debug" option causes the generation of run-time error checking code. This includes subscript and substring bounds checking as well as code that allows a run-time traceback to be issued when an error occurs. During the execution of the application, if an error occurs, the code added with the "debug" option will halt the program and provide an informative error message.

1

16-bit 123
 16-bit DLL 151
 16-bit DOS applications 5
 16-bit far pointer 123
 16-bit near pointer 123
 16-bit OS/2 1.x applications 209
 16-bit Windows 3.x applications 103
 16-bit Windows 3.x non-GUI applications 107
 _16xxx functions 183

3

32-bit 123
 32-bit DLL 139, 151
 32-bit DOS/4GW applications 13
 32-bit far pointer 123
 32-bit gates 40
 32-bit near pointer 123
 32-bit OS/2 applications 213
 32-bit Phar Lap 386/DOS-Extender applications 9
 32-bit Windows 3.x applications 113
 32-bit Windows 3.x non-GUI applications 117
 386enh 130

4

4GWPRO.EXE 95

8

8042 auxiliary processor 33

A

A20 line 33-34
 address line 20 34
 AllocAlias16 **160**, 132-133, 161, 168, 178
 ALLOCATE 124, 132, 134, 144, 164
 AllocHugeAlias16 **161**, 132, 161, 169, 178
 answers to general problems 249
 API special functions 183
 application development 1
 array 124-125, 132
 autopassup range 40

B

beginthread function 195, 217
 binary data 251
 binding 32-bit applications 114, 119
 binding a 32-bit DLL 115, 120
 BINP directory 211
 BINW directory 115, 120
 building 386/DOS-Extender applications 10
 building DOS applications 5
 building DOS/4GW applications 13
 building OS/2 1.x applications 210
 building OS/2 applications 214
 building Windows 3.x applications 104, 114
 building Windows NT applications 192

C

_Call16 **162**, 137, 172, 181-182
 cdecl 137, 162, 166
 class 131
 common questions 249
 DOS/4GW 85
 Compaq 386 memory 32
 CONFIG.SYS 211
 converting to Open Watcom F77 249
 common problems 249
 what you need to know 249
 cstart 258

D

- d1 250
- d1 versus d2 264
- d2 250
- d2 versus debug 264
- DEALLOCATE 134
- DEBUG option 259
- debug versus d2 264
- debugger option
 - d1 264
 - d2 264
- debugging 250
 - optimized programs 250
 - stack overflow 258
 - techniques 258
- debugging 386|DOS-Extender applications 10
- debugging DOS applications 6
- debugging DOS/4GW applications 14
- debugging information
 - global variables 262
 - line numbering 262
 - local variables 262
 - Open Watcom Debugger 263
 - Open Watcom F77 262
 - source file 262
 - types 262
 - WLINK 262
- debugging Non-GUI 16-bit Windows 3.x applications 109
- debugging Non-GUI 32-bit Windows 3.x applications 120
- debugging OS/2 1.x applications 210
- debugging OS/2 applications 214
- debugging Windows 3.x applications 104, 115
- debugging Windows NT applications 192
- default type 129
- default windowing library functions 110, 121
- DefineDLLEntry **164**
- DefineUserProc16 **166**, 174
- DELETESWAP virtual memory option 35, 93-95
- distribution rights 130
- DLL
 - 16-bit 151
 - 16-bit calls into 32-bit DLLs 141
 - 32-bit 139, 151
 - 32-bit calls into 32-bit DLLs 143
 - 32-bit Windows example 140
 - creating 147-148
 - debugging 147
 - debugging example 149
 - installing example 148
 - OS/2 2.x 221
 - passing information in a structure 144
 - running example 149
 - summary 150
 - Windows NT 199
- DLL access
 - OS/2 2.x 223
 - Windows NT 202
- DLL creation
 - OS/2 2.x 221
 - Windows NT 199
- DLL directory 211
- DLL initialization
 - OS/2 2.x 225
 - Windows NT 204
- DLL sample
 - OS/2 2.x 222
 - Windows NT 200
- DLL termination
 - OS/2 2.x 225
 - Windows NT 204
- DLL_CHAR 164
- DLL_DWORD 164
- DLL_ENDLIST 164
- DLL_PTR 164
- DLL_WORD 164
- DOS extenders
 - common problems 17
- DOS file I/O 92
- DOS memory management 49
- DOS Protected-Mode Interface 43
- DOS/4GW
 - 4GWPRO.EXE 95
 - address line 20 34
 - asynchronous interrupts 89
 - cannot lock stack 100
 - chaining handlers 41
 - code and data addresses 88
 - common questions 85
 - Ctrl-Break handling 90
 - debugger version 95
 - debugging bound applications 95
 - demand-loading 94
 - differences with DOS/4G 86
 - differences with Professional version 85
 - documentation 85
 - DOS file I/O 92
 - DOSX.EXE 100
 - DPMI support 87
 - EMM386.EXE 99
 - error messages 80
 - extender messages 77
 - extra memory 32

- int 70h-77h 90
- interrupt handler address 41
- interrupt handlers 41, 90
- kernel error messages 77
- linear vs physical addresses 89
- locking memory 90
- Lotus 1-2-3 99
- low memory access 88
- memory addressability 92
- memory control 31
- memory range 31
- memory use 26
- mouse support 92
- NULL pointer references 89
- OS/2 bug 99
- out of memory 93
- pointers vs linear addresses 88
- realloc 92
- register dump 97
- runtime options 33
- spawning 92
- switch mode setting 30
- TCPIP.EXE 100
- transfer stack overflow 97
- TSR not supported 23
- unexpected interrupt 96
- utilities 69
- VESA support 93
- VM configuration 94
- VMM 93
- VMM instability 93
- VMM restriction 23
- Windows NT bug 100
- DOS/4GW DOS extender 23
- DOS16M
 - + option 32
 - A20 option 34
 - loops option 34
 - runtime options 33
- DOS16M environment variable 29-34, 73
- DOS4G
 - NULLP option 29, 89
 - QUIET option 29
 - VERBOSE option 29, 97
- DOS4G environment variable 29
- DOS4GPATH environment variable 25
- DOS4GVM
 - DELETESWAP 93-95
 - MAXMEM 95
 - MINMEM 95
 - SWAPINC 93, 95
 - SWAPMIN 93, 95
 - SWAPNAME 94
 - VIRTUALSIZE 93, 95
- DOS4GVM environment variable 35-36
- DOS4GVM.SWP 35
- DOS4GW 70
- DOS4GW.EXE 25
- DOSX.EXE 100
- DPMI 32, 40, 43
 - allocate DOS memory block 49
 - allocate memory block 62
 - allocate real-mode callback address 56
 - demand paging 63
 - discard page 64
 - free DOS memory block 50
 - free memory block 62
 - free physical address mapping 65
 - free real-mode callback address 60
 - function calls 44
 - get and disable virtual interrupt state 66
 - get and enable virtual interrupt state 66
 - get API entry point 67
 - get coprocessor status 67
 - get DPMI version 60
 - get exception handler vector 51
 - get free memory information 61
 - get page size 63
 - get protected-mode interrupt vector 52
 - get real-mode interrupt vector 51
 - get virtual interrupt state 66
 - lock linear region 63
 - mark page 63
 - physical address mapping 64
 - resize DOS memory block 50
 - resize memory block 62
 - set coprocessor emulation 68
 - set exception handler vector 51
 - set protected-mode interrupt vector 52
 - set real-mode interrupt vector 51
 - simulate real-mode far call 55
 - simulate real-mode interrupt 54
 - simulate real-mode iret call 56
 - unlock linear region 63
 - vendor extensions 67
 - virtual interrupt state 65
- DPMI host
 - 386Max 43
 - OS/2 VDM 43
 - QEMM QDPMI 43
 - Windows 3.1 43
- DPMI specification 17, 85
- DPMI_MEMORY_LIMIT
 - DOS setting 100
- dragon book 250
- dwfDeleteOnClose 110, 121
- dwfSetAboutDlg 110, 121
- dwfSetAppTitle 111, 122

dwfSetConTitle 111, 122
dwfShutDown 111, 122
dwfYield 111, 122
DWORD 137
dynamic link libraries 211
 OS/2 2.x 221
 Windows NT 199
dynamic link library 139, 151
dynamic link library access
 OS/2 2.x 223
 Windows NT 202
dynamic link library creation
 OS/2 2.x 221
 Windows NT 199
dynamic link library initialization
 OS/2 2.x 225
 Windows NT 204
dynamic link library sample
 OS/2 2.x 222
 Windows NT 200
dynamic link library termination
 OS/2 2.x 225
 Windows NT 204
dynamic linking 199, 221

E

EMM386.EXE 99
endthread subroutine 196, 218
EnumChildWindows 174
EnumFonts 174
EnumMetaFile 174
EnumObjects 174
EnumProps 174
EnumTaskWindows 174
EnumWindows 174
environment variables
 DOS16M 29-34, 73
 DOS4G 29
 DOS4GPATH 25
 DOS4GVM 35-36
 FINCLUDE 151, 198, 219, 257
 PATH 115, 119-120, 151
 WATCOM 115, 120, 198, 219, 257-258
 WINDOWS_INCLUDE 157
error messages
 DOS/4GW 80
 kernel 77
executable
 linear 25

 segmented 25
executable file 6, 10, 14, 104, 109, 114, 119, 192, 210,
 214
EXPLICIT option 130
extended memory 29
extender messages
 DOS/4GW 77

F

far 123-125, 132, 144, 164, 175, 246
far pointer 123
__fdll_initialize_ 204, 225
__fdll_terminate_ 204, 225
files
 more than 20 261
 unable to find 257
FINCLUDE environment variable 151, 198, 219, 257
fixed record type 251
floating-point options 260
formatted 251
FORTRAN 77
 Extensions 249
fpc option 260
fpi option 260
fpi87 option 260
free 172
free memory 19
FreeAlias16 **168**, 133
FreeHugeAlias16 **169**, 169
FreeIndirectFunctionHandle **170**, 162, 172, 181-182
FreeLibrary 201
FreeProcInstance 135
fsignal 251
__fthrd_initialize_ 204
__fthrd_terminate_ 204
Fujitsu FMR-70 switch mode setting 30
FWinMain 139

G

GetIndirectFunctionHandle **172**, 162, 170, 176,
 181-182
GetProcAddress **174**, 132, 134, 166, 180, 182
GETPROC_ABORTPROC 174
GETPROC_CALLBACK 174
GETPROC_ENUMCHILDDWINDOWS 174

GETPROC_ENUMFONTS 174
 GETPROC_ENUMMETAFILE 174
 GETPROC_ENUMOBJECTS 174
 GETPROC_ENUMPROPS_FIXED_DS 174
 GETPROC_ENUMPROPS_MOVEABLE_DS 174
 GETPROC_ENUMTASKWINDOWS 174
 GETPROC_ENUMWINDOWS 174
 GETPROC_GLOBALNOTIFY 174
 GETPROC_GRAYSTRING 174
 GETPROC_LINEDDA 174
 GETPROC_SETRESOURCEHANDLER 174
 GETPROC_SETTIMER 174
 GETPROC_SETWINDOWSHOOK 174
 GETPROC_USERDEFINED_1 166
 GETPROC_USERDEFINED_32 166
 GETPROC_USERDEFINED_x 174
 GetProcAddr 182
 GetProcAddress 137, 162, 181
 GlobalAlloc 134
 GlobalLock 182
 GlobalNotify 174
 GMEM_DDESHARE 134
 GrayString 174
 GROWHANDLES 261
 GWL_WNDPROC 136

H

HIMEM.SYS 33
 Hitachi B32 switch mode setting 30

I

IBM PS/55 switch mode setting 30
 IDT 40
 import definitions 199, 221
 import library 203, 223
 INDIR_CDECL 172
 INDIR_CHAR 172
 INDIR_DWORD 172
 INDIR_ENDLIST 172
 INDIR_PTR 172, 176
 INDIR_WORD 172
 initialization
 OS/2 2.x dynamic link library 225
 Windows NT dynamic link library 204
 initializing

 variables 258
 Instant-D 25
 INT 21H 37
 INT 31H 43
 int 31H function calls 44
 inter-language calls 239
 interrupt handling 40
 interrupt services 51
 interrupts
 using DOS/4GW 18
 InvokeIndirectFunction **176**, 172, 181-182
 InvokeIndirectFunctionHandle 162

K

kernel error messages 77
 keyboard status 33

L

LDT 44
 LE format 25
 LibMain 200-202, 222
 library 258
 library functions
 default windowing 110, 121
 line number information 250
 linear executable 25
 LineDDA 174
 linker
 undefined references 257
 LoadLibrary 181, 201
 local descriptor table 44
 LocalAlloc 134
 LOCATION 124, 132, 144, 164
 LOCATION= 124, 132
 Lotus 1-2-3 99

M

MakeProcInstance 134-135
 malloc 172
 MapAliasToFlat **178**

- math errors 251
- MAXMEM virtual memory option 35, 95
- memory management services 61
- memory transfer rate 71
- memory wait states 71
- message
 - include files 257
 - no starting address found 257
 - stack segment not found 257
 - unable to find files 257
 - undefined references 257
- messages
 - DOS/4GW 77
- MINMEM virtual memory option 35, 95
- mixed-language programming 239
 - argument passing 240
 - common blocks 245
 - integer type 241
 - linking issues 241
 - memory models 240
 - parameter passing 240
 - passing integers 241-242
 - passing strings 243-244
 - symbol names 239
 - variable number of arguments 246
- mode switching
 - basis 74
 - performance 71
- multi-threaded applications 195, 217
 - OS/2 2.x 217
 - Windows NT 195
- multi-threading issues
 - OS/2 2.x 217
 - Windows NT 195

N

- NE format 25
- near 123
- near pointer 123
- NEC 98-series switch mode setting 30
- NLM
 - debugging 235
 - header files 235
 - import libraries 235
 - libraries 235
 - sampler 235
- NLM support
 - version 4.0 235
 - version 4.1 235

- no starting address found 257
- NOAUTOPROCS 135-136
- NOCOVERSENDS 133
- Novell
 - TCPIP.EXE 100
- NT development 189
- NULLP 29

O

- object file 6, 10, 14, 104, 109, 114, 119, 192, 210, 214
- OKI if800 switch mode setting 30
- Open Watcom F77
 - converting to 249
- Open Watcom F77 debugging
 - d2 264
 - debug 264
- Open Watcom F77 options
 - d1 250, 262
 - d2 250, 262
- Open Watcom Strip Utility 250
- opening more than 20 files 261
- optimization
 - suggested reading 250
 - what you should know 250
- OS/2
 - fullscreen application 209, 213
 - PM-compatible application 209, 213
 - Presentation Manager application 209, 213
- OS/2 PM
 - API calls 228
 - non-GUI applications 227
 - non-GUI example 227
- OS/2 Presentation Manager 227

P

- page locking services 62
- page tuning services 63
- PASCAL 137, 162, 166
- PASS_WORD_AS_POINTER **179**
- PATH environment variable 115, 119-120, 151
- performance 72
- Phar Lap TNT 189
- PMINFO 31, 71
- pointers
 - 16-bit 123

32-bit 123
 far 123
 near 123
 pragma 250
 PRINT 195, 217
 private memory pool 73
 PRIVATXM 32, 73, 99
 PROCPTR 182
 protected mode 33
 PS/2 switch mode setting 30

Q

questions 249
 QUIET 29

R

real mode 33
 ReleaseProc16 **180**
 resource compiler 115, 120
 RMINFO 74

S

SAVE 249, 258
 segmented executable 25
 SendDlgItemMessage 133
 SendMessage 133
 sequential 251
 SetResourceHandler 174
 SetTimer 174
 setvbuf 92
 SetWindowLong 136
 SetWindowsHook 174
 sieve 103, 108, 113, 118, 191
 sieve program 5, 9, 13, 209, 213
 STACK option 259
 stack overflow 258
 stack segment not found 257
 static linking 199, 221
 stub program 25, 70
 supervisor 114, 119

SWAPINC virtual memory option 35, 93, 95
 SWAPMIN virtual memory option 35, 93, 95
 SWAPNAME virtual memory option 35, 94
 switch mode setting
 Fujitsu FMR-70 30
 Hitachi B32 30
 IBM PS/55 30
 NEC 98-series 30
 OKI if800 30
 PS/2 30
 switching modes
 performance 71
 symbolic information 250
 system 258
 system configuration file 211
 SYSTEM.INI 130

T

TCPIP.EXE 100
 termination
 OS/2 2.x dynamic link library 225
 Windows NT dynamic link library 204
 thread creation
 OS/2 2.x 217
 Windows NT 195
 thread example
 OS/2 2.x 218
 Windows NT 196
 thread identifier
 OS/2 2.x 218
 Windows NT 196
 thread limits
 OS/2 2.x 220
 thread termination
 OS/2 2.x 218
 Windows NT 196
 threadid function 196, 218
 threads of execution 195, 217
 TNT 189
 transfer rate
 memory 71
 translation services 53

U

UDP16_CDECL 166

UDP16_CHAR 166
UDP16_DWORD 166
UDP16_ENDLIST 166
UDP16_PTR 166
UDP16_WORD 166
unable to find files 257
undefined references 257
 cstart 258
Unexpected Interrupt 259

V

variables
 set to zero 258
VCPI 32
VERBOSE 29
video memory 17
virtual memory manager 35, 93
VIRTUALSIZE virtual memory option 35, 93, 95
Visual Basic 151
 16-bit DLL 154-155
 32-bit DLL 154
 building examples 156
 example 152
 Version 3.0 151
VMC extension 36
VMM 35, 93

W

W386DLL.EXT 115, 120
WATCOM environment variable 115, 120, 198, 219,
 257-258
WBIND 114-115, 119, 181
WBIND.EXE 114, 119
WDEBUG.386 130
WEMU387.386 130
WFL 6-7, 104-105, 109-110, 210-211
WFL386 10-11, 14-15, 115-116, 120-121, 192-193,
 214-215
WIN386 library routines 159
WIN386.EXT 114-115, 119-120
Win386LibEntry 151
WINAPI.FI 129, 133
WINDLG.FI 129
windowed applications
 default windowing environment 107, 117

Windows
 binding 32-bit applications 114, 119
Windows 3.x extender 123
 _16xxx functions 182-183
 32-bit callback routines 182
 calling 16-bit code 181
 components 126
 creating applications 127
 floating-point 130
 function pointers 182
 multiple instances 130
 overview 124
 pointer conversion 132
 pointer handling 131
 pointers 123
 programming notes 129
 questions 181
 resources 181
 special functions 183
 structure 125
Windows API 129
Windows NT 189
 character-mode applications 189
 GUI applications 189
 programming overview 189
Windows NT Character-mode applications 191
Windows supervisor 114, 119
WINDOWS.FI 129
__WINDOWS_386__ 157
__WINDOWS__ 157
WINDOWS_INCLUDE environment variable 157
WINFONT.FI 130
WSTUB.C 25

X

XMS 33